


Huge Ensembles of Weather Extremes using NVIDIA's Fourier Forecasting Neural Network (FourCastNet)



William D. Collins
and Ankur Mahesh
Berkeley Lab and UC Berkeley
NERSC
NVIDIA

Outline

- Low-likelihood High Impact Extremes (LLHIs)
- Background: data-driven forecasting
- FourCastNet model
 - Forecast results on extremes
 - Architecture & training strategy
- Ongoing/future work
 - Construction of Huge Ensembles
 - Operationalizing data-driven weather & climate Prediction



Low-Likelihood High Impact Extremes

IPCC AR6:


... In summary, the future occurrence of LLHIs events linked to climate extremes is generally associated with low confidence, but cannot be excluded, especially at global warming levels above 4°C.



Types of LLHIs to be investigated


- Heatwaves
- Atmosphere rivers
- Tropical cyclones and hurricanes
- Extreme downpours and flooding
- Co-occurring extreme events
- Western US hydroclimate extreme events (droughts, floods, mountain snow, wildfires)

Low-Likelihood High Impact Extremes




Anthropogenic Contributions to the 2021 Pacific Northwest Heatwave

Emily Bercos-Hickey¹ , Travis A. O'Brien^{1,2} , Michael F. Wehner³ , Likun Zhang^{1,4} , Christina M. Patricola^{1,5} , Huanping Huang^{1,6} , and Mark D. Risser¹



Low-Likelihood High Impact Extremes



Multiple ARs landing in California Dec-Jan, 2023




Why LLHIs? From the IPCC AR6 WG1 SPM....



- ▶ “**Low-likelihood outcomes**, such as *ice-sheet collapse, abrupt ocean circulation changes, some compound extreme events, and warming substantially larger than the assessed very likely range of future warming*, **cannot be ruled out and are part of risk assessment.**”
- ▶ “**Low-likelihood, high-warming outcomes are associated with potentially very large impacts**, such as *through more intense and more frequent heatwaves and heavy precipitation, and high risks for human and ecological systems, particularly for high GHG emissions scenarios.*”
- ▶ “*Low-likelihood, high-impact outcomes could occur at global and regional scales even for global warming within the very likely range for a given GHG emissions scenario. The probability of low-likelihood, high-impact outcomes increases with higher global warming levels (high confidence).*”


Numerical Weather Prediction (NWP)

- Urgency under climate change:
 - Reducing uncertainties, predicting extremes, disaster mitigation, etc....
- Entire earth system: complex phenomena across wide range of physical scales
- Traditional NWP consumes substantial HPC resources
 - Dedicated machines running physical models + data assimilation 24/7





Ingredients of NWP

- PDEs w/ hundreds of variables, complex parameterizations for subgrid and multi-physics processes
- High resolution grids (computational cost scales as \sim fourth power)
- Large ensembles for UQ and long-term forecasting



Deterministic Numerical Weather Prediction (NWP)


Solving equations of motion for an incompressible fluid on a rotating sphere



HPC-intensive: European Center for Medium Range Weather Forecasting

The slow but steady evolution of classical NWP


Data assimilation and other advances have slowly revolutionized the quality & accuracy over decades




Bauer et al., *Nature*, 2015

Caveat: These advances are the foundation for assimilated state estimates that ML approaches rely on.



Climate Phenomena and Models are Inherently Multiscale



©The COMET Program




Faster computers have bridged 4 orders of magnitude, out of 16



V. Balaji / GFDL

Enabling Technology: DOE's GPU Exascale Systems



2021: 100 petaFLOPS/s (NERSC)





2023: ~1 exaFLOPS/s (ALCF)


2022: ~1.7 exaFLOPS/s (OLCF)




Gains in Resolution for Climate Modeling



Climate Phenomena and Models are Inherently Multiscale




©The COMET Program




Fully Data-Driven Weather Prediction (DDWP)

A modern alternative, intriguingly complementary




- Mining the data assimilated states from operational NWP
- # training samples = length of satellite record (~ 15k days)
- Can be stood up by small teams within tech companies.
- Is producing skill gains rapidly.

FourCastNet: DDWP at very high resolution



FourCastNet, Pathak et al. (2022), 0.25° , ~1,000,000 Pixels, ViT+AFNO

GNN, Keisler et al. (2022), 1° , 64,000 Pixels, Graph Neural Networks



DLWP, Weyn et al. (2020). 2° , 16K pixels, Deep CNN on Cubesphere/(2021) ResNet



Weyn et al. (2019), 2.5° N.H only, 72x36, 2.6k pixels, ConvLSTM




WeatherBench, Rasp et al. (2020). 5.625° , 64x32, 2K pixels, CNN




Deuben & Bauer (2018), 6° , 60x30, 1.8K pixels, MLP

FourCastNet pushing the frontier of AI-Driven Digital Twins



- Scope Global
- Model Type Full-Atmosphere AI Surrogate
- Architecture Fourier Neural Operator
- Resolution 25km, 6-hourly
- Training Data ERA5 Reanalysis
- Initial Condition ERA5 / GFS / UFS
- Training Time 1000 GPU-hrs
- Inference Time 3 sec (2-week forecast)
- Calibration IC + Bayesian model uncertainty
- Speedup vs NWP O(10,000 – 100,000)
- Power Savings O(10,000)
- Max Stable Rollout Years
- Project Type Open-source

FourCastNet (Fourier foreCasting Network)



- Purely data-driven ML surrogate weather model
- Fourier transform for global convolution
- Learns solution operator, mesh and resolution invariant
- Training data: ERA5 reanalysis
- Autoregressive time interval: 6 hours
- 73 state variables selected:
 - Temperatures, winds, geopotential & humidity (surface & 12 vertical levels)
 - Surface pressure, column water vapor, ...

Extending to include radiation processes, vapor transport, clouds

Training set: 1979 to 2015
Validation set: 2016, 2017
Held out: 2018 onwards

FourCastNet and other DDWP models available online


Modulus-Makani, Earth2-MIP, ECMWF AI Lab



Makani: Massively parallel training of machine-learning based weather and climate models

[Overview](#) | [Getting started](#) | [More information](#) | [Known issues](#) | [Contributing](#) | [Further reading](#) | [References](#)

tests passing


Makani (the Hawaiian word for wind 🌬️) is an experimental library designed to enable the research and development of machine-learning based weather and climate models in PyTorch. Makani is used for ongoing research. Stable features are regularly ported to the [NVIDIA Modulus](#) framework, a framework used for training Physics-ML models in Science and Engineering.






- Open sourcing and rapid training are indispensable for academic partnerships
 - Our collaboration have actively contributed to these code bases.

FourCastNet predicts many atmospheric fields accurately




FCN has impressive skill on forecasting extremes

Including
tropical
cyclones, extra-
tropical
cyclones, and
atmospheric
rivers.




FourCastNet predicts hurricane paths and intensities



FCN accurately emulates Atmospheric Rivers


- Backbone forecasts moisture vars well (e.g., total column water vapor has ACC>0.6 out to 10 days)
- Key to capturing atmospheric river dynamics



ARs approach West coasts of N. America, Europe

FourCastNet's speed enables massive ensembles.

To capture low-likelihood high-impact extreme events more accurately – far into the long tails of distributions



For simulating LLHIs, DDWP is a “killer app”

Characterizing LLHIs requires examination of many realizations of them.

The only way to produce a sufficient number of realizations is via simulation.

Given the low frequency of LLHIs, the only way to conduct these simulations is with an emulation of NWP/ESM codes that runs orders of magnitude faster.

This is why DDWP is valuable for characterizing rare climate extremes.

"Must-have" use case for DDWP for Extremes

Proposal: Generate statistics on simulated LLHIs that could have occurred under historical conditions, as well as their drivers, by generating

HENS: Huge Ensembles of 10^N members,
where $N \geq 4$ required to converge statistics

The ensemble will consist of short (2-week long) hindcasts.



Hypothesis: Ergodicity of climate system means we can "trade"
increasing ensemble size with increasing length of sampling time.

Data-Driven Forecasts from FourCastNet

FourCastNet is a machine learning architecture used for weather and climate prediction.

It can produce a 2-week, 0.25 degree forecast in 0.5 seconds. It's ~4 orders of magnitude faster than traditional models.

This enables the creation of **huge** ensembles (1,000-10,000 members). This ensemble size is vital for capturing the long-tails of climate distributions.




NVIDIA: <https://www.youtube.com/watch?v=FUUT6lrQjo4>

Huge Ensembles (HENs)

Constructing HENS: We construct ensembles with FourCastNet using the same ensembling techniques as operational weather centers


Validating HENS: We validate these ensembles on extremes using the same techniques as NWP

LLHIs in HENS: Summer 2023 was the hottest summer on record. We will study and quantify near-miss LLHIs in ultra-large counterfactuals of summer 2023.



Constructing Huge Ensembles with FourCastNet


1. Perturb the initial conditions



2. Perturb the model




3. TBD: Inject noise during model integration



FCN scales efficiently up to ~ 4000 GPUs on 3 HPC systems


NVIDIA's full-stack AI + HPC expertise shines thru... training on large amounts of the world's petabytes of past weather data.



Peak performance is 140.8 petaFLOPS in mixed precision (averaged over a full epoch). Time to solution decreased from 24+ hours to 67 minutes with model and data parallelism

Another key scaling property: training in <1 day


Model and data parallelism reduces training time from ~ 24 hours to 67.4 minutes



Given Future Data, FCN's Speed Yields Fast Tethering


AI nimbly generates details between "checkpoints" saved only infrequently from physics-based climate simulations

-- Bjorn Stevens, GTC 2021





Game changer: Regenerate data rather than moving it

For how long can full-AI models like FourCastNet be trusted to "tether" between climate checkpoints?



Spherical harmonics enable stable, high-fidelity long rollouts.

Respecting spherical geometry key to maintaining robustness for auto-regressive rollout.



* Stable one-year rollout (1460 autoregressive steps) computed in 13 minutes on a single NVIDIA RTX A6000


harmonics
<https://github.com/NVIDIA/torch-harmonics>

As of January 2024, our newest **spherical FNO** is making convincing S2S predictions and stable up to decades.


While not yet ready for real-world S2S forecasting this is a community milestone.

NVIDIA

Looking ahead: Filling the gap between weather (short-term) and climate (long-term).



A critical timescale with immense value: sub-seasonal ensemble forecasting




- Iterating on & resolving current pathologies.
- In partnership with operational NWP centers:
 - ECMWF & NOAA
- Exploring data assimilation, optimal initialization.
- Generative post-processing.

Source: National Academies, "Next-generation Earth System Prediction", <https://nap.nationalacademies.org/catalog/21873/next-generation-earth-system-prediction-strategies-for-subseasonal-to-seasonal>

Towards stable, accurate, high-fidelity rollouts for decadal climate

Enabled by spherical harmonics, equivariance, and physical constraints



Physics > Atmospheric and Oceanic Physics


[Submitted on 3 Oct 2023 ([v1](#)), last revised 6 Dec 2023 (this version, v2)]

ACE: A fast, skillful learned global atmospheric model for climate prediction

Oliver Watt-Meyer, Gideon Dresdner, Jeremy McGibbon, Spencer K. Clark, Brian Henn, James Duncan, Noah D. Brenowitz, Karthik Kashinath, Michael S. Pritchard, Boris Bonev, Matthew E. Peters, Christopher S. Bretherton



Temperature
(boundary
layer average)

Total Water Path
(vertically integrated
moisture)



FCN will sample future extreme-weather for decades

For more accurate statistics, to capture rare extremes, and identify significant forecast-bifurcations



Towards kilometer-scale emulation

To meet the needs of society:

- (i) predicting impacts at scales that matter
- (ii) interacting with data at low-latency


Model: FourCastNet ~ 2.5B Parameter FCN (5km Resolution)

System: 4K H100

Training: Full State Vector, 10 Years > 5 PB > ~3 Days


Inference: 30 Days, 1000-member ensemble > ~1 Hour

O(1,000)X Speed-Up Versus Simulation



Next Steps and Future Directions

- Extending emulators to all components of the Earth system
- Developing extrapolation methods to “no analog” climates
- Introducing data-driven climate prediction to the IPCC (Intergovernmental Panel on Climate Change)
- Integrating data-driven weather prediction with existing operational forecasting
- Demonstrating the real-world benefits from emulation



Thanks to all collaborators!



RICE



Ankur Mahesh



Jaideep P.
NVIDIA



Shashank S.
LBL



Peter H.
LBL



Sanjeev R.
U Michigan



Ashesh C.
Rice U.




Morteza M.
NVIDIA



Thorsten K.
NVIDIA



David H.
NVIDIA



Zongyi L.
Caltech



Kamyar A.
Purdue



Pedram H.
Rice U.



Karthik K.
NVIDIA



Anima A.
NVIDIA / Caltech



Mike Pritchard

Data-Driven Models Revolutionize Environmental Prediction

At the dawn of an entirely new way to simulate the global atmosphere



Questions?


Feedback?

Paths Forward?

BACKUP SLIDES


The spread of an ensemble prediction should grow at the same rate as its error

The European Center for Medium Range Weather Forecasting (ECMWF) is respected for its well-calibrated ensembles




FCN medium-range forecast skill comparable to IFS

And can be initialized with GFS initial conditions, proving zero-shot transfer learning.




t2m	3 days 00:00:00	0.02
	7 days 00:00:00	-0.00
	14 days 00:00:00	0.00
t850	3 days 00:00:00	-0.03
	7 days 00:00:00	-0.04
	14 days 00:00:00	-0.03
u10m	3 days 00:00:00	-0.09
	7 days 00:00:00	-0.06
	14 days 00:00:00	-0.06
v10m	3 days 00:00:00	-0.09
	7 days 00:00:00	-0.06
	14 days 00:00:00	-0.05
z500	3 days 00:00:00	0.09
	7 days 00:00:00	0.01
	14 days 00:00:00	-0.02



Modeled on the ECMWF scorecard
Green is better, Blue is worse.

FourCastNet passes key tests for extreme forecasts

FCN's ensembles calibrated using initial condition uncertainty and model uncertainty drawing on NWP wisdom



Ankur Mahesh