

# Accelerating NetworkX: The Future of Easy Graph Analytics



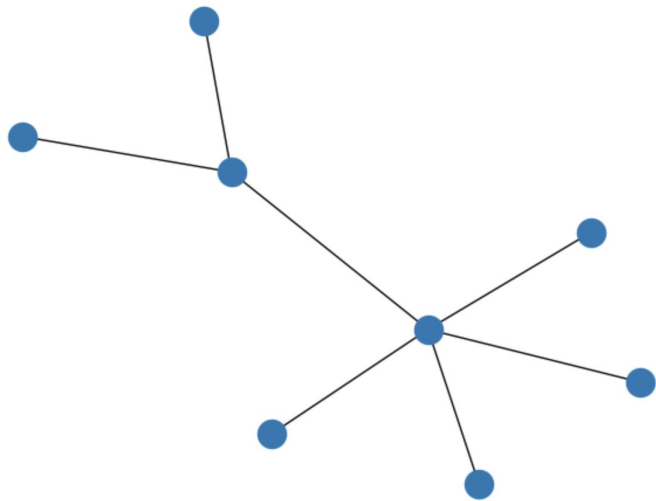
**Mridul Seth - Core Developer, NetworkX**

Erik Welch - Sr. System Software Engineer, NVIDIA

Rick Ratzel - Sr. System Software Engineer, NVIDIA

NetworkX and backend developers

# Why do we care about Graphs?



Do you have data?

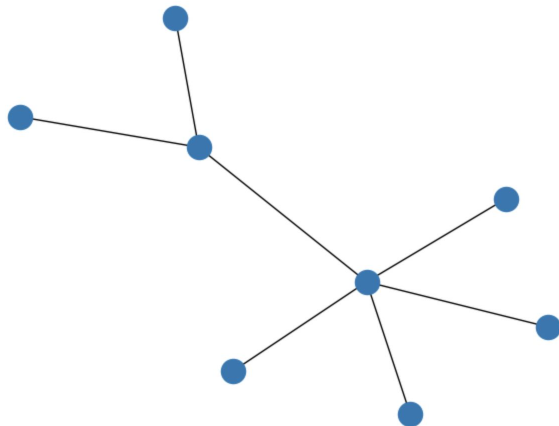
Does your data have *relationships*? (hint: probably!)

Do you want to gain insights from your data?

- Most “important”, “central” or “unusual” entities
- Detect communities and cliques
- Fraud detection
- Recommendation systems
  - e.g., movie, product, or LinkedIn recommendation
- Connectivity, shortest paths, route planning
- Much, much more (including GNNs)
- Used by virtually all science and engineering fields!

# What do you need to work with graphs?

- A good graph analytics library should be:
  - Easy to use
  - Well-documented
  - Robust, reliable, and well-tested
  - Popular, with a large, friendly community backing it
  - Able to support a wide breadth of algorithms
    - ...and plotting would also be very nice
  - Easy to install without exotic dependencies
  - Permissive to use with an open-source license
  - Performant



# ...this describes NetworkX



**NetworkX**  
Network Analysis in Python

- Considered the de facto standard

R

You

What nodes in a Graph representation of the Zachary Karate Club dataset have the highest jaccard coefficient? Please provide the answer as node IDs in the graph.

ChatGPT

ChatGPT

To perform these steps, you can use a programming language like Python and a graph processing library like NetworkX. Here's a simple example using NetworkX:

python

Copy code

```
import networkx as nx

# Load the Zachary Karate Club dataset
G = nx.karate_club_graph()

# Compute Jaccard coefficients
jaccard_coefficients = nx.jaccard_coefficient(G)

# Sort nodes based on Jaccard coefficients in descending order
sorted_nodes = sorted(jaccard_coefficients, key=lambda x: x[2], reverse=True)

# Print nodes with the highest Jaccard coefficients
for u, v, coef in sorted_nodes:
    print(f"Nodes {u} and {v} have Jaccard coefficient: {coef}")
```

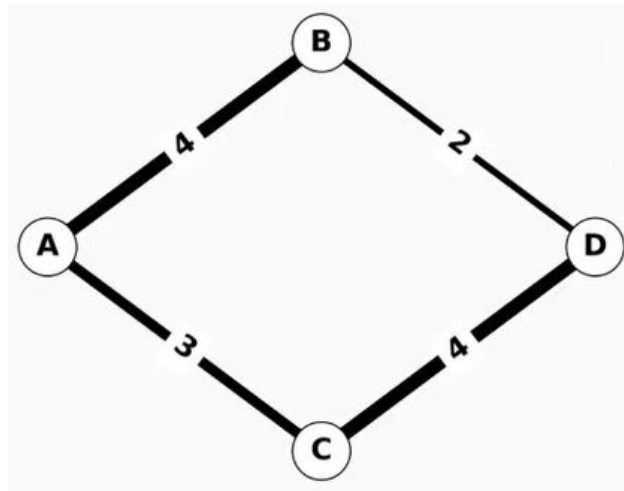
- Easy to use - API supports familiar usage and flexible types to accommodate real-world data
- Pure python, runs anywhere python is supported
- Easy install provided by all major package managers
- Well documented: excellent API docs, examples, notebooks
- Permissive open source (BSD license) and community governed
- Mature, battle-tested, and bug-vetted
  - First public release was 19 years ago!
- Very popular with a large, friendly community
  - ~14K stars and 3K forks on GitHub
  - ~47M PyPI and ~360k conda downloads per month
- Hundreds of graph algorithms and functions
  - Exploring rich graph data is rarely limited by the algorithms available

# In NetworkX, simple things stay simple

```
>>> import networkx as nx

>>> G = nx.Graph()
>>> G.add_edge("A", "B", my_cost=4)
>>> G.add_edge("B", "D", my_cost=2)
>>> G.add_edge("A", "C", my_cost=3)
>>> G.add_edge("C", "D", my_cost=4)

>>> nx.shortest_path(G, "A", "D", weight="my_cost")
['A', 'B', 'D']
```



# ...but what about *less simple* things?

*"Um, did you say 'pure python'?"*

- NetworkX is a data tool... and data is getting bigger!
- Pure python dictionary-of-dictionaries implementation does not scale well to larger problems
  - Maximum reasonable graph size varies by algo and options used
    - Eg. `nx.betweenness_centrality()` on 3.7M nodes, 16.5M edge graph,  $k=500$ , takes approx. 80 minutes\*
  - Single Processor bound to CPU



\* NetworkX 3.2 on Intel(R) Xeon(R) Gold 6128 CPU @ 3.40GHz, 45GB RAM

How do we fix this without losing all the good things about NetworkX?

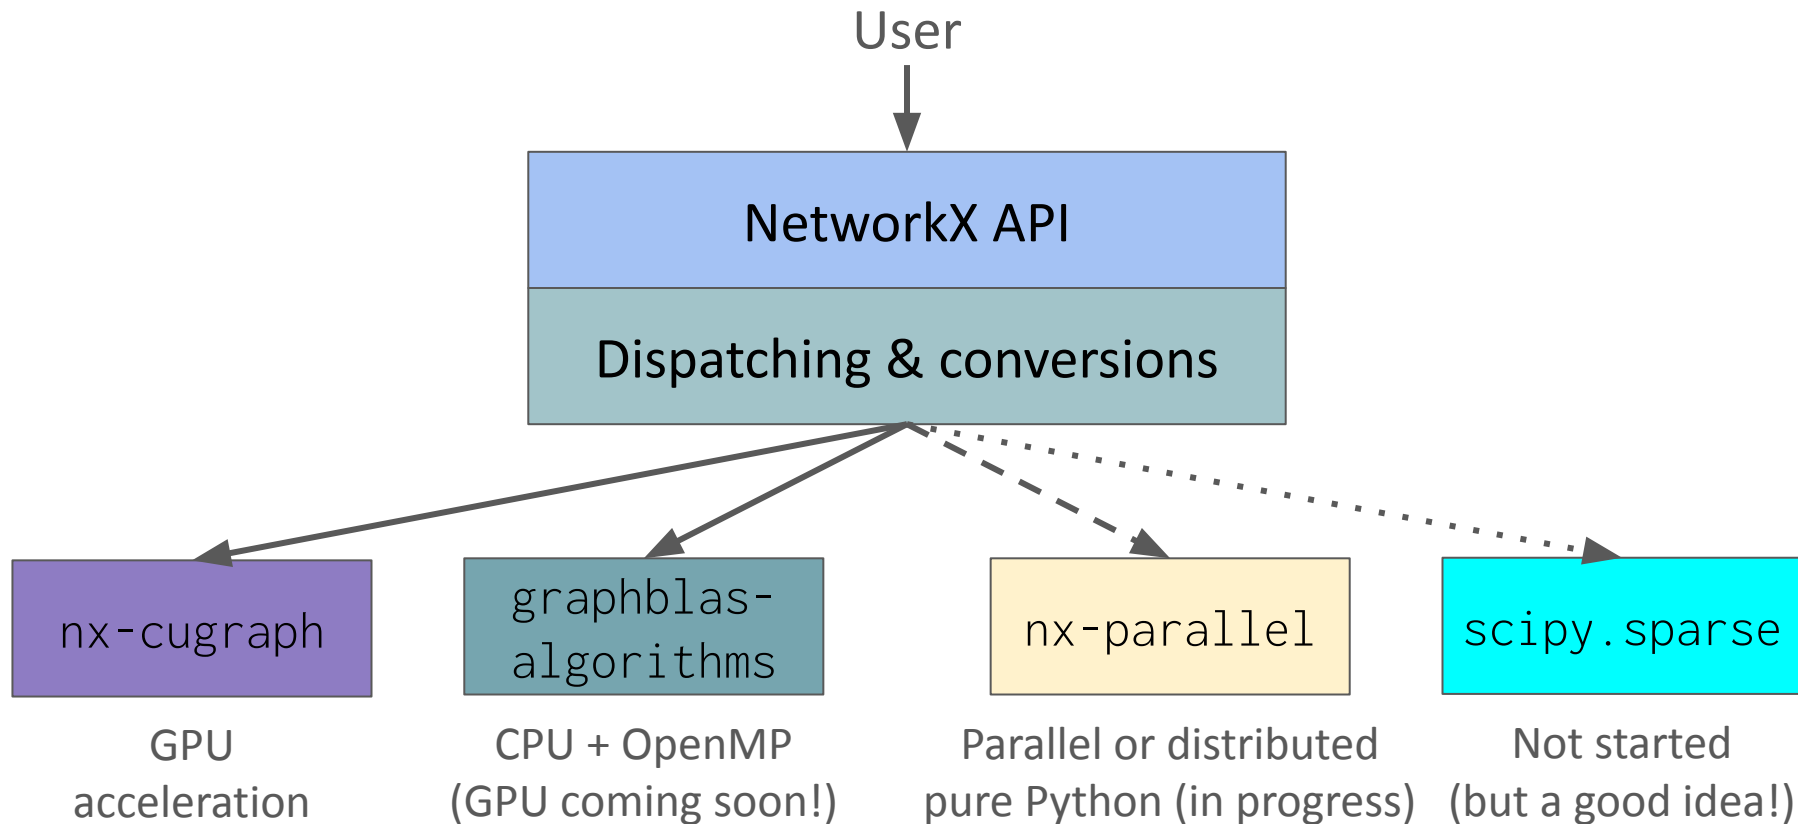
Let's rewrite everything in Rust.



~~Let's rewrite everything in Rust.~~

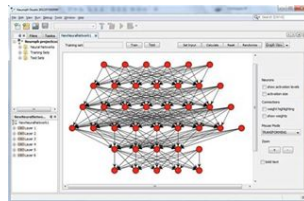
Dispatching.

# NetworkX as a user-facing API + pluggable backends



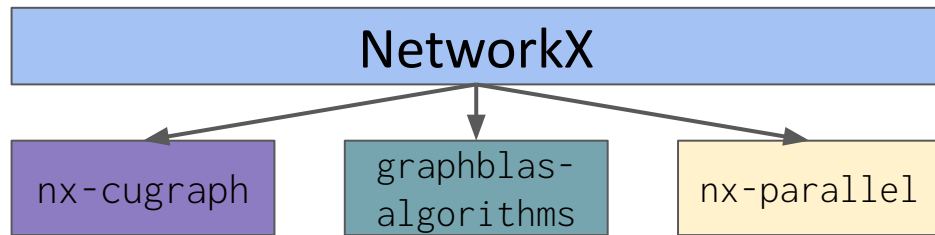
# Some other benefits of dispatching to backends

- NetworkX becomes a “standard” graph API
  - Learn a single library, scale as needed with backends.
- Share your graph algorithms with the NetworkX community - write your own backend



```
def test_pagerank(G):  
    nx.pagerank(G)
```

```
def bench_pagerank(G):  
    nx.pagerank(G)
```



- Don't pay the open-source tax
  - Good OSS needs good docs, examples, benchmarks, etc. - these aren't cheap!
- Backends leverage the NetworkX API and everything written to it
  - Run against the NetworkX test suite - tests for free!
  - Reference your backend in the NetworkX docs - docs for free!
  - Run with NetworkX benchmarks - benchmarks for free!

# This is GTC! Tell me more about that GPU-accelerated NetworkX backend!

**RAPIDS**

- **nx-cugraph** :  
NetworkX  
backend using  
RAPIDS cuGraph
- No-code-change  
GPU acceleration
- Run on GPU if  
available, fallback  
to CPU if not

bc\_demo.ipynb

```
import pandas as pd
import networkx as nx

url = "https://data.rapids.ai/cugraph/datasets/cit-Patents.csv"
df = pd.read_csv(url, sep=" ", names=["src", "dst"], dtype="int32")
G = nx.from_pandas_edgelist(df, source="src", target="dst")

%time result = nx.betweenness centrality(G, k=10)
```

```
user@machine:/# ipython bc_demo.ipynb
CPU times: user 7min 38s, sys: 5.6 s, total: 7min 44s
Wall time: 7min 44s
```

```
user@machine:/# NETWORKX_BACKEND_PRIORITY=cugraph ipython bc_demo.ipynb
CPU times: user 18.4 s, sys: 1.44 s, total: 19.9 s
Wall time: 20 s
```

NetworkX 3.2, CPU: Intel(R) Xeon(R) Gold  
6128 CPU @ 3.40GHz 45GB RAM, GPU:  
NVIDIA Quadro RTX 8000 80GB RAM

# nx-cugraph – GPU-accelerated NetworkX



- Included in the RAPIDS cuGraph github repo
  - <https://github.com/rapidsai/cugraph/tree/branch-24.04/python/nx-cugraph>
  - Leverages RAPIDS' established release process and CI infrastructure
  - **nx-cugraph CI runs NetworkX unit tests using the nx-cugraph backend**
- Install instructions
  - conda: `conda install -c rapidsai -c conda-forge -c nvidia nx-cugraph`
  - pip: `pip install --extra-index-url=https://pypi.nvidia.com nx-cugraph-cu12`
  - from source: <https://github.com/rapidsai/cugraph>
  - for more information: <https://rapids.ai/#quick-start>

# nx-cugraph – supported algorithms

centrality	core	reciprocity	traversal
betweenness_centrality	core_number	overall_reciprocity	bfs_edges
edge_betweenness_centrality	k_truss	reciprocity	bfs_layers
degree_centrality			bfs_predecessors
in_degree_centrality	<b>dag</b>	<b>shortest_paths</b>	bfs_successors
out_degree_centrality	ancestors	has_path	bfs_tree
eigenvector_centrality	descendants	shortest_path	descendants_at_distance
katz_centrality		shortest_path_length	generic_bfs_edges
	<b>isolate</b>	all_pairs_shortest_path	<b>tree</b>
<b>cluster</b>	is_isolate	all_pairs_shortest_path_length	is_arborescence
average_clustering	isolates	bidirectional_shortest_path	is_branching
clustering	number_of_isolates	single_source_shortest_path	is_forest
transitivity		single_source_shortest_path_length	is_tree
triangles	<b>link_analysis</b>	single_target_shortest_path	
	hits	single_target_shortest_path_length	
<b>community</b>	pagerank	all_pairs_bellman_ford_path	
louvain_communities		all_pairs_bellman_ford_path_length	
	<b>operators</b>	bellman_ford_path	
<b>components</b>	complement	bellman_ford_path_length	
connected_components	reverse	single_source_bellman_ford	
is_connected		single_source_bellman_ford_path	
node_connected_component		single_source_bellman_ford_path_length	
number_connected_components			
is_weakly_connected			
number_weakly_connected_components			
weakly_connected_components			

- 60 graph algorithms
- 42 accelerated graph generators (not shown)
- More added with every release

# nx-cugraph benchmarks - betweenness centrality

**Dataset:** U.S. patent dataset, National Bureau of Economic Research (<https://snap.stanford.edu/data/cit-Patents.html>) : directed graph, **3.7M nodes**, **16.5M edges**

**CPU:** Intel(R) Xeon(R) Gold 6128 CPU @ 3.40GHz, 45GB

**GPU:** NVIDIA Quadro RTX 8000, 48GB

```
bash $> NETWORKX_BACKEND_PRIORITY="cugraph" python
```

```
>>> import networkx as nx
```

```
>>> G = nx.DiGraph()
```

```
...
```

```
>>> nx.betweenness centrality(G, k=k)
```

	k=10	k=20	k=50	k=100	k=500	k=1000
NetworkX	100.6 s	206.01 s	494.78 s	1007.71 s	5132.08 s	10233.25 s
NetworkX + nx-cugraph (cold graph conversion cache)	10.22 s	10.28 s	10.81 s	11.45 s	16.43 s	23.07 s
<i>speedup</i>	9.84 X	20.04 X	45.77 X	88.0 X	312.36 X	443.57 X
NetworkX + nx-cugraph (warm graph conversion cache)	1.16 s	1.19 s	1.65 s	2.19 s	7.24 s	13.57 s
<i>speedup</i>	86.7 X	173.11 X	299.87 X	460.14 X	708.85 X	754.11 X



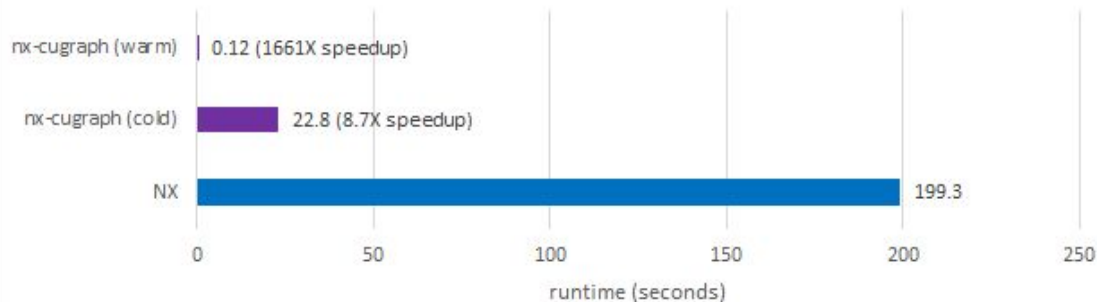
# nx-cugraph benchmarks

## overall\_reciprocity - runtime (*shorter is better*)

dataset: soc-livejournal1 (LiveJournal social network) 4.8M nodes, 69M edges

CPU: Intel Xeon Platinum 8480CL, 2TB

GPU: NVIDIA H100, 80GB

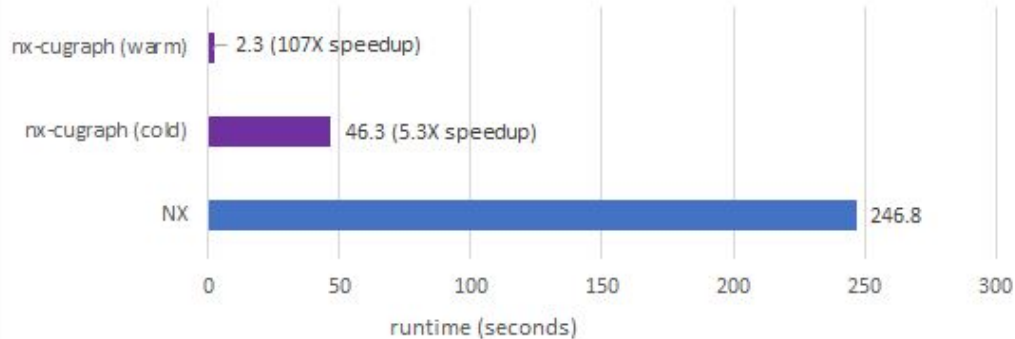


## pagerank (personalized) - runtime (*shorter is better*)

dataset: soc-livejournal1 (LiveJournal social network) 4.8M nodes, 69M edges

CPU: Intel(R) Xeon(R) Gold 6128 CPU @ 3.40GHz, 45GB

GPU: NVIDIA Quadro RTX 8000, 50GB



# nx-cugraph benchmarks

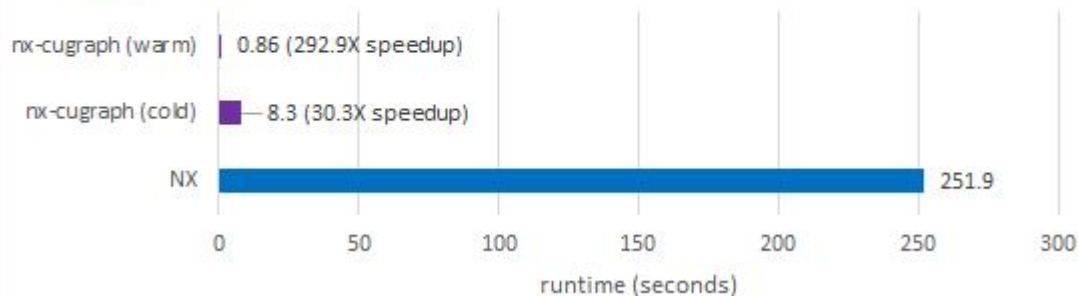
## katz centrality - runtime (*shorter is better*)

dataset: cit-Patents (citation graph of patent citations) 3.7M nodes, 16.5M

edges

CPU: Intel Xeon Platinum 8480CL, 2TB

GPU: NVIDIA H100, 80GB

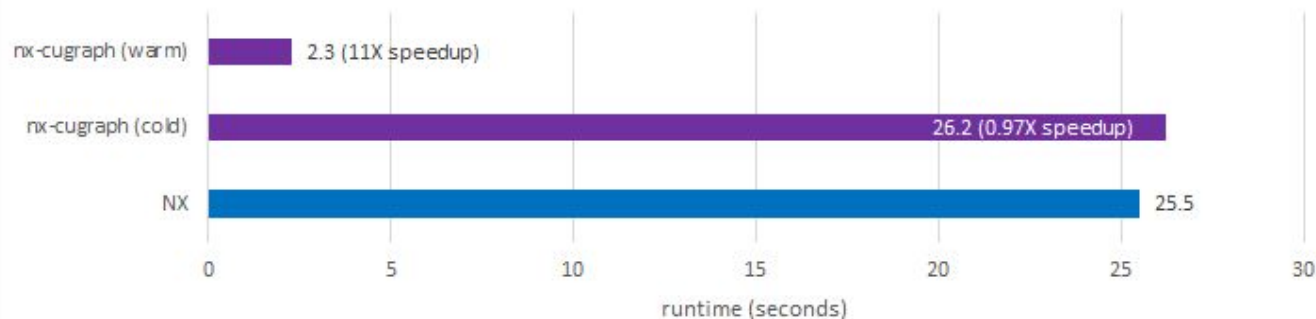


## weakly\_connected\_components - runtime (*shorter is better*)

dataset: soc-livejournal1 (LiveJournal social network) 4.8M nodes, 69M edges

CPU: Intel Xeon Platinum 8480CL, 2TB

GPU: NVIDIA H100, 80GB



# Next Steps

*Something not here you'd like to see?*

*Leave a suggestion...or a pull-request!*

<https://github.com/networkx/networkx>

<https://github.com/rapidsai/cugraph>

## NetworkX dispatching

- Configuration API
- Introspection and Logging
- More dispatchable APIs (plotting? 3rd-party-only APIs?)

## nx-cugraph

- Many more algorithms (*suggestions welcome*)
- Multi-GPU support

**Thank You!**

BEGIN EXTRA SLIDES

# nx-cugraph benchmarks - betweenness centrality

**Dataset:** U.S. patent dataset maintained by the National Bureau of Economic Research (<https://snap.stanford.edu/data/cit-Patents.html>) : directed graph, 3.7M nodes, 16.5M edges

**CPU:** Intel(R) Xeon(R) Platinum 8480CL, 2TB

**GPU:** NVIDIA H100, 80GB

```
bash $> NETWORKX_AUTOMATIC_BACKENDS="cugraph" python
```

```
>>> import networkx as nx
```

```
>>> G = nx.DiGraph()
```

```
...
```

```
>>> nx.betweenness centrality(G, k=k)
```

	k=10	k=20	k=50	k=100	k=500	k=1000
NetworkX	97.28 s	184.77 s	463.15 s	915.84 s	4,585.96 s	9,125.48 s
nx-cugraph	8.71 s	8.26 s	8.91 s	8.67 s	11.31 s	14.37 s
speedup	11.17 X	22.37 X	51.96 X	105.58 X	405.59 X	634.99 X