Accelerating NetworkX: The Future of Easy Graph Analytics

Mridul Seth - Core Developer, NetworkX

Erik Welch - Sr. System Software Engineer, NVIDIA

Rick Ratzel - Sr. System Software Engineer, NVIDIA

NetworkX and backend developers

Why do we care about Graphs?

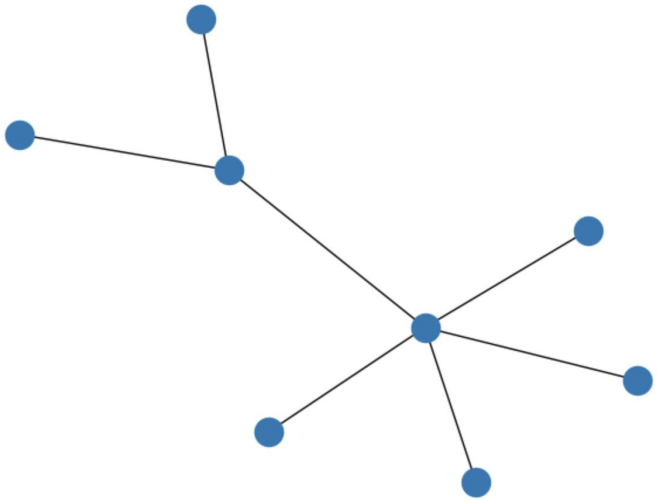

Do you have data?

Does your data have *relationships*? (hint: probably!)

Do you want to gain insights from your data?

- Most “important”, “central” or “unusual” entities
- Detect communities and cliques
- Fraud detection
- Recommendation systems
  - e.g., movie, product, or LinkedIn recommendation
- Connectivity, shortest paths, route planning
- Much, much more (including GNNs)
- Used by virtually all science and engineering fields!

What do you need to work with graphs?

- A good graph analytics library should be:
  - Easy to use
  - Well-documented
  - Robust, reliable, and well-tested
  - Popular, with a large, friendly community backing it
  - Able to support a wide breadth of algorithms
    - ...and plotting would also be very nice
  - Easy to install without exotic dependencies
  - Permissive to use with and open-source license
  - Performant

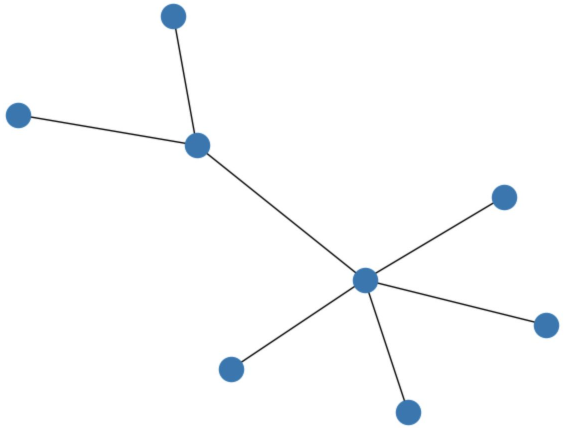

...this describes NetworkX

- Considered the de facto standard

You

What nodes in a Graph representation of the Zachary Karate Club dataset have the highest jaccard coefficient? Please provide the answer as node IDs in the graph.

ChatGPT

To perform these steps, you can use a programming language like Python and a graph processing library like NetworkX. Here's a simple example using NetworkX:

python — Copy code

```python
import networkx as nx

# Load the Zachary Karate Club dataset
G = nx.karate_club_graph()

# Compute Jaccard coefficients
jaccard_coefficients = nx.jaccard_coefficient(G)

# Sort nodes based on Jaccard coefficients in descending order
sorted_nodes = sorted(jaccard_coefficients, key=lambda x: x[2], reverse=True)

# Print nodes with the highest Jaccard coefficients
for u, v, coef in sorted_nodes:
    print(f"Nodes {u} and {v} have Jaccard coefficient: {coef}")
```

- Easy to use - API supports familiar usage and flexible types to accommodate real-world data
- Pure python, runs anywhere python is supported
- Easy install provided by all major package managers
- Well documented: excellent API docs, examples, notebooks
- Permissive open source (BSD license) and community governed
- Mature, battle-tested, and bug-vetted
  - First public release was 19 years ago!
- Very popular with a large, friendly community
  - ~14K stars and 3K forks on GitHub
  - ~47M PyPI and ~360k conda downloads per month
- Hundreds of graph algorithms and functions
  - Exploring rich graph data is rarely limited by the algorithms available

In NetworkX, simple things stay simple

```
>>> import networkx as nx

>>> G = nx.Graph()
>>> G.add_edge("A", "B", my_cost=4)
>>> G.add_edge("B", "D", my_cost=2)
>>> G.add_edge("A", "C", my_cost=3)
>>> G.add_edge("C", "D", my_cost=4)

>>> nx.shortest_path(G, "A", "D", weight="my_cost")
['A', 'B', 'D']
```

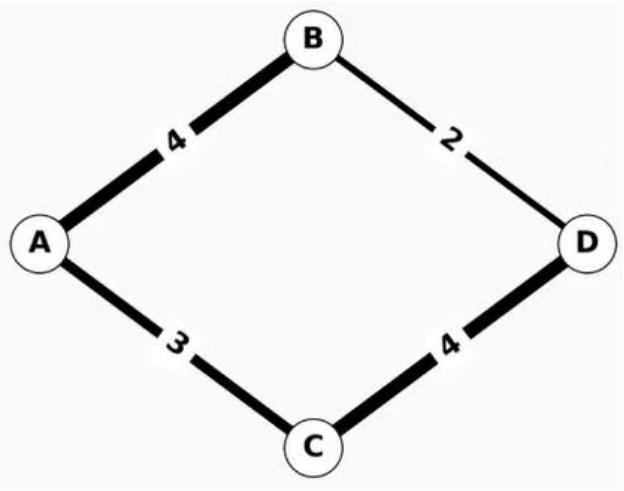

...but what about *less simple* things?

"Um, did you say 'pure python'"?

- NetworkX is a data tool... and data is getting bigger!
- Pure python dictionary-of-dictionaries implementation does not scale well to larger problems
  - Maximum reasonable graph size varies by algo and options used
    - Eg. `nx.betweenness_centrality()` on 3.7M nodes, 16.5M edge graph, k=500, takes approx. 80 minutes*
  - Single Processor bound to CPU

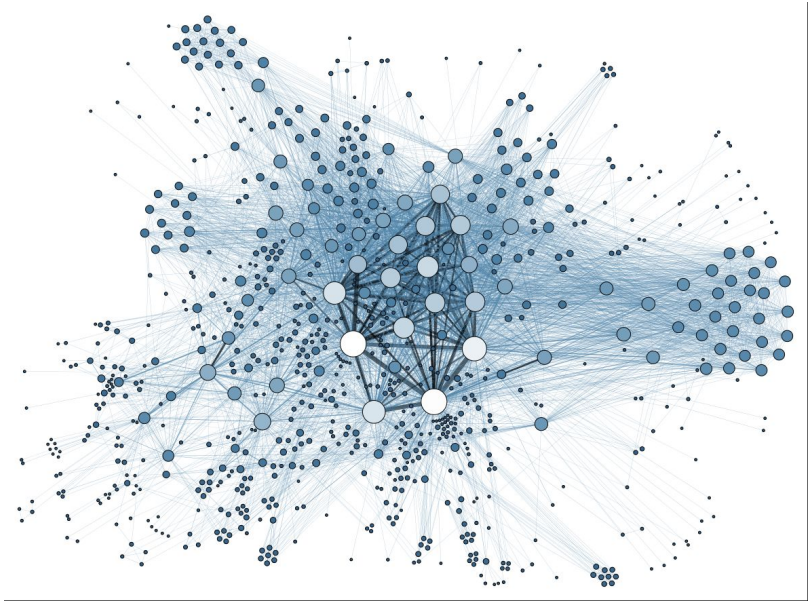

* NetworkX 3.2 on Intel(R) Xeon(R) Gold 6128 CPU @ 3.40GHz, 45GB RAM

How do we fix this without losing all the good things about NetworkX?

Let’s rewrite everything in Rust.

~~Let’s rewrite everything in Rust.~~

Dispatching.

NetworkX as a user-facing API + pluggable backends

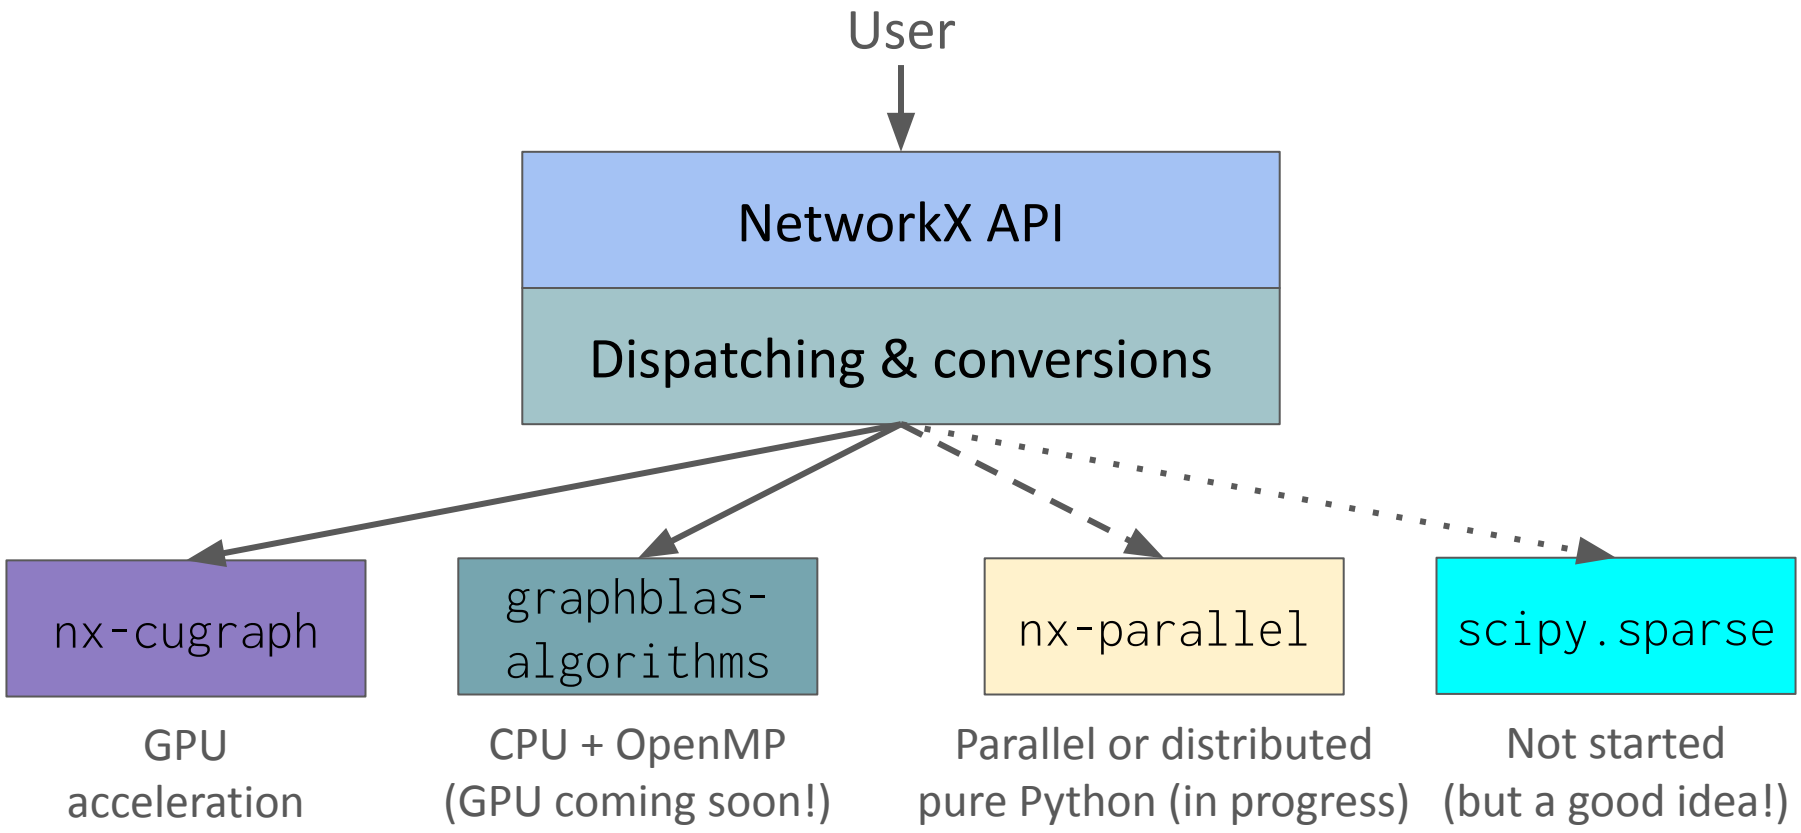

Some other benefits of dispatching to backends

- NetworkX becomes a “standard” graph API
  - Learn a single library, scale as needed with backends.
- Share your graph algorithms with the NetworkX community - write your own backend

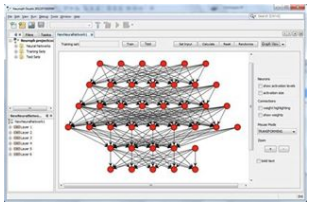

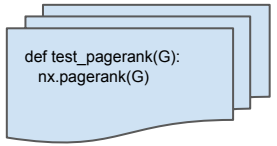

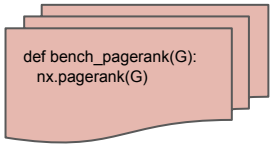

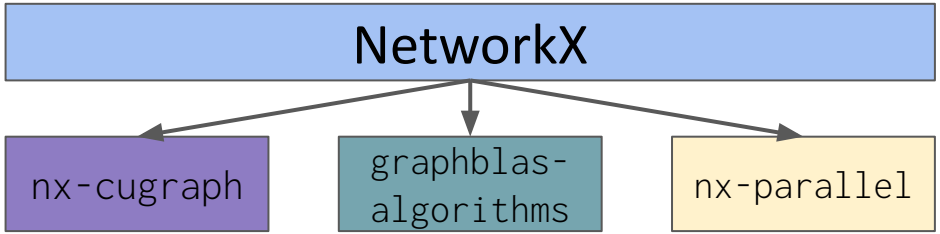

  - Don’t pay the open-source tax
    - Good OSS needs good docs, examples, benchmarks, etc. - these aren’t cheap!
  - Backends leverage the NetworkX API and everything written to it
    - Run against the NetworkX test suite - tests for free!
    - Reference your backend in the NetworkX docs - docs for free!
    - Run with NetworkX benchmarks - benchmarks for free!

This is GTC! Tell me more about that GPU-accelerated NetworkX backend!

- **nx-cugraph** : NetworkX backend using RAPIDS cuGraph
- No-code-change GPU acceleration
- Run on GPU if available, fallback to CPU if not

bc_demo.ipy

```
import pandas as pd
import networkx as nx

url = "https://data.rapids.ai/cugraph/datasets/cit-Patents.csv"
df = pd.read_csv(url, sep=" ", names=["src", "dst"], dtype="int32")
G = nx.from_pandas_edgelist(df, source="src", target="dst")

%time result = nx.betweenness_centrality(G, k=10)
```

```
user@machine:/# ipython bc_demo.ipy
CPU times: user 7min 38s, sys: 5.6 s, total: 7min 44s
Wall time: 7min 44s

user@machine:/# NETWORKX_BACKEND_PRIORITY=cugraph ipython bc_demo.ipy
CPU times: user 18.4 s, sys: 1.44 s, total: 19.9 s
Wall time: 20 s
```

NetworkX 3.2, CPU: Intel(R) Xeon(R) Gold 6128 CPU @ 3.40GHz 45GB RAM, GPU: NVIDIA Quadro RTX 8000 80GB RAM

nx-cugraph – GPU-accelerated NetworkX

- Included in the RAPIDS cuGraph github repo
  - https://github.com/rapidsai/cugraph/tree/branch-24.04/python/nx-cugraph
  - Leverages RAPIDS' established release process and CI infrastructure
  - **nx-cugraph CI runs NetworkX unit tests using the nx-cugraph backend**

- Install instructions
  - conda: `conda install -c rapidsai -c conda-forge -c nvidia nx-cugraph`
  - pip: `pip install --extra-index-url=https://pypi.nvidia.com nx-cugraph-cu12`
  - from source: https://github.com/rapidsai/cugraph
  - for more information: https://rapids.ai/#quick-start

nx-cugraph – supported algorithms

centrality
betweenness_centrality
edge_betweenness_centrality
degree_centrality
in_degree_centrality
out_degree_centrality
eigenvector_centrality
katz_centrality

cluster
average_clustering
clustering
transitivity
triangles

community
louvain_communities

components
connected_components
is_connected
node_connected_component
number_connected_components
is_weakly_connected
number_weakly_connected_components
weakly_connected_components

core
core_number
k_truss

dag
ancestors
descendants

isolate
is_isolate
isolates
number_of_isolates

link_analysis
hits
pagerank

operators
complement
reverse

reciprocity
overall_reciprocity
reciprocity

shortest_paths
has_path
shortest_path
shortest_path_length
all_pairs_shortest_path
all_pairs_shortest_path_length
bidirectional_shortest_path
single_source_shortest_path
single_source_shortest_path_length
single_target_shortest_path
single_target_shortest_path_length
all_pairs_bellman_ford_path
all_pairs_bellman_ford_path_length
bellman_ford_path
bellman_ford_path_length
single_source_bellman_ford
single_source_bellman_ford_path
single_source_bellman_ford_path_length

traversal
bfs_edges
bfs_layers
bfs_predecessors
bfs_successors
bfs_tree
descendants_at_distance
generic_bfs_edges

tree
is_arborescence
is_branching
is_forest
is_tree

- 60 graph algorithms
- 42 accelerated graph generators (not shown)
- More added with every release

nx-cugraph benchmarks - betweenness_centrality

Dataset: U.S. patent dataset, National Bureau of Economic Research (https://snap.stanford.edu/data/cit-Patents.html) : directed graph, **3.7M nodes, 16.5M edges**
CPU: Intel(R) Xeon(R) Gold 6128 CPU @ 3.40GHz, 45GB
GPU: NVIDIA Quadro RTX 8000, 48GB

```
bash $> NETWORKX_BACKEND_PRIORITY="cugraph" python
>>> import networkx as nx
>>> G = nx.DiGraph()
...
>>> nx.betweenness_centrality(G, k=k)
```

	k=10	k=20	k=50	k=100	k=500	k=1000
NetworkX	100.6 s	206.01 s	494.78 s	1007.71 s	5132.08 s	10233.25 s
NetworkX + nx-cugraph (**cold** graph conversion cache)	10.22 s	10.28 s	10.81 s	11.45 s	16.43 s	23.07 s
speedup	*9.84 X*	*20.04 X*	*45.77 X*	*88.0 X*	*312.36 X*	*443.57 X*
NetworkX + nx-cugraph (**warm** graph conversion cache)	1.16 s	1.19 s	1.65 s	2.19 s	7.24 s	13.57 s
speedup	*86.7 X*	*173.11 X*	*299.87 X*	*460.14 X*	*708.85 X*	*754.11 X*

https://github.com/rapidsai/cugraph/blob/branch-24.04/benchmarks/nx-cugraph/pytest-based/bench_algos.py

nx-cugraph benchmarks

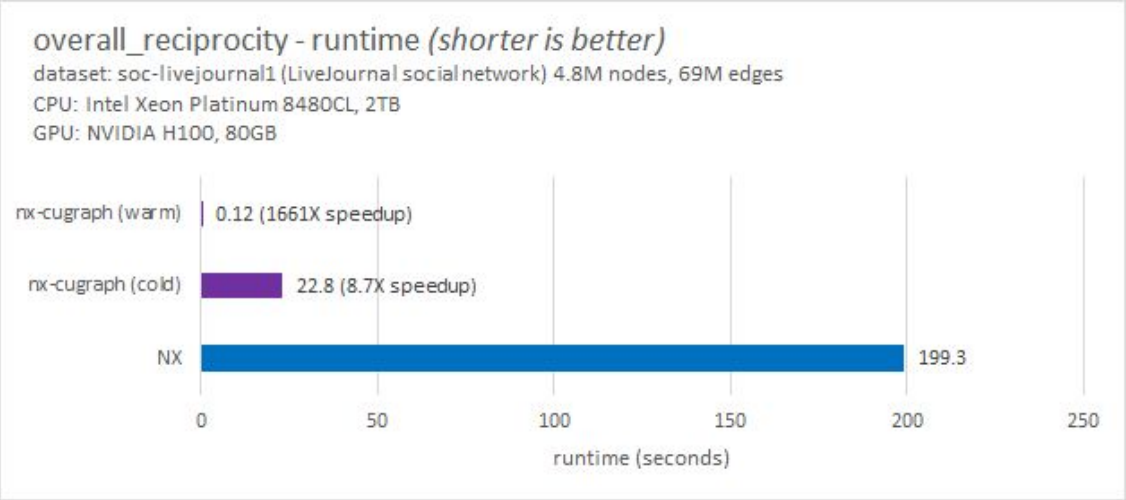

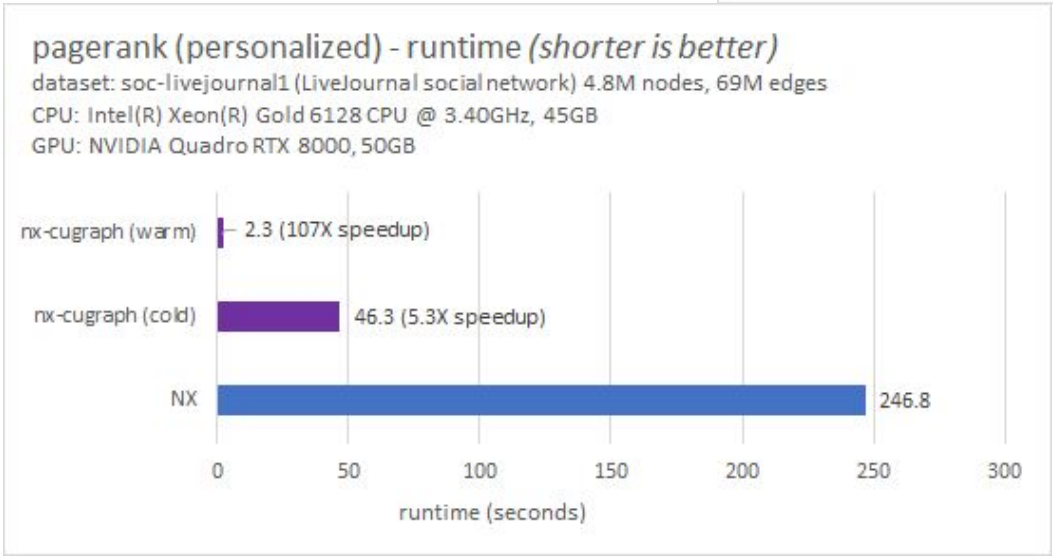

nx-cugraph benchmarks

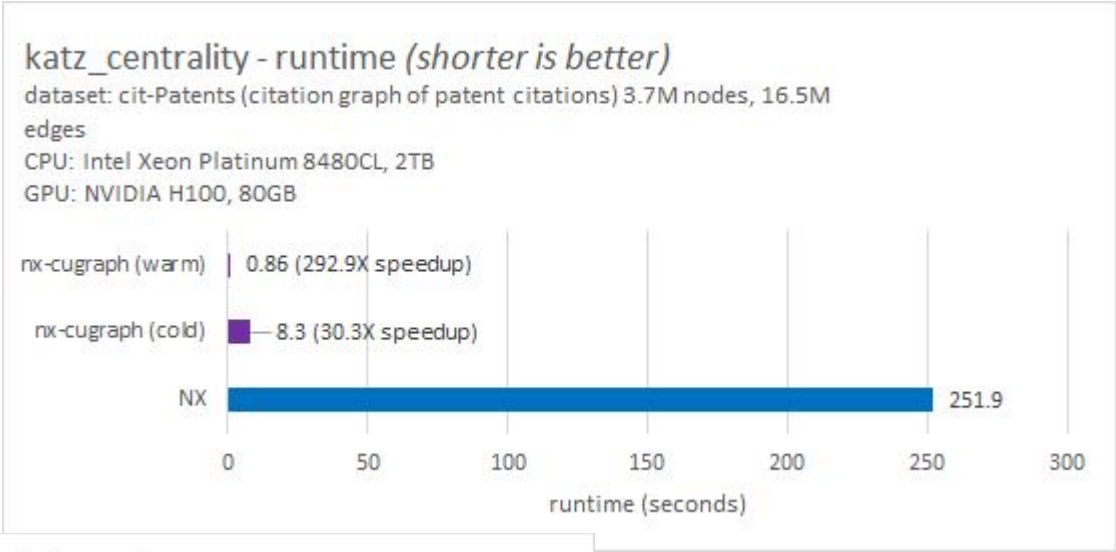

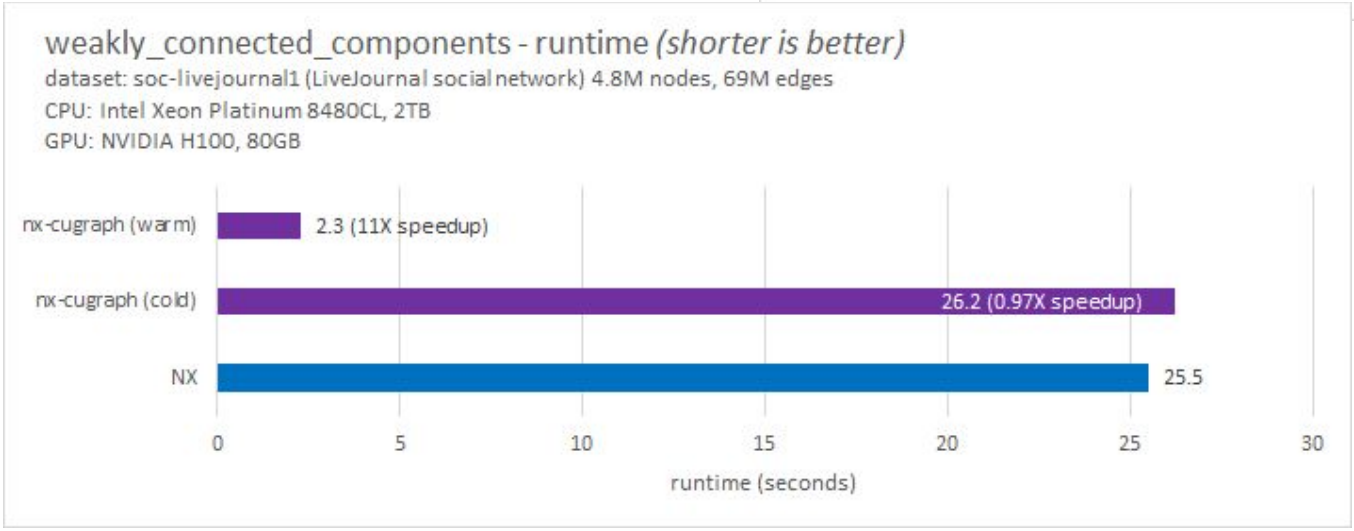

Next Steps

Something not here you'd like to see?

Leave a suggestion…or a pull-request!

https://github.com/networkx/networkx
https://github.com/rapidsai/cugraph

NetworkX dispatching

- Configuration API
- Introspection and Logging
- More dispatchable APIs (plotting? 3rd-party-only APIs?)

nx-cugraph

- Many more algorithms *(suggestions welcome)*
- Multi-GPU support

Thank You!

BEGIN EXTRA SLIDES

nx-cugraph benchmarks - betweenness_centrality

Dataset: U.S. patent dataset maintained by the National Bureau of Economic Research (https://snap.stanford.edu/data/cit-Patents.html) : directed graph, 3.7M nodes, 16.5M edges
CPU: Intel(R) Xeon(R) Platinum 8480CL, 2TB
GPU: NVIDIA H100, 80GB

```
bash $> NETWORKX_AUTOMATIC_BACKENDS="cugraph" python
>>> import networkx as nx
>>> G = nx.DiGraph()

...
>>> nx.betweenness_centrality(G, k=k)
```

	k=10	k=20	k=50	k=100	k=500	k=1000
NetworkX	97.28 s	184.77 s	463.15 s	915.84 s	4,585.96 s	9,125.48 s
nx-cugraph	8.71 s	8.26 s	8.91 s	8.67 s	11.31 s	14.37 s
speedup	11.17 X	22.37 X	51.96 X	105.58 X	405.59 X	634.99 X

https://github.com/rapidsai/cugraph/blob/branch-24.04/benchmarks/nx-cugraph/pytest-based/bench_algos.py