



# Advanced Performance Optimization in CUDA [S62192]

Igor Terentyev\*, NVIDIA DevTech Compute  
GPU Technology Conference / March 18<sup>th</sup>, 2024

\* With Guillaume Thomas-Collignon & Athena Elafrou



# Agenda

- Cooperative Grid Arrays

---
- Memory Model

---
- Asynchronous Barriers

---
- Asynchronous Data Copies

---
- Compressible Memory

---
- CUDA Graphs

# Nomenclature

**CTA** (Cooperative Thread Array) == Thread Block

**CGA** (Cooperative Grid Array) == Thread Block Cluster

Code snippets:

```
namespace cg = cooperative_groups;
```







# **Cooperative Grid Arrays**



# Cooperative Grid Arrays

Teaser

Increased SMEM saves GMEM trips in many algorithms...



# Cooperative Grid Arrays


## Properties

New level in the hierarchy:

$T < \text{CTA} < \text{CGA} < \text{GRID}$

### Clusters:

- Support guaranteed for up to 8 CTAs / CGA  
need to opt-in for more than 8 (max 16 on Hopper)
- CTAs can be logically organized in 3D blocks:
  - linear rank: `cg.this_cluster().block_rank()`
  - 3D rank : `cg.this_cluster().block_index()`
- CTAs in cluster are co-scheduled
- Thread synchronization:
  - CTA  $\rightarrow$  `cg.this_block().sync()`
  - CGA  $\rightarrow$  `cg.this_cluster().sync()`
- Possible resource under-utilization:  
cluster size > 2 may lead to unused SMs



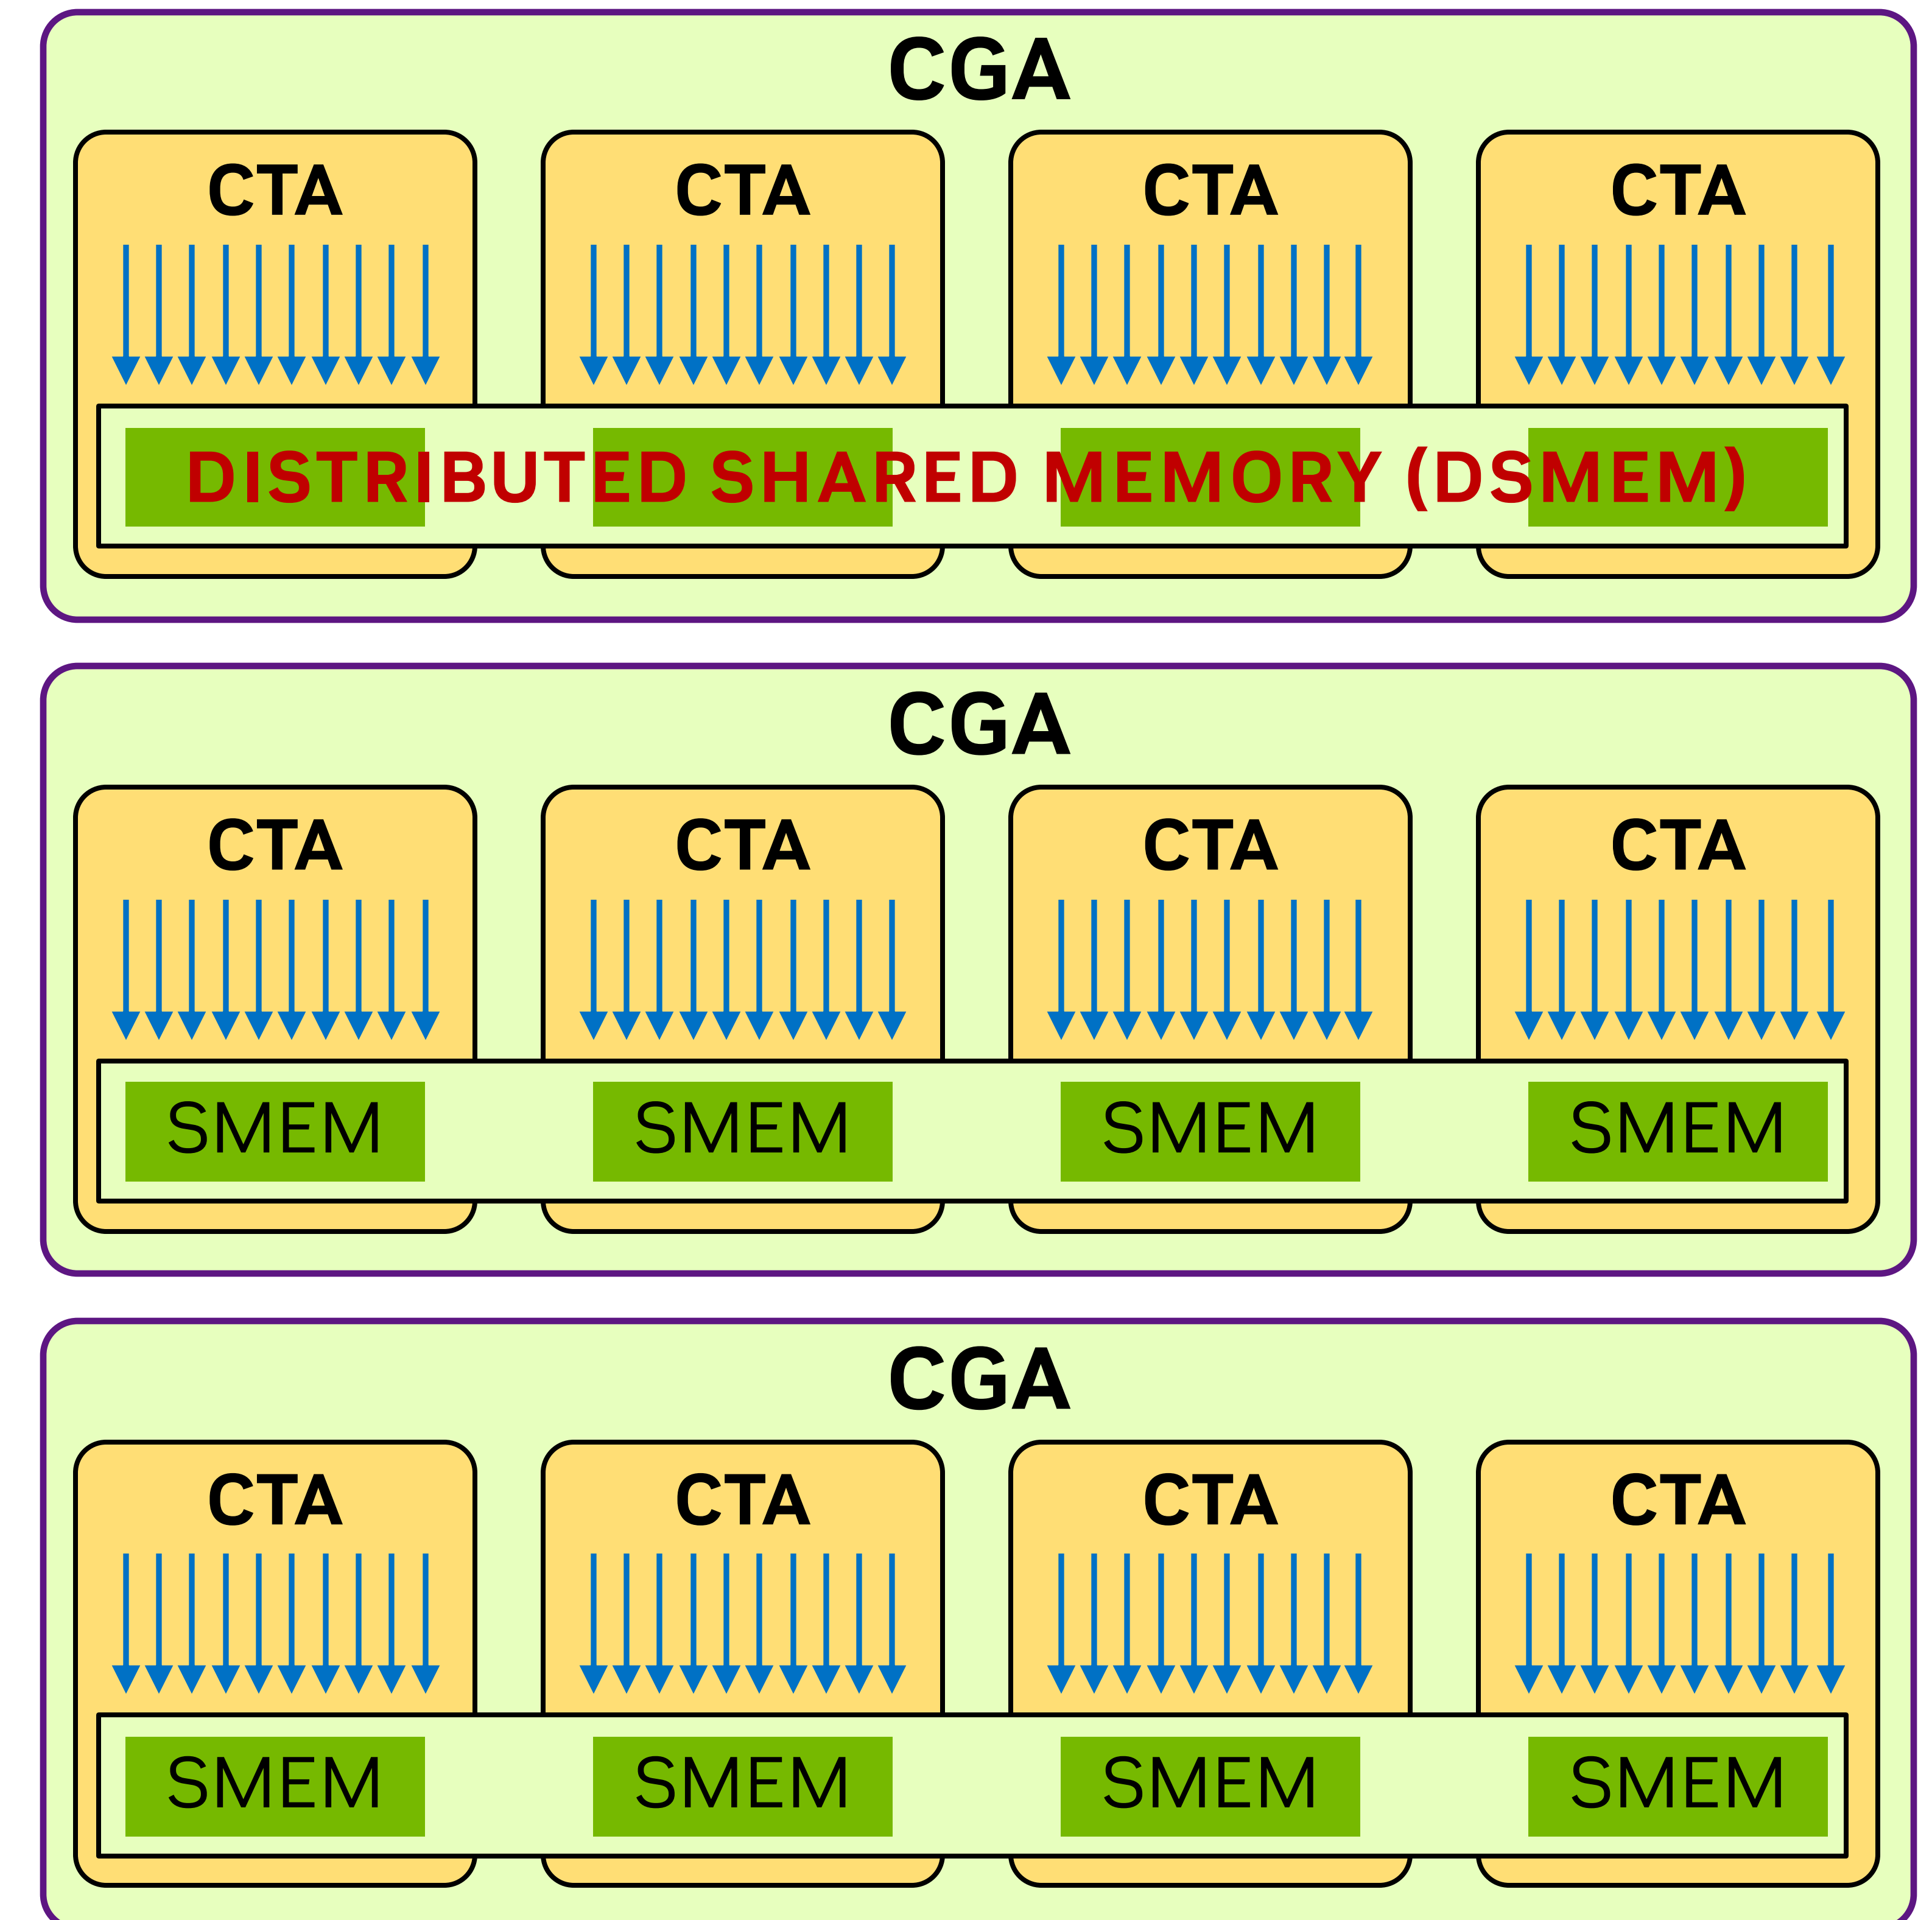
# Cooperative Grid Arrays

## DSMEM properties

### Distributed SMEM (DSMEM):

- SMEM within cluster is accessible by all CTAs → DSMEM  
E.g., H100 total DSMEM size ↗ 3.6 MB
- DSMEM allows:
  - Access: LD / ST / ATOM
  - Async copies into remote DSMEM
  - Multicasting
- Note: remote DSMEM is slower than local:  
low latency, but relatively low bandwidth
- Remote DSMEM has to be mapped:

```
__shared__ T smem;  
auto dsmem_ptr =  
    cg::this_cluster().map_shared_rank(&smem, rank);
```
- **Important:**  
**make sure remote CTAs are alive while accessing DSMEM !**  
(use `cluster.sync()` accordingly)



# Cooperative Grid Arrays

Launching kernels with CGA support

## Compile-time CGA:

```
constexpr int X = 2;
constexpr int Y = 2;
constexpr int Z = 2;
__global__ __cluster_dims__(X, Y, Z) void kernel ();
dim3 grid = {...};
assert(grid.x % X == 0 && grid.y % Y == 0 && grid.z % Z == 0); // Launch failure otherwise
kernel<<<grid_size,...>>>();
```

## Run-time CGA:

```
__global__ void kernel (); // as usual
cudaLaunchAttribute attribute[1];
attribute[0].id = cudaLaunchAttributeClusterDimension;
attribute[0].val.clusterDim.x = X;
attribute[0].val.clusterDim.y = Y;
attribute[0].val.clusterDim.z = Z;
assert(grid.x % attribute[0].val.clusterDim.x == 0 && ...); // Launch failure otherwise
cudaLaunchConfig_t config = {0};
config.attrs = attribute;
config.numAttrs = 1;
cudaLaunchKernelEx(&config, kernel, ...);
```





# Memory Model



# Memory Model

Why memory model?

Why do we need Memory (Consistency) Model?

Memory operations behave “strangely” in a multithreaded world...





# Memory Model

Weak memory model

## Thread 0:

```
data = 42;  
flag = 1;
```

## Thread 1:

```
while (!flag) {}  
assert(data == 42); // not guaranteed !!
```

CUDA uses **Weak Memory Model**:

- Writes by one thread may not be observed in the same order by other threads



# Memory Model

Weak memory model

**Thread 0:**

```
data = 42;  
flag = 1;
```

**Thread 1:**

```
while (!flag) {}  
assert(data == 42);
```

// not guaranteed !!

Moreover, Undefined Behavior (UB)!

CUDA uses **Weak Memory Model**:

- Writes by one thread may not be observed in the same order by other threads
- Concurrent access of same memory with at least one 'write' is a **CONFLICT (UNDEFINED BEHAVIOR)**
- Compiler assumes no-conflict code, **it's a programmer's responsibility to synchronize**



# Memory Model

Weak memory model

Thread 0:	Thread 1:
<pre>data = 42; flag = 1;</pre>	<pre>while (!flag) {} assert(data == 42);</pre>

// not guaranteed !!

Moreover, Undefined Behavior (UB)!

CUDA uses **Weak Memory Model**:

- Writes by one thread may not be observed in the same order by other threads
- Concurrent access of same memory with at least one 'write' is a **CONFLICT (UNDEFINED BEHAVIOR)**
- Compiler assumes no-conflict code, **it's a programmer's responsibility to synchronize**

Synchronization between threads requires:

1. **Strong ordering: "happens before" relation**
2. **Reaching correct Point of Coherency**

# Memory Model

Weak memory model

**Thread 0:**

```
data = 42;  
flag = 1;
```

**Thread 1:**

```
while (!flag) {}  
assert(data == 42);
```

// not guaranteed !!

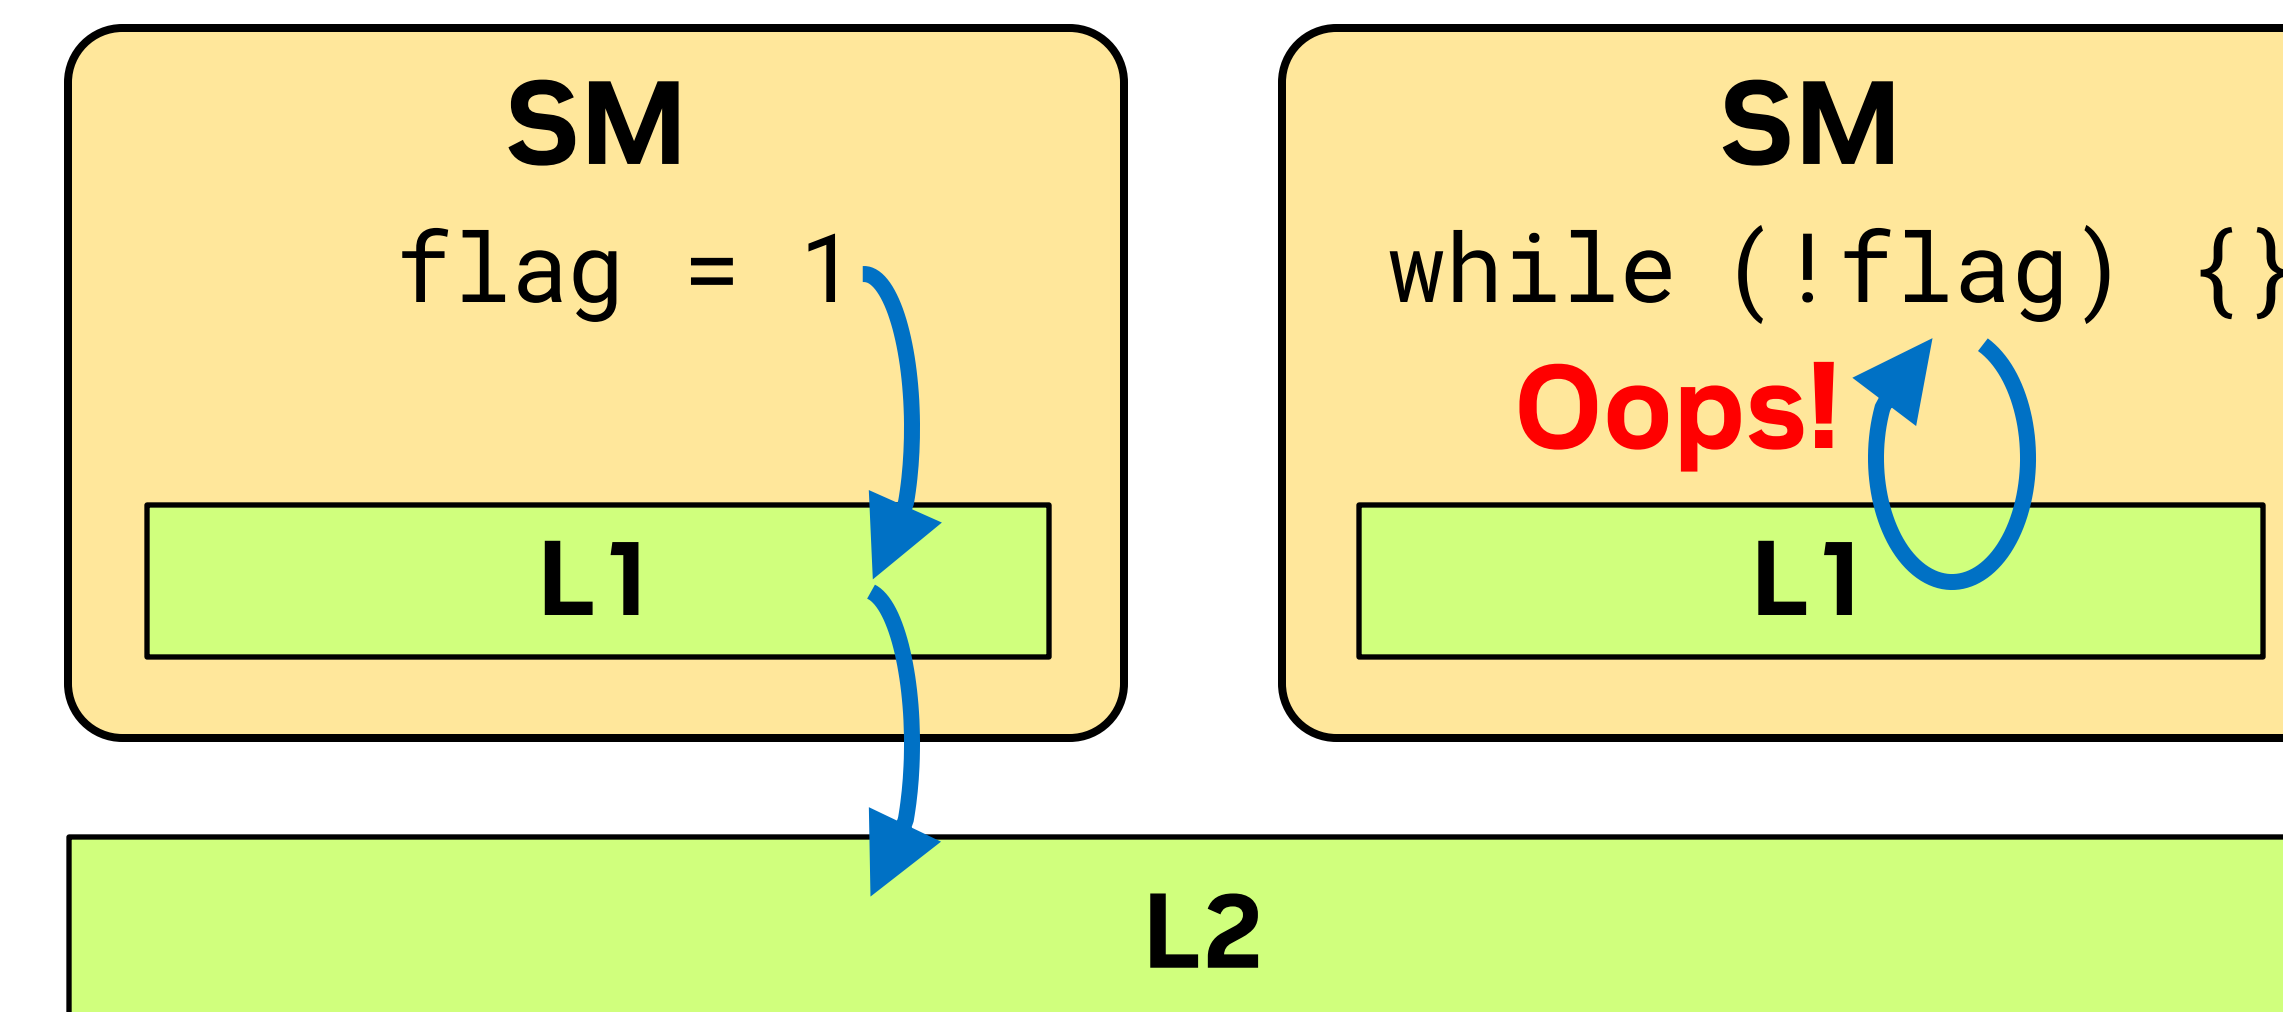
Moreover, Undefined Behavior (UB)!

CUDA uses **Weak Memory Model**:

- Writes by one thread may not be observed in the same order by other threads
- Concurrent access of same memory with at least one 'write' is a **CONFLICT (UNDEFINED BEHAVIOR)**
- Compiler assumes no-conflict code, **it's a programmer's responsibility to synchronize**

Synchronization between threads requires:

1. **Strong ordering: "happens before" relation**
2. **Reaching correct Point of Coherency**





# Memory Model

Common synchronization using shared memory

A very common communication pattern:

```
__shared__ T buf[1024];
```

```
buf[ind1] = value1;      // Threads writing to shared memory
```

```
__syncthreads();        // Barrier
```

```
value2 = buf[ind2];      // Threads reading values written by other threads
```

# Memory Model

Common synchronization using shared memory

A very common communication pattern:

```
__shared__ T buf[1024];
```

```
buf[ind1] = value1;      // Threads writing to shared memory
```

```
__syncthreads();         // Barrier + makes all previous writes visible to all threads in the block
```

```
value2 = buf[ind2];      // Threads reading values written by other threads
```

- ✓ 1. Strong ordering: “happens before” relation
- ✓ 2. Reaching correct Point of Coherency



# Memory Model

Common synchronization using global memory

A less common communication pattern:

```
__global__ void kernel (T* buf)
{
    buf[ind1] = value1;          // Threads writing to global memory

    __syncthreads();             // Barrier + makes all previous writes visible to all threads in block
                                // Including global memory !

    value2 = buf[ind2];          // Threads reading values written by other threads in the same block
}
```

- ✓ 1. Strong ordering: “happens before” relation
- ✓ 2. Reaching correct Point of Coherency

# Memory Model

Common synchronization using global memory

A less common communication pattern:

```
__global__ void kernel (T* buf)
{
    buf[ind1] = value1;          // Threads writing to global memory

    cg::this_grid().sync();    // Barrier + makes all previous writes visible to all threads in block grid
                                Including global memory !

    value2 = buf[ind2];          // Threads reading values written by other threads in the same block!
}
```

requires cooperative launch

- ✓ 1. Strong ordering: “happens before” relation
- ✓ 2. Reaching correct Point of Coherency



# Memory Model

Common synchronization using global memory

A less common communication pattern:

```
__global__ void kernel (T* buf)
{
```

```
    buf[ind1] = value1;        // Threads writing to global memory
```

```
    cg::this_grid().sync();    // Barrier + makes all previous writes visible to all threads in block grid  
                                Including global memory !
```

```
    value2 = buf[ind2];        // Threads reading values written by other threads in the same block!
```

```
}
```

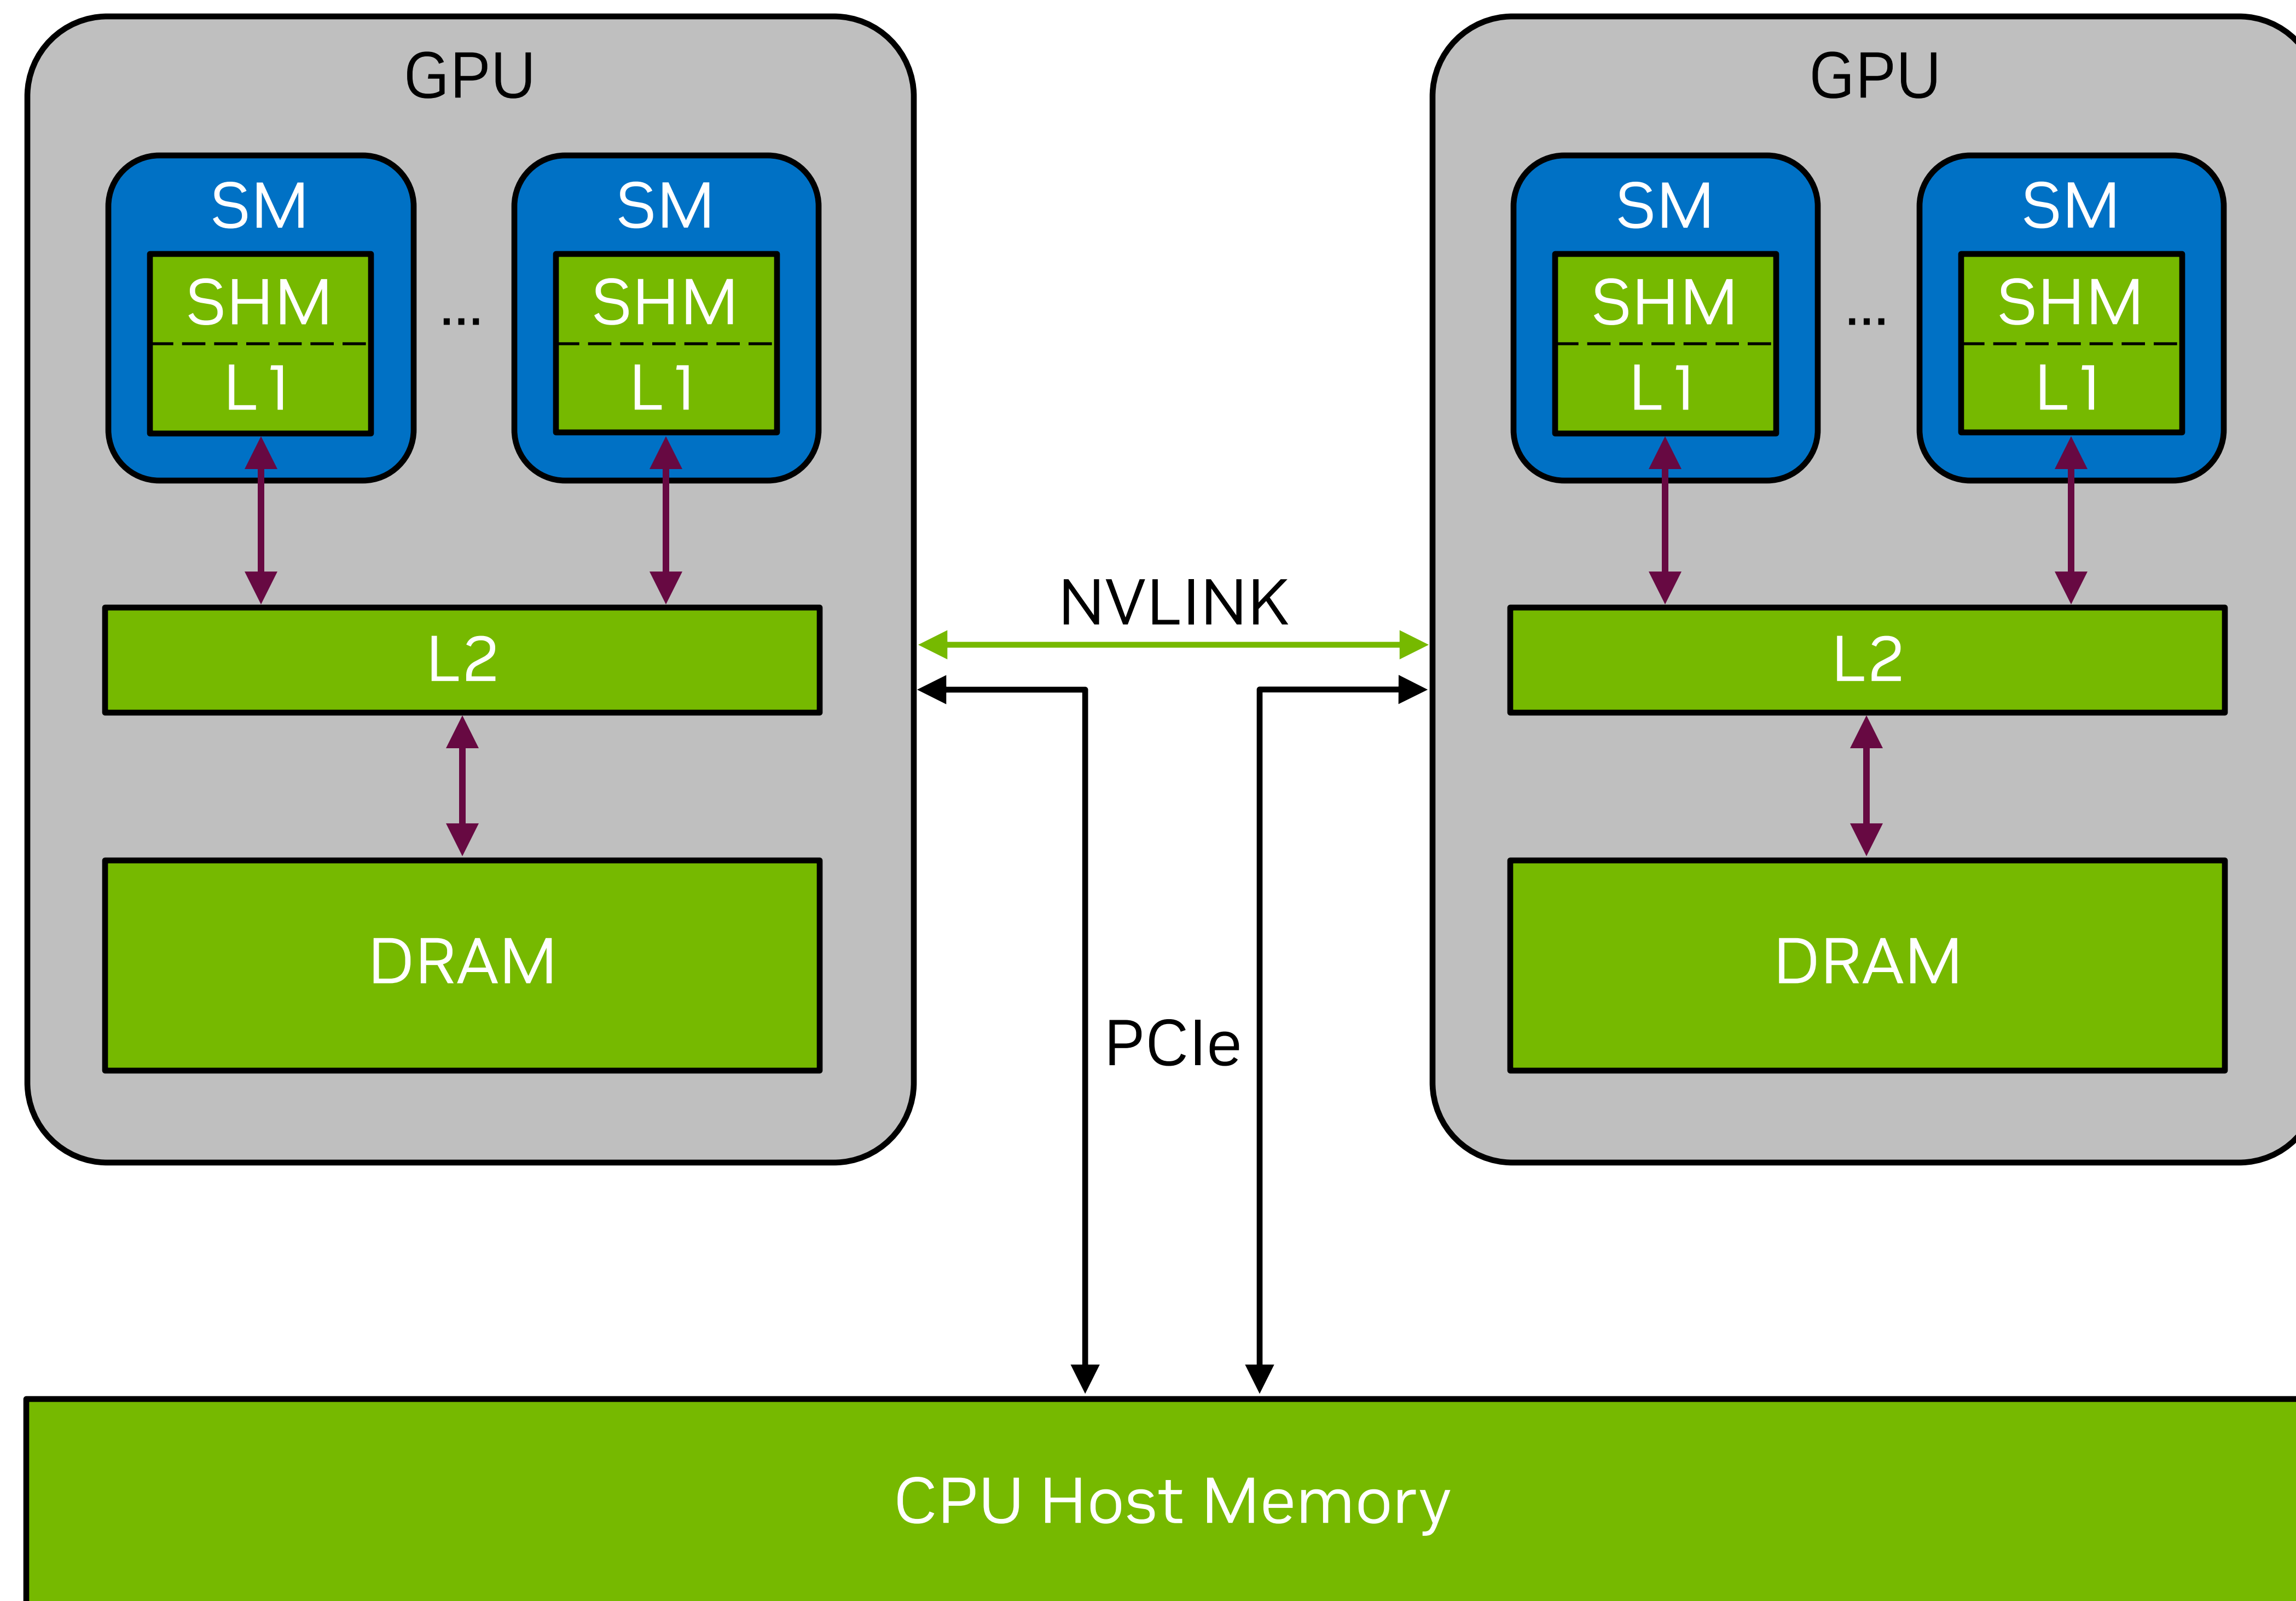
requires cooperative launch

What's really happening here?



# Memory Model

GPU memory hierarchy

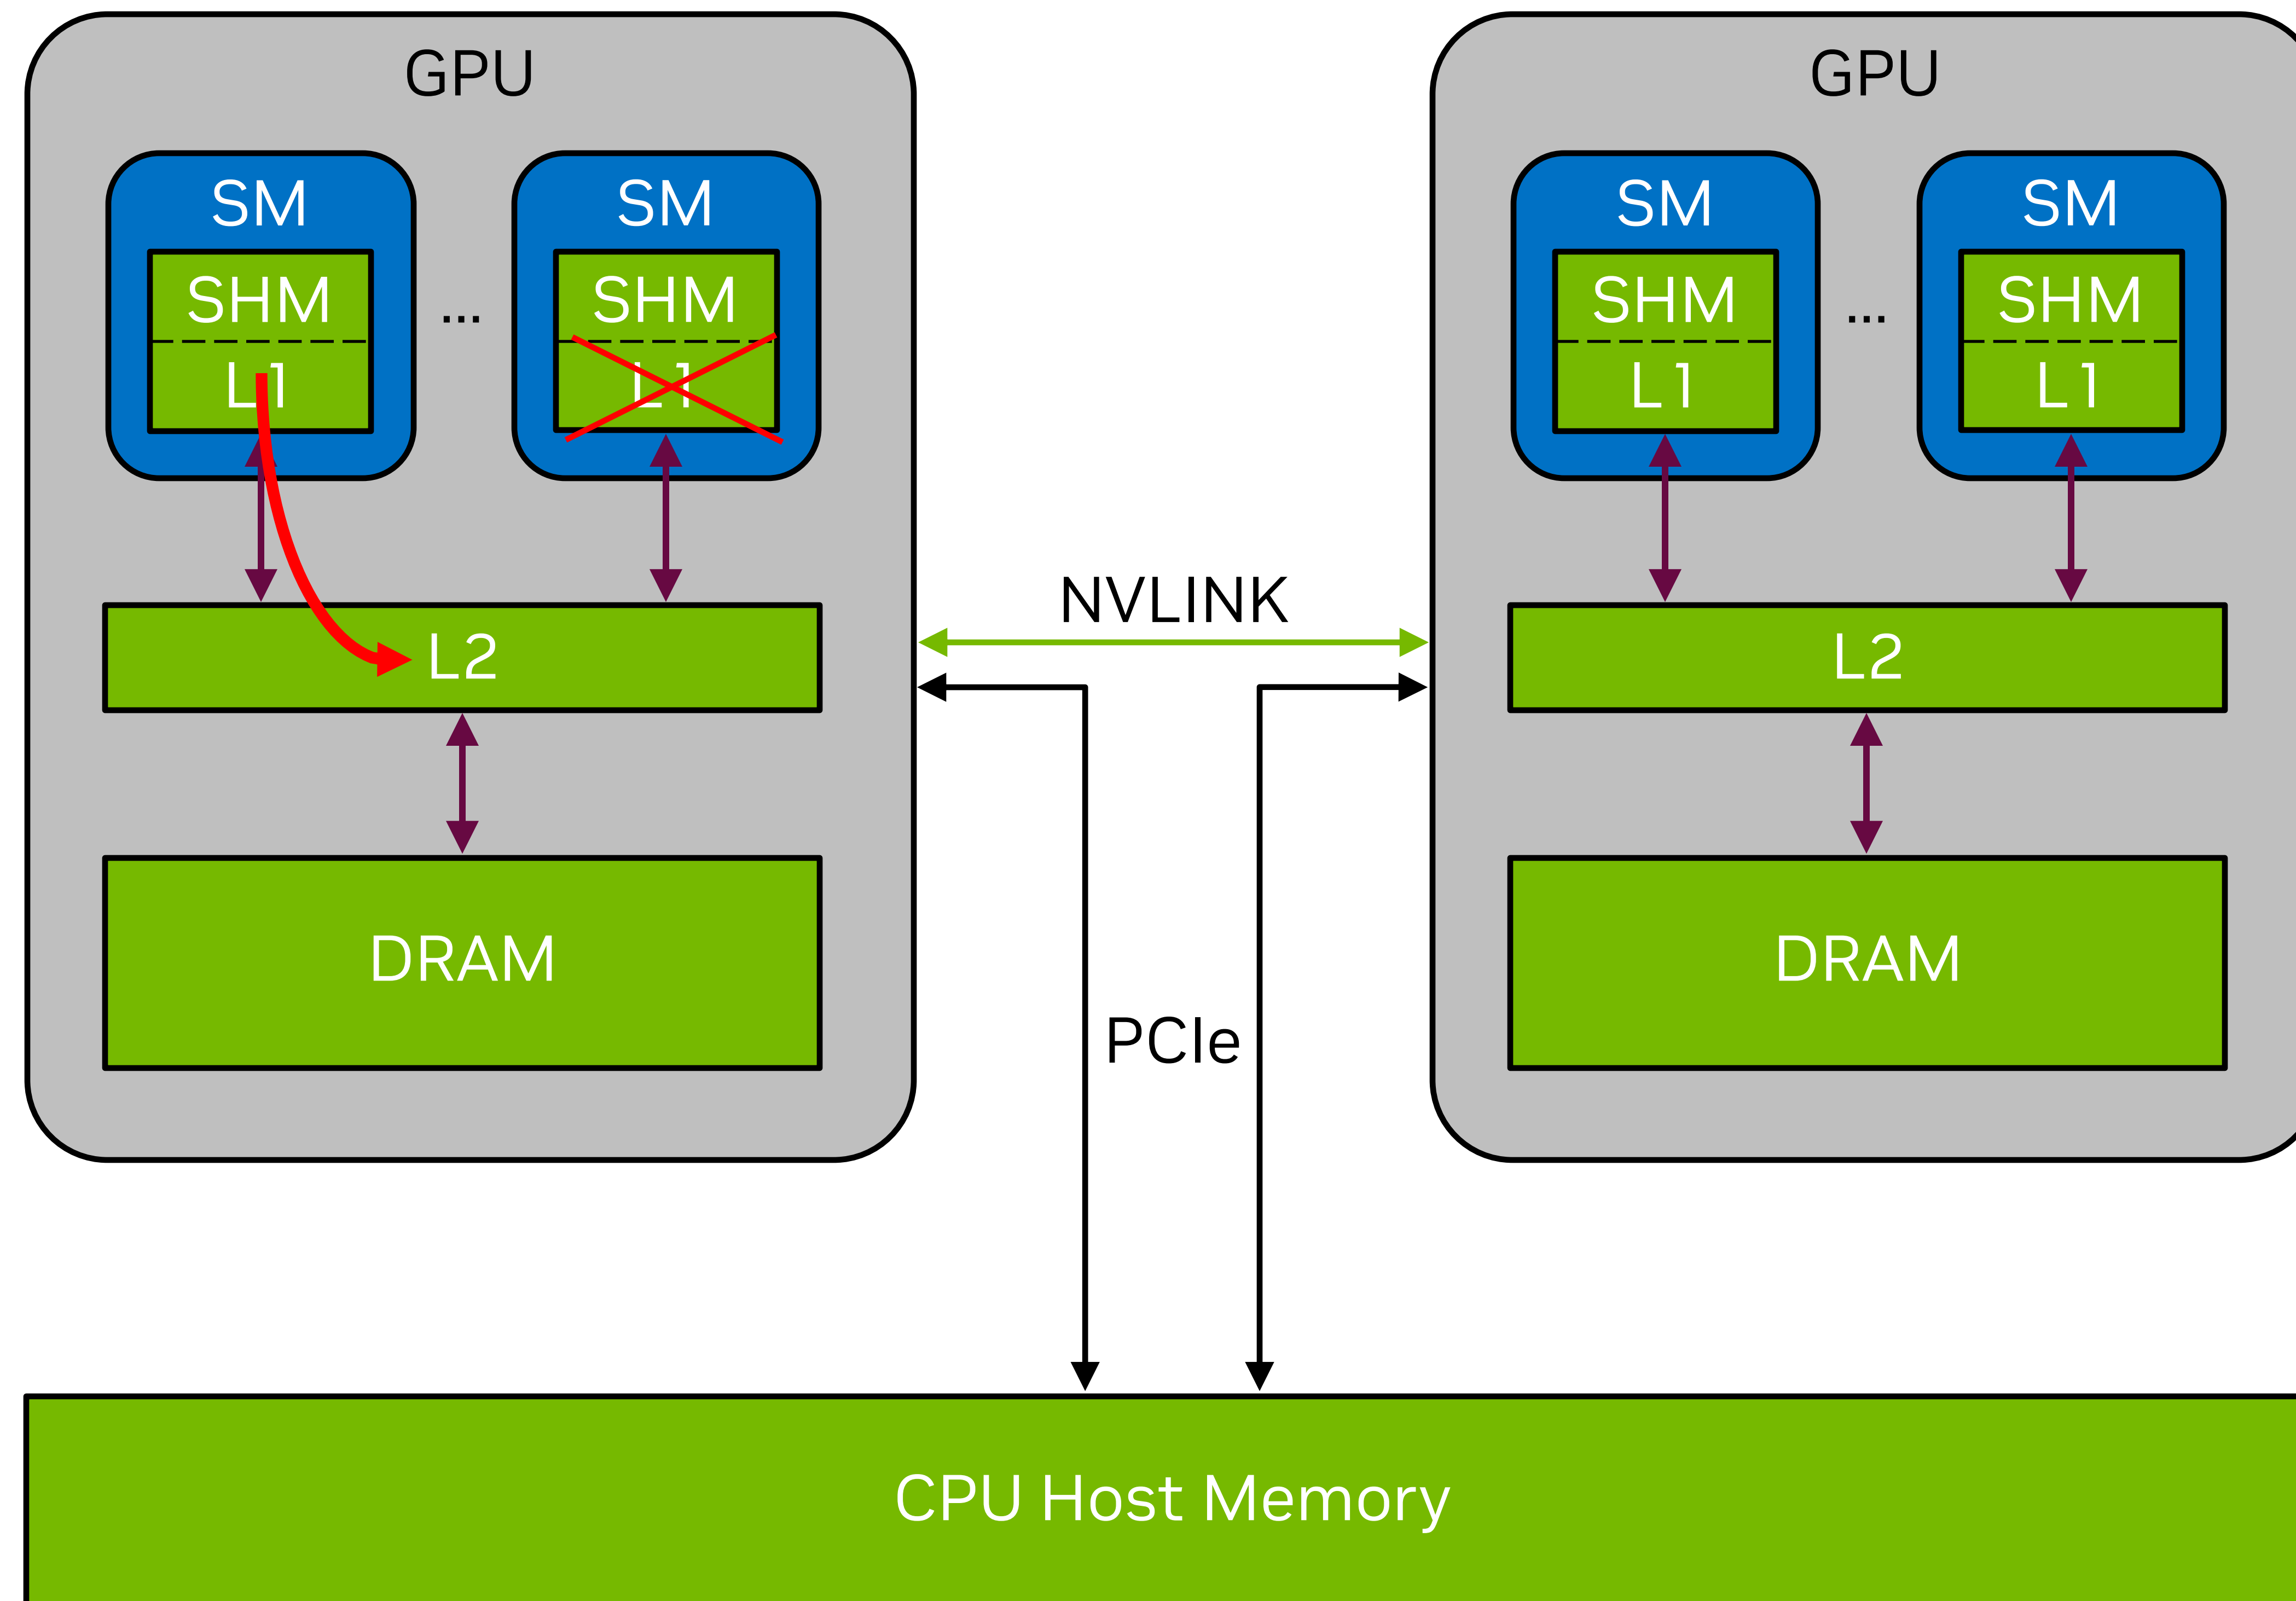


**Point of Coherency**  
depends on the level at which  
threads are communicating



# Memory Model

GPU memory hierarchy



**Point of Coherency**  
depends on the level at which  
threads are communicating

# Memory Model

Common synchronization using global memory

A less common communication pattern:

```
__global__ void kernel (T* buf)
{
```

```
    buf[ind1] = value1;        // Threads writing to global memory
```

```
    cg::this_grid().sync();    // Barrier + makes all previous writes visible to all threads in block  
                                Including global memory !
```

```
    value2 = buf[ind2];        // Threads reading values written by other threads in the same block!
```

```
}
```

requires cooperative launch

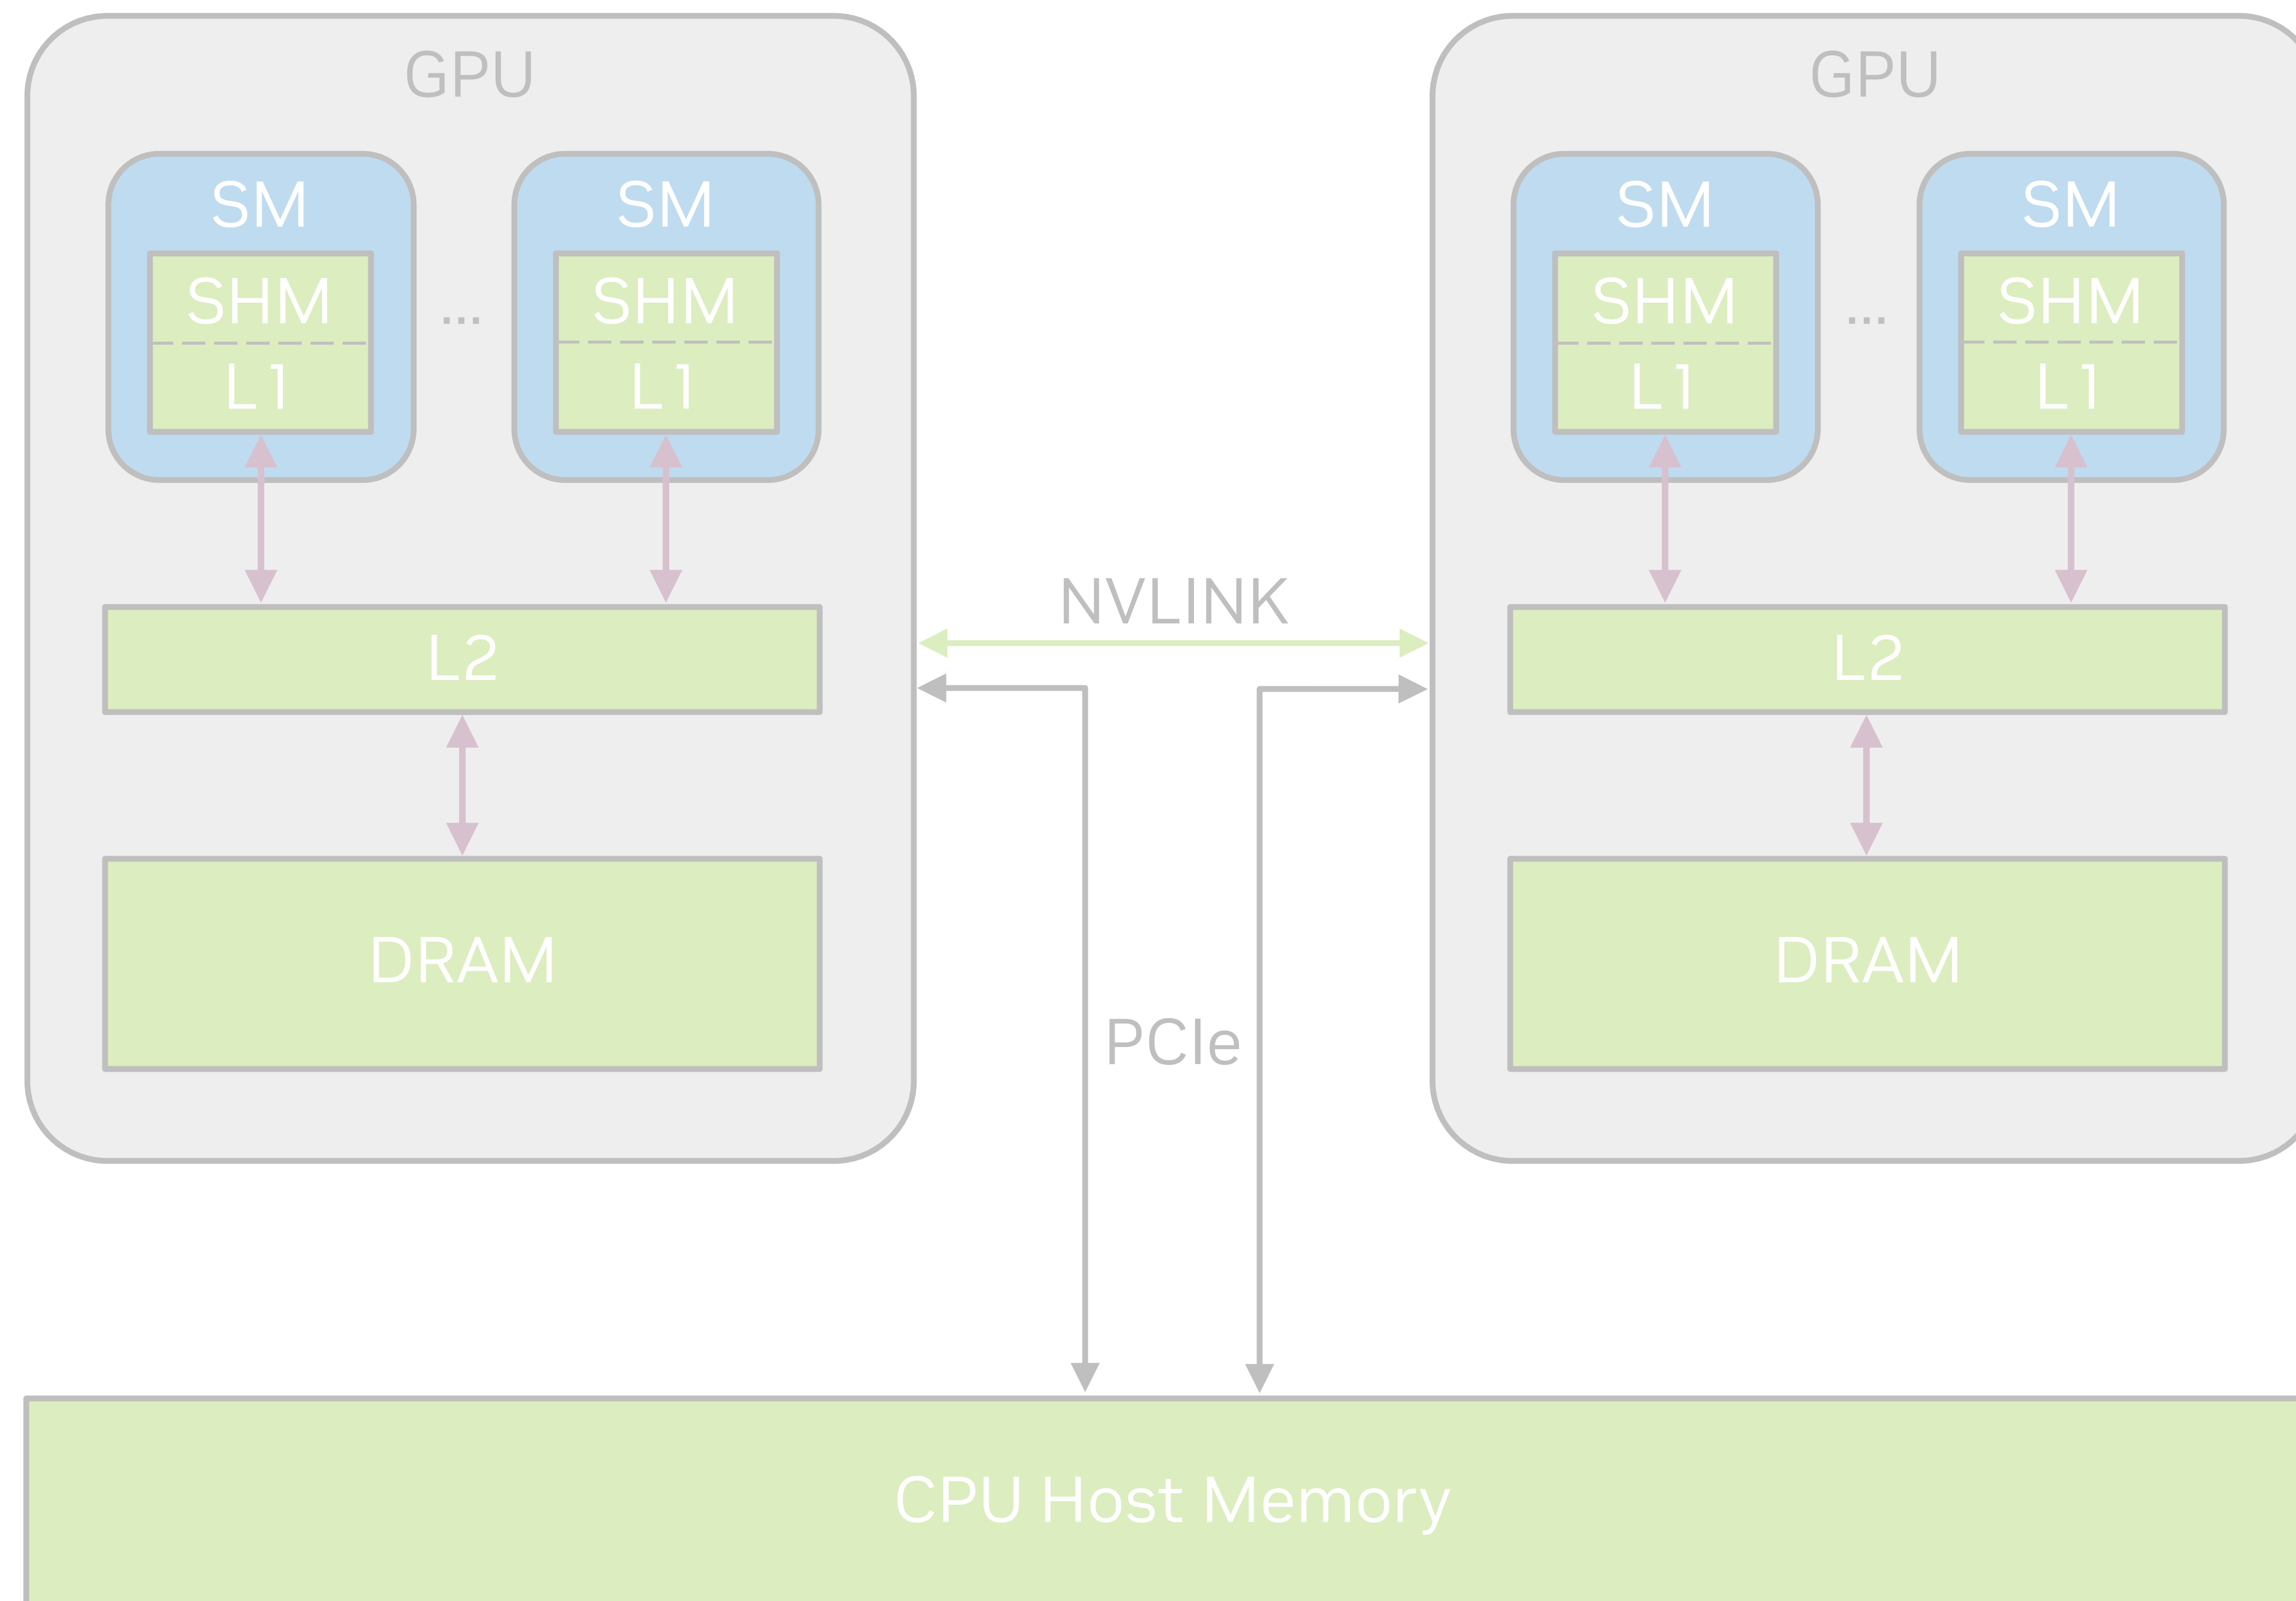
What's really happening here?

**Memory fence + barrier (device level) +  
L1 cache invalidate**




# Memory Model

GPU memory hierarchy



**Point of Coherency**  
depends on the level at which  
threads are communicating:



**All you really need to know  
about POC**

# Memory Model

Fine-grained communication

Producer/consumer, different threads running on different GPUs

```
__global__ void producer (  
    T* data, int* flag)
```

```
*data = value;           // Write data
```

```
*flag = 1;               // Then flag
```

```
__global__ void consumer (  
    T* data, int* flag)
```

```
while (*flag == 0)       // Wait for flag  
{
```

```
    value = *data;       // Then read data  
}
```

Does this work? **NO**

- ✗ 1. Strong ordering: “happens before” relation
- ✗ 2. Reaching correct Point of Coherency



# Memory Model

Fine-grained communication

Producer/consumer, different threads running on different GPUs

```
__global__ void producer (  
    T* data, int* flag)  
  
*data = value;           // Write data  
  
__threadfence_system(); // Memory fence  
  
*flag = 1;               // Then flag
```

```
__global__ void consumer (  
    T* data, int* flag)  
  
while (*flag == 0)       // Wait for flag  
{  
    __threadfence_system(); // Memory fence  
}  
  
value = *data;           // Then read data
```

Does this work? **NO**

- ✓ 1. Strong ordering: “happens before” relation
- ✗ 2. Reaching correct Point of Coherency

# Memory Model

Fine-grained communication

Producer/consumer, different threads running on different GPUs

```
__global__ void producer (  
    T* data, int* volatile flag)
```

```
*data = value;           // Write data
```

```
__threadfence_system(); // Memory fence
```

```
*flag = 1;               // Then flag
```

```
__global__ void consumer (  
    T* data, int* volatile flag)
```

```
while (*flag == 0)       // Wait for flag  
{
```

```
    __threadfence_system(); // Memory fence
```

```
    value = *data;        // Then read data  
}
```

Does this work? **YES (in practice)**

- ✓ 1. Strong ordering: “happens before” relation
- ✓ 2. Reaching correct Point of Coherency – **forced by volatile**

Is it right? **NO**

- **volatile does not imply atomicity**

Concurrent access of same memory with at least one ‘write’ is a **CONFLICT (UNDEFINED BEHAVIOR)** – remember?!



# Memory Model

Fine-grained communication

Producer/consumer, different threads running on different GPUs

```
__global__ void producer (  
    T* data, int* flag)  
  
*data = value;           // Write data  
  
__threadfence_system(); // Memory fence  
  
atomicExch_system((int*)flag, 1); // Then flag
```

```
__global__ void consumer (  
    T* data, int* flag)  
  
while (atomicAdd_system(flag, 0) == 0) // Wait for flag  
{  
    __threadfence_system(); // Memory fence  
  
    value = *data;           // Then read data  
}
```

Does this work? **YES**

- ✓ 1. Strong ordering: “happens before” relation
- ✓ 2. Reaching correct Point of Coherency

Is it right? **YES**

But...

- Not the most elegant/efficient
- Cannot communicate with a CPU – no compatible atomic

# Memory Model

Fine-grained communication

Producer/consumer, different threads running on different GPUs

```
__global__ void producer (  
    T* data,  
    cuda::atomic<int, thread_scope_system>* flag)
```

```
*data = value;           // Write data
```

```
*flag = 1;               // Then flag  
flag->notify_all();
```

```
__global__ void consumer (  
    T* data,  
    cuda::atomic<int, thread_scope_system>* flag)
```

```
flag->wait(0);           // Wait for flag
```

```
value = *data;           // Then read data
```

C++ atomics have a built-in memory fence

Does this work? **YES**

- ✓ 1. Strong ordering: “happens before” relation
- ✓ 2. Reaching correct Point of Coherency

Is it right? **YES, elegant, C++ atomics can communicate with CPU**

But...

- Still not the most efficient



# Memory Model

## Memory order

*“Memory order specifies how memory accesses, including regular (non-atomic) accesses, are to be ordered around an atomic operation”*

Default behavior of `std::atomic` is **sequentially-consistent**:

- Enforces strict ordering not only between atomics but between all memory ops  
Think about it as **“slow for HW to do”**

# Memory Model

Fine-grained communication

Producer/consumer, different threads running on different GPUs

```
__global__ void producer (  
    T* data,  
    cuda::atomic<int, thread_scope_system>* flag)
```

```
*data = value;           // Write data
```

```
*flag = 1;               // Then flag  
flag->notify_all();
```

Only need preceding writes not reordered after flag-write

```
__global__ void consumer (  
    T* data,  
    cuda::atomic<int, thread_scope_system>* flag)
```

```
flag->wait(0);           // Wait for flag
```

```
value = *data;           // Then read data
```



# Memory Model

Fine-grained communication

Producer/consumer, different threads running on different GPUs

```
__global__ void producer (  
    T* data,  
    cuda::atomic<int, thread_scope_system>* flag)
```

```
*data = value;           // Write data
```

```
*flag = 1;               // Then flag  
flag->notify_all();
```

Only need preceding writes not reordered after flag-write

```
__global__ void consumer (  
    T* data,  
    cuda::atomic<int, thread_scope_system>* flag)
```

```
flag->wait(0);           // Wait for flag
```

```
value = *data;           // Then read data
```

Only need subsequent reads not reordered before flag-read

# Memory Model

Acquire / release

*“Memory order specifies how memory accesses, including regular (non-atomic) accesses, are to be ordered around an atomic operation”*

Default behavior of `std::atomic` is **sequentially consistent**:

- Enforces strict ordering not only between atomics but between all memory ops  
Think about it as **“slow for HW to do”**

Fine-grained ordering around atomic ops:

- `memory_model::acquire`

No reads or writes in the current thread can be moved before the acquire

Makes writes from other threads (that use same atomic) visible to the current thread

- `memory_model::release`

No reads or writes in the current thread can be moved after the release

Makes writes from the current thread visible to the other threads (that use same atomic)

- Other orders



# Memory Model

Fine-grained communication

Producer/consumer, different threads running on different GPUs

```
__global__ void producer (  
    T* data,  
    cuda::atomic<int, thread_scope_system>* flag)
```

```
*data = value;           // Write data
```

```
flag->write(1, memory_order_release); // Then flag  
flag->notify_all();
```

```
__global__ void consumer (  
    T* data,  
    cuda::atomic<int, thread_scope_system>* flag)
```

```
flag->wait(0, memory_order_acquire); // Wait for flag
```

```
value = *data;           // Then read data
```

**CORRECT, ELEGANT, and EFFICIENT !!!**

# Memory Model

Cumulativity

## Thread 1

```
std::atomic<int, thread_scope_device>* flag_dev;  
*data = 42;  
flag_dev->write(1, memory_order_release);
```

## Thread 2

```
std::atomic<int, thread_scope_system>* flag_sys;  
flag_dev->wait(0, memory_order_acquire);  
assert(*data == 42);  
flag_sys->write(1, memory_order_release);
```

DEVICE 1

## Thread 3

```
flag_sys->wait(0, memory_order_acquire);  
assert(*data == 42);
```

DEVICE 2





# **Asynchronous Barriers**

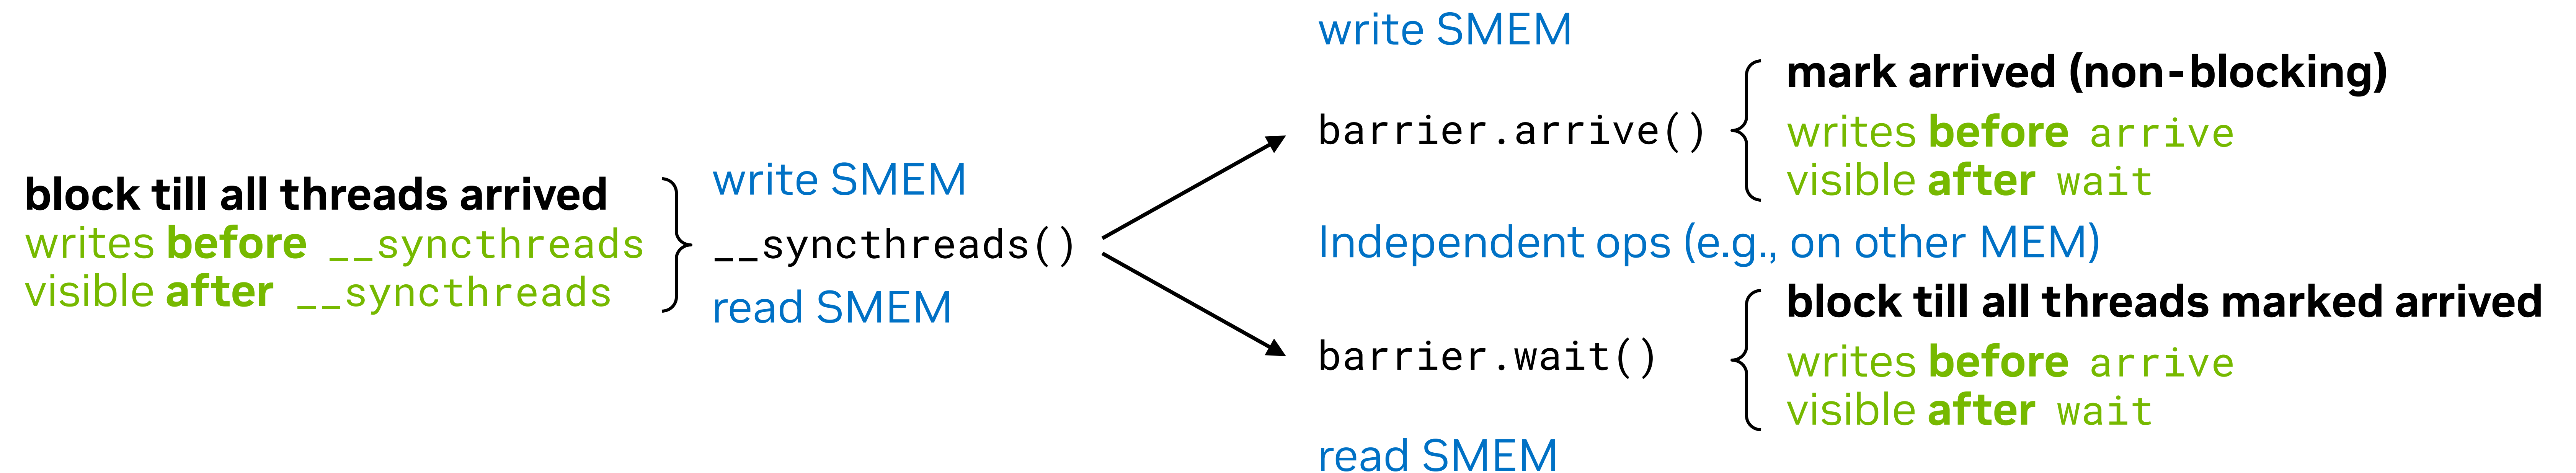


# Asynchronous Barriers

## Barrier intro

**CUDA barriers provide (within scope!):**

- Synchronization points (including async mem transactions)
- Memory ordering





# Asynchronous Barriers

## Barrier details

### Barrier

#### Tracked quantities:

- expected\_arrival\_count ( $\geq 0$ )
- transaction\_count ( $\geq 0$ ) – optional

#### State:

- arrival\_count
- phase (0/1)

#### Methods:

- arrive
- wait

### Initialization:

- expected\_arrival\_count := user-provided value
  - arrival\_count := expected\_arrival\_count
  - phase := 0
- 
- transaction\_count – incremented by a user-provided value;  
(Hopper++) decremented (by HW) on completion of memory transactions  
(allows to synchronize async copies)

# Asynchronous Barriers

## Barrier details

### Barrier

#### Tracked quantities:

- expected\_arrival\_count ( $\geq 0$ )
- transaction\_count ( $\geq 0$ ) – optional

#### State:

- arrival\_count
- phase (0/1)

#### Methods:

- arrive
- wait

**Arrive** (single atomic non-blocking operation):

1. arrival\_count -= 1 (or -= n)
2. IF arrival\_count == 0  
    phase ^= 1 **WHEN TRANSACTION COUNT REACHES 0**  
    arrival\_count := expected\_arrival\_count (barrier reset)

**Wait** (two ways to wait for phase flip):

- Token-based. Typical code:  
    auto token = bar.arrive();  
    bar.wait(std::move(token)); **// more expensive than explicit tracking**
- Explicitly tracking the phase. Typical code:  
    while (!mbarrier\_try\_wait\_parity(bar, curr\_phase, 1000)) {}  
    curr\_phase ^= 1; **// need to explicitly keep track of phase**

#### Initialization:

- expected\_arrival\_count := user-provided value
  - arrival\_count := expected\_arrival\_count
  - phase := 0
- 
- transaction\_count – incremented by a user-provided value;  
    **(Hopper++)** decremented (by HW) on completion of memory transactions  
    (allows to synchronize async copies)

# Asynchronous Barriers

## Summary

HW-accelerated  
on Hopper

Scope	Synchronizes	Typical Memory Location	Initialization	Arrive	Wait
Block	CTA	Local SMEM	Single-thread initialization (followed by CTA sync)	ALLOWED	ALLOWED
Cluster	CGA	Local SMEM (owner CTA)	Single-thread initialization (followed by CGA sync)	ALLOWED	ALLOWED
		Remote DSMEM (after mapping)	NOT ALLOWED	ALLOWED	NOT ALLOWED
Device	GPU	GMEM	Single-thread initialization (followed by device sync#)	ALLOWED	ALLOWED
System	NODE	Accessible by all actors: <ul style="list-style-type: none"><li>• GMEM*</li><li>• Uniform (managed)</li></ul>	Single-thread initialization	ALLOWED*	ALLOWED*

# – cooperative launch required

\* – to access from non-owning GPU, peer access must be enabled



# Asynchronous Barriers

Barrier location

## Producer:

```
Loop {  
  Produce data for Consumer  
  Arrive at Data barrier  
  (notifies Consumer that data is ready)  
  Wait at Copy barrier  
  (waits for next iteration notification)  
}
```

## Consumer:

```
Loop {  
  Wait at Data barrier  
  (waits for data ready notification)  
  Use produced data  
  Arrive at Copy barrier  
  (notifies Consumer that ready for next iteration)  
}
```

GPU-GPU Timings

		Data barrier location	
		Producer	Consumer
Copy barrier location	Consumer	0.035	0.065
	Producer	0.065	0.032


Recommended barrier location:  
where waiting

CPU-GPU Timings

		Data barrier location	
		Producer	Consumer
Copy barrier location	Consumer	0.27	0.27
	Producer	3.39	3.39

Recommended barrier location:  
local to most threads (GPU)





# **Asynchronous Data Copies**



# Asynchronous Data Copies

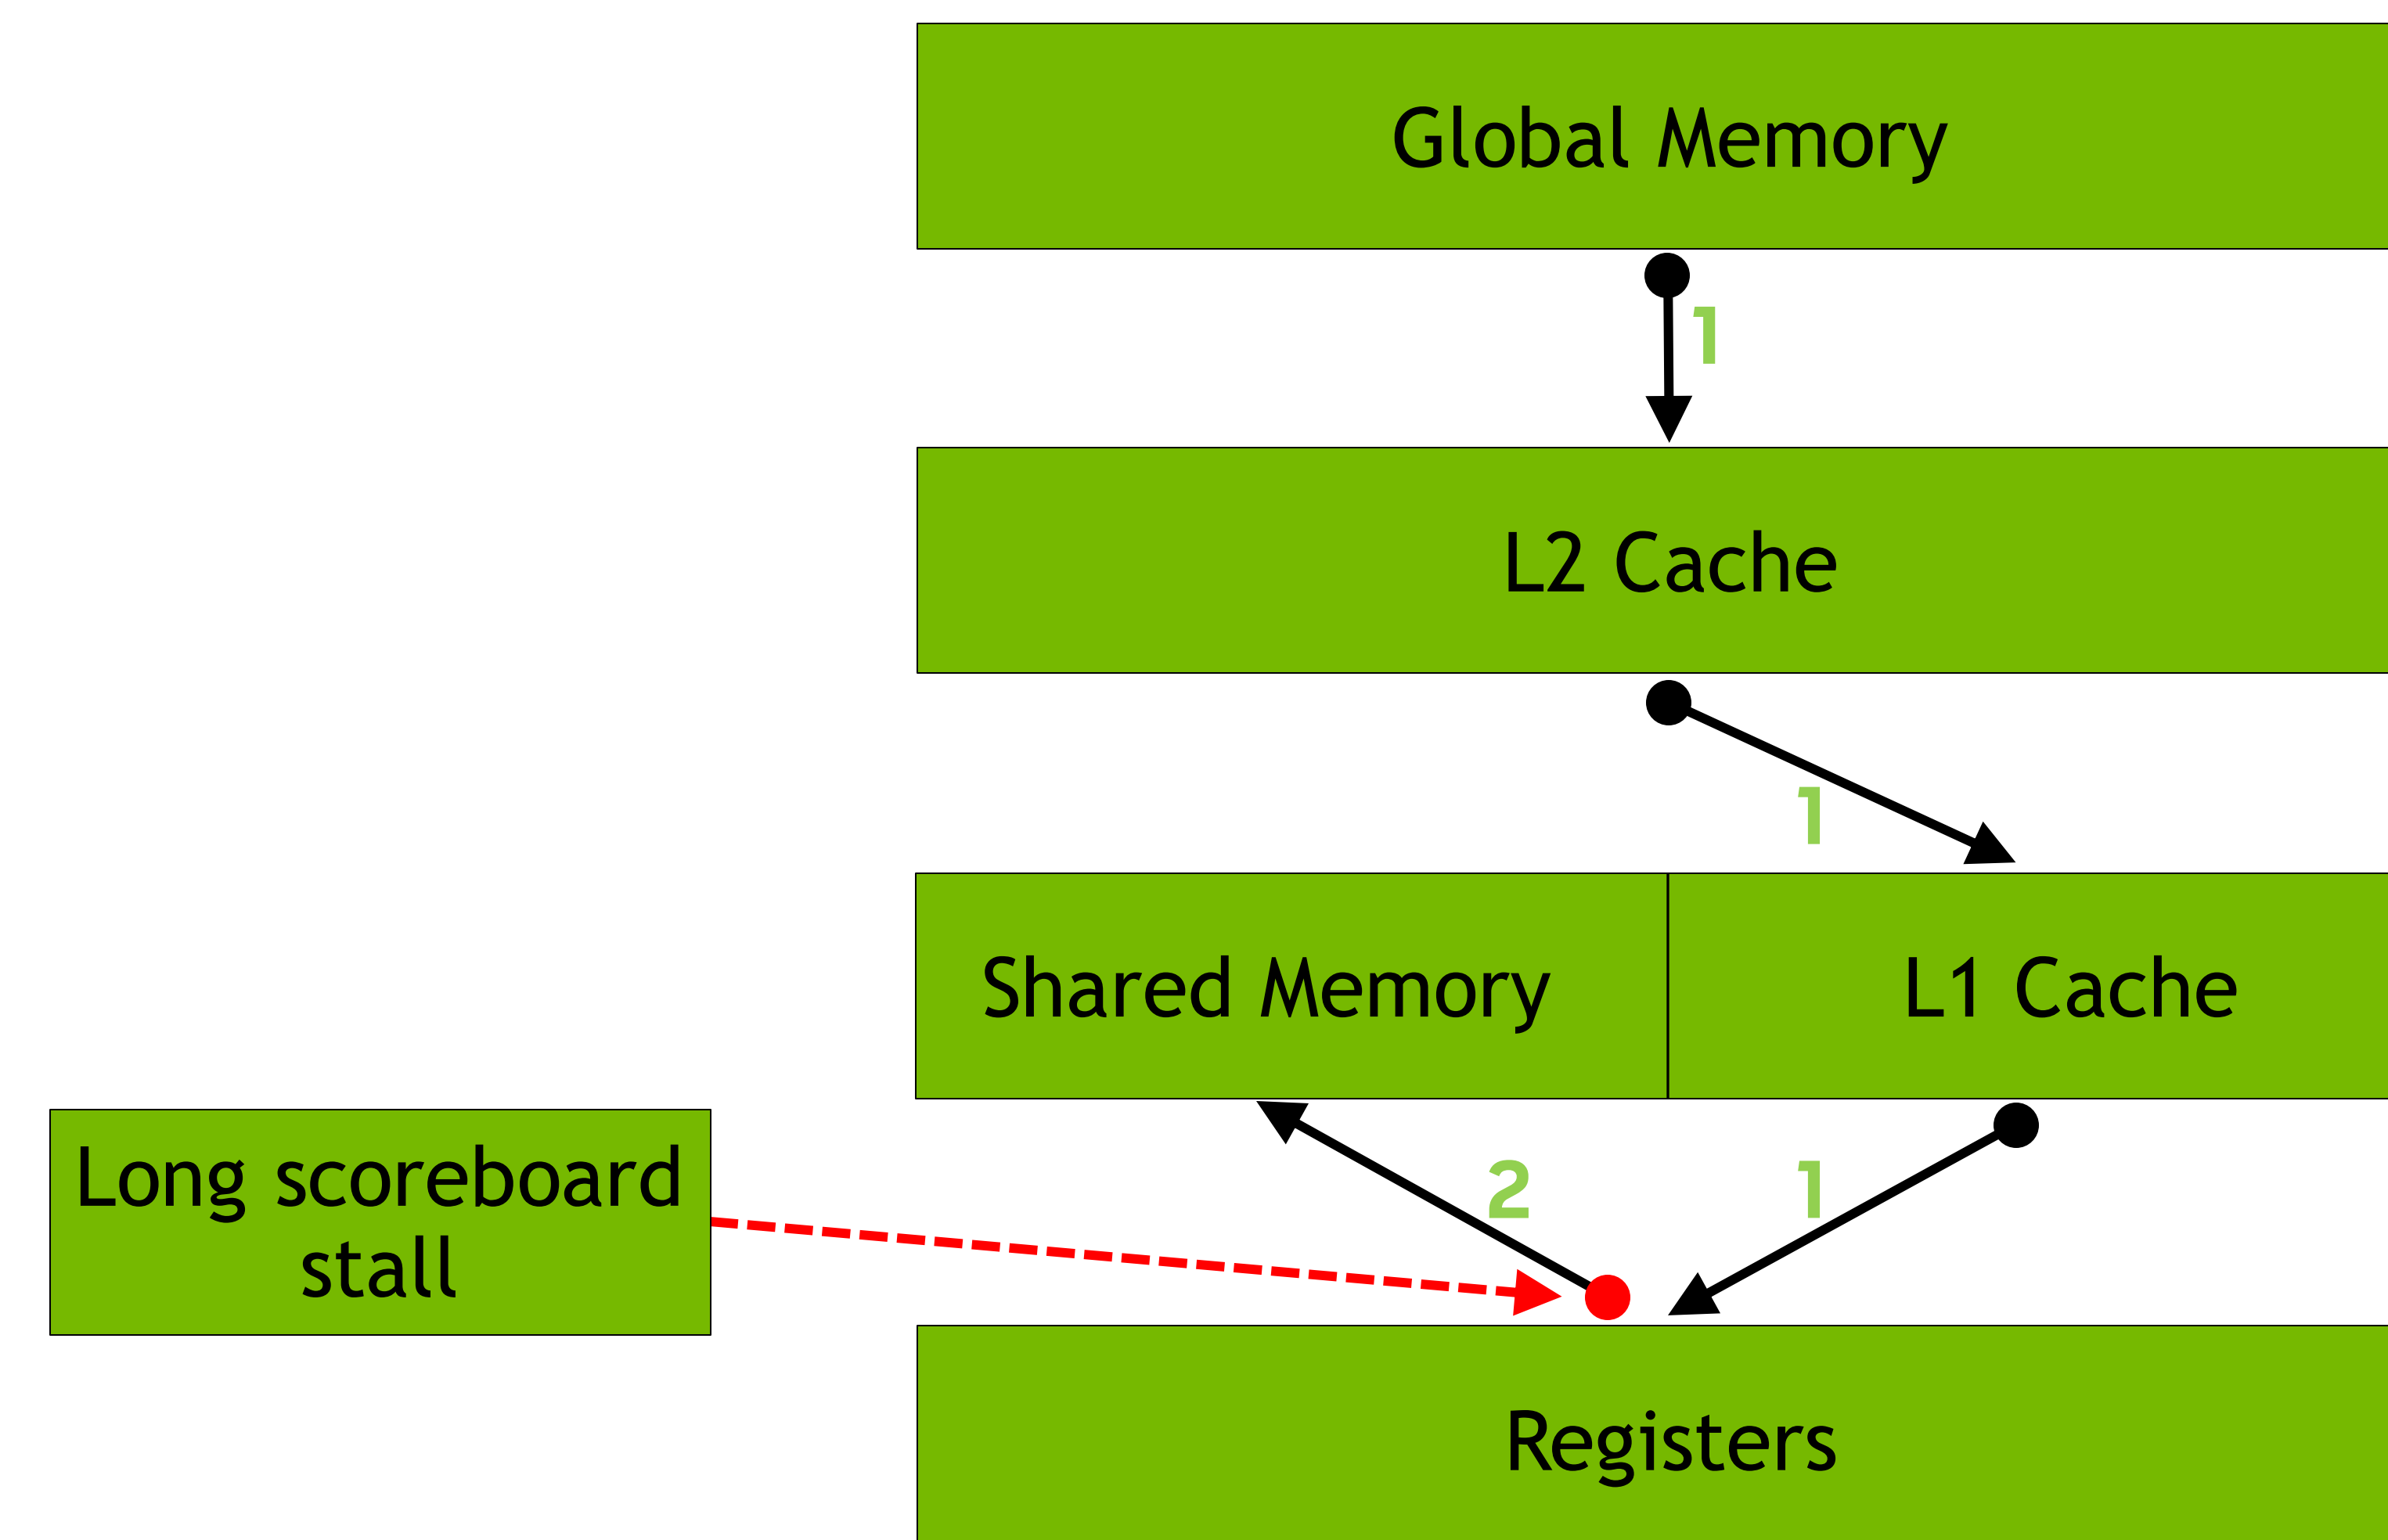
Synchronous copies dissected

```
T* gmem;  
__shared__ T smem[ ];  
smem[sind] = gmem[gind];
```

Involves:

1. Copying data from global memory to register
2. Copying data from register to shared memory

- **STALL** [Scoreboard wait]
- Wasting registers and L1 bandwidth





# Asynchronous Data Copies

Synchronous copies dissected

```
T* gmem;  
__shared__ T smem[ ];  
smem[sind] = gmem[gind];
```

Involves:

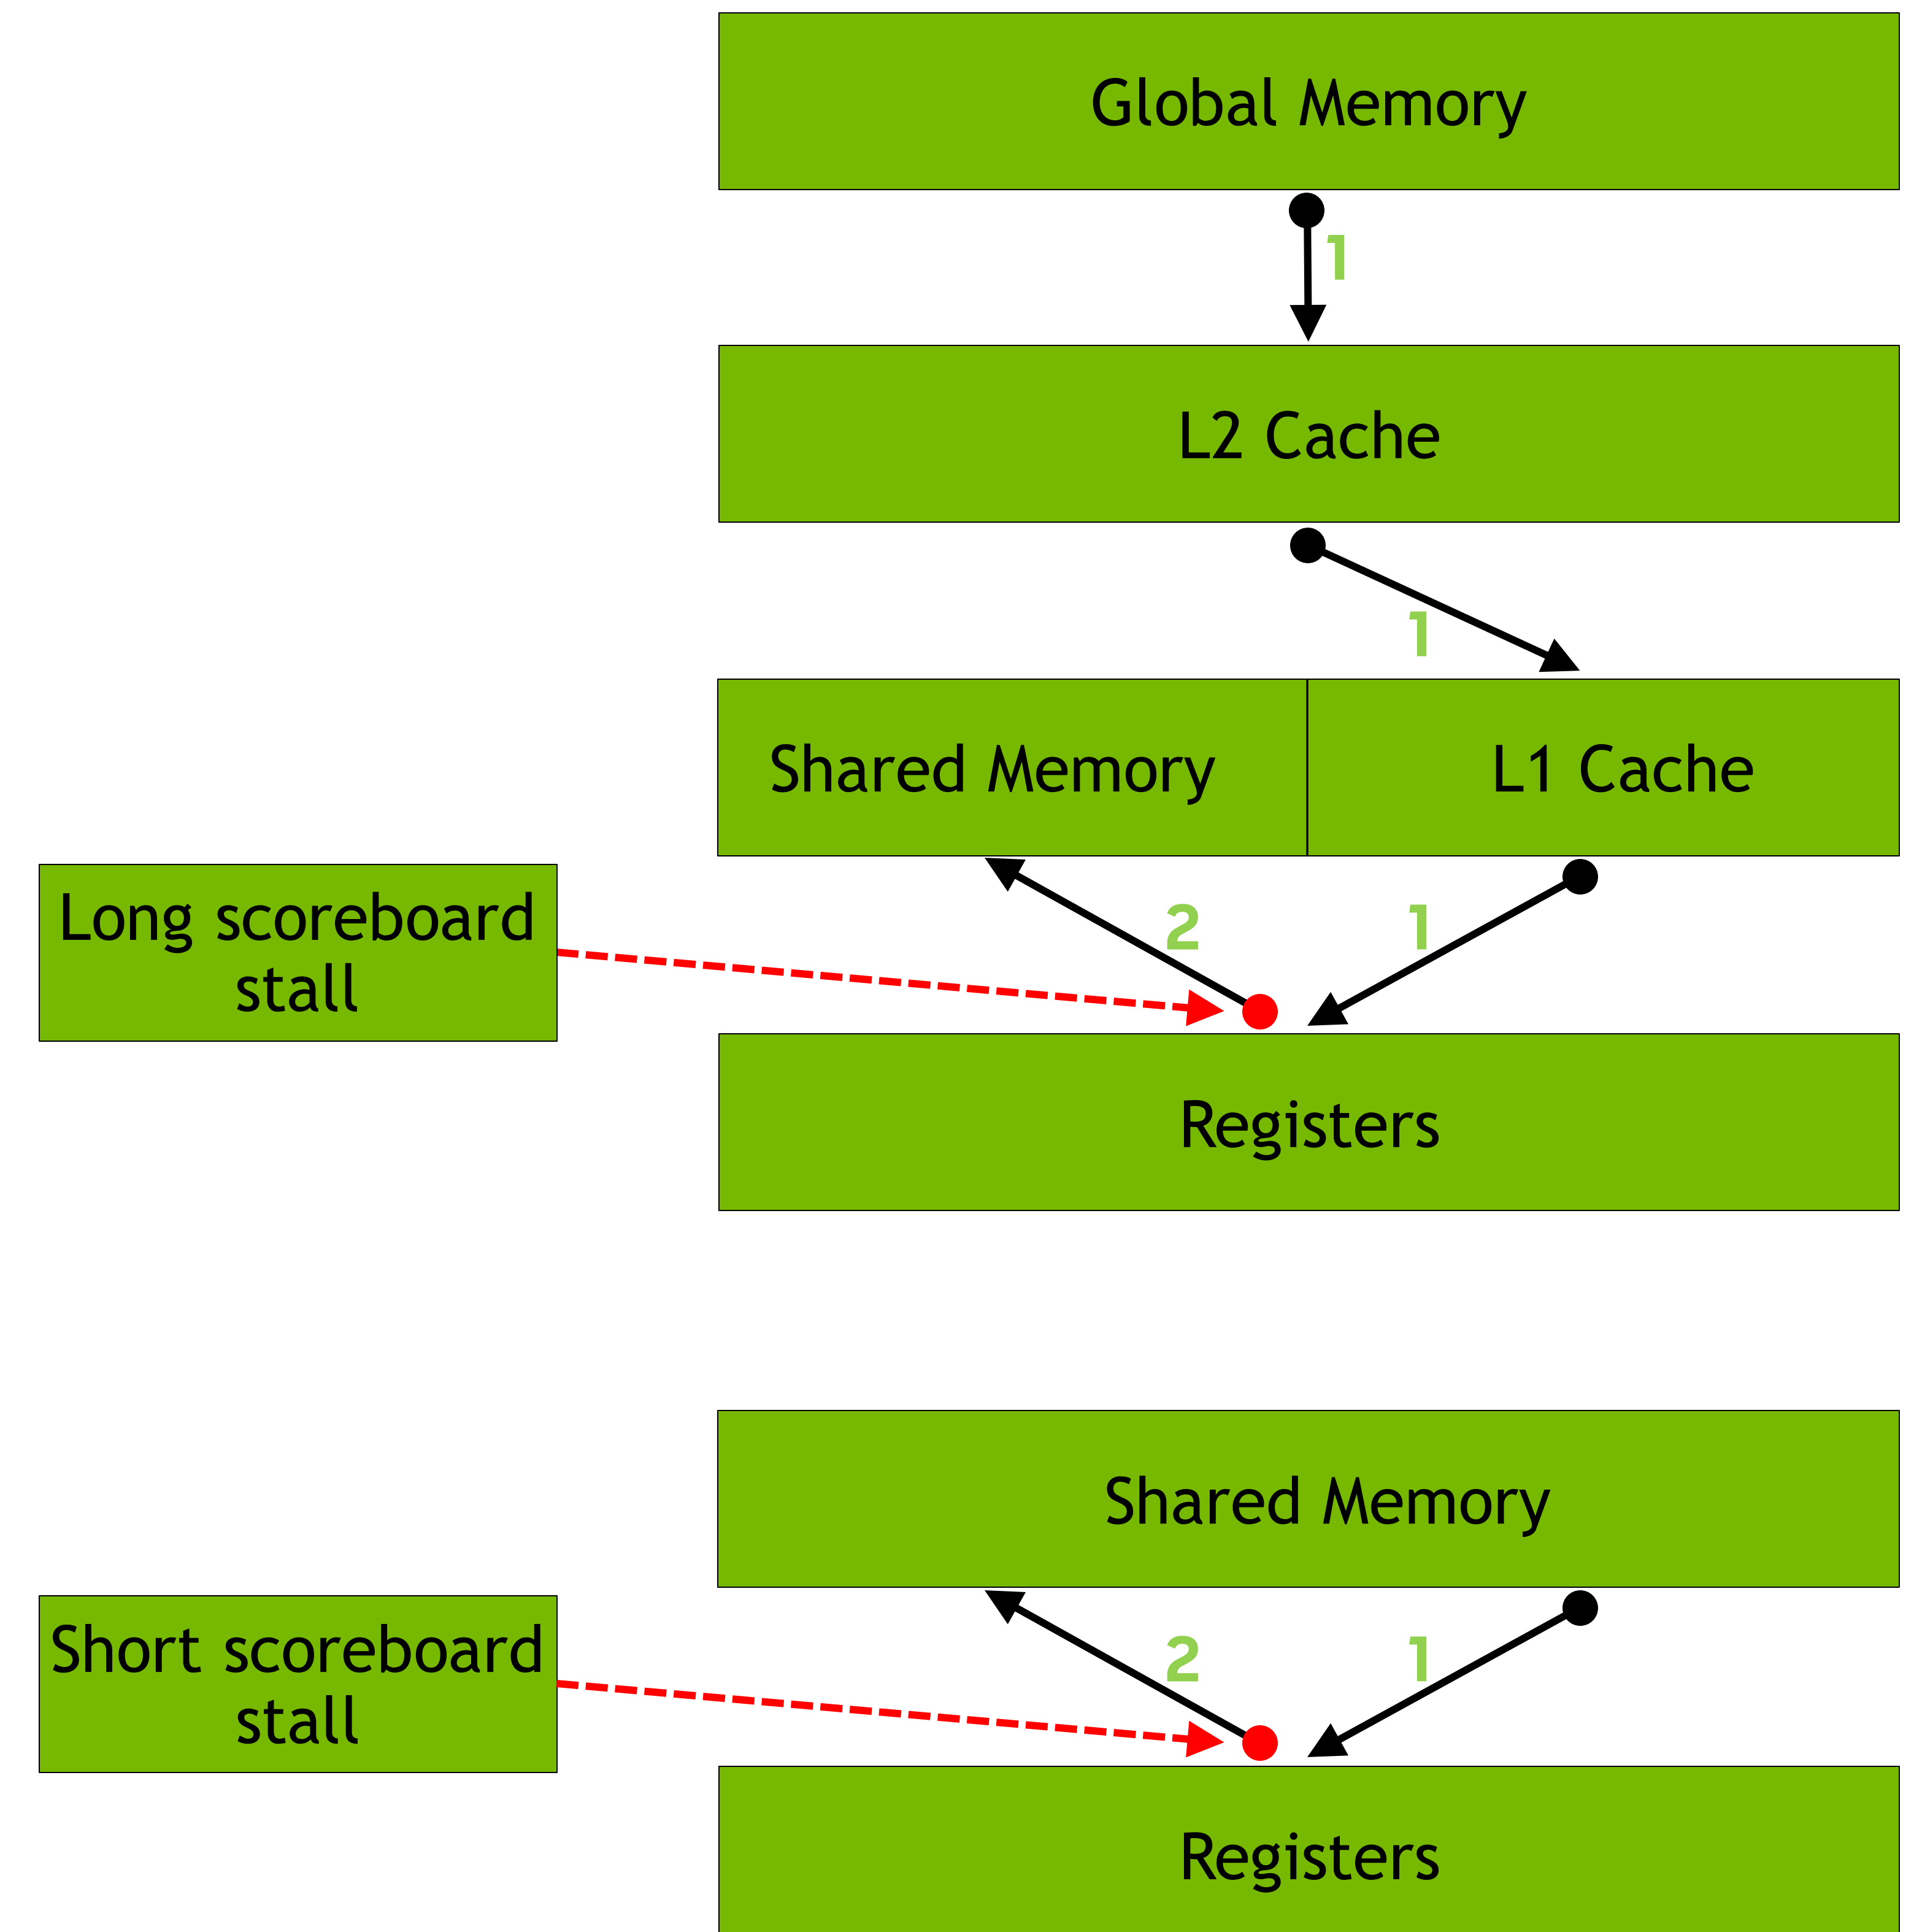
1. Copying data from global memory to register
2. Copying data from register to shared memory

- **STALL** [Scoreboard wait]
- Wasting registers and L1 bandwidth

```
smem[sind1] = smem[sind2];
```

Involves:

1. Copying data from shared memory to register
2. Copying data from register to shared memory



# Asynchronous Data Copies

Synchronous copies – less bytes in flight ☹️

Representative code:

```
T* gmem;
__shared__ T smem[..];
for (int i = 0; i < n; ++i) {
    smem[sind] = g_mem[gind];
    gind += stride;
    __syncthreads();
    // compute with smem[]
    __syncthreads();
}
```



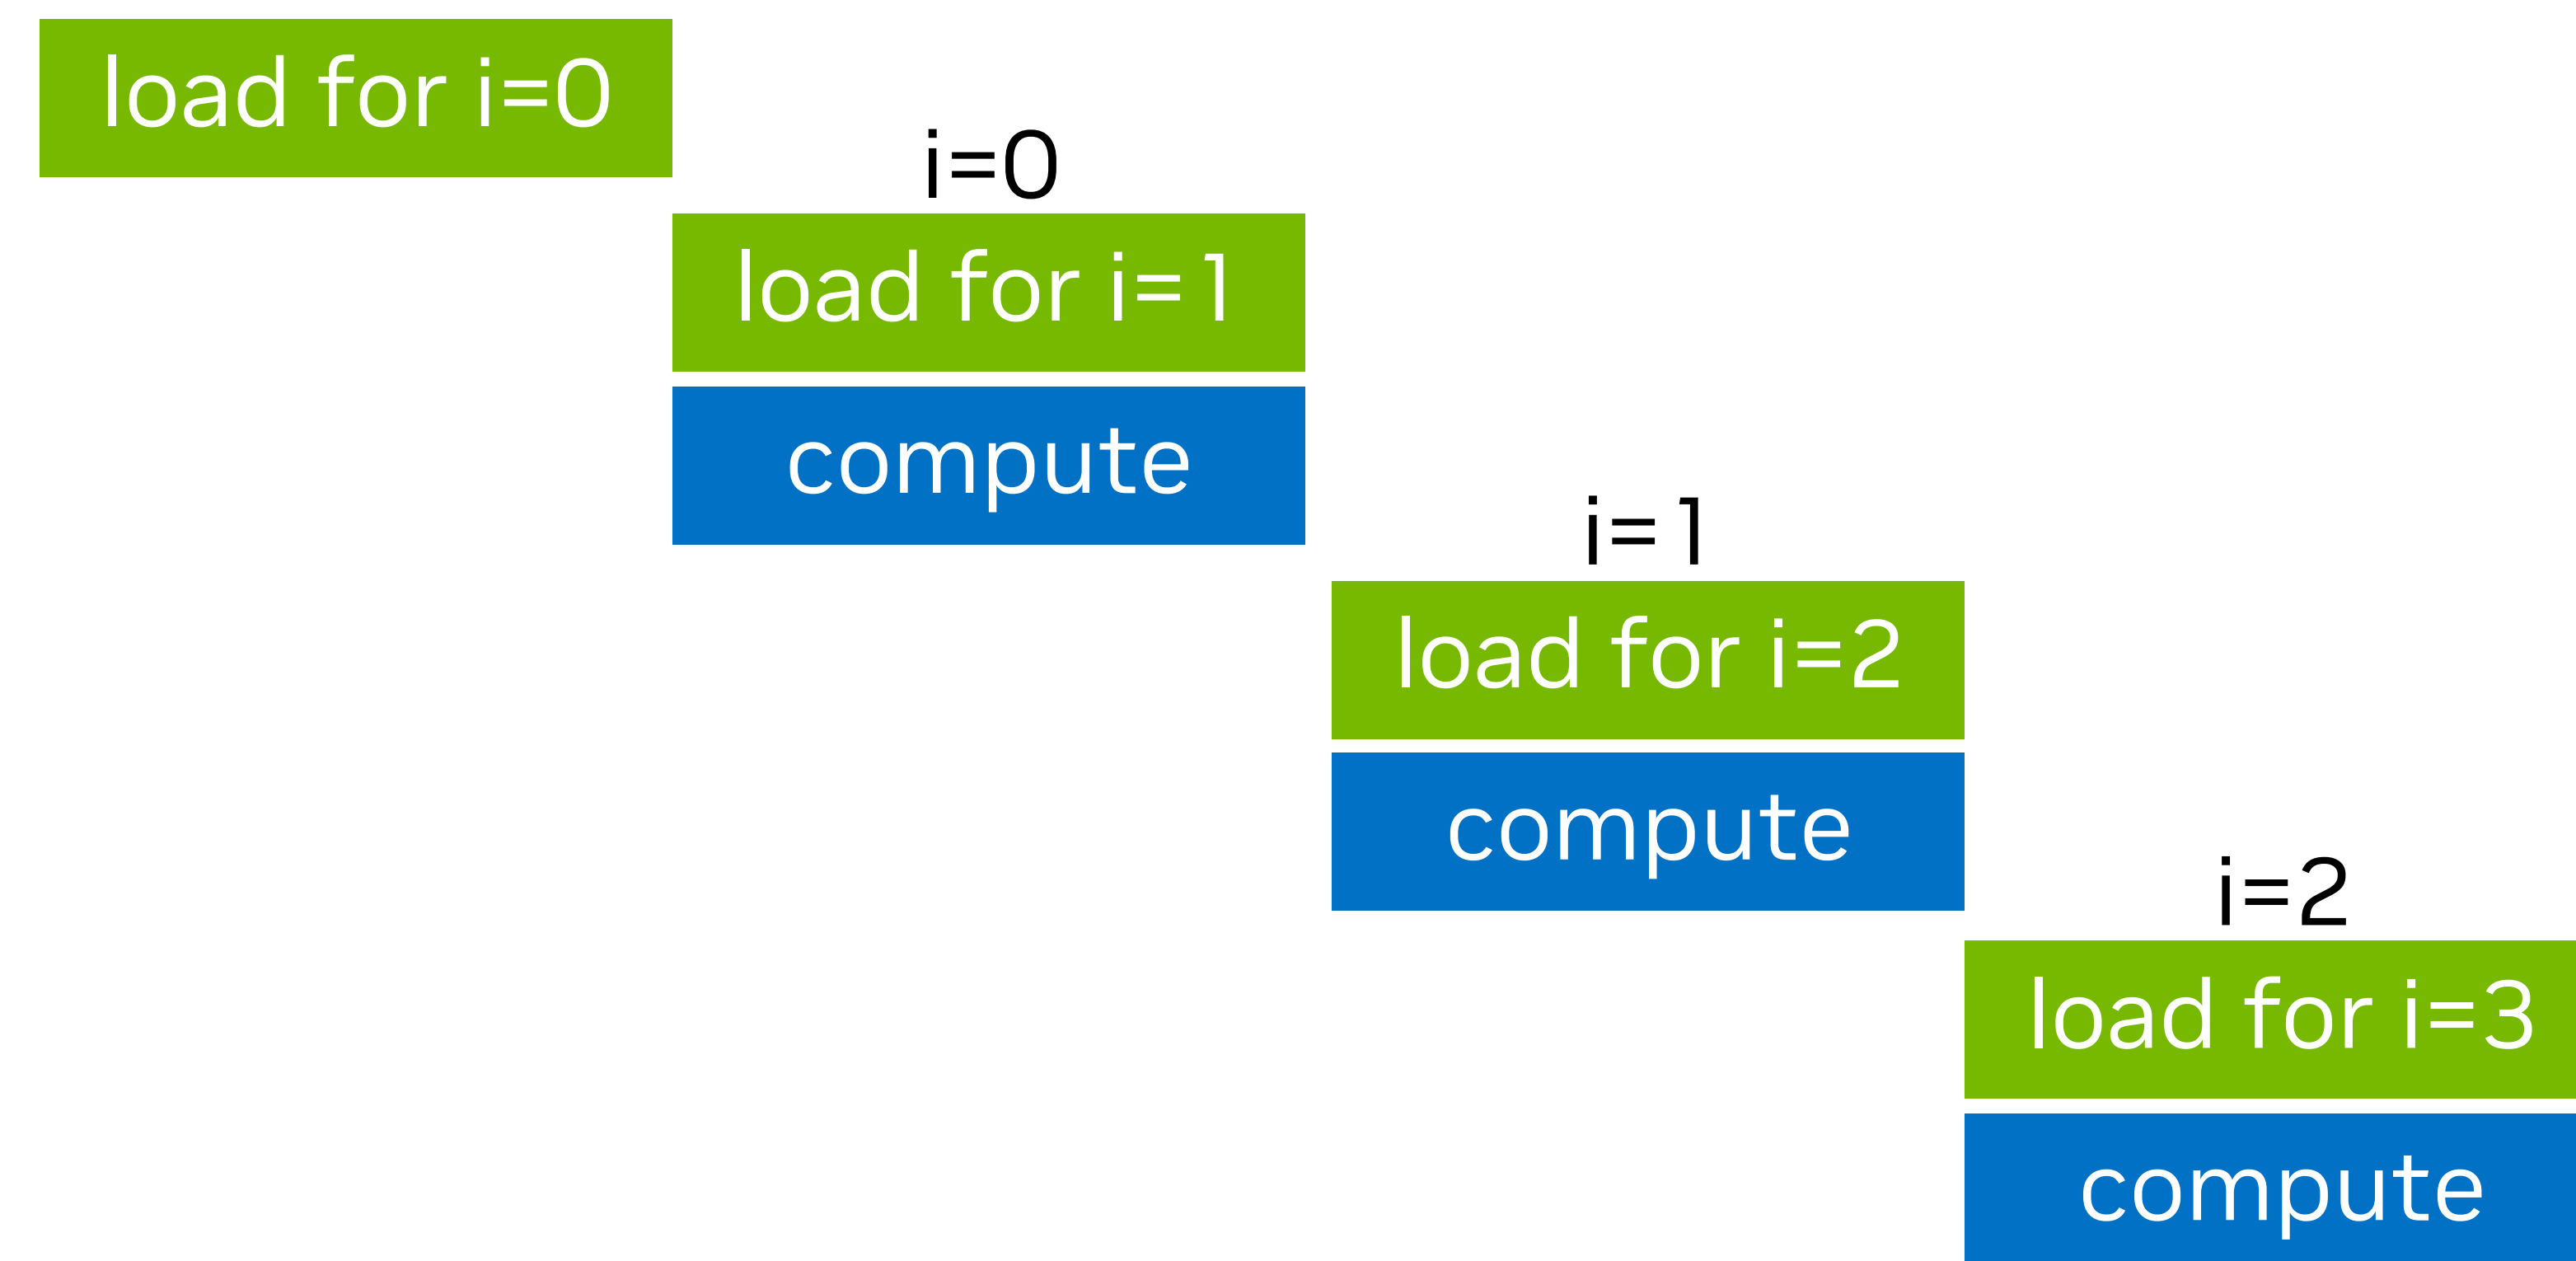
Copy and compute are sequenced.  
Cannot fire SMEM prefetch during compute...

# Asynchronous Data Copies

Asynchronous copies – more bytes in flight 😊

## Asynchronous copy benefits:

- **Increases bytes in flight.**
- Other benefits:
  - Register bypass
  - Less L1 traffic
  - Less MIO pressure (fewer instructions)
  - Better dependency management



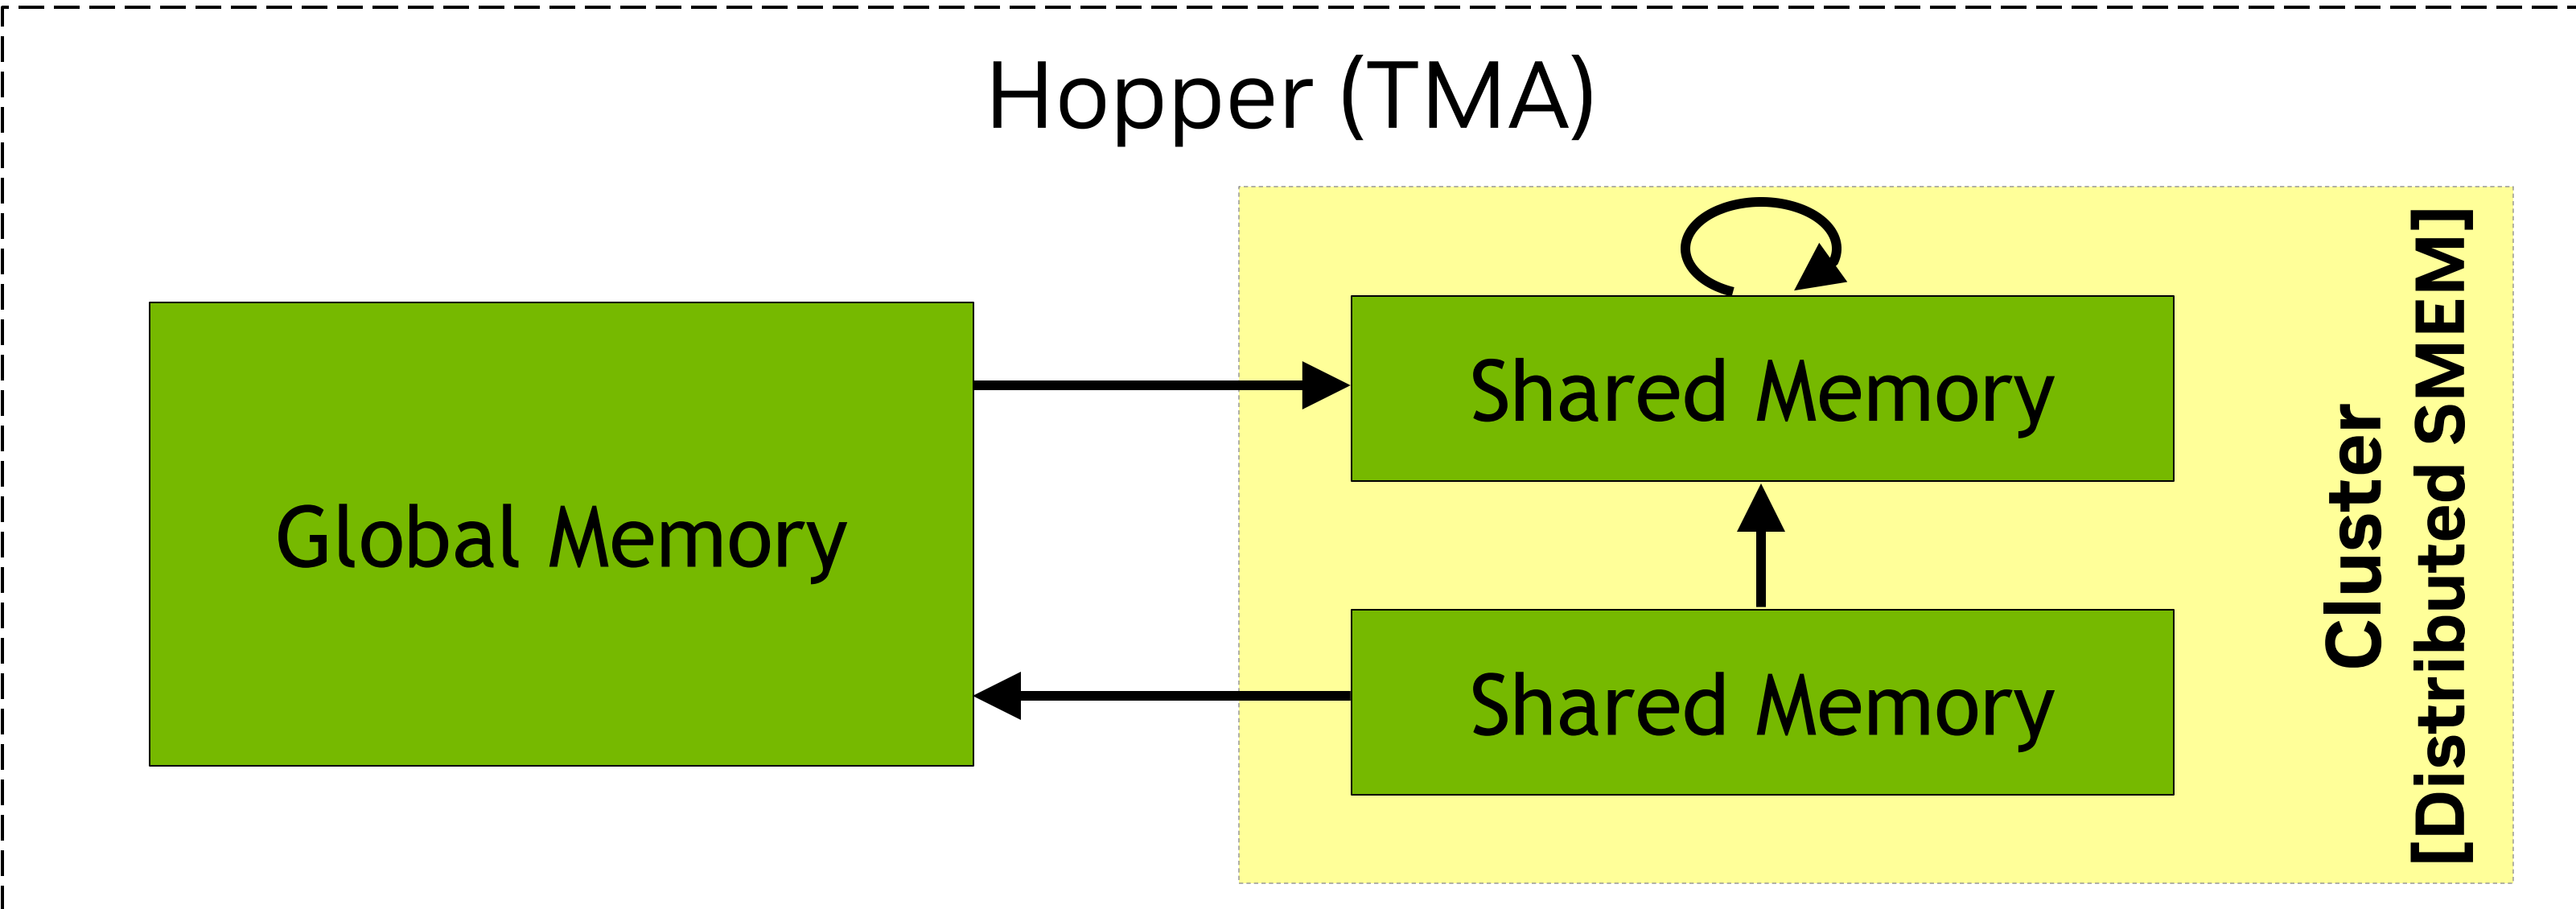
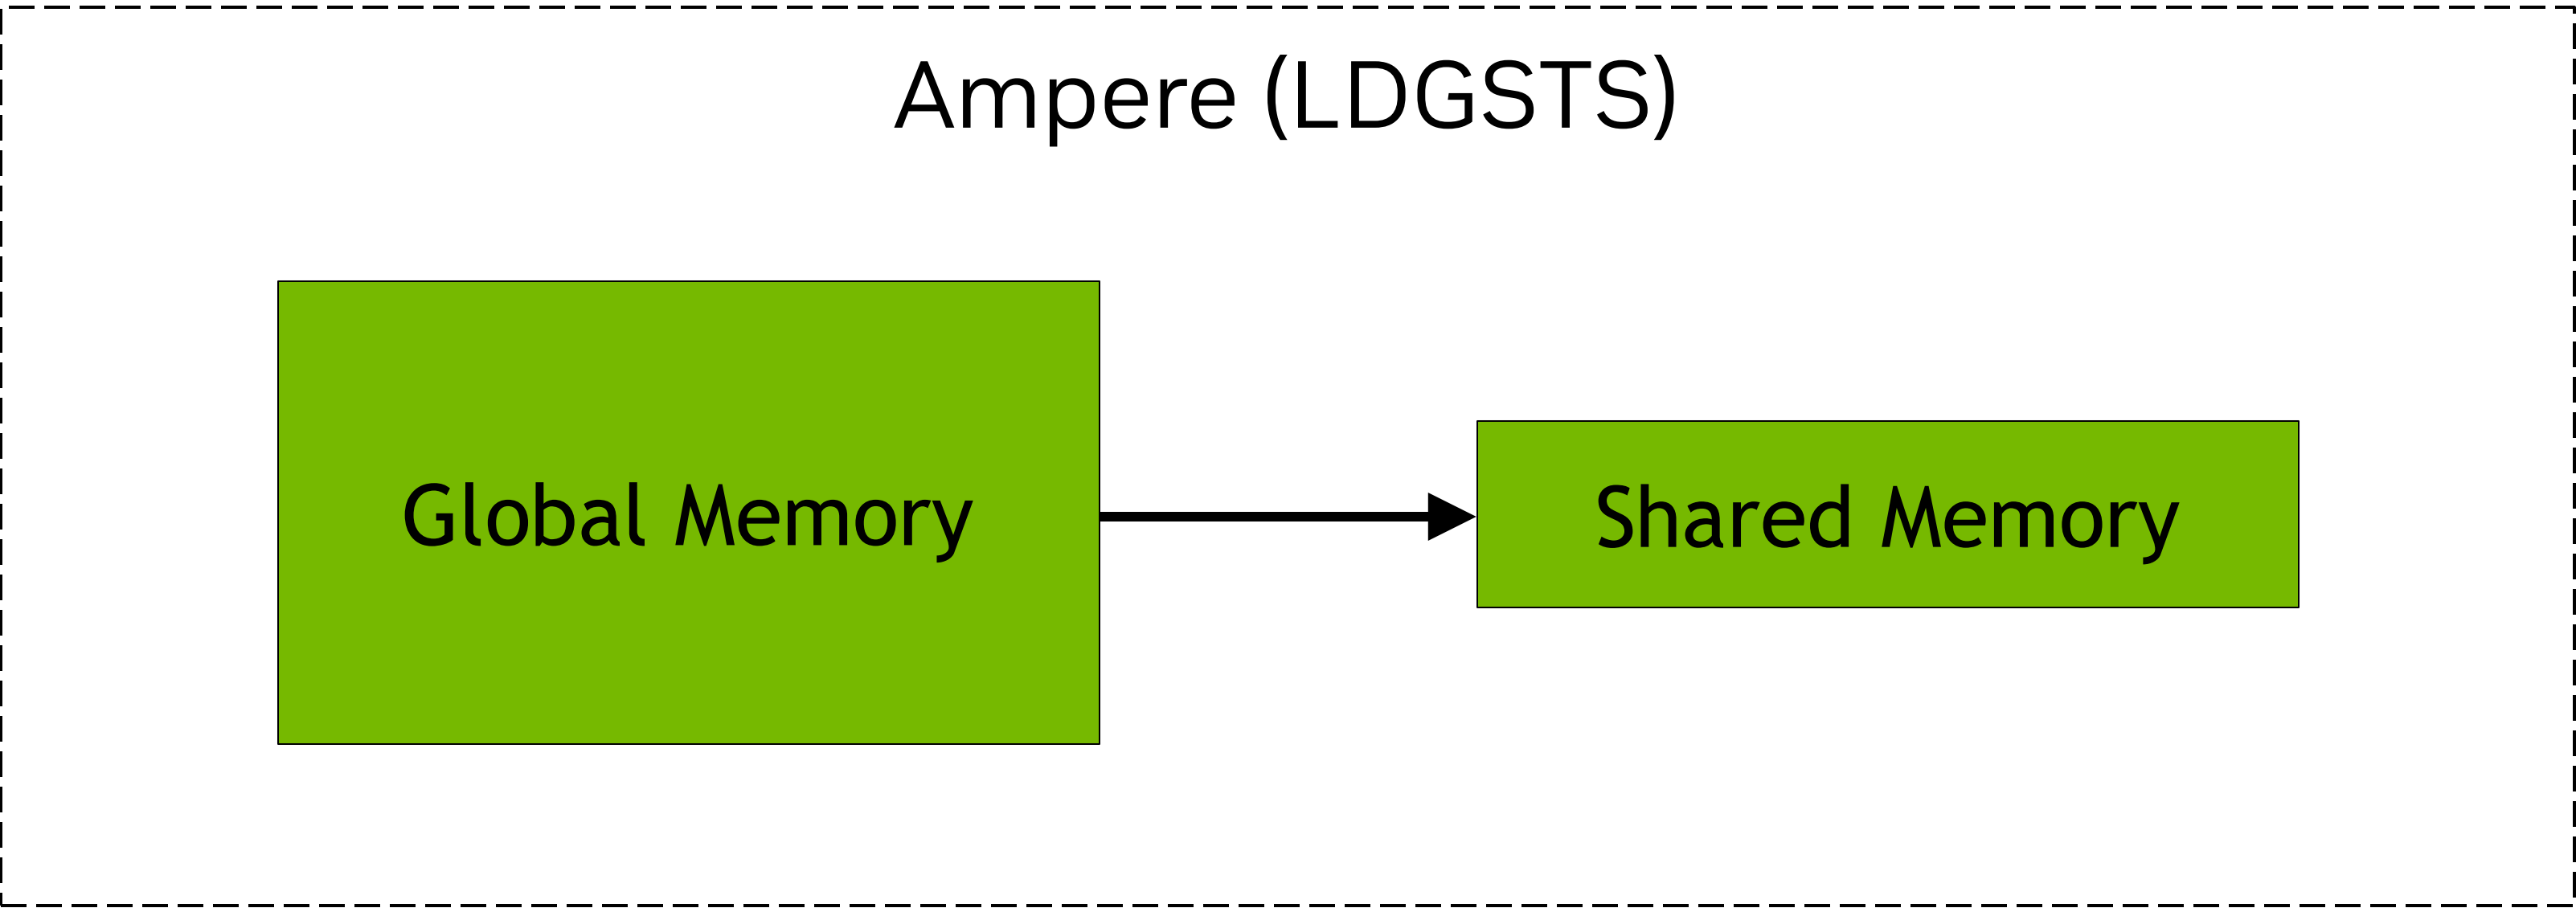


# Asynchronous Data Copies

Asynchronous copies summary

HW	Direction	Operation	Synchronization	C++	SASS
<div><div>Ampere</div><div>Hopper</div><div>TMA</div></div>	GMEM → SMEM	Asynchronous version of smem[sind] := gmem[gind]	Thread-local SMEM barrier	__pipeline_memcpy_async	LDGSTS
	REG → DSMEM	Asynchronous version of dsmem[sind] := var	DSMEM barrier	cuda::ptx::st_async	STAS
	GMEM → (D)SMEM	1D – 5D block copy Uniform	(D)SMEM barrier	cuda::device::experimental:: cp_async_bulk_tensor_Nd_global_to_shared cp_async_bulk_tensor_Nd_shared_to_global	UTMALDG
	SMEM → GMEM		Thread-local	UTMASTG	
	GMEM → SMEM	1D block copy Uniform	SMEM barrier	cuda::device::experimental:: cp_async_bulk_global_to_shared cp_async_bulk_shared_to_global	UBLKCP
	SMEM → GMEM		Thread-local		
	SMEM → (D)SMEM		(D)SMEM barrier	cuda::ptx::cp_async_bulk*	

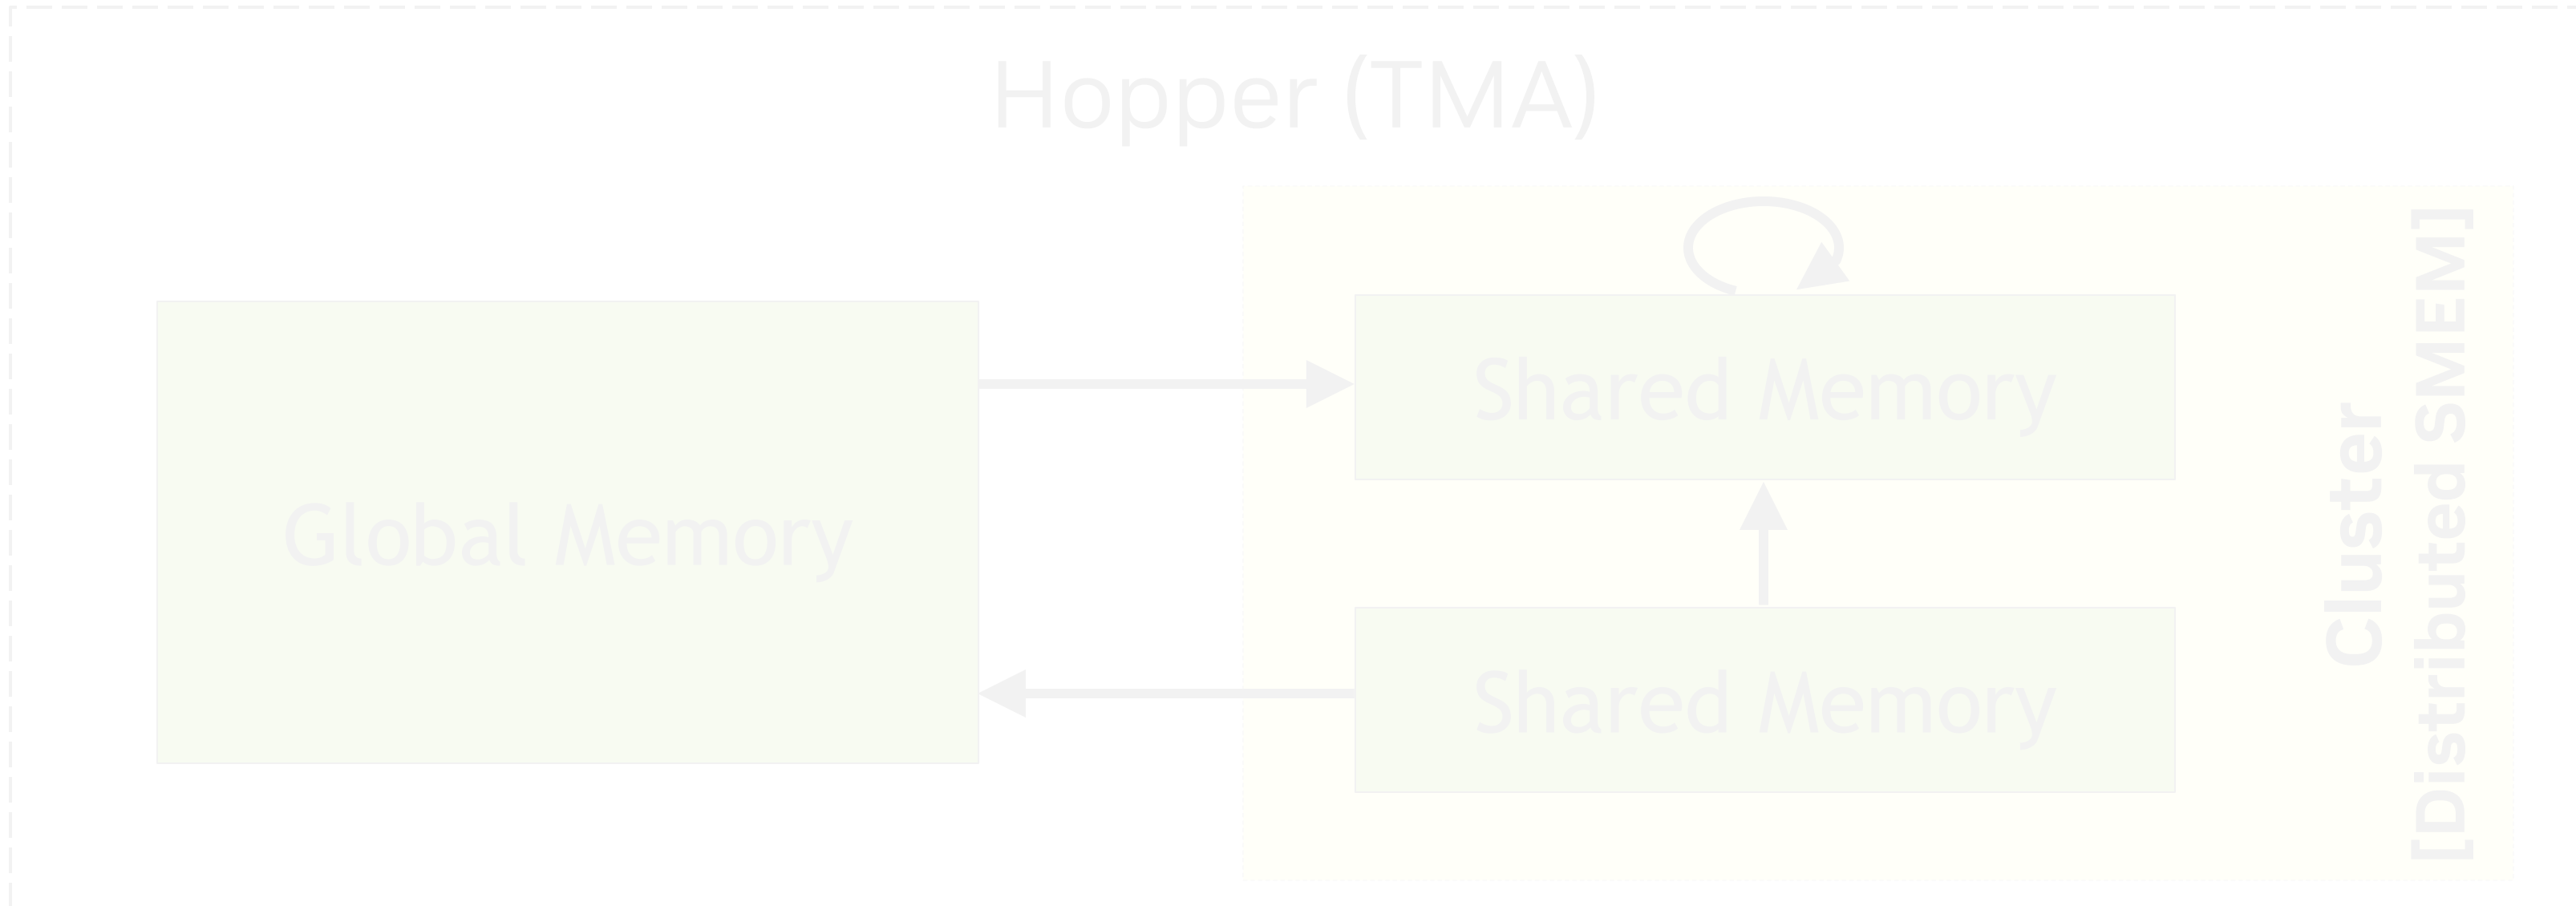
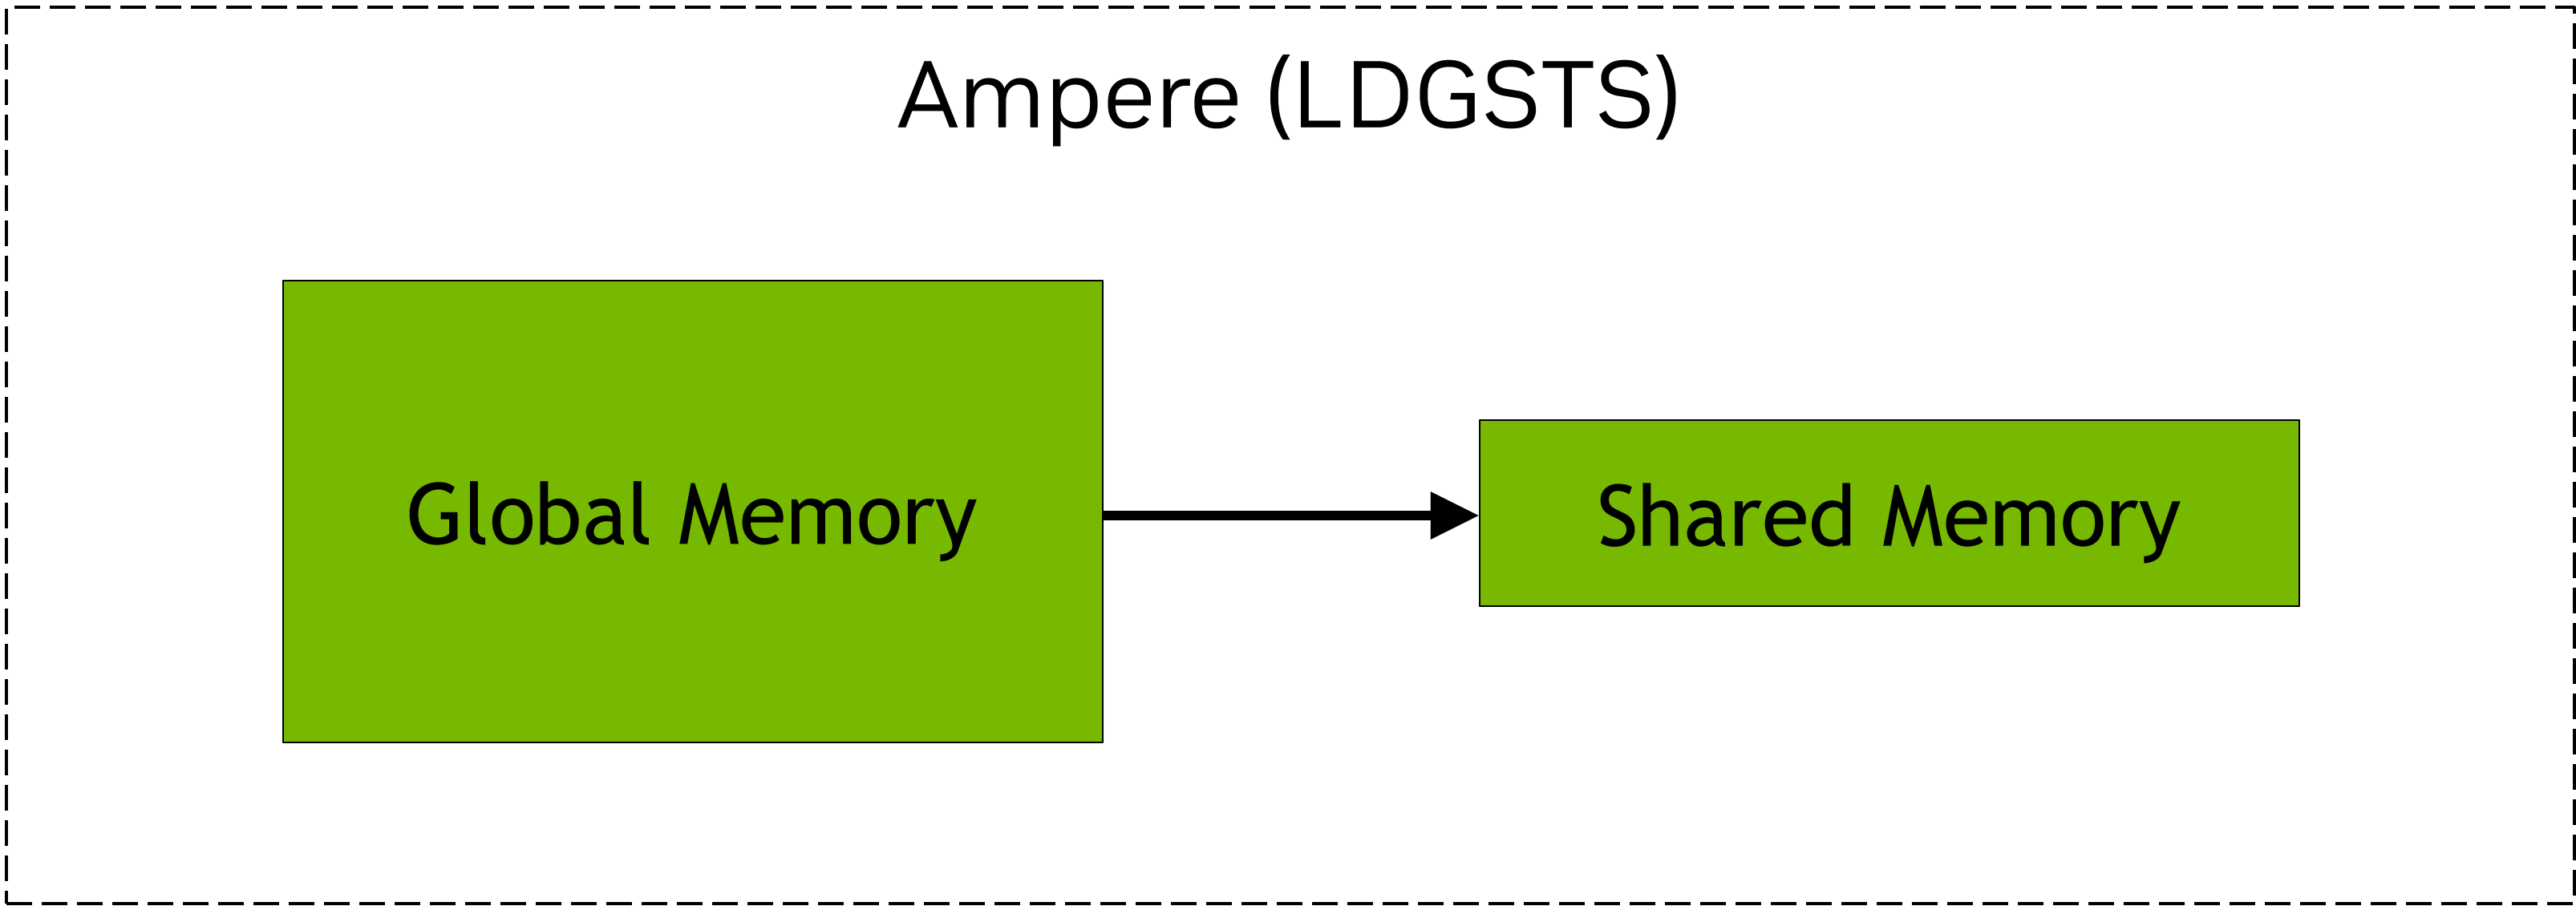
\* Available in latest github CCCL



# Asynchronous Data Copies

Asynchronous copies – LDGSTS

HW	Direction	Operation	Synchronization	C++	SASS
Ampere    Hopper  TMA	GMEM → SMEM	Asynchronous version of smem[sind] := gmem[gind]	Thread-local SMEM barrier	--pipeline_memcpy_async	LDGSTS
	REG → DSMEM	Asynchronous version of dsmem[sind] := var	DSMEM barrier	cuda::ptx::st_async	STAS
	GMEM → (D)SMEM	1D – 5D block copy Uniform	(D)SMEM barrier	cuda::device::experimental:: cp_async_bulk_tensor_Nd_global_to_shared	UTMALDG
	SMEM → GMEM		Thread-local	cp_async_bulk_tensor_Nd_shared_to_global	UTMASTG
	GMEM → SMEM	1D block copy Uniform	SMEM barrier	cuda::device::experimental:: cp_async_bulk_global_to_shared	UBLKCP
	SMEM → GMEM		Thread-local	cp_async_bulk_shared_to_global	
	SMEM → (D)SMEM		(D)SMEM barrier	cuda::ptx::cp_async_bulk*	



# Asynchronous Data Copies

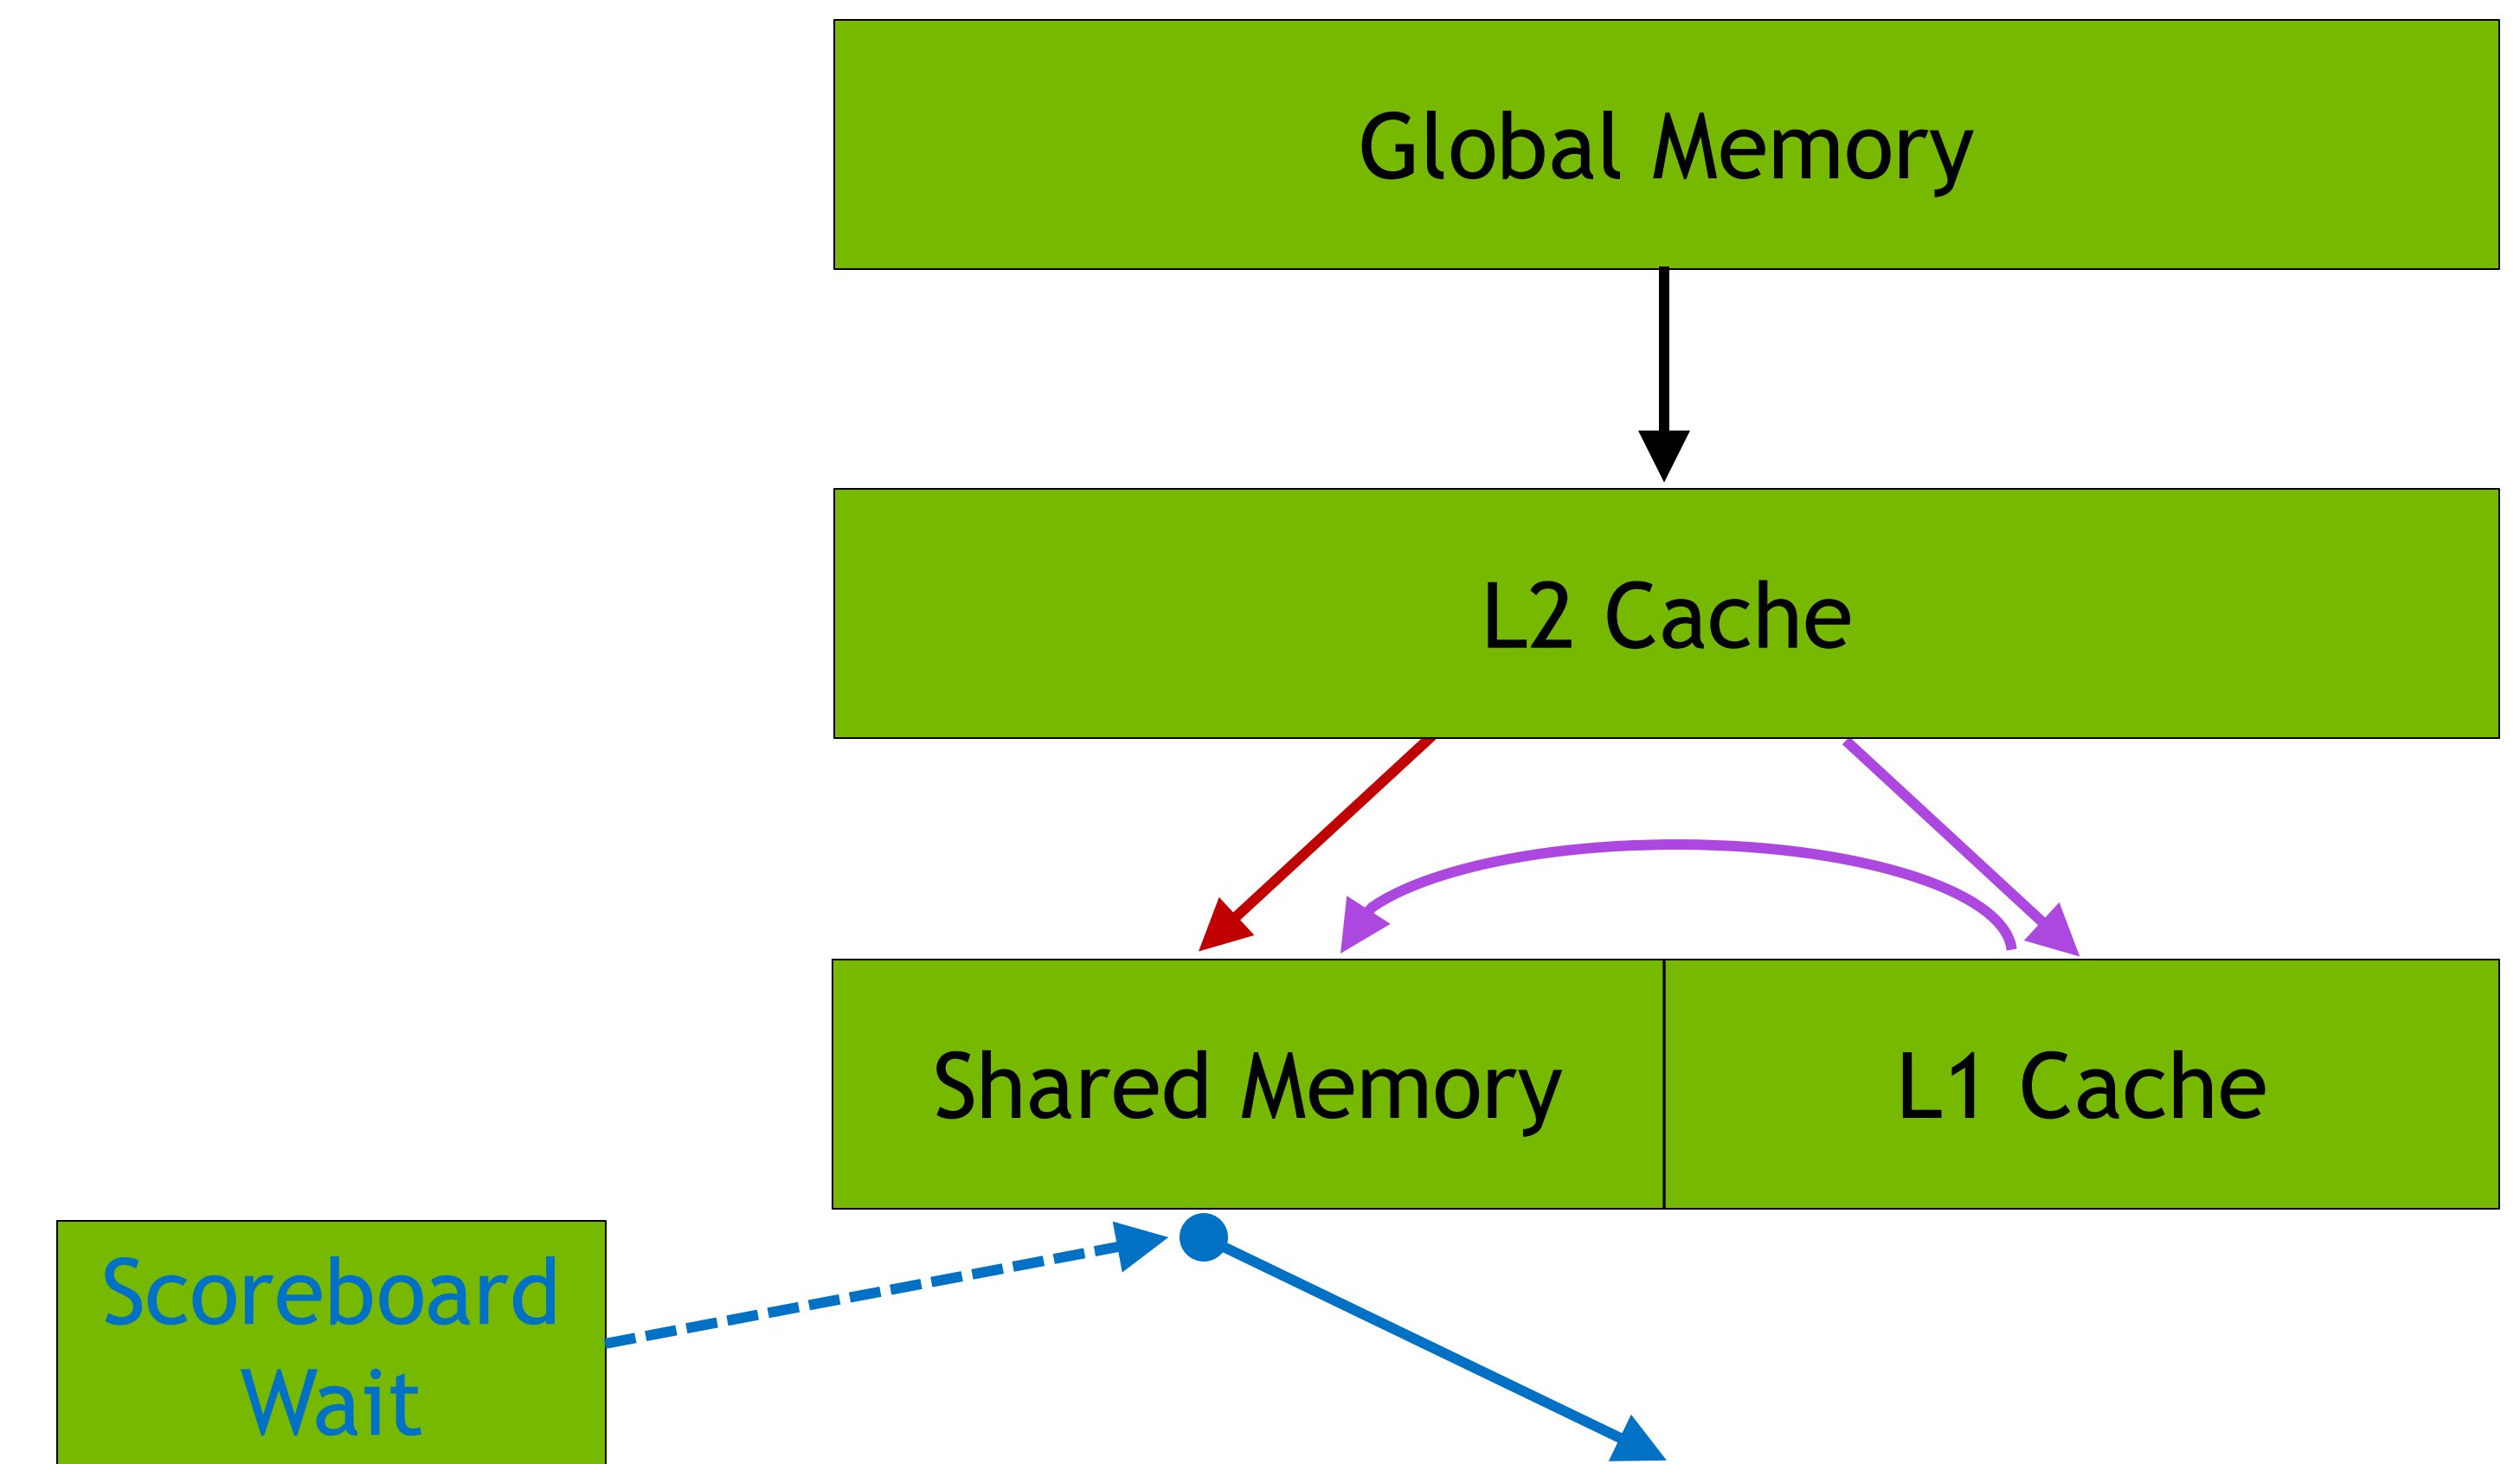
LDGSTS dissected

Asynchronous version of `smem[sind] = gmem[gind];`  
`__pipeline_memcpy_async(&smem[sind], &gmem[gind], sizeof(T));`

Features:

1. Bypassing registers  
(less registers may lead to better occupancy / reduced spillage)
2. Two modes:
  - L1 bypass mode: kicks in when `sizeof(T)` and alignment are 16 B  
(benefit: no L1 pollution)
  - L1 access mode: when `sizeof(T)` and alignment are 4 or 8 B

**Hopper note:** shared memory must be local SMEM when CTA is part of CGA.





# Asynchronous Data Copies

LDGSTS copies and synchronization

```
T* gmem;
```

```
__shared__ T smem[..];
```

```
__pipeline_commit();
```

```
__pipeline_memcpy_async(...);
```

```
...
```

```
__pipeline_memcpy_async(...);
```

```
__pipeline_commit();
```

```
__pipeline_wait_prior(N);
```

– issues async load instructions

– (non-blocking) batch-counting dependency barrier for previous loads released when data have written to SHMEM (for the executing thread)

– blocks until all previous commits except last N are “ready”

# Asynchronous Data Copies

LDGSTS copies and synchronization

```
T* gmem;
```

```
__shared__ T smem[..];
```

```
__pipeline_memcpy_async(...);  
__pipeline_memcpy_async(...);  
__pipeline_commit();
```

```
__pipeline_memcpy_async(...);  
__pipeline_memcpy_async(...);  
__pipeline_memcpy_async(...);  
__pipeline_commit();
```

```
__pipeline_memcpy_async(...);  
__pipeline_commit();
```

```
__pipeline_wait_prior(2);
```

Async copies  
ready after wait

```
__pipeline_memcpy_async(...);  
__pipeline_memcpy_async(...);  
__pipeline_commit();
```

```
__pipeline_memcpy_async(...);  
__pipeline_memcpy_async(...);  
__pipeline_memcpy_async(...);  
__pipeline_commit();
```

```
__pipeline_memcpy_async(...);  
__pipeline_commit();
```

```
__pipeline_wait_prior(0);
```

Async copies  
ready after wait

# Asynchronous Data Copies

LDGSTS copies and synchronization

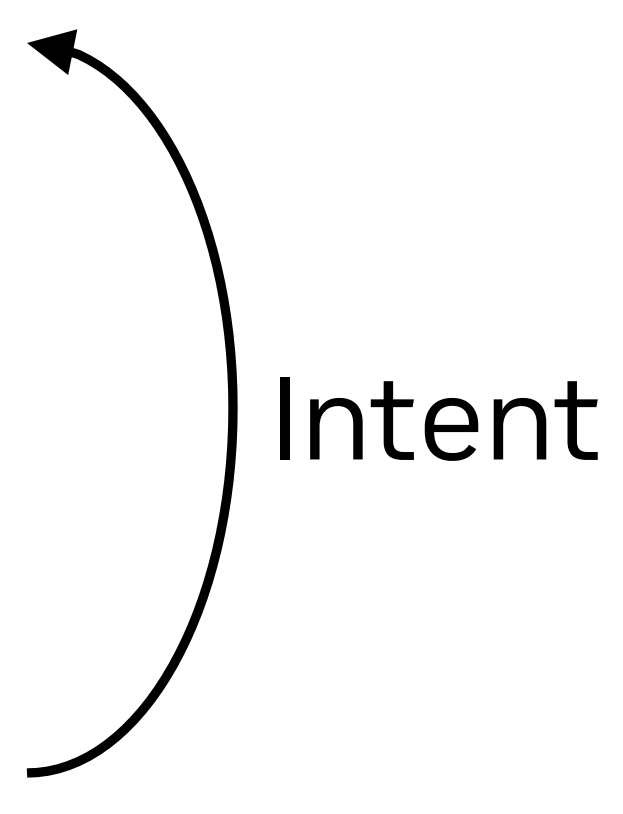
`__pipeline_commit` behaves as a full-warp level instruction regardless of warp divergence!

The code below is equivalent to 2 commits in case of divergence (and up to 32 for fully diverged warp):

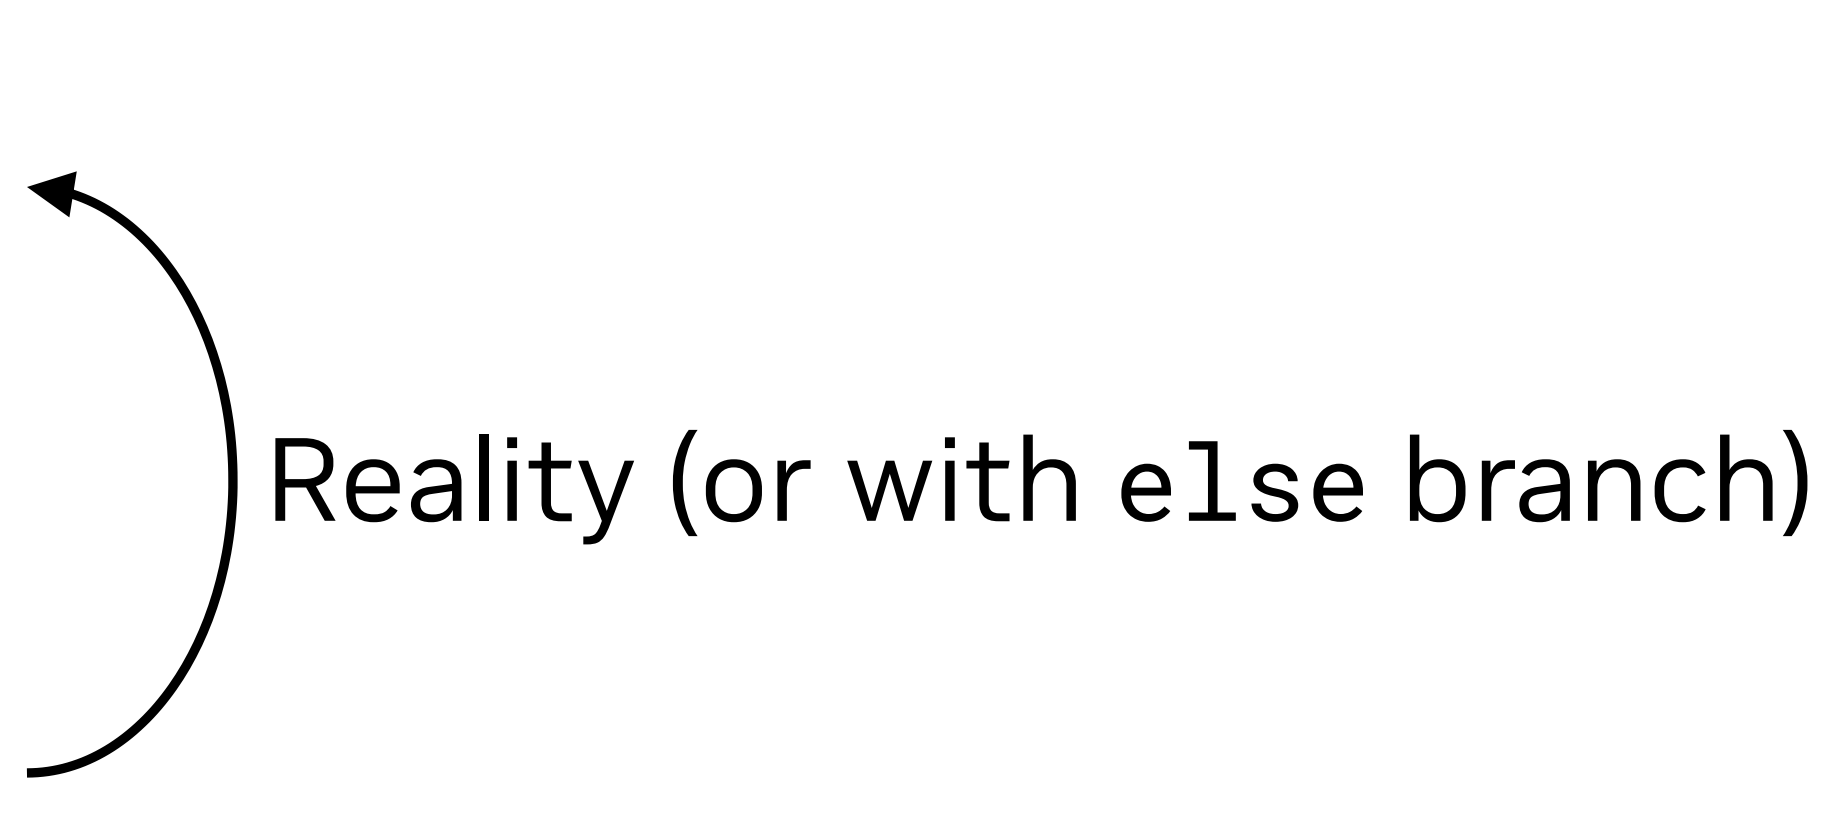
```
if (cond)
    __pipeline_commit();
else
    __pipeline_commit();
```

**Consequence:** overwait for 1<sup>st</sup> batch in case of warp divergence:

```
__pipeline_commit();
if (cond)
    __pipeline_commit();
else
    __pipeline_commit();
__pipeline_wait_prior(1);
```



```
__pipeline_commit();
if (cond)
    __pipeline_commit();
else
    __pipeline_commit();
__pipeline_wait_prior(1);
```



**Recommendation:** do not commit in divergent code!



# Asynchronous Data Copies

~~LDGSTS~~ use case: batching loads in conditional code

```
T const* __restrict__ left;
T const* __restrict__ right;
T const* __restrict__ center;

// Halo-Center-Halo.
__shared__ T smem[8 + 32 + 8];
auto const tx = threadIdx.x;

if (tx < 8)
    smem[tx] = left[tx];

else if (tx >= 32 - 8)
    smem[tx + 16] = right[tx];

smem[tx + 8] = center[tx];

__syncthreads();


// COMPUTE STENCIL
```

Compiler generates (pseudocode):

```
BRA LABEL1
LDG R1, [left]
STS [..] R1      ← STALL
GOTO LABEL2

LABEL1: BRA LABEL2
LDG R2, [right]
STS [..] R2      ← STALL

LABEL2: LDG R3, [center]
STS [..] R3      ← STALL
```



# Asynchronous Data Copies

~~LDGSTS~~ use case: batching loads in conditional code

```
T const* __restrict__ left;
T const* __restrict__ right;
T const* __restrict__ center;

// Halo-Center-Halo.
__shared__ T smem[8 + 32 + 8];
auto const tx = threadIdx.x;


if (tx < 8)
    smem[tx] = left[tx];

else if (tx >= 32 - 8)
    smem[tx + 16] = right[tx];

smem[tx + 8] = center[tx];

__syncthreads();

// COMPUTE STENCIL
```



Compiler generates (pseudocode):

```
BRA LABEL1
LDG R1, [left]
STS [..] R1      ← STALL
GOTO LABEL2

LABEL1: BRA LABEL2
LDG R2, [right]
STS [..] R2      ← STALL

LABEL2: LDG R3, [center]
STS [..] R3      ← STALL
```

Ideally, compiler could use predicates (but it may not):

```
PREDICATE1 LDG R1, [left]
PREDICATE2 LDG R2, [right]
            LDG R3, [center] } ALL IN FLIGHT

PREDICATE1 STS [..] R1
PREDICATE2 STS [..] R2
            STS [..] R3
```

# Asynchronous Data Copies

LDGSTS use case: batching loads in conditional code

```
T const* __restrict__ left;
T const* __restrict__ right;
T const* __restrict__ center;

// Halo-Center-Halo.
__shared__ T smem[8 + 32 + 8];
auto const tx = threadIdx.x;

if (tx < 8)
    __pipeline_memcpy_async(
        &smem[tx], &left[tx], sizeof(T));

else if (tx >= 32 - 8)
    __pipeline_memcpy_async(
        &smem[tx + 16], &right[tx], sizeof(T));

__pipeline_memcpy_async(
    &smem[tx + 8], &center[tx], sizeof(T));

__pipeline_commit(); ← single commit, all in flight!

__pipeline_wait_prior<0> ();
__syncthreads();

// COMPUTE STENCIL
```

Compiler generates (pseudocode):

```
BRA LABEL1
LDGSTS [...], [left] ← NO STALL
GOTO LABEL2

LABEL1: BRA LABEL2
LDGSTS [...], [right] ← NO STALL

LABEL2: LDGSTS [...], [center] ← NO STALL

LDGDEPBAR

DEPBAR ← SINGLE STALL
```



# Asynchronous Data Copies

LDGSTS use case: batching loads in conditional code

```
__shared__ float smem[8 + 32 + 8][32];  
  
int const tx = threadIdx.x;  
int const ty = threadIdx.y;  
int gind = (ty * stride_y) + (blockIdx.x * 32 + tx);  
  
for (int i = 0; i < nz; ++i) {  
    if (ty < 8)  
        __pipeline_memcpy_async(&smem[ty][tx], &gl[gind], sizeof(float));  
    else if (ty >= 24)  
        __pipeline_memcpy_async(&smem[ty + 16][tx], &gr[gind], sizeof(float));  
    __pipeline_memcpy_async(&smem[ty + 8][tx], &gc[gind], sizeof(float));  
  
    __pipeline_commit();  
    __pipeline_wait_prior(0);  
    __syncthreads();  
  
    // "STENCIL":  
    float d = smem[ty + 8][tx];  
    #pragma unroll  
    for (int r = 1; r <= 8; ++r)  
        d += (smem[ty + 8 + r][tx] - smem[ty + 8 - r][tx]) * float(r + 0.1);  
  
    out[gind] = d;  
    gind += stride_z;  
    __syncthreads();  
}
```

Extra dimension added to code from previous slide  
Note: condition is on ty => no divergence within warp

# Asynchronous Data Copies

LDGSTS use case: batching loads in conditional code

```
__shared__ float smem[8 + 32 + 8][32];  
int const tx = threadIdx.x;  
int const ty = threadIdx.y;  
int gind = (ty * stride_y) + (blockIdx.x * 32 + tx);  
for (int i = 0; i < nz; ++i) {  
    if (ty < 8)  
        __pipeline_memcpy_async(&smem[ty][tx], &gl[gind], sizeof(float));  
    else if (ty >= 24)  
        __pipeline_memcpy_async(&smem[ty + 16][tx], &gr[gind], sizeof(float));  
    __pipeline_memcpy_async(&smem[ty + 8][tx], &gc[gind], sizeof(float));  
  
    __pipeline_commit();  
    __pipeline_wait_prior(0);  
    __syncthreads();  
  
    // "STENCIL":  
    float d = smem[ty + 8][tx];  
    #pragma unroll  
    for (int r = 1; r <= 8; ++r)  
        d += (smem[ty + 8 + r][tx] - smem[ty + 8 - r][tx]) * float(r + 0.1);  
  
    out[gind] = d;  
    gind += stride_z;  
    __syncthreads();  
}
```

H200 experiment.

Speedup over sync-copy:

- Un-staged: 1.3X

Memory throughput:

- Sync copy: 46%
- Un-staged: 67%

Extra dimension added to code from previous slide  
Note: condition is on ty => no divergence within warp

# Asynchronous Data Copies

LDGSTS use case: batching loads in conditional code

```
__shared__ float smem[8 + 32 + 8][32];  
int const tx = threadIdx.x;  
int const ty = threadIdx.y;  
int gind = (ty * stride_y) + (blockIdx.x * 32 + tx);  
for (int i = 0; i < nz; ++i) {  
    if (ty < 8)  
        __pipeline_memcpy_async(&smem[ty][tx], &gl[gind], sizeof(float));  
    else if (ty >= 24)  
        __pipeline_memcpy_async(&smem[ty + 16][tx], &gr[gind], sizeof(float));  
    __pipeline_memcpy_async(&smem[ty + 8][tx], &gc[gind], sizeof(float));  
  
    __pipeline_commit();  
    __pipeline_wait_prior(0);  
    __syncthreads();  
  
    // "STENCIL":  
    float d = smem[ty + 8][tx];  
    #pragma unroll  
    for (int r = 1; r <= 8; ++r)  
        d += (smem[ty + 8 + r][tx] - smem[ty + 8 - r][tx]) * float(r + 0.1);  
  
    out[gind] = d;  
    gind += stride_z;  
    __syncthreads();  
}
```

H200 experiment.

Speedup over sync-copy:

- Un-staged: 1.3X
- Two-stage: 1.7X

Memory throughput:

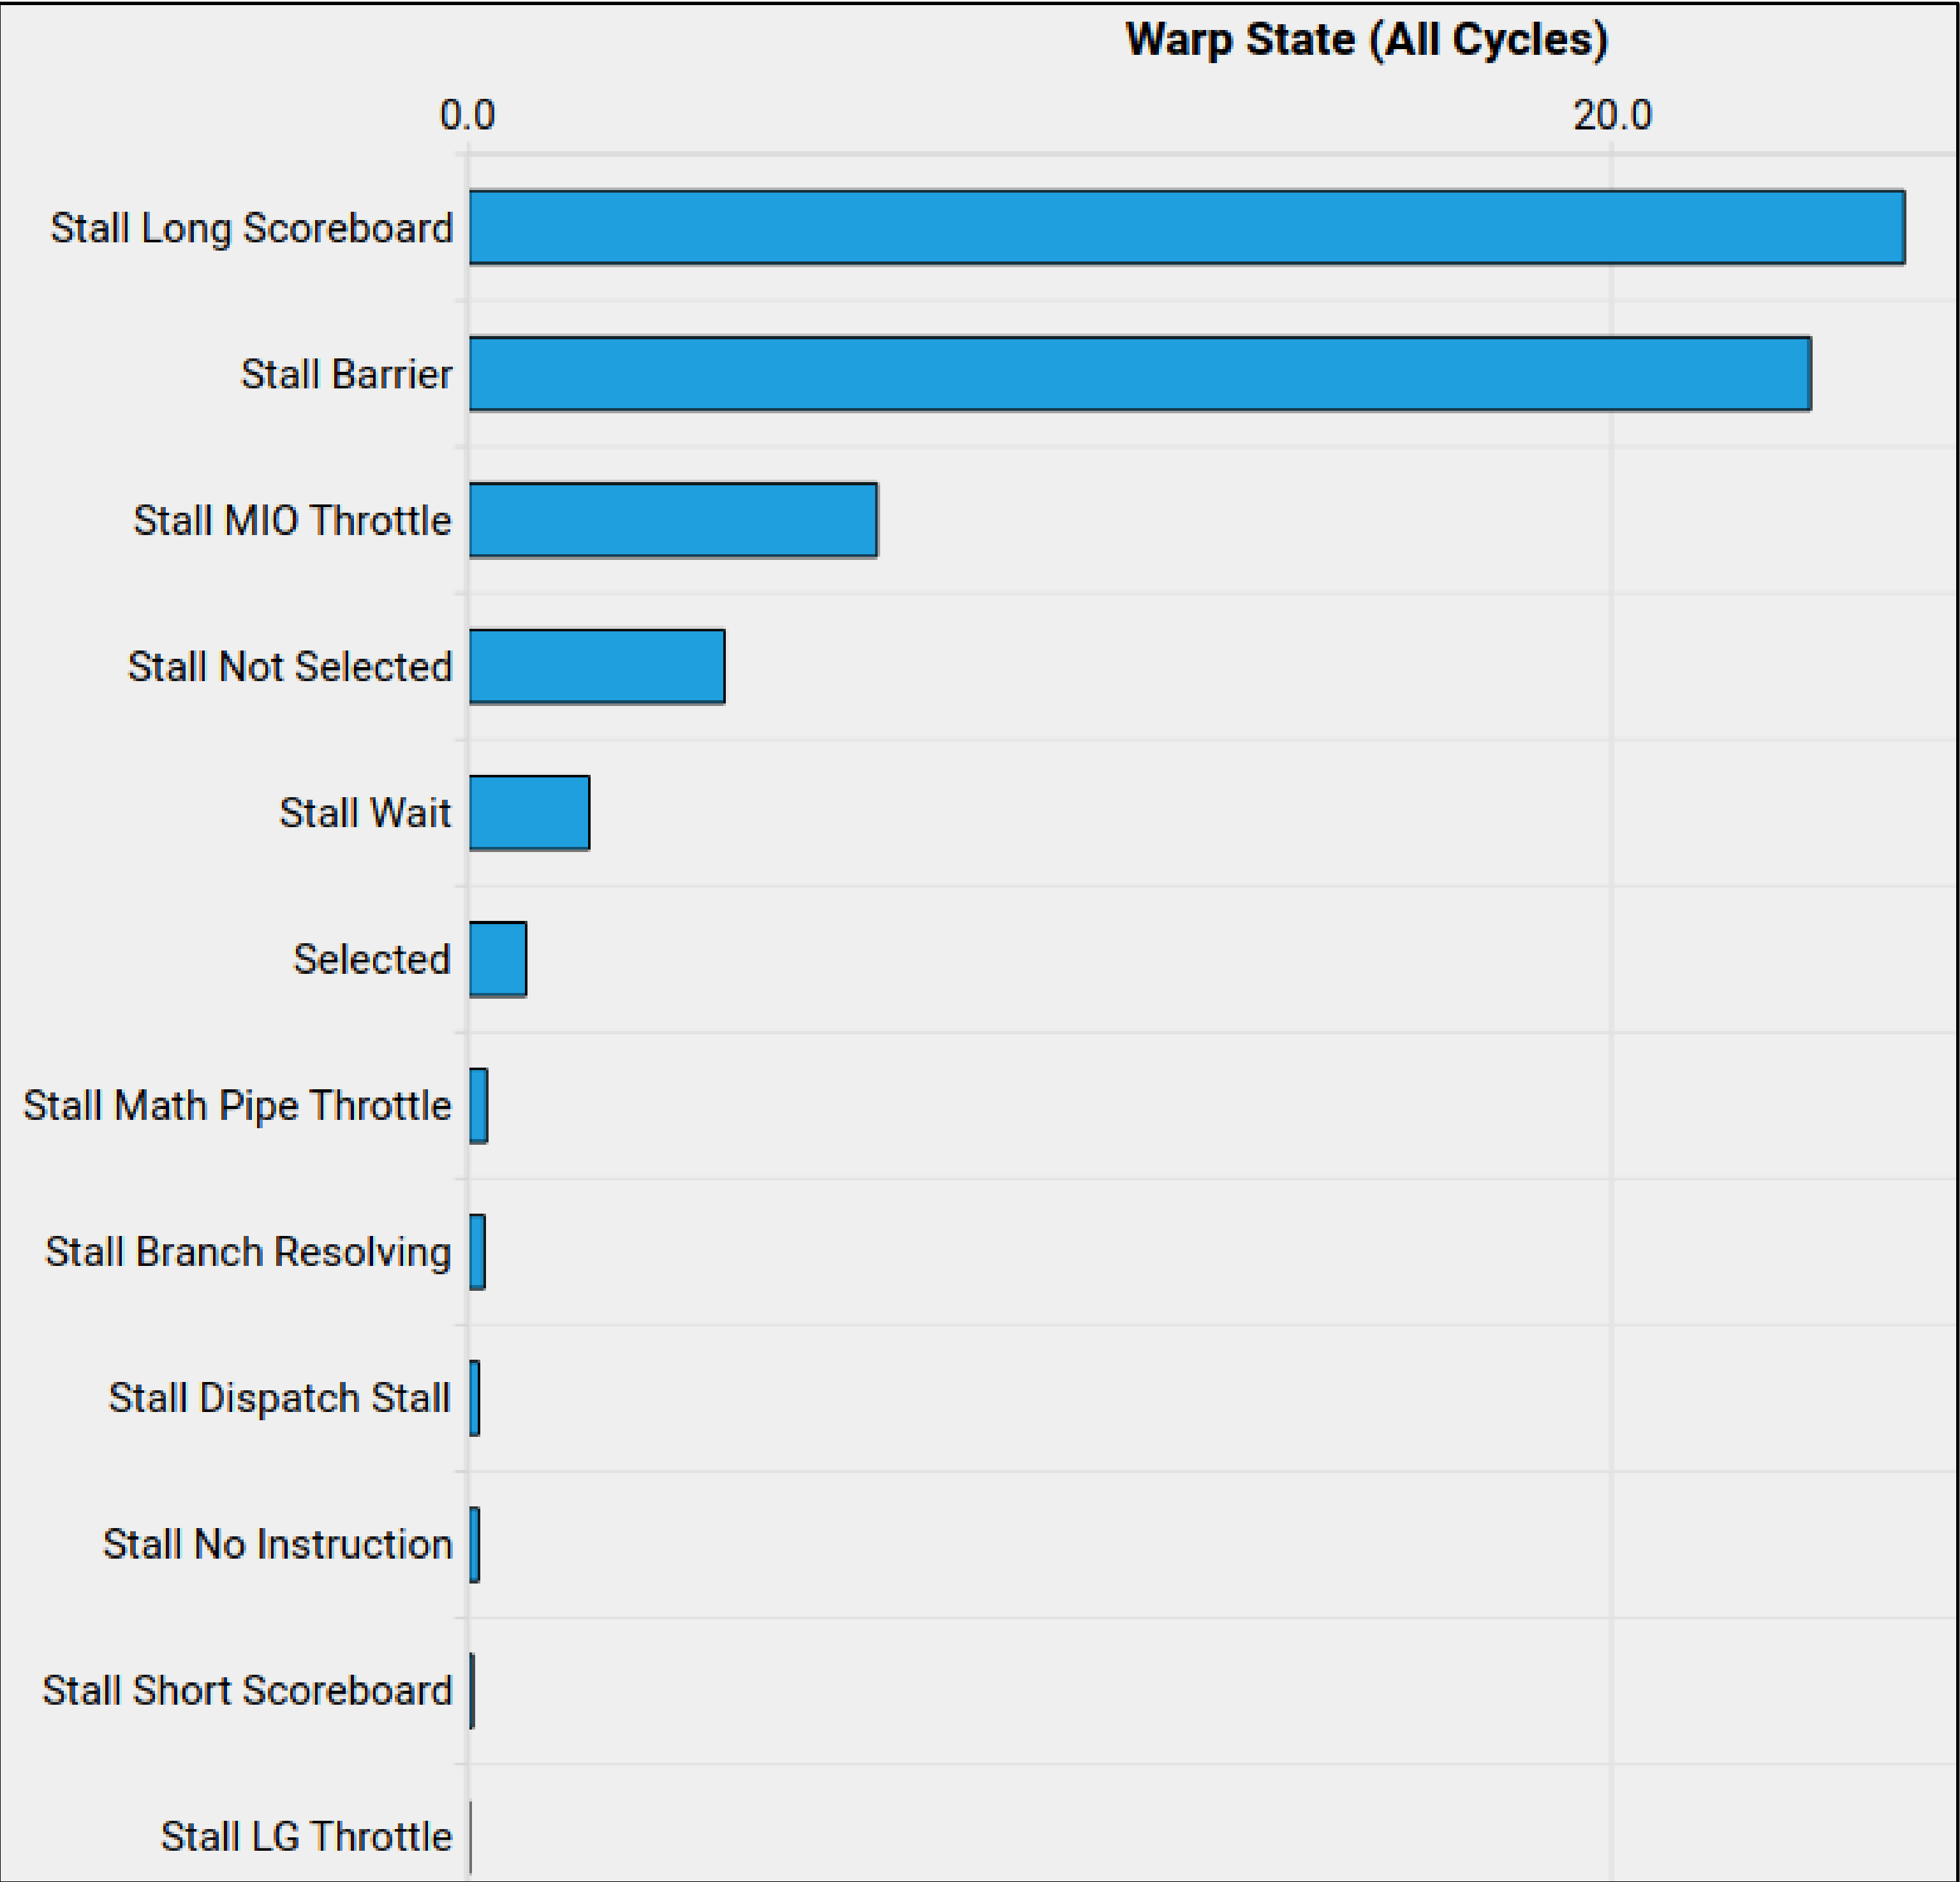
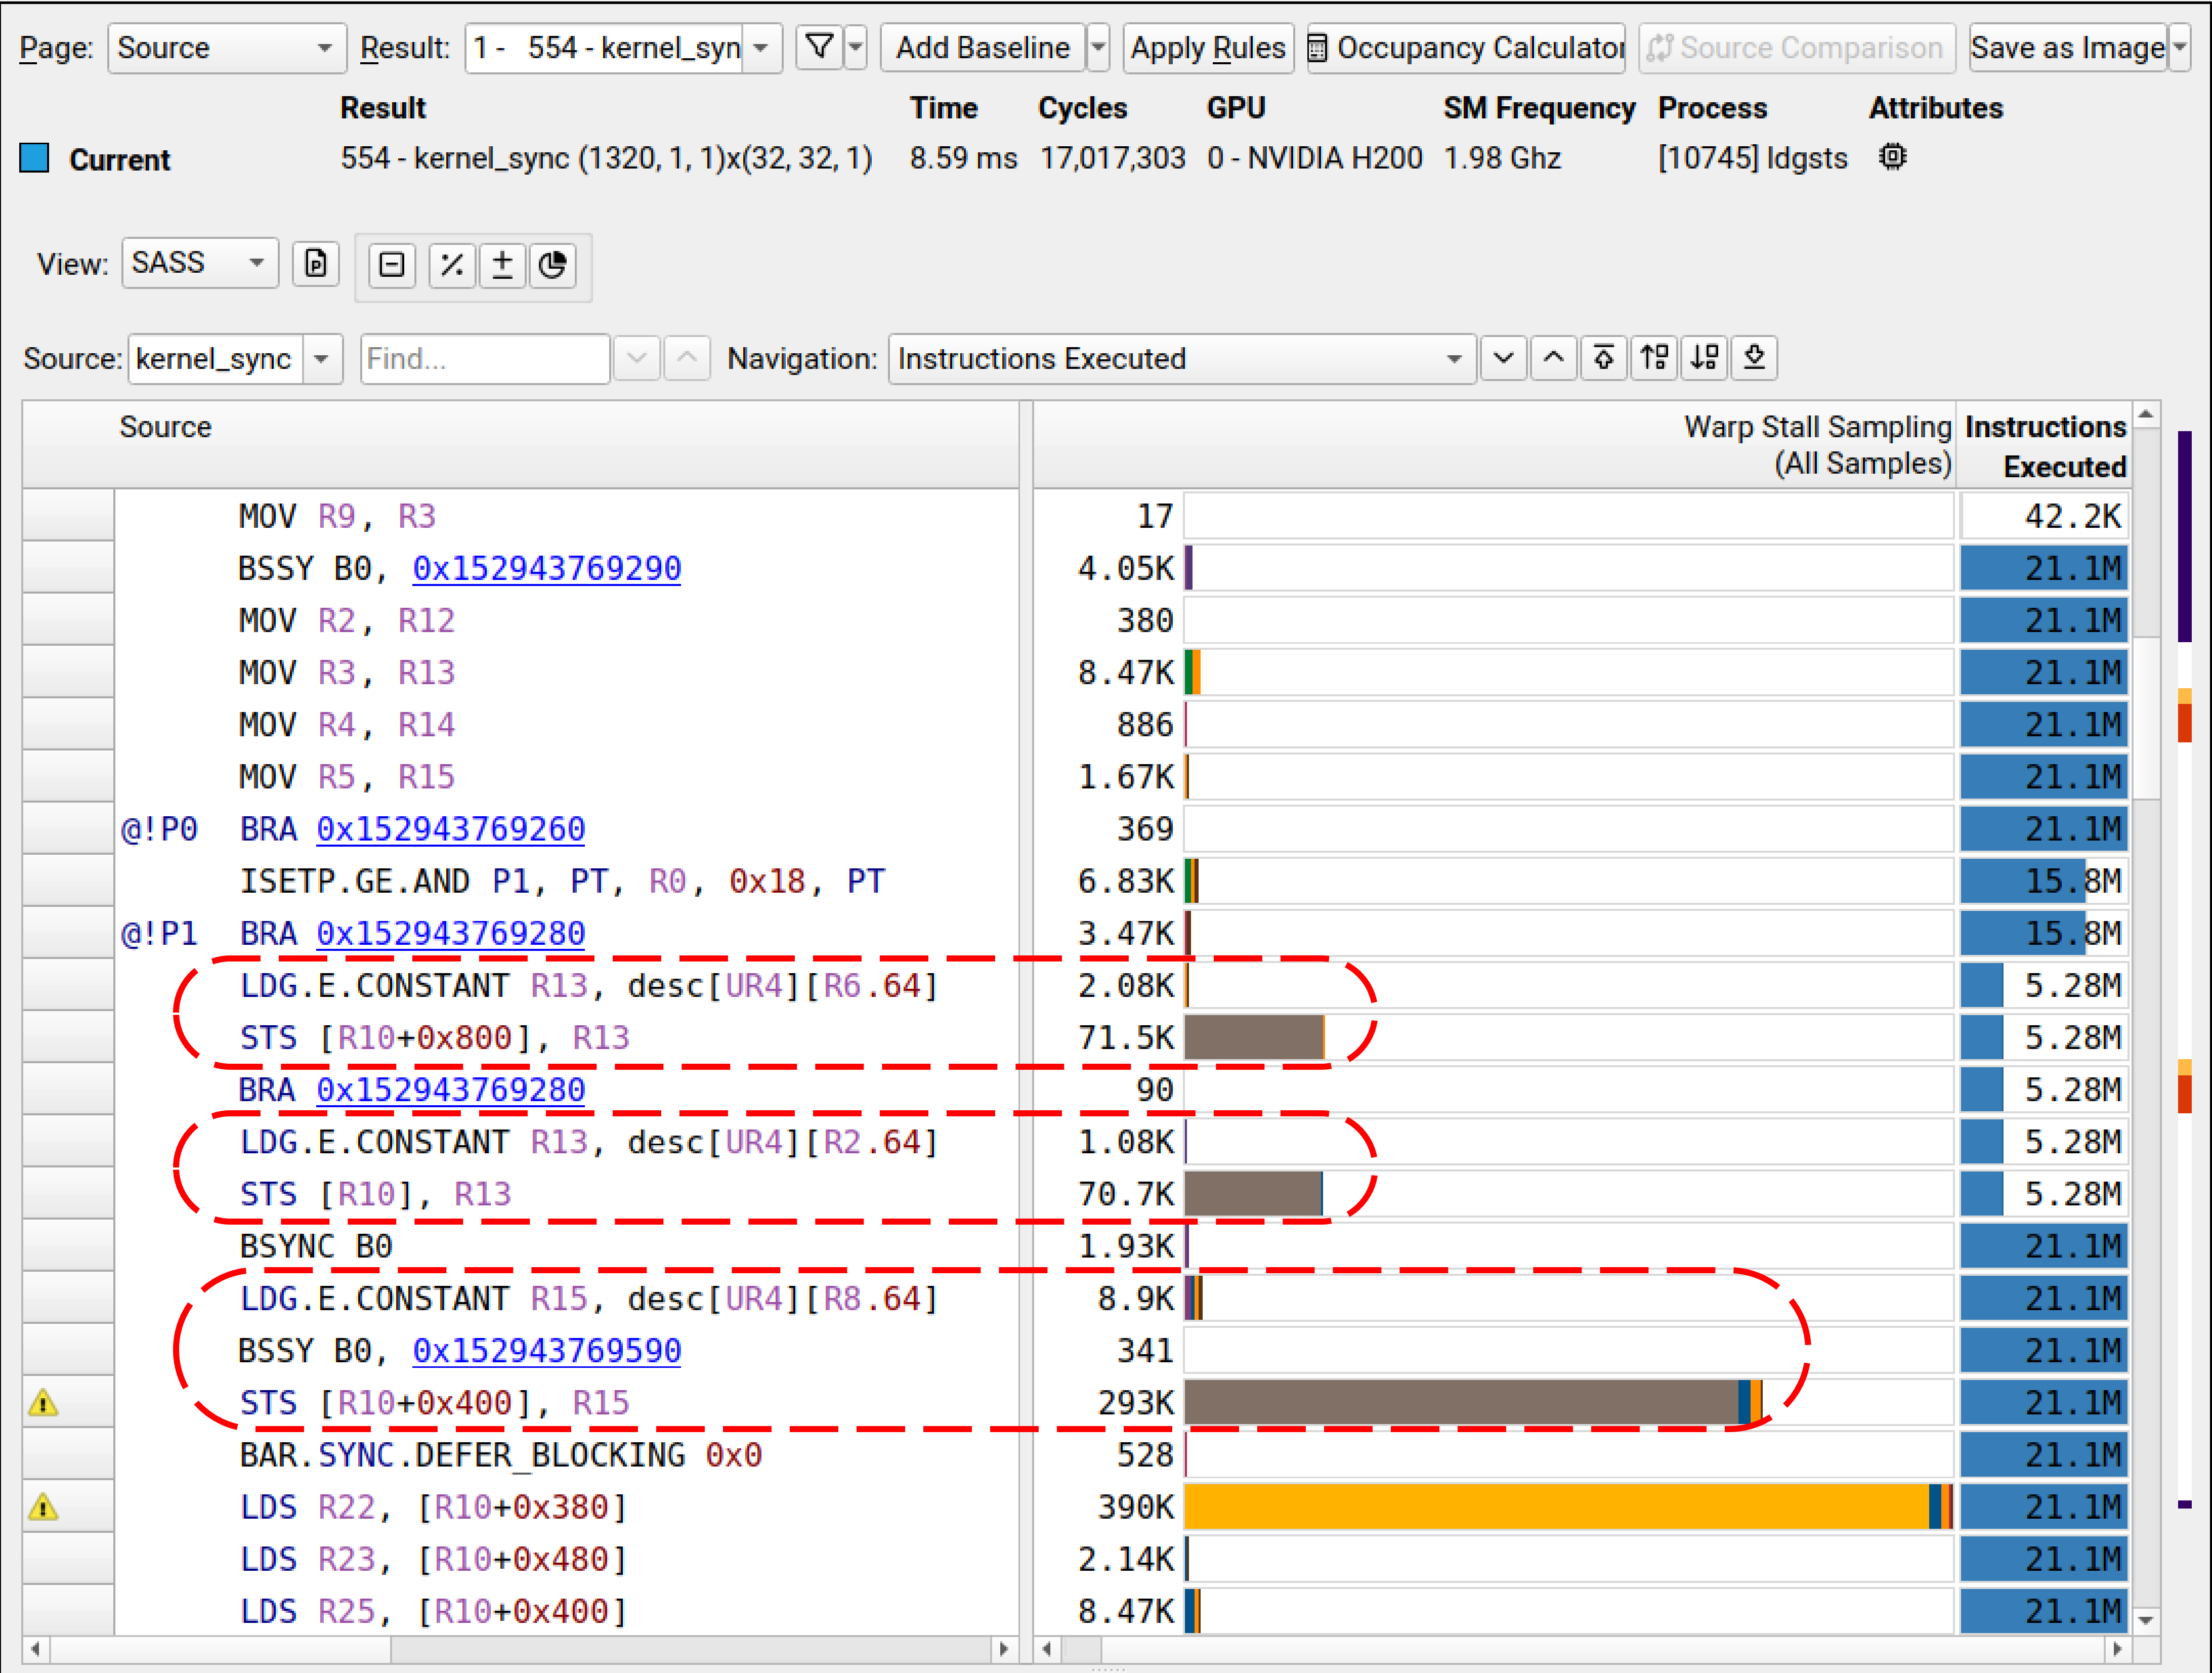
- Sync copy: 46%
- Un-staged: 67%
- Two-stage: 84%

Extra dimension added to code from previous slide  
Note: condition is on ty => no divergence within warp



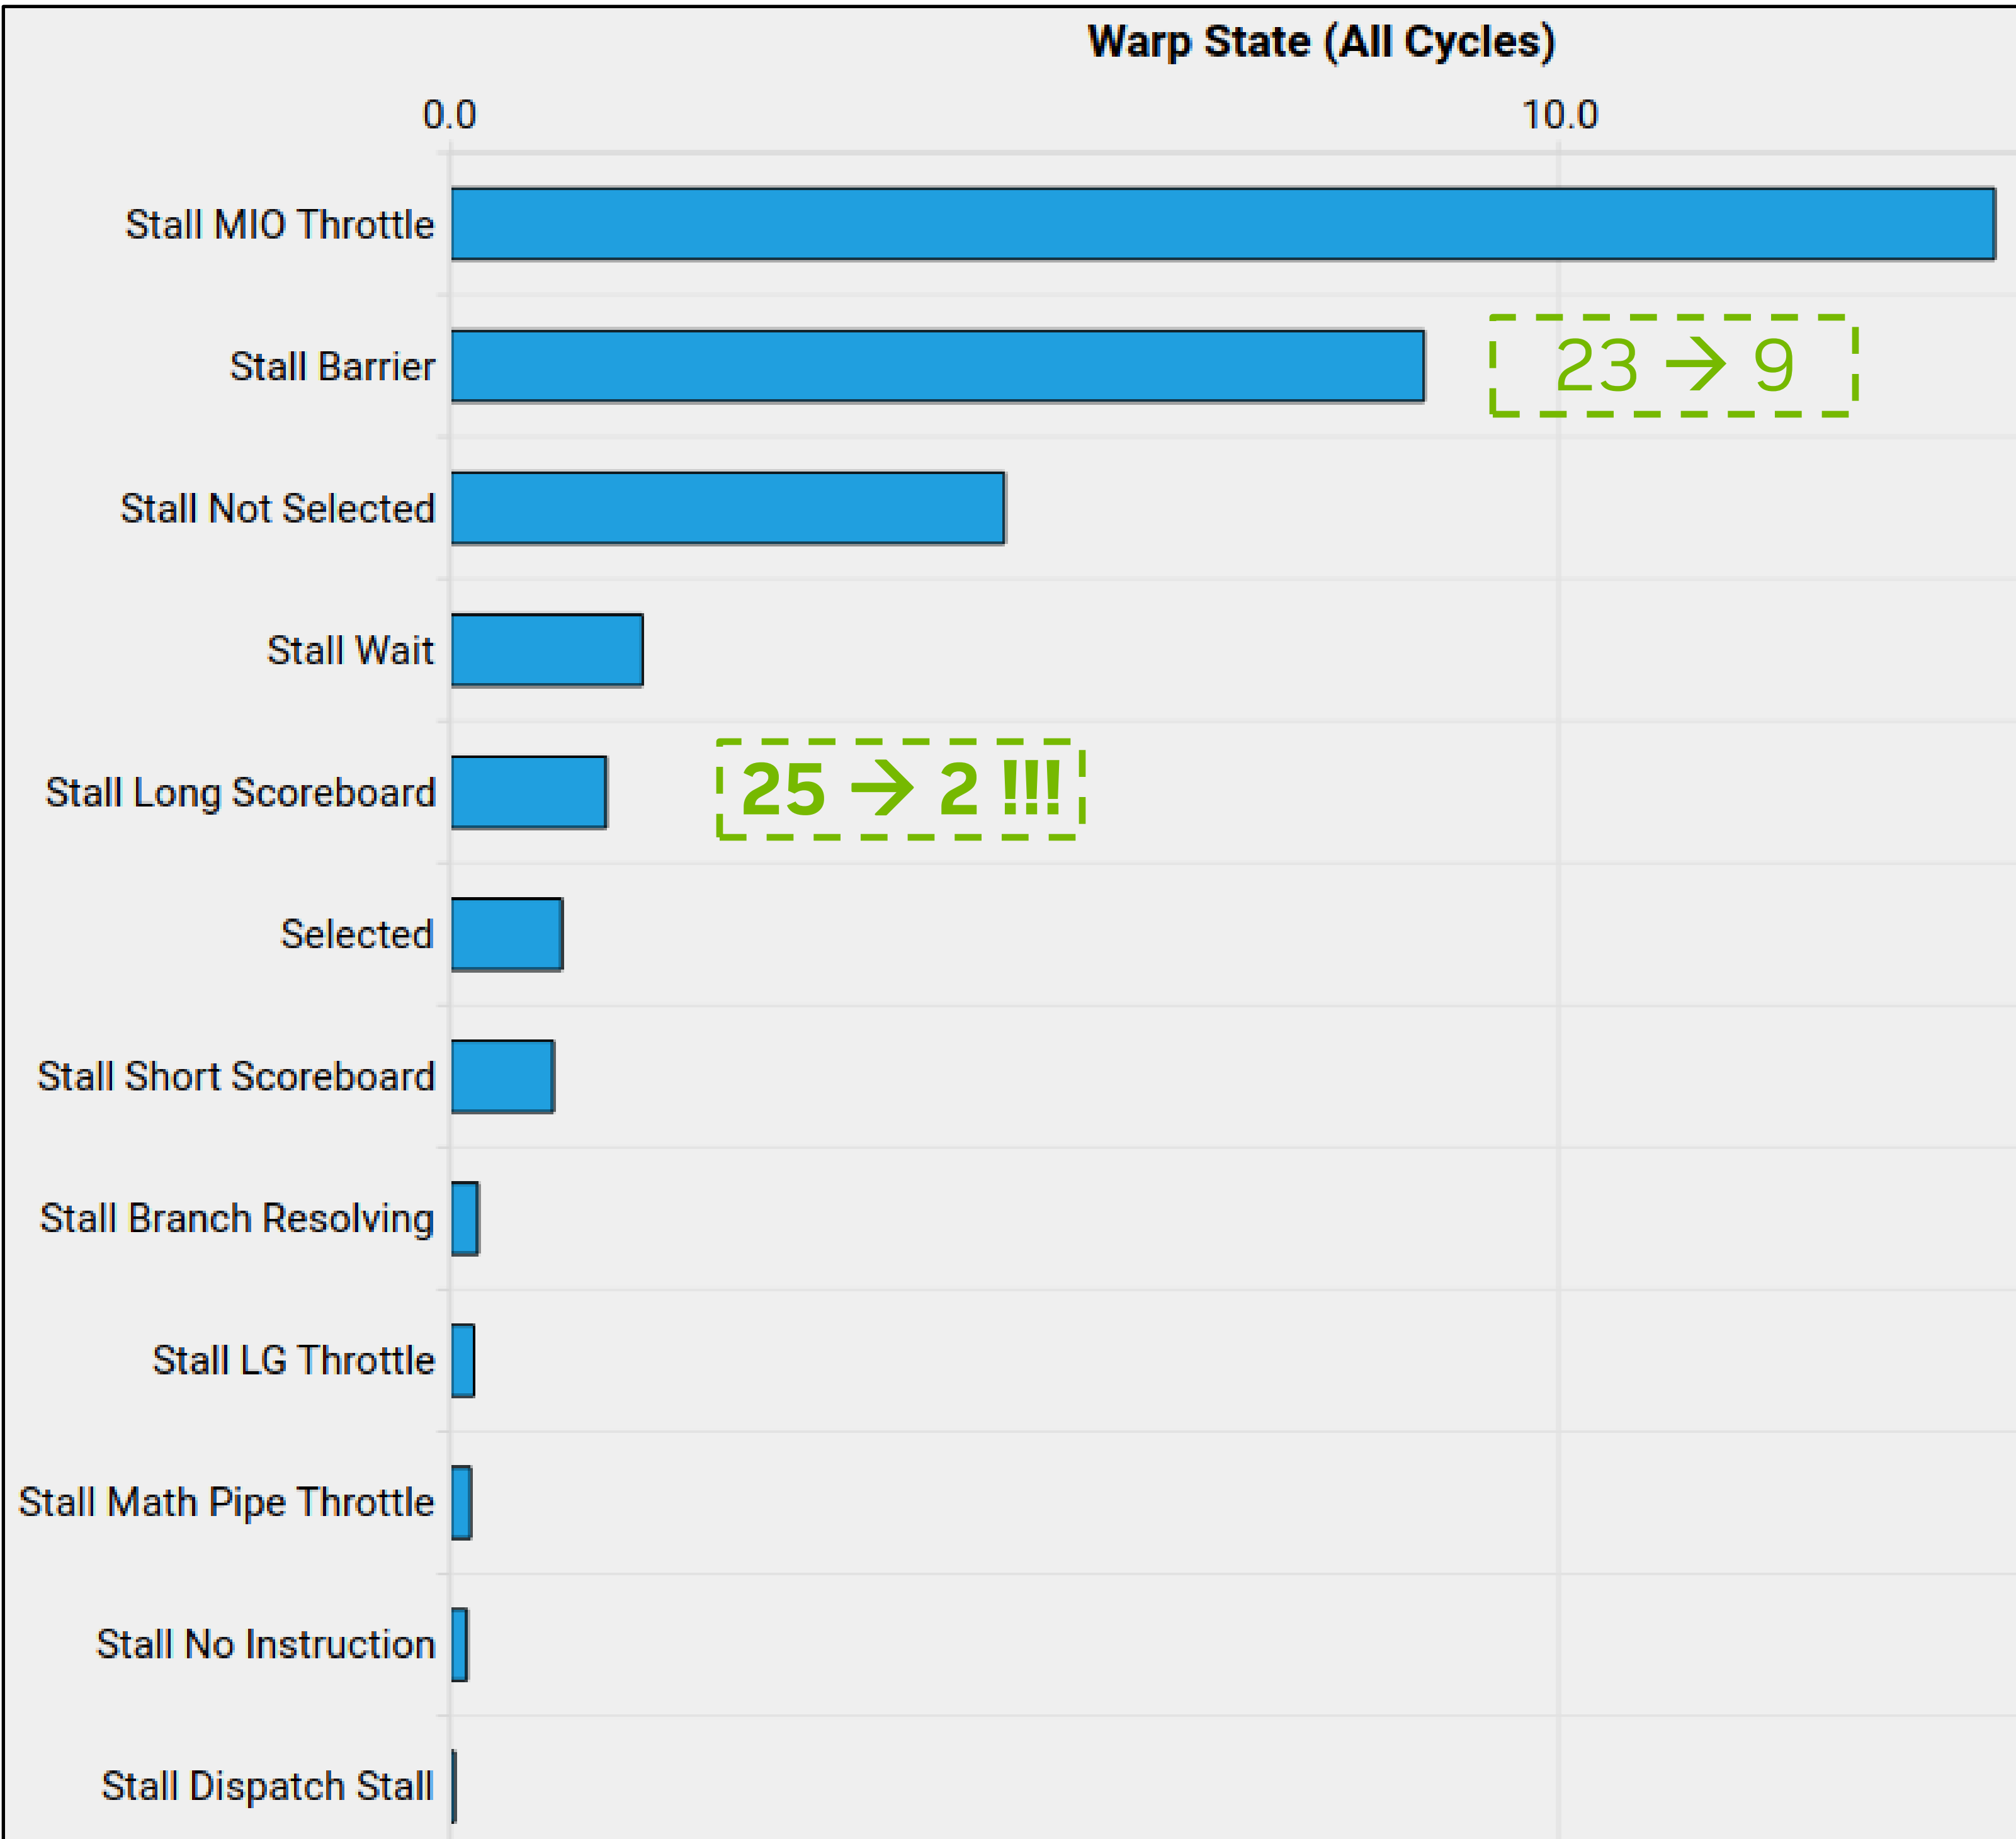
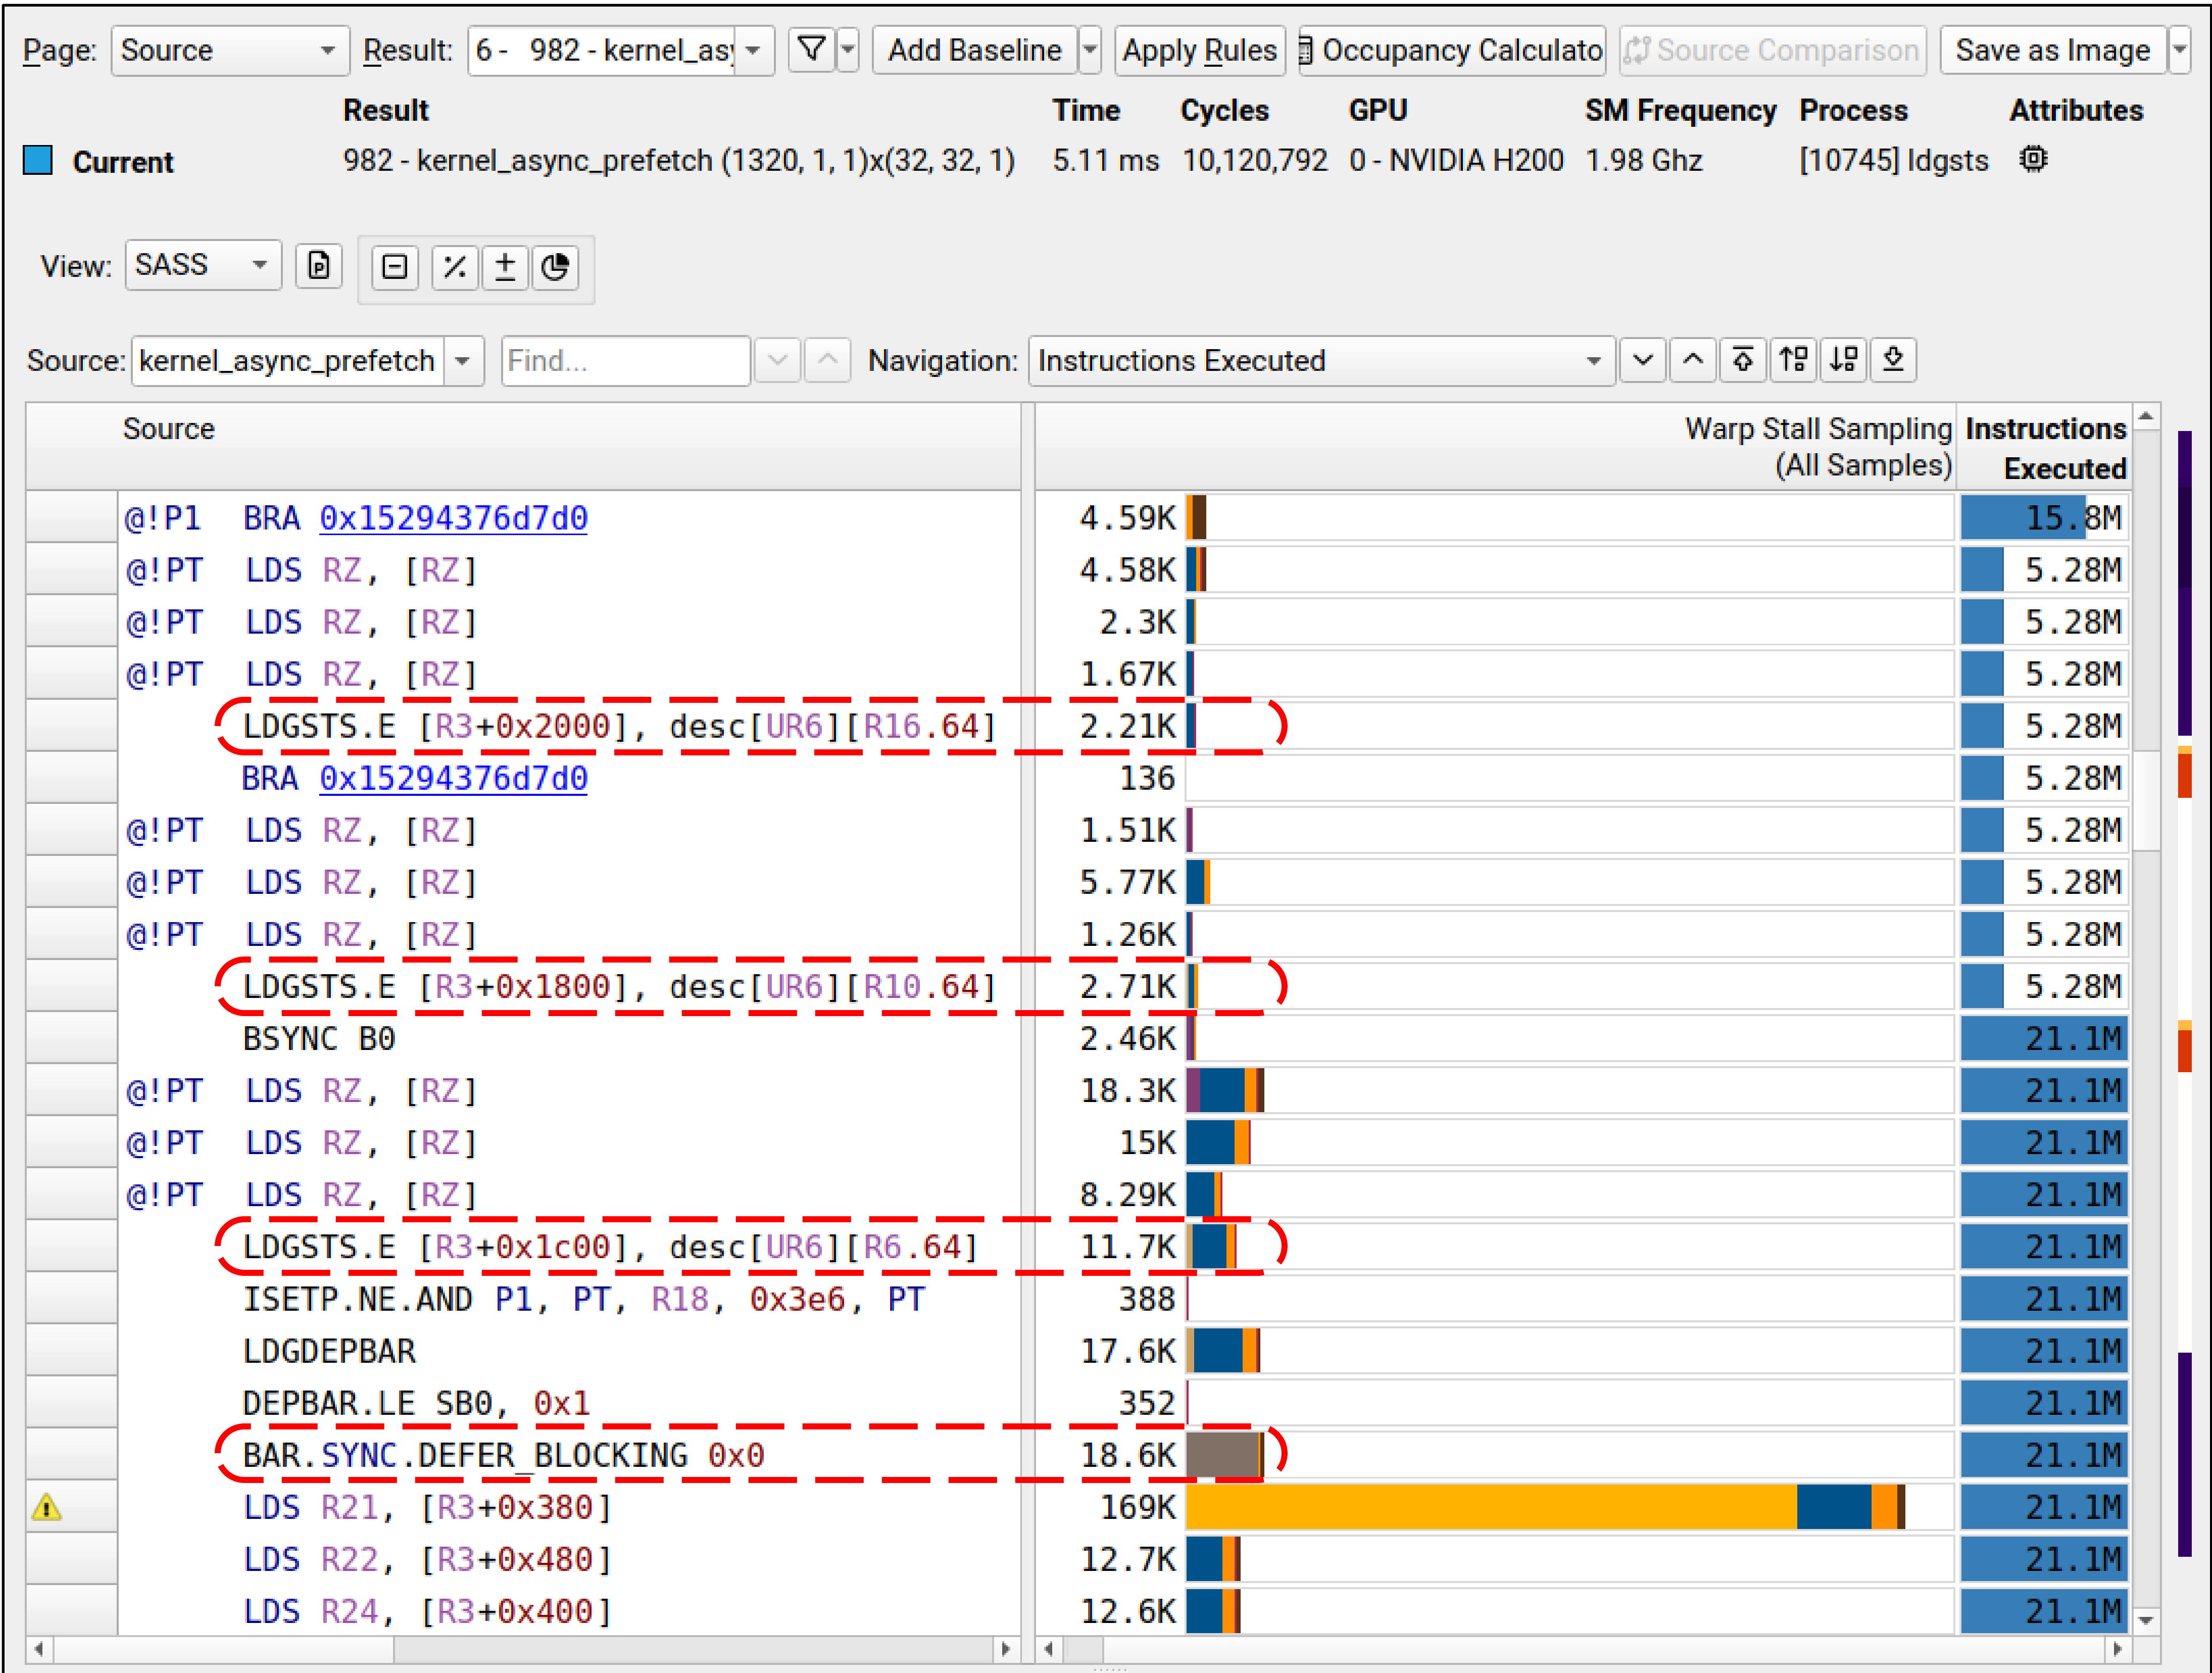
# Asynchronous Data Copies

~~LDGSTS~~ use case: batching loads in conditional code – NCU



# Asynchronous Data Copies

LDGSTS use case: batching loads in conditional code – NCU



# Asynchronous Data Copies

LDGSTS use case: staging/prefetching

```
T* gmem;
__shared__ T smem[2][..];
int is = 0;
__pipeline_memcpy_async(&smem[is][sind], &gmem[gind], sizeof(T));
__pipeline_commit();
is ^= 1;          gind += stride;

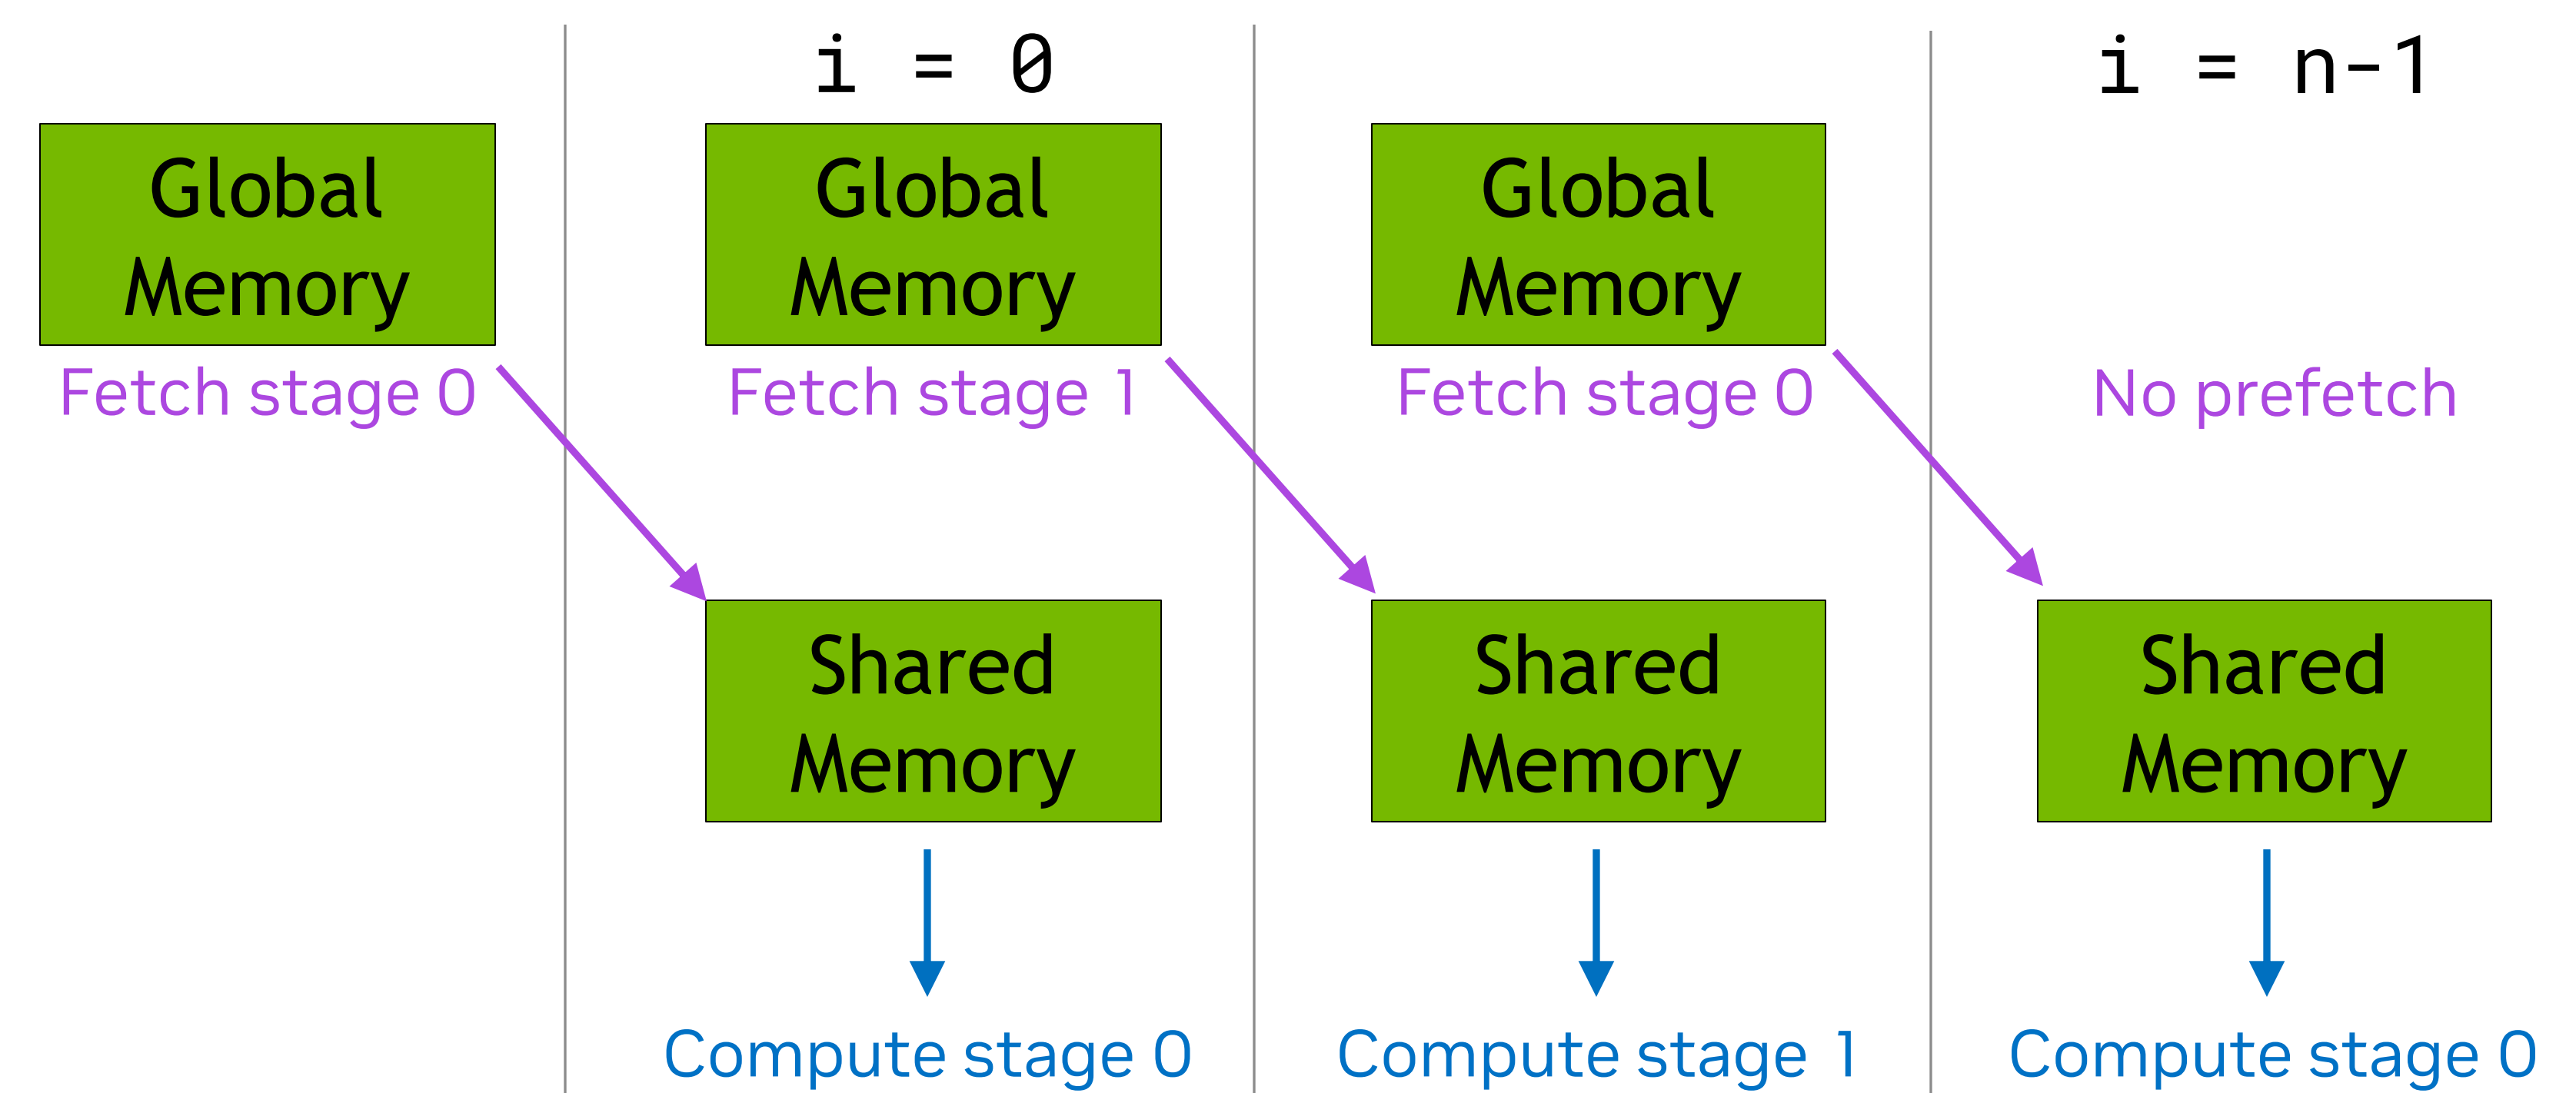
for (int i = 0; i < n - 1; ++i) {
    __pipeline_memcpy_async(&smem[is][sind], &gmem[gind], sizeof(T));
    __pipeline_commit();
    is ^= 1;          gind += stride;

    __pipeline_wait_prior(1);
    __syncthreads();

    // compute with smem[is][]
    __syncthreads();
}

__pipeline_wait_prior(0);
__syncthreads();
// compute with smem[is][]
```

```
// Synchronous copy code:
for (int i = 0; i < n; ++i) {
    smem[sind] = g_mem[gind];
    gind += stride;
    __syncthreads();
    // compute with smem[]
    __syncthreads();
}
```





# Asynchronous Data Copies

LDGSTS use case: staging/prefetching

```
T* gmem;
__shared__ T smem[2][..];
int is = 0;
__pipeline_memcpy_async(&smem[is][sind], &gmem[gind], sizeof(T));
__pipeline_commit();
is ^= 1;          gind += stride;

for (int i = 0; i < n - 1; ++i) {
    __pipeline_memcpy_async(&smem[is][sind], &gmem[gind], sizeof(T));
    __pipeline_commit();
    is ^= 1;          gind += stride;


    __pipeline_wait_prior(1);
    __syncthreads();

    // compute with smem[is][]
    __syncthreads();
}

__pipeline_wait_prior(0);
__syncthreads();

// compute with smem[is][]
```

```
// Synchronous copy code:
for (int i = 0; i < n; ++i) {
    smem[sind] = g_mem[gind];
    gind += stride;
    __syncthreads();
    // compute with smem[]
    __syncthreads();
}
```

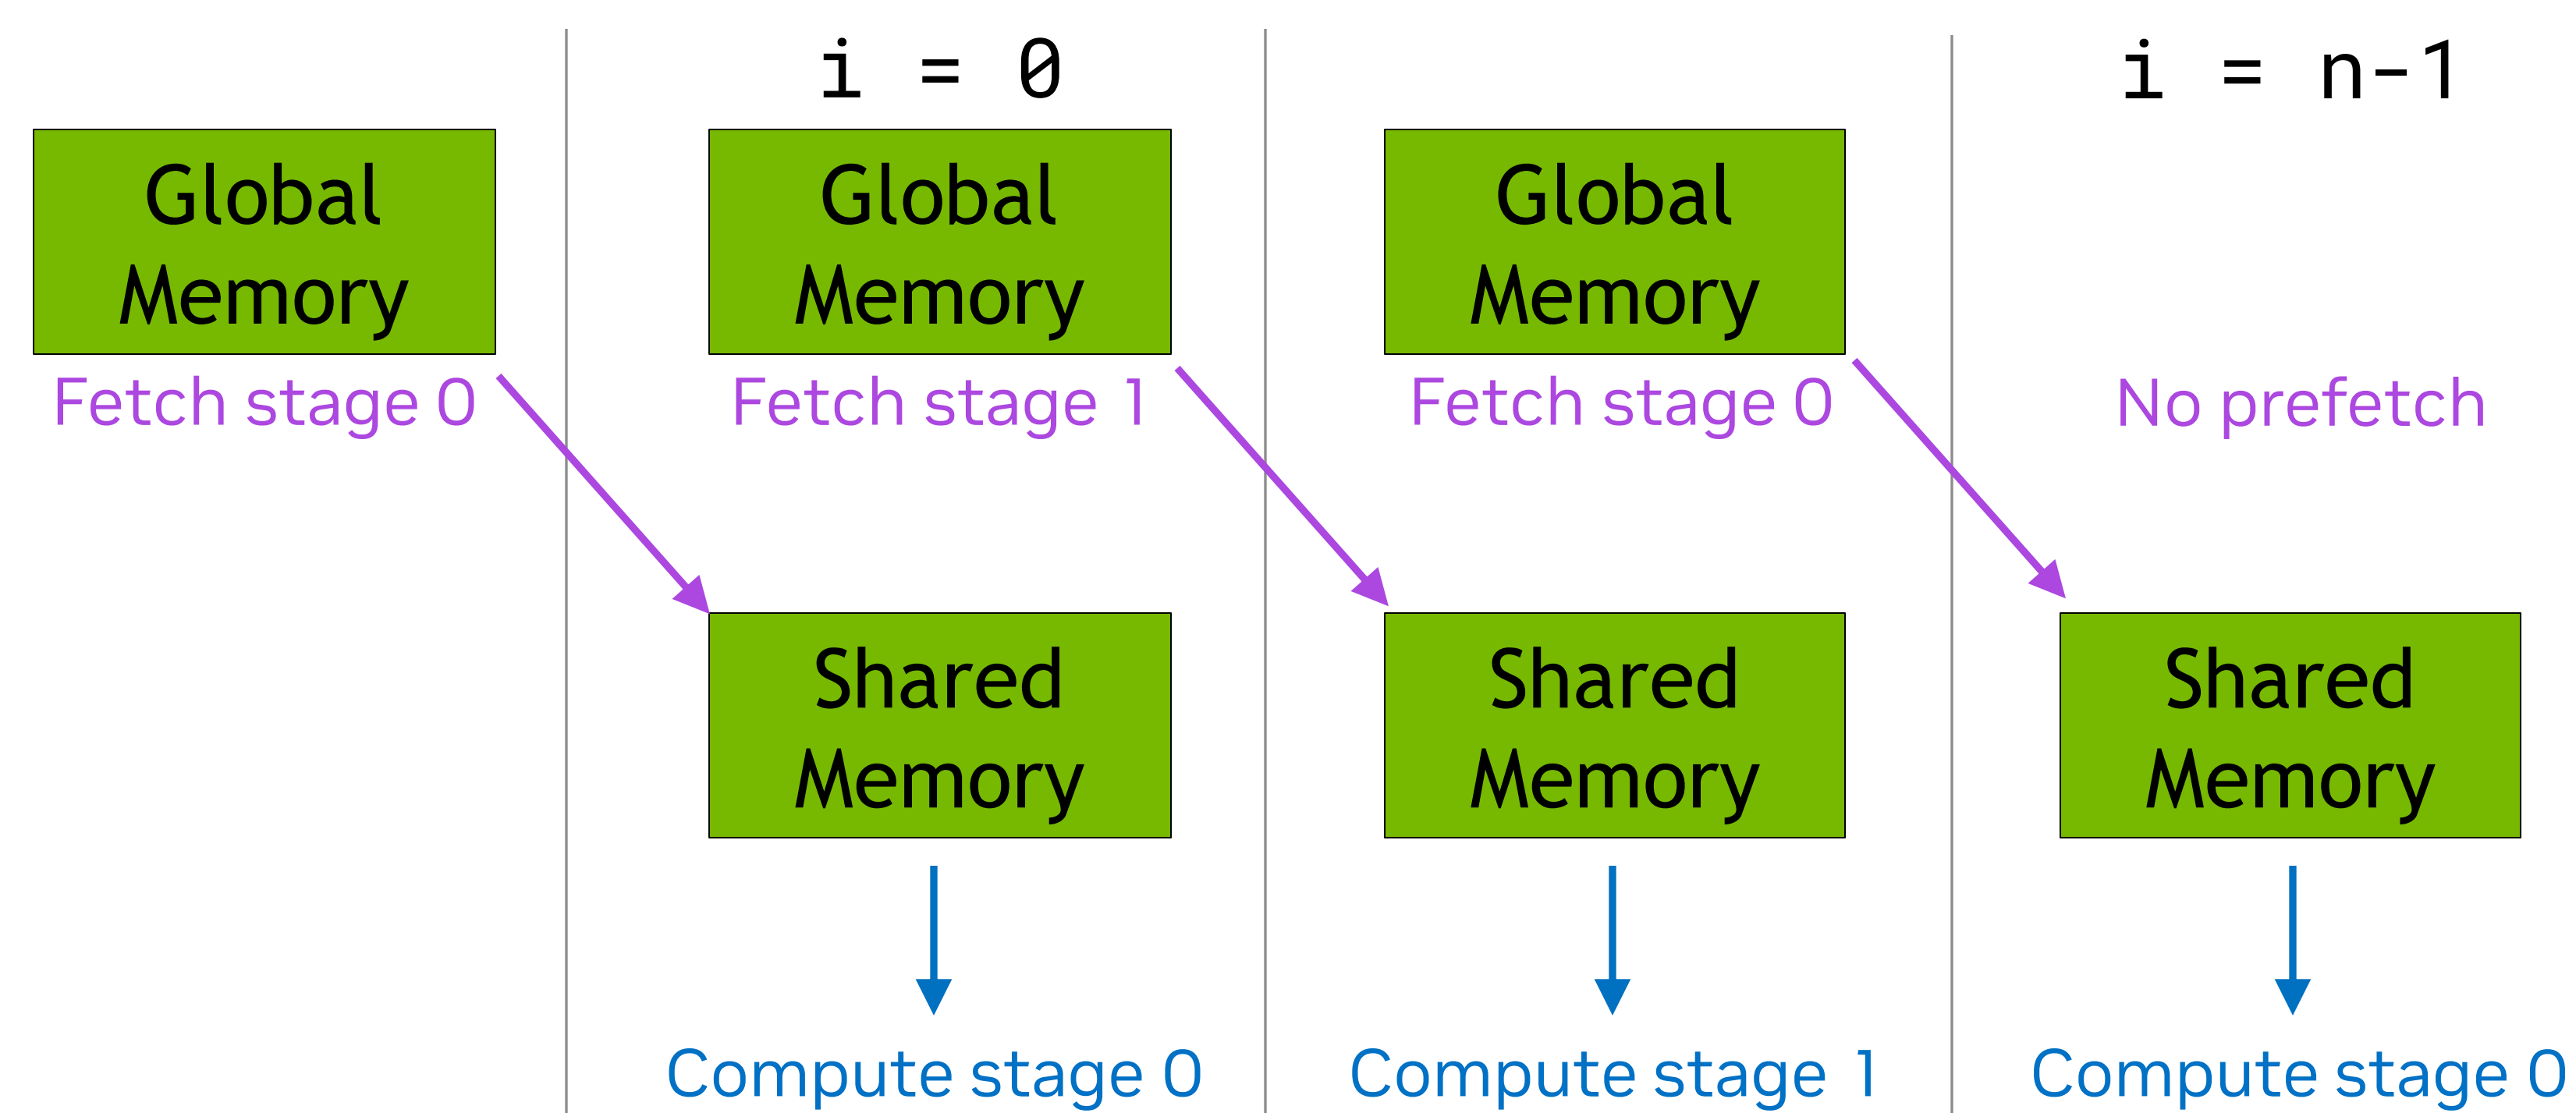


# Asynchronous Data Copies

LDGSTS use case: staging/prefetching

```
T* gmem;
__shared__ T smem[3][..];
int is = 0;
__pipeline_memcpy_async(&smem[is][sind], &gmem[gind], sizeof(T));
__pipeline_commit();
is = (is + 1) % 3; gind += stride;
for (int i = 0; i < n - 1; ++i) {
    __pipeline_memcpy_async(&smem[is][sind], &gmem[gind], sizeof(T));
    __pipeline_commit();
    is = (is + 1) % 3; gind += stride;
    __pipeline_wait_prior(1);
    __syncthreads();
    // compute with smem[is][]
    __syncthreads();
}
__pipeline_wait_prior(0);
__syncthreads();
// compute with smem[is][]
```

```
// Synchronous copy code:
for (int i = 0; i < n; ++i) {
    smem[sind] = g_mem[gind];
    gind += stride;
    __syncthreads();
    // compute with smem[]
    __syncthreads();
}
```



# Asynchronous Data Copies

LDGSTS use case: staging/prefetching

**More stages  
(prefetch multiple iterations)  
=> more Bytes in flight!**

```
T* gmem;
__shared__ T smem[3][..];

int is = 0;
__pipeline_memcpy_async(&smem[is][sind], &gmem[gind], sizeof(T));
__pipeline_commit();
is = (is + 1) % 3; gind += stride;

for (int i = 0; i < n - 1; ++i) {
    __pipeline_memcpy_async(&smem[is][sind], &gmem[gind], sizeof(T));
    __pipeline_commit();
    is = (is + 1) % 3; gind += stride;


    __pipeline_wait_prior(1);
    __syncthreads();

    // compute with smem[is][]
    __syncthreads();
}

__pipeline_wait_prior(0);
__syncthreads();

// compute with smem[is][]
```

```
// Synchronous copy code:
for (int i = 0; i < n; ++i) {
    smem[sind] = g_mem[gind];
    gind += stride;
    __syncthreads();
    // compute with smem[]
    __syncthreads();
}
```





# Asynchronous Data Copies

## Ways to code LDGSTS

```
// primitives <cuda_pipeline.h>:
__pipeline_memcpy_async(&smem[sind], &gmem[gind], sizeof(T)); // LDGSTS
__pipeline_commit(); // LDGDEPBAR
__pipeline_wait_prior(N); // DEPBAR.LE SB,N

// libcu++ <cuda/pipeline.h>:
auto pipeline = cuda::make_pipeline();
pipeline.producer_acquire();
cuda::memcpy_async(&smem[sind], &gmem[gind], sizeof(T), pipeline); // LDGSTS
pipeline.producer_commit(); // LDGDEPBAR
cuda::pipeline_consumer_wait_prior<N>(pipeline); // DEPBAR.LE SB,N
pipeline.consumer_release();

// cooperative_groups <cooperative_groups/memcpy_async.h>:
auto block = cg::this_thread_block();
cg::memcpy_async(block, smem, gmem, sizeof(T) * block.size()); // LDGSTS + LDGDEPBAR
cg::wait_prior<N>(block); // DEPBAR.LE SB,N
```


# Asynchronous Data Copies

LDGSTS with barrier

## Producer/consumer example:

Loop, at each iteration:

1. Producer threads copy area of GMEM to SMEM
2. Consumer threads access SMEM after copy completion




# Asynchronous Data Copies

LDGSTS with barrier

## Producer/consumer example:

Loop, at each iteration:

1. Producer threads copy area of GMEM to SMEM
2. Consumer threads access SMEM after copy completion





# Asynchronous Data Copies

LDGSTS with barrier

```
__global__ void kernel (T const* __restrict__ gmem, int n)
{
    __shared__ T smem[SMEM_SIZE];
    __shared__ __mbarrier_t bar;
    bool is_producer = (threadIdx.x < N_PRODUCERS);
    if (threadIdx.x == 0)
        __mbarrier_init(&bar, N_PRODUCERS); // count := N_PRODUCERS (reset)
    __syncthreads();
    for (int it = 0, phase = 0; it < n; ++it) {
        if (is_producer) {
            for (int i = threadIdx.x; i < SMEM_SIZE; i += N_PRODUCERS)
                __pipeline_memcpy_async(&smem[i], &gmem[it * SMEM_SIZE + i], sizeof(T));
            In this order! { __pipeline_arrive_on(&bar); // count += N_PRODUCERS and count -= N_PRODUCERS by HW on copy completion
                           __mbarrier_arrive(&bar); // count -= N_PRODUCERS
                           }
        }
    }
}
```

# Asynchronous Data Copies

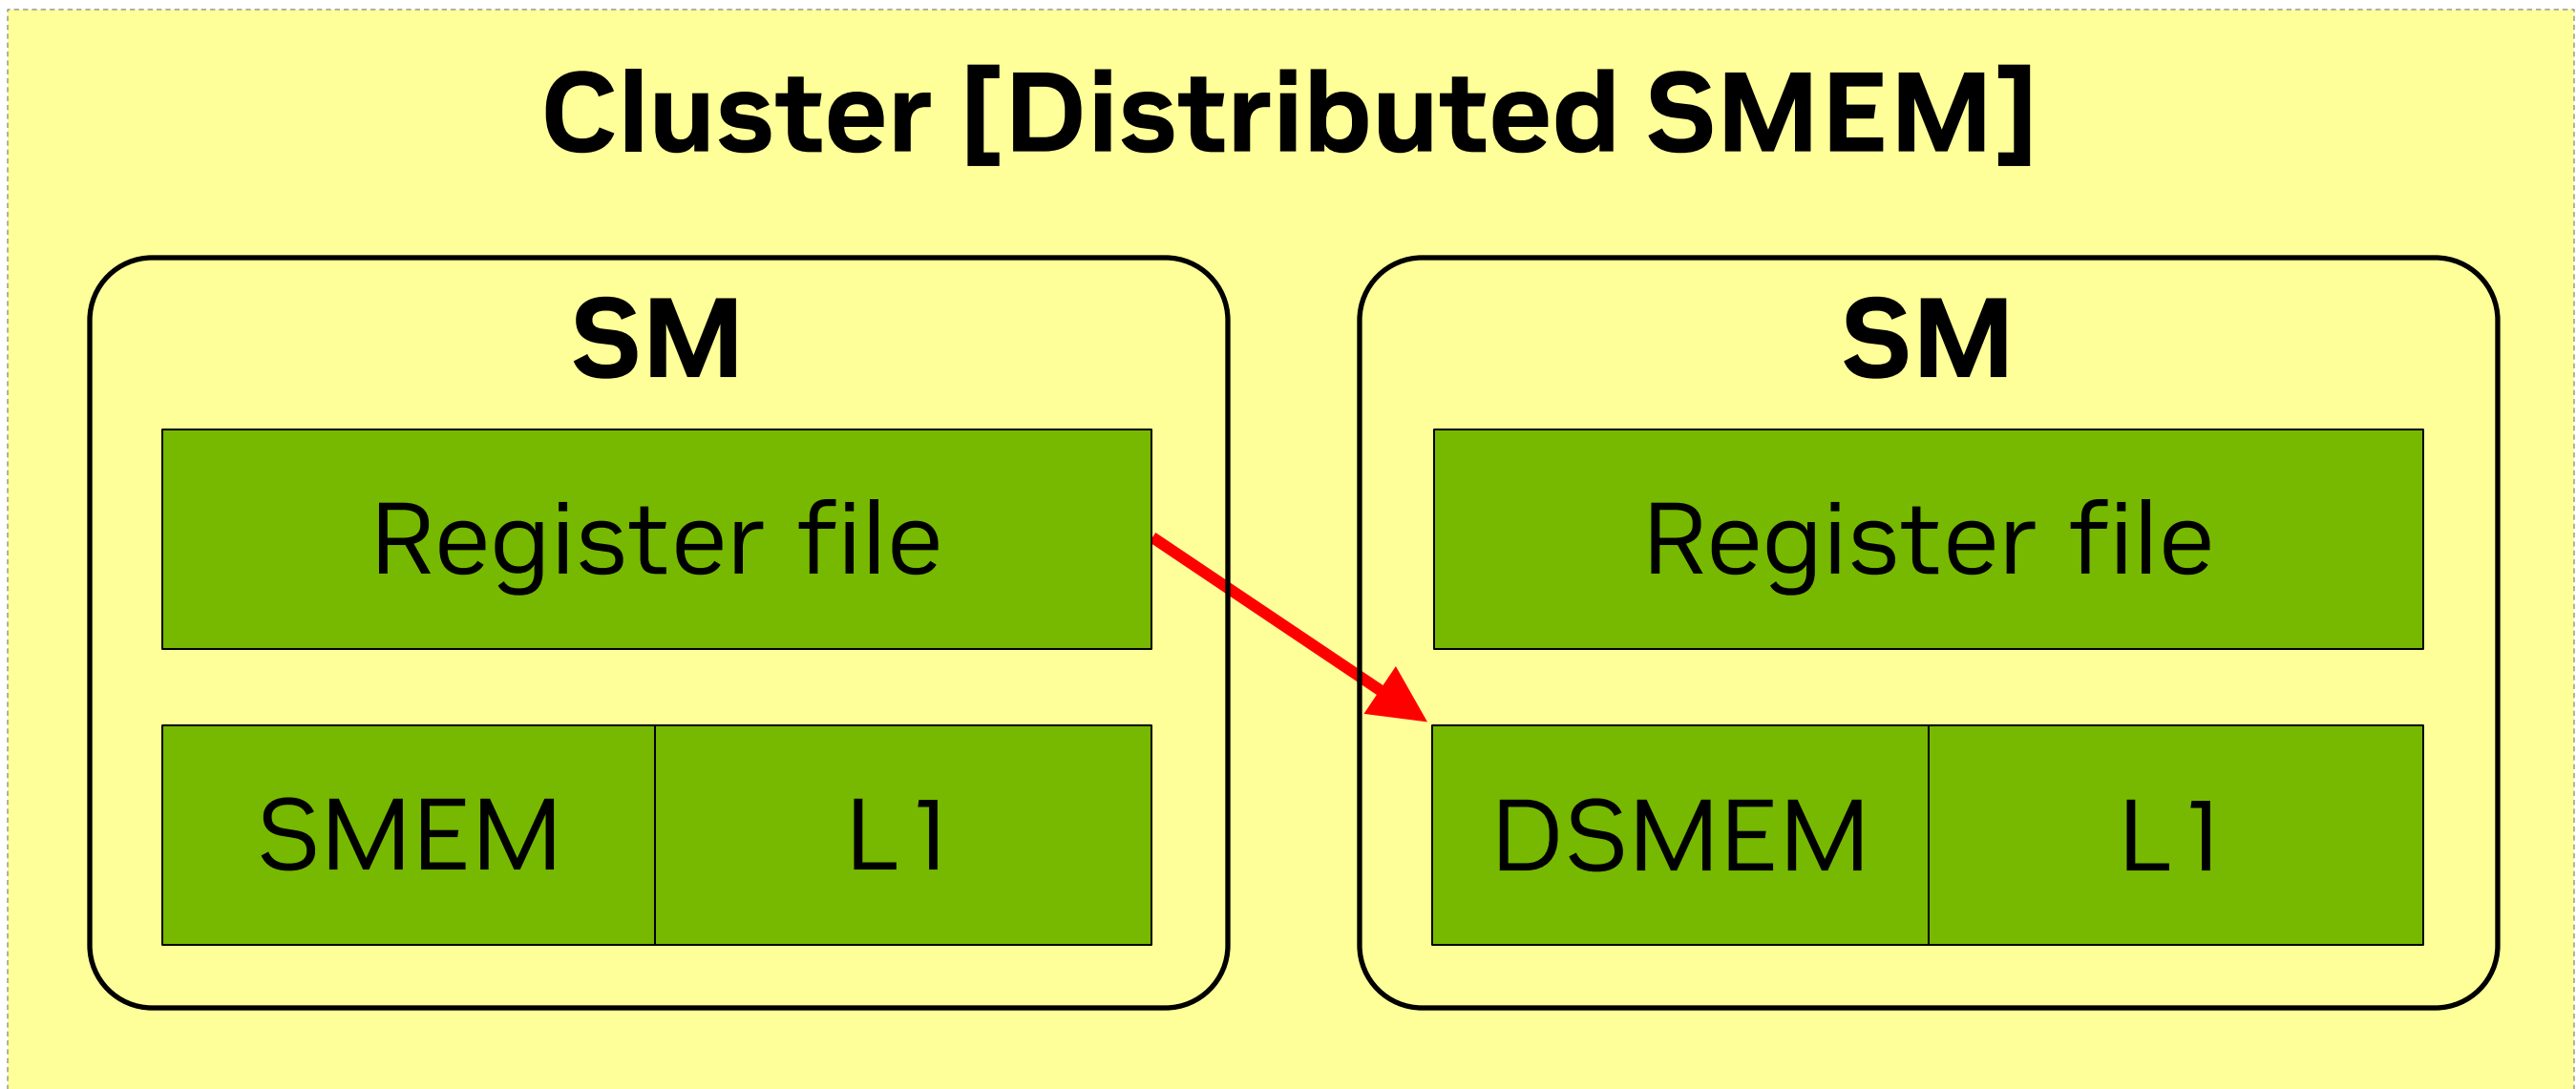
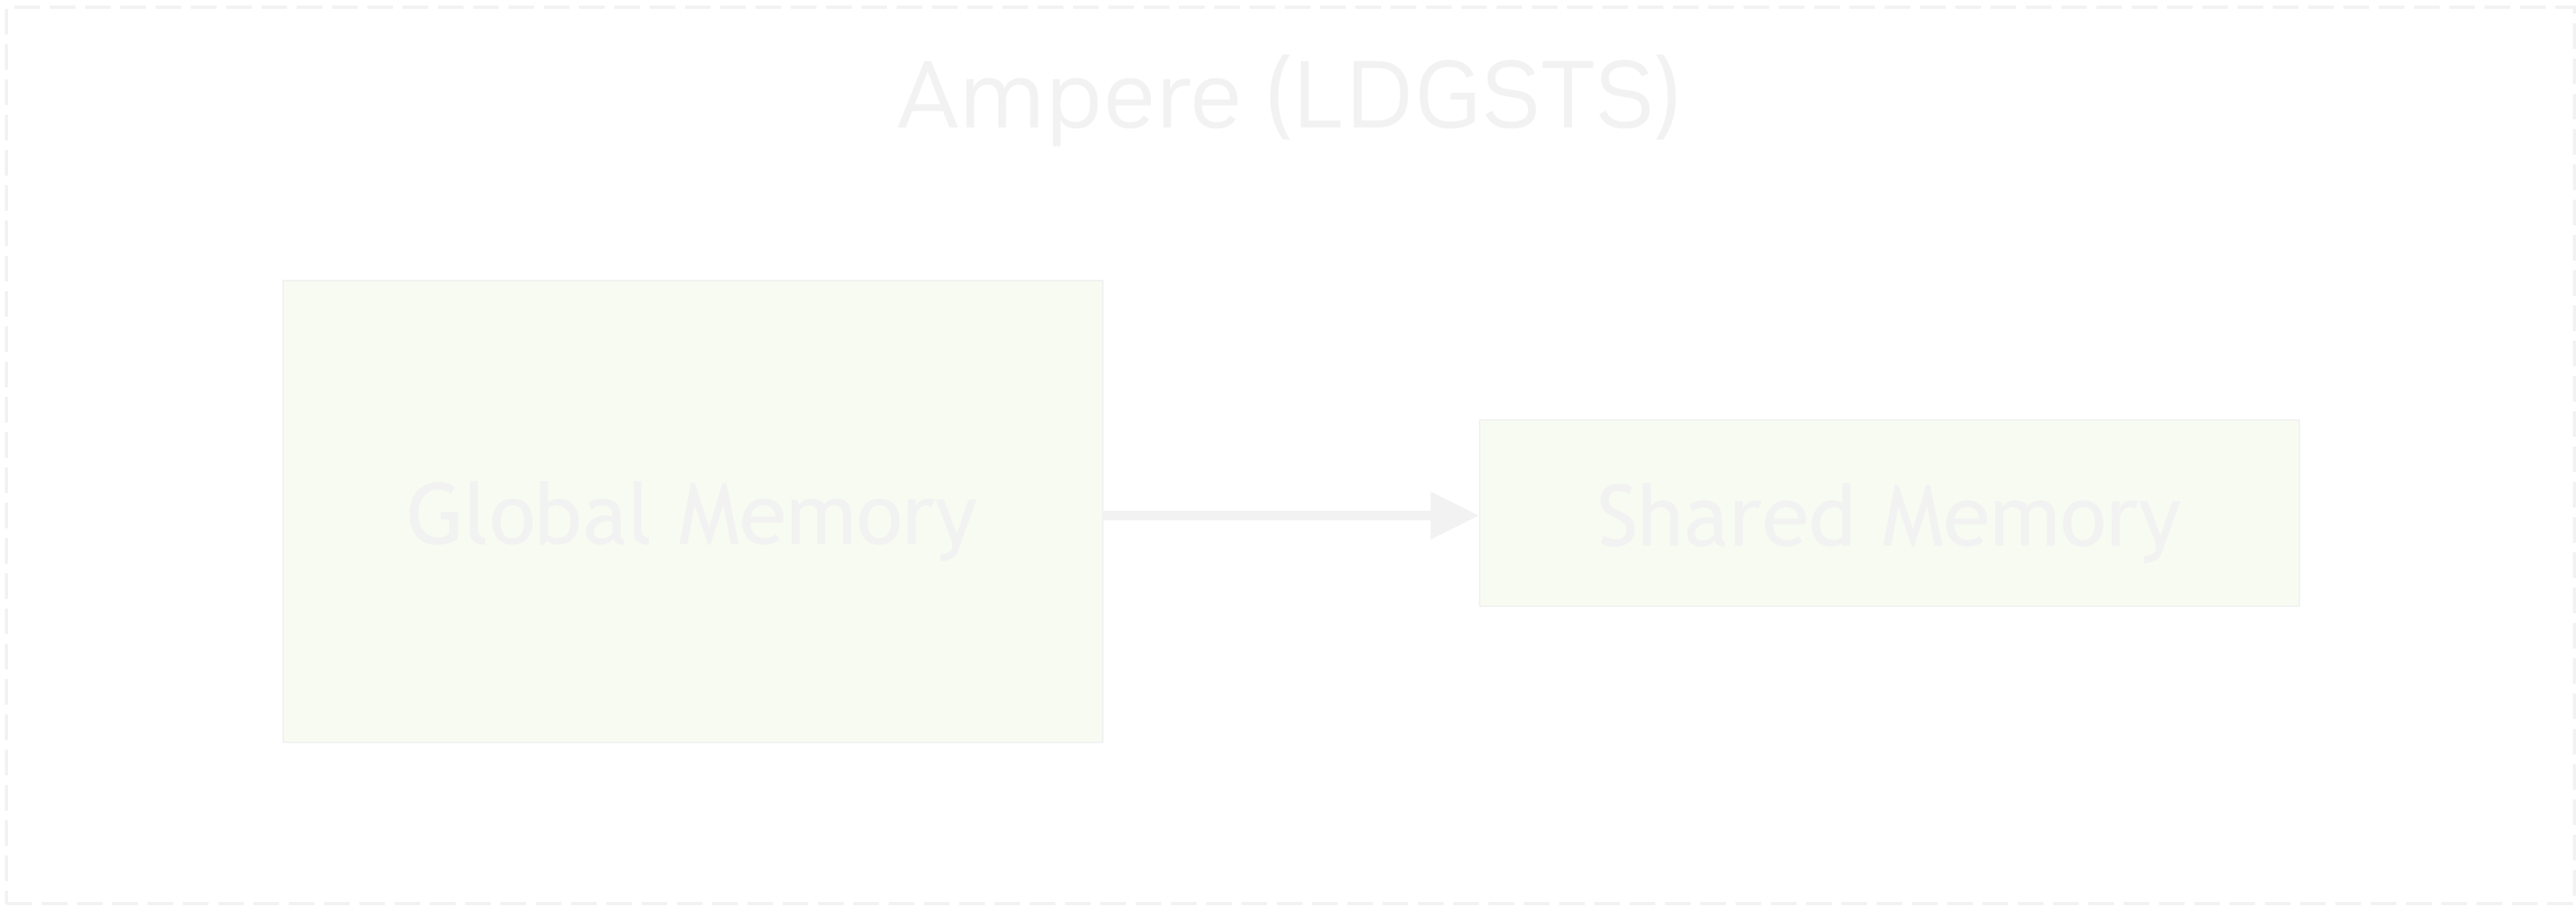
LDGSTS with barrier

```
__global__ void kernel (T const* __restrict__ gmem, int n)
{
    __shared__ T smem[SMEM_SIZE];
    __shared__ __mbarrier_t bar;
    bool is_producer = (threadIdx.x < N_PRODUCERS);
    if (threadIdx.x == 0)
        __mbarrier_init(&bar, N_PRODUCERS); // count := N_PRODUCERS (reset)
    __syncthreads();
    for (int it = 0, phase = 0; it < n; ++it) {
        if (is_producer) {
            for (int i = threadIdx.x; i < SMEM_SIZE; i += N_PRODUCERS)
                __pipeline_memcpy_async(&smem[i], &gmem[it * SMEM_SIZE + i], sizeof(T));
        }
        In this order! { __pipeline_arrive_on(&bar); // count += N_PRODUCERS and count -= N_PRODUCERS by HW on copy completion
                        __mbarrier_arrive(&bar); // count -= N_PRODUCERS
        } else {
            while (!__mbarrier_try_wait_parity(&bar, phase, 1000))
                {}
            // Use loaded smem here
        }
        phase ^= 1;
        __syncthreads();
    }
}
```

# Asynchronous Data Copies

Asynchronous copies – STAS

HW	Direction	Operation	Synchronization	C++	SASS
Hopper	GMEM → SMEM	Asynchronous version of smem[sind] := gmem[gind]	Thread-local	__pipeline_memcpy_async	LDGSTS
			SMEM barrier		
	REG → DSMEM	Asynchronous version of dsmem[sind] := var	DSMEM barrier	cuda::ptx::st_async	STAS
	GMEM → (D)SMEM	1D – 5D block copy Uniform	(D)SMEM barrier	cuda::device::experimental:: cp_async_bulk_tensor_Nd_global_to_shared cp_async_bulk_tensor_Nd_shared_to_global	UTMALDG
	SMEM → GMEM		Thread-local		UTMASTG
	GMEM → SMEM	1D block copy Uniform	SMEM barrier	cuda::device::experimental:: cp_async_bulk_global_to_shared cp_async_bulk_shared_to_global	UBLKCP
	SMEM → GMEM		Thread-local		
	SMEM → (D)SMEM		(D)SMEM barrier	cuda::ptx::cp_async_bulk*	





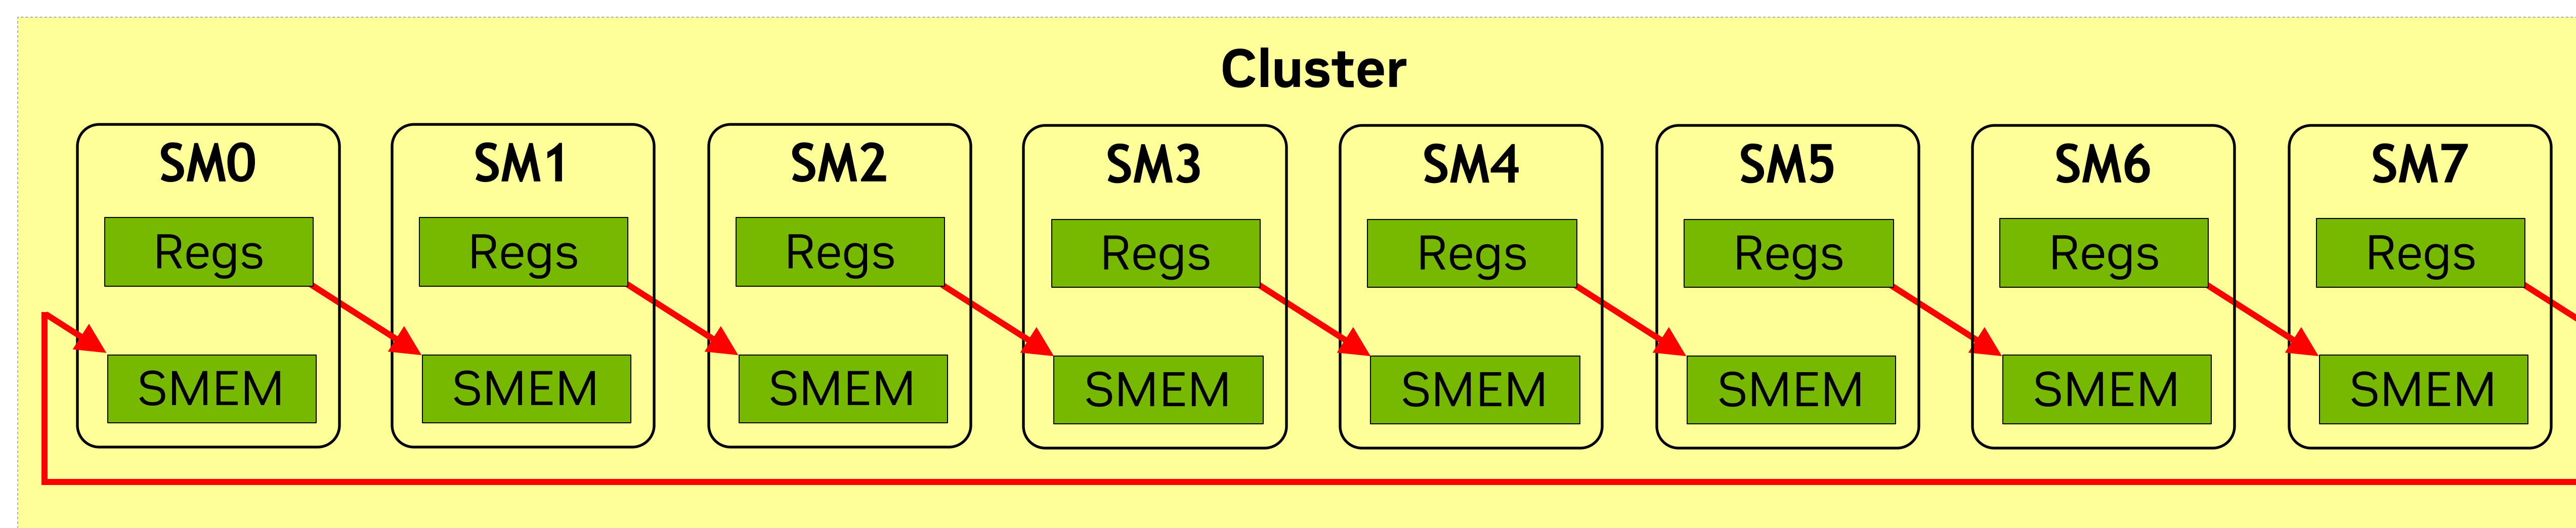
# Asynchronous Data Copies

Hopper: cluster asynchronous copies (REG -> DSMEM)

**Example (code in next slide):**

In a loop:

CTA #i copies data to SMEM of CTA #i+1, last CTA – to SMEM of CTA #0



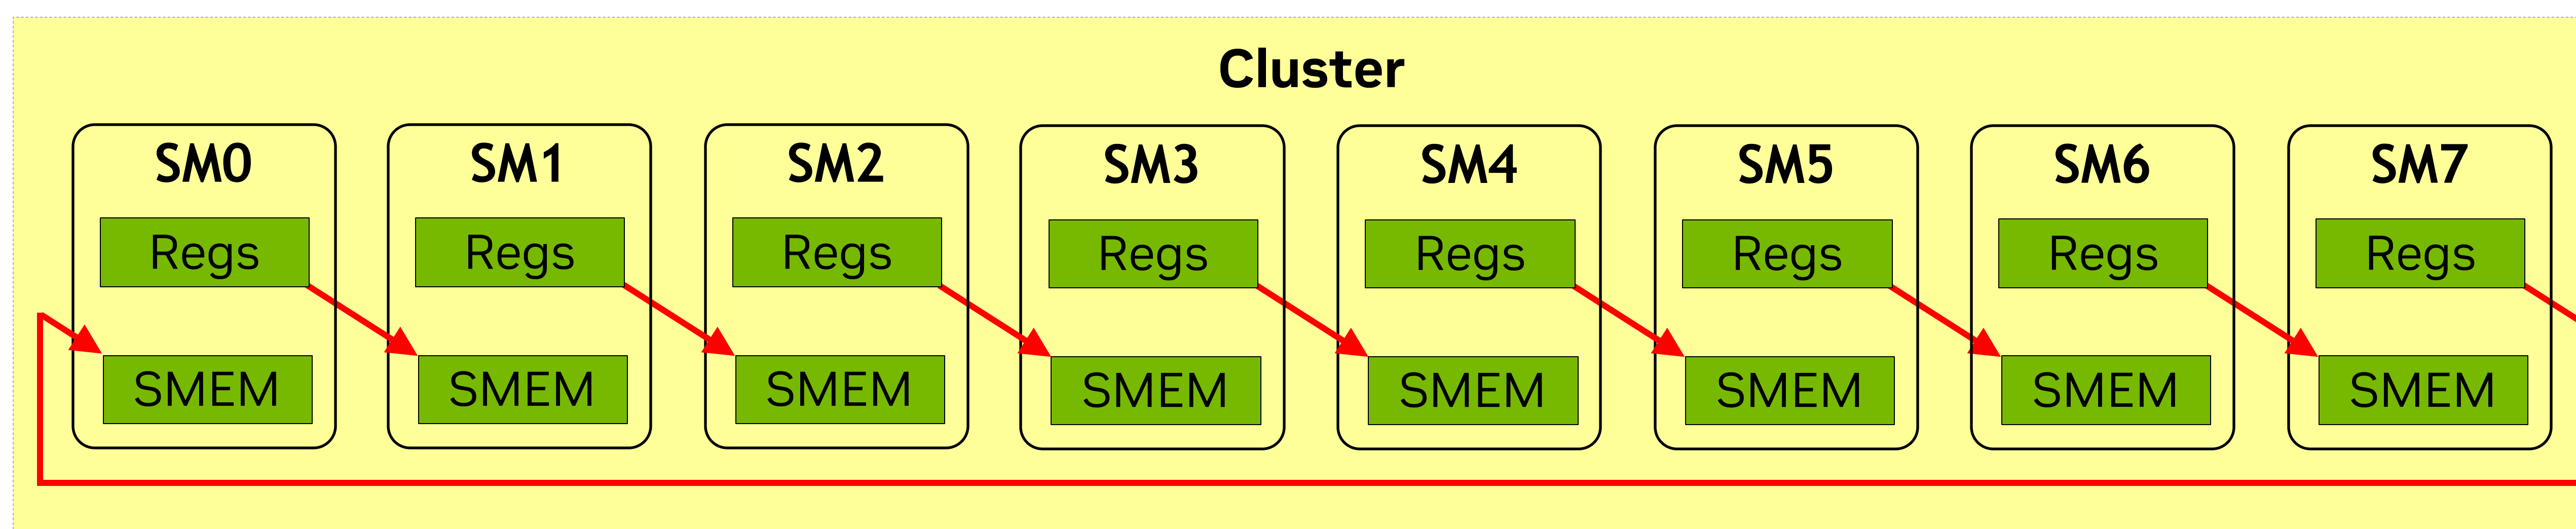
# Asynchronous Data Copies

Hopper: cluster asynchronous copies (REG -> DSMEM)

**Example (code in next slide):**

In a loop:

CTA #i copies data to SMEM of CTA #i+1, last CTA – to SMEM of CTA #0



Two barriers on each CTA:

- `bar` – notifies **consumer CTA** that data copy from **producer CTA** is complete
- `bar_done` – notifies **producer CTA** that **consumer CTA** is ready to receive data

# Asynchronous Data Copies

Hopper: cluster asynchronous copies (REG -> DSMEM)

```
__global__ __cluster_dims__(8, 1, 1)
void kernel ()
{
    using namespace cooperative_groups;
    using namespace cuda::device;
    using namespace cuda::ptx;

    __shared__ int4 smem[1024];
    __shared__ cuda::barrier<thread_scope_block> bar;
    __shared__ cuda::barrier<thread_scope_block> bar_done;
    if (threadIdx.x == 0) {
        init(&bar, 1);
        init(&bar_done, 1024);
    }

    auto cluster = this_cluster();
    cluster.sync();

    int rk = cluster.block_rank();
    int rk_next = (rk + 1) % 8;
    int rk_prev = (rk + 7) % 8; // Same as '-1'.

    auto smem_next = cluster.map_shared_rank(smem, rk_next);
    auto bar_next = cluster.map_shared_rank(
        barrier_native_handle(bar), rk_next);
    auto bar_prev = cluster.map_shared_rank(
        barrier_native_handle(bar_done), rk_prev);
```



# Asynchronous Data Copies

Hopper: cluster asynchronous copies (REG -> DSMEM)

```
__global__ __cluster_dims__(8, 1, 1)
void kernel ()
{
    using namespace cooperative_groups;
    using namespace cuda::device;
    using namespace cuda::ptx;

    __shared__ int4 smem[1024];
    __shared__ cuda::barrier<thread_scope_block> bar;
    __shared__ cuda::barrier<thread_scope_block> bar_done;
    if (threadIdx.x == 0) {
        init(&bar, 1);
        init(&bar_done, 1024);
    }

    auto cluster = this_cluster();
    cluster.sync();

    int rk = cluster.block_rank();
    int rk_next = (rk + 1) % 8;
    int rk_prev = (rk + 7) % 8; // Same as '-1'.

    auto smem_next = cluster.map_shared_rank(smem, rk_next);
    auto bar_next = cluster.map_shared_rank(
        barrier_native_handle(bar), rk_next);
    auto bar_prev = cluster.map_shared_rank(
        barrier_native_handle(bar_done), rk_prev);
```

// kernel cont-d:

```
    for (int it = 0, phase = 0; it < 1000; ++it) {
        int4 r = {rk, rk_next, (int)threadIdx.x, it};

        st_async(&smem_next[1023 - threadIdx.x].x,
                {r.x, r.y, r.z, r.w}, bar_next);
        if (threadIdx.x == 0)
            mbarrier_arrive_expect_tx(
                sem_release, scope_cluster, space_shared,
                barrier_native_handle(bar), sizeof(smem));

        while (!mbarrier_try_wait_parity(
            barrier_native_handle(bar), phase, 1000)) {}

        r = smem[threadIdx.x]; // done reading
        mbarrier_arrive(sem_release, scope_cluster, space_cluster,
            bar_prev);

        while (!mbarrier_try_wait_parity(
            barrier_native_handle(bar_done), phase, 1000)) {}

        phase ^= 1;
    }
} // void kernel ()
```

Any order

tx\_count -= 16 by each of 1024 threads on completion

tx\_count += 16384

phase flip when tx\_count gets to 0

# Asynchronous Data Copies

Hopper: cluster asynchronous copies (REG -> DSMEM)

```
__global__ __cluster_dims__(8, 1, 1)
void kernel ()
{
    using namespace cooperative_groups;
    using namespace cuda::device;
    using namespace cuda::ptx;

    __shared__ int4 smem[1024];
    __shared__ cuda::barrier<thread_scope_block> bar;
    __shared__ cuda::barrier<thread_scope_block> bar_done;

    if (threadIdx.x == 0) {
        init(&bar, 1);
        init(&bar_done, 1024);
    }

    auto cluster = this_cluster();
    cluster.sync();

    int rk = cluster.block_rank();
    int rk_next = (rk + 1) % 8;
    int rk_prev = (rk + 7) % 8; // Same as '-1'.

    auto smem_next = cluster.map_shared_rank(smem, rk_next);
    auto bar_next = cluster.map_shared_rank(
        barrier_native_handle(bar), rk_next);
    auto bar_prev = cluster.map_shared_rank(
        barrier_native_handle(bar_done), rk_prev);
}
```

// kernel cont-d:

```
for (int it = 0, phase = 0; it < 1000; ++it) {
    int4 r = {rk, rk_next, (int)threadIdx.x, it};

    st_async(&smem_next[1023 - threadIdx.x].x,
        {r.x, r.y, r.z, r.w}, bar_next);
    if (threadIdx.x == 0)
        mbarrier_arrive_expect_tx(
            sem_release, scope_cluster, space_shared,
            barrier_native_handle(bar), sizeof(smem));

    while (!mbarrier_try_wait_parity(
        barrier_native_handle(bar), phase, 1000)) {}

    r = smem[threadIdx.x]; // done reading
    mbarrier_arrive(sem_release, scope_cluster, space_cluster,
        bar_prev);

    while (!mbarrier_try_wait_parity(
        barrier_native_handle(bar_done), phase, 1000)) {}

    phase ^= 1;
} // void kernel ()
```

Any order

tx\_count -= 16 by each of 1024 threads on completion

tx\_count += 16384

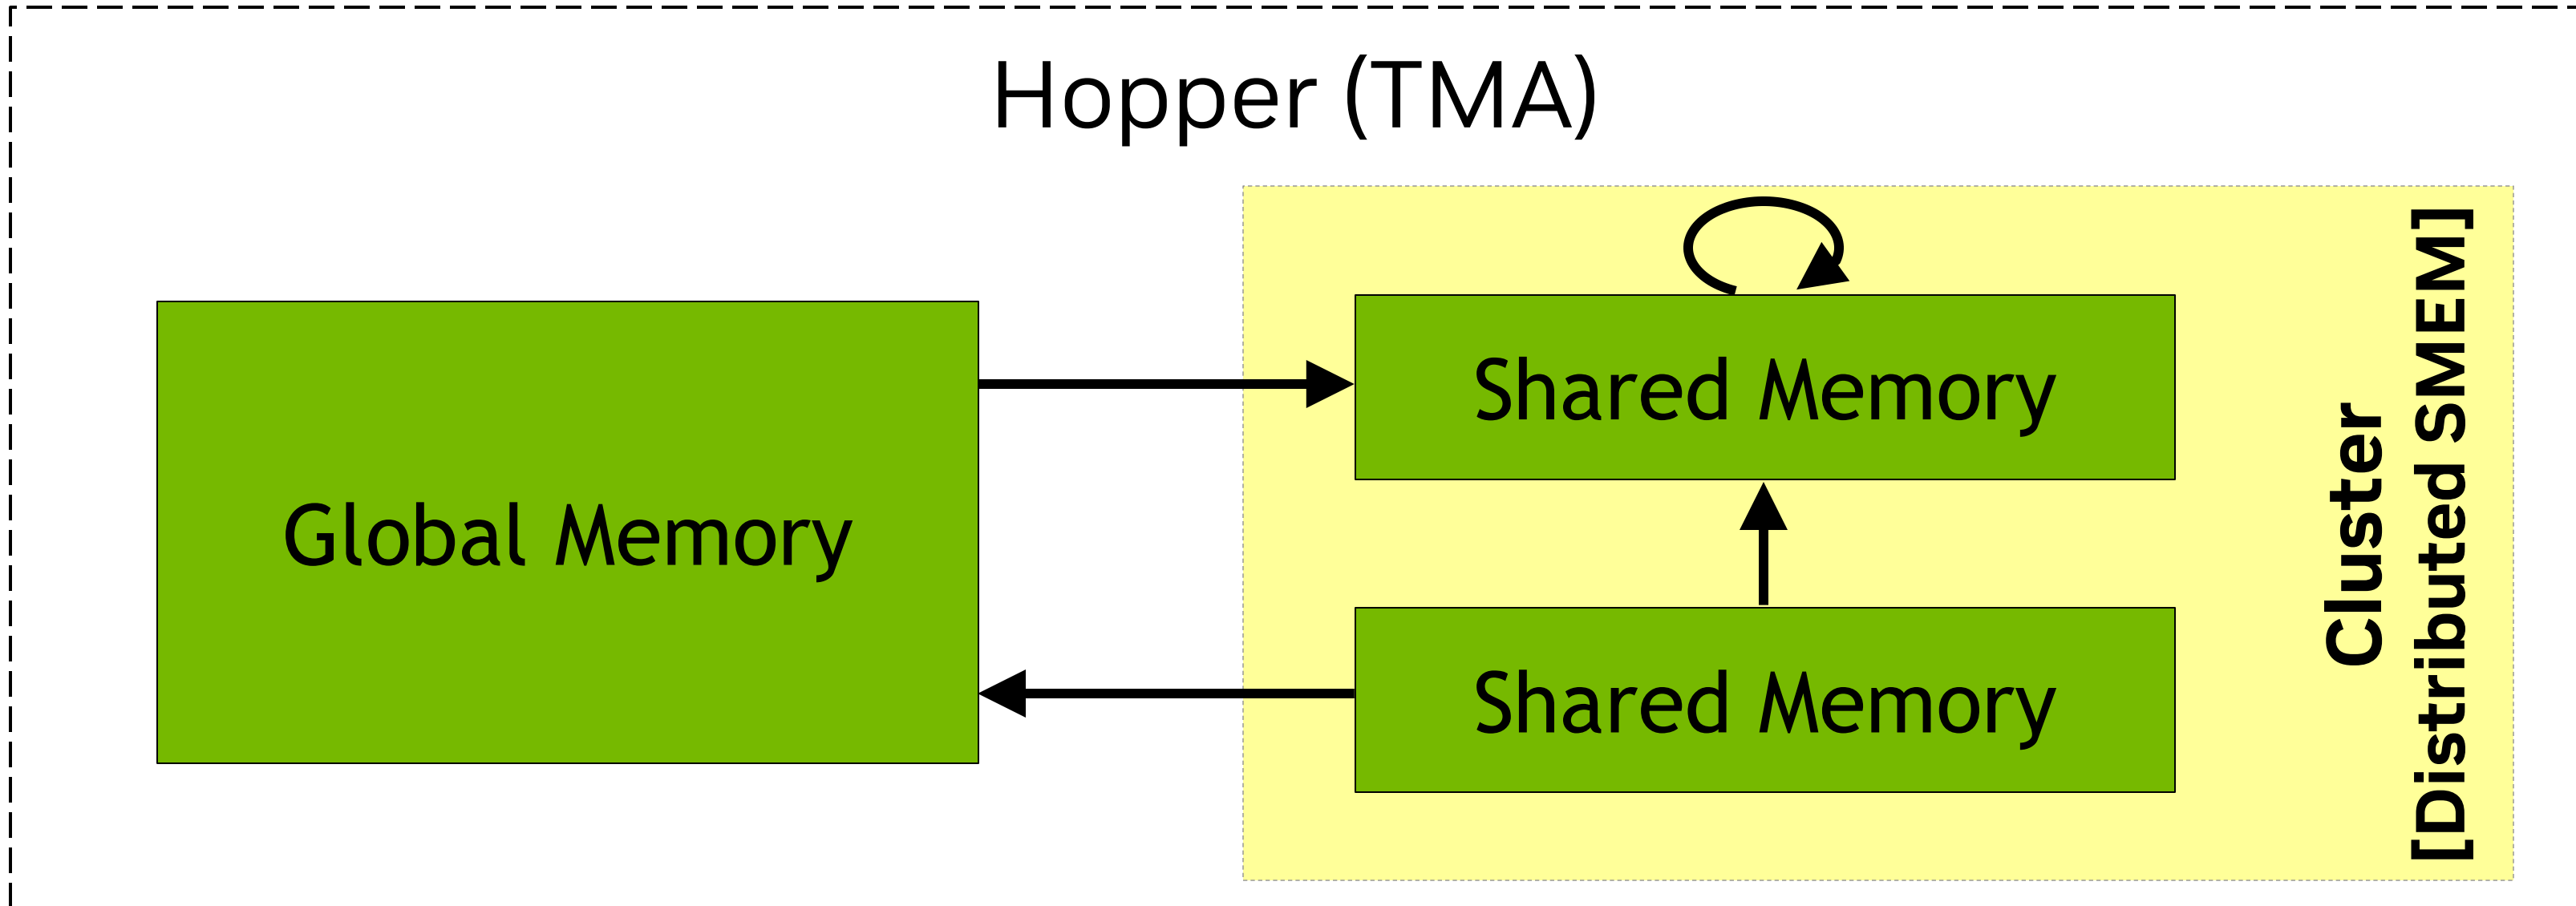
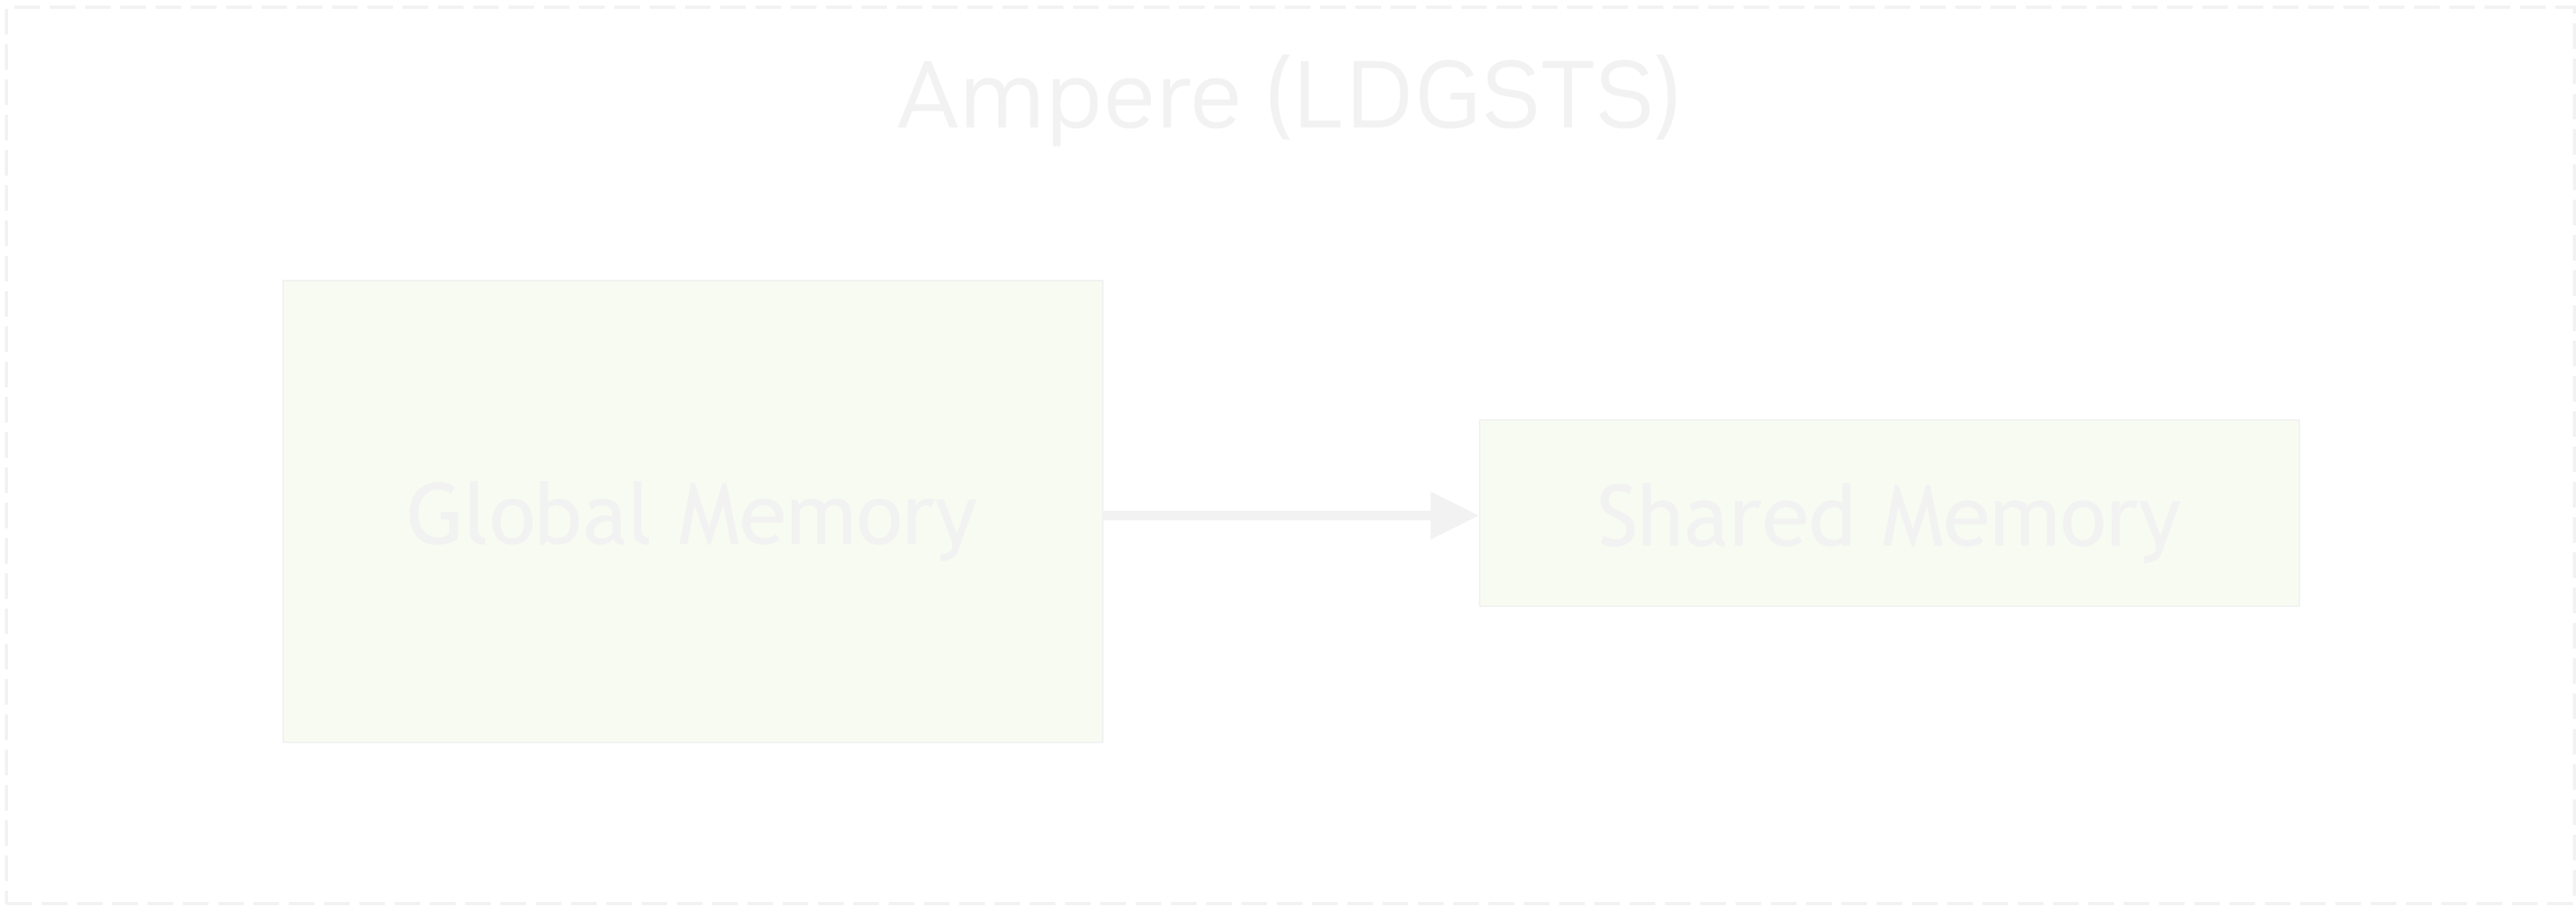
phase flip when tx\_count gets to 0

Space matches barrier location!

# Asynchronous Data Copies

Asynchronous copies – UTMA\*

HW	Direction	Operation	Synchronization	C++	SASS
Hopper	GMEM → SMEM	Asynchronous version of smem[sind] := gmem[gind]	Thread-local	__pipeline_memcpy_async	LDGSTS
			SMEM barrier		
	REG → DSMEM	Asynchronous version of dsmem[sind] := var	DSMEM barrier	cuda::ptx::st_async	STAS
	GMEM → (D)SMEM	1D – 5D block copy Uniform	(D)SMEM barrier	cuda::device::experimental:: cp_async_bulk_tensor_Nd_global_to_shared	UTMALDG
	SMEM → GMEM		Thread-local	cp_async_bulk_tensor_Nd_shared_to_global	UTMASTG
	GMEM → SMEM	1D block copy Uniform	SMEM barrier	cuda::device::experimental:: cp_async_bulk_global_to_shared cp_async_bulk_shared_to_global	UBLKCP
	SMEM → GMEM		Thread-local		
	SMEM → (D)SMEM		(D)SMEM barrier	cuda::ptx::cp_async_bulk*	





# Asynchronous Data Copies

TMA nD – properties

## Tensor copies functionality (not comprehensive list):

- Single thread can fire 1D-5D block copy
- Destination modes:
  - Tile – source layout preserved
  - Im2col – source copy-box elements rearranged into columns
- Out-of-bounds fill modes:
  - Zeros
  - NaNs
- SMEM swizzling

## Constraints (not comprehensive list):

- SMEM address must be 128B aligned
- GMEM address must be 16B aligned
- Strides must be multiples of 16B
- Copy-box fast-dimension size must be multiple of 16B
- Copy-box start address must be 16B aligned

## Tensor copy geometry is described by a **CUtensorMap** descriptor:

- Can be a `__grid_constant__` kernel parameter or `__constant__` object (64B aligned)

# Asynchronous Data Copies



TMA nD – tensor descriptor

GMEM

## Tensor descriptor (constant):

- nD (1, 2, 3, 4, or 5)
- base\_pointer (GMEM)
- tensor\_stride[nD - 1]
- tensor\_size[nD]
- box\_size[nD]
- element\_stride[nD]

Box copied to SMEM/DSMEM:



# Asynchronous Data Copies


TMA nD (**GMEM**->**SMEM**) – tensor descriptor

GMEM

## Instruction args (variable):

- dst\_pointer (SMEM/DSMEM)
  - coords[nD]
- 
- Tensor descriptor
  - Synchronization barrier

Box copied to SMEM/DSMEM:





# Asynchronous Data Copies


TMA nD (**GMEM->SMEM**) – tensor descriptor

GMEM

## Instruction args (variable):

- dst\_pointer (SMEM/DSMEM)
  - coords[nD]
- 
- Tensor descriptor
  - Synchronization barrier

Box copied to SMEM/DSMEM:



# Asynchronous Data Copies

TMA nD (GMEM->SMEM) – tensor descriptor


GMEM

## Tensor descriptor (constant):

- nD
- base\_pointer (GMEM)
- tensor\_stride[nD - 1]
- tensor\_size[nD]
- box\_size[nD]
- element\_stride[nD] := {1, 2}

← must be 1

## Box copied to SMEM/DSMEM:



# Asynchronous Data Copies


TMA nD (**SMEM->GMEM**) – tensor descriptor

GMEM

## Instruction args (variable):

- src\_pointer (**SMEM**/~~DSMEM~~)
  - coords[nD]
- 
- Tensor descriptor
  - ~~Synchronization barrier~~  
-> thread-local sync

Box copied from SMEM:







# Asynchronous Data Copies

TMA nD (**SMEM->GMEM**) – tensor descriptor


## Instruction args (variable):

- src\_pointer (**SMEM**/~~DSMEM~~)
  - coords[nD]
- 
- Tensor descriptor
  - ~~Synchronization barrier~~  
-> thread-local sync

Box copied from SMEM:



base\_pointer  
(coords == {0,0})



# Asynchronous Data Copies

## TMA nD vs LDGSTS

### Example:

✕-shaped 16-th order FD stencil.

Thread block size: 32×8, shared memory size: 48×24 (32×8 center + 8-point halo on each side)

### LDGSTS (pseudocode):

```
int tx = threadIdx.x;
int ty = threadIdx.y;

int ix = blockIdx.x * 32 + tx;
int iy = blockIdx.y * 32 + ty;

if (tx < 8 && ty < 8)
    async_load_0(offset0);
if (tx < 8 && ty < 8 && iy < ny+8)
    async_load_1(offset1);
...
if (ix < nx && iy < ny)
    async_load_4(offset2);
...
commit_async_loads();
```


### TMA (pseudocode):

```
if (tid == 0)
    async_load_tma(...);
```

Single-thread, unconditional  
(wrt. domain bounds) load

Multiple (up to 9) conditional loads  
by each thread with different offset  
indices:

- MIO pressure
- Register pressure
- Predicate pressure



# Asynchronous Data Copies

## TMA nD vs LDGSTS

### Example:

✕-shaped 16-th order FD stencil.

Thread block size:  $32 \times 8$ , shared memory size:  $48 \times 24$  ( $32 \times 8$  center + 8-point halo on each side)

### LDGSTS (pseudocode):

```
int tx = threadIdx.x;
int ty = threadIdx.y;

int ix = blockIdx.x * 32 + tx;
int iy = blockIdx.y * 32 + ty;

if (tx < 8 && ty < 8)
    async_load_0(offset0);

if (tx < 8 && ty < 8 && iy < ny+8)
    async_load_1(offset1);
...

if (ix < nx && iy < ny)
    async_load_4(offset2);
...

commit_async_loads();
```

### TMA (pseudocode):


```
if (tid == 0)
    async_load_tma(...);
```

Single-thread, unconditional  
(wrt. to domain bounds) load

Multiple (up to 9) conditional loads  
by each thread with different offset  
indices:

- MIO pressure
- Register pressure
- Predicate pressure

*Even worse for higher order stencil.  
E.g., 12-point halos: regions #3 and #5 are not  
covered by TB and require 2 loads per region!*





# Asynchronous Data Copies


## TMA details

TMA programming model is Uniform:

all active threads in a warp are executing TMA operation “with same arguments”.

### Code:

```
async_load_tma(args_per_thread);
```



```
Loop over active thread i in {0,...,31}  
{  
    Issue_TMA_Instruction(args_of_thread #i);  
}
```

# Asynchronous Data Copies


## TMA details

TMA programming model is Uniform:

all active threads in a warp are executing TMA operation “with same arguments”.

### Code:

```
async_load_tma(args_per_thread);
```



```
Loop over active thread i in {0,...,31}
{
    Issue_TMA_Instruction(args_of_thread #i);
}
```

### Erroneous code (race):

```
async_load_tma(single_thread_args);
```

### Not optimal code:

```
if (tid == 0)
    async_load_tma(single_thread_args);
```

Compiler doesn't know that `tid == 0` for only one thread!

Compiler still creates a peeling loop over the active threads:

```
Loop over active thread i in {0,...,31}
{
    if (tid(i) == 0) // uniform condition
        Issue_TMA_Instruction(single_thread_args);
}
```

# Asynchronous Data Copies


## TMA details

TMA programming model is Uniform:

all active threads in a warp are executing TMA operation “with same arguments”.

### Code:

```
async_load_tma(args_per_thread);
```



```
Loop over active thread i in {0,...,31}
{
    Issue_TMA_Instruction(args_of_thread #i);
}
```

### Erroneous code (race):

```
async_load_tma(single_thread_args);
```

### Not optimal code:

```
if (tid == 0)
    async_load_tma(single_thread_args);
```

Compiler doesn't know that `tid == 0` for only one thread!

Compiler still creates a peeling loop over the active threads:

```
Loop over active thread i in {0,...,31}
{
    if (tid(i) == 0) // uniform condition
        Issue_TMA_Instruction(single_thread_args);
}
```

### Better code:

```
if (tid == 0)
    cg::invoke_one(cg::coalesced_threads(),
        async_load_tma(single_thread_args));
```



# Asynchronous Data Copies

## TMA nD main steps summary

### GMEM->SMEM:

```
INIT { if (threadIdx.x == 0) {
      __mbarrier_init(&bar, 1);
      cuda::device::experimental::fence_proxy_async_shared_cta();
    }
    __syncthreads(); // cluster.sync() if GMEM->DSMEM

FIRE { if (threadIdx.x == 0) {
      invoke_one(coalesced_threads(), cp_async_bulk_tensor_2d_global_to_shared, smem, &map,
                coord[0], coord[1], bar); // mapped_remote_bar if GMEM->DSMEM
      mbarrier_arrive_expect_tx(sem_release, scope_cta, space_shared, &bar, SMEM_SIZE);
    }

WAIT FOR COMPLETION { while (!mbarrier_try_wait_parity(&bar, phase, 1000))
                     {}
```

# Asynchronous Data Copies

TMA nD main steps summary

## SMEM->GMEM:

```
__syncthreads(); // After all threads use SMEM

if (threadIdx.x == 0) {
    FIRE {
        invoke_one(coalesced_threads(),
            [] () { cp_async_bulk_tensor_2d_shared_to_global(&map, coordA[0], coordA[1], smemA);
                    cp_async_bulk_tensor_2d_shared_to_global(&map, coordB[0], coordB[1], smemB); }
        );
    COMMIT GROUP { cp_async_bulk_commit_group();
        FIRE {
            invoke_one(coalesced_threads(), cp_async_bulk_tensor_2d_shared_to_global,
                &map, coordD[0], coordD[1], smemD);
        COMMIT GROUP { cp_async_bulk_commit_group();
            cp_async_bulk_wait_group_read<N>(); } WAIT FOR COMPLETION EXCEPT LAST N GROUPS
        }
    }


__syncthreads(); // For all threads to use SMEM again
```

# Asynchronous Data Copies

## SMEM bank conflicts

Given 2D `array[][NX]`, index `[y][x]` maps to linear index `ind` via:

- **Conventional:**  $\text{ind} := y * NX + x$






# Asynchronous Data Copies

SMEM bank conflicts – padding

Given 2D array  $array[][\overset{NX+1}{NX}]$ , index  $[y][x]$  maps to linear index  $ind$  via:

- **Conventional:**  $ind := y * \overset{(NX+1)}{NX} + x$



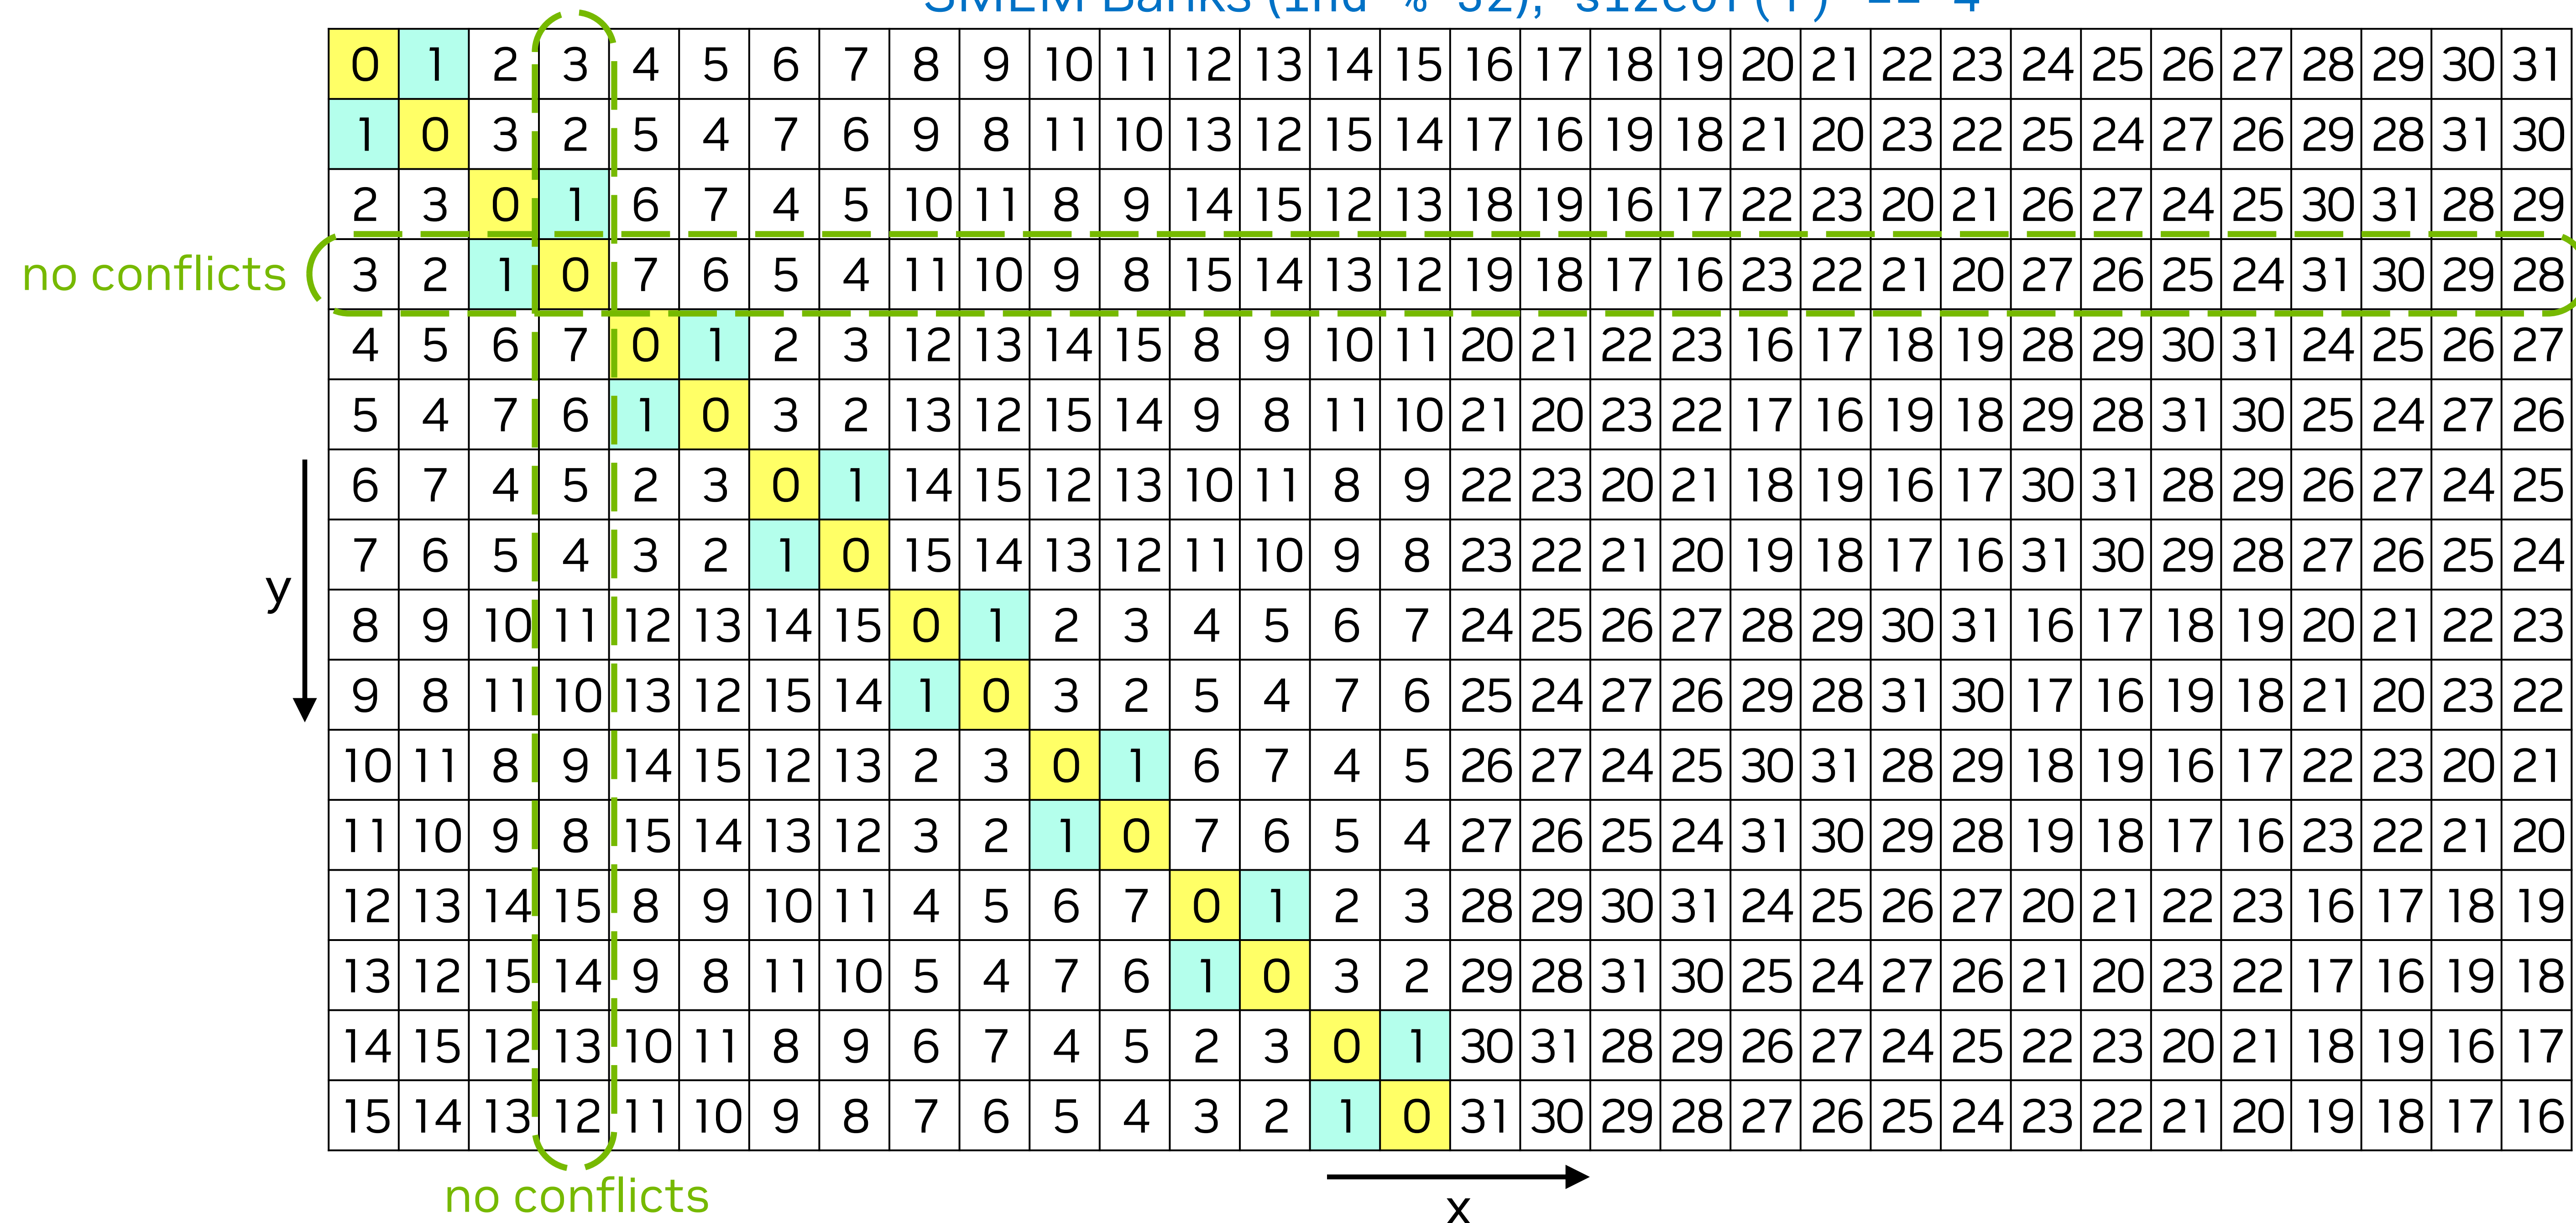
# Asynchronous Data Copies

SMEM bank conflicts – swizzling

Given 2D `array[][NX]`, index `[y][x]` maps to linear index `ind` via:

- **Conventional:** `ind := y * NX + x`
- **Swizzling:** `ind := y * NX + (y ^ x)`

SMEM Banks (`ind % 32`), `sizeof(T) == 4`





# Asynchronous Data Copies

## TMA nD – SMEM swizzling

TMA swizzles eight 16-byte chunks within 128-byte segments.

### Constraints:

- $NX * \text{sizeof}(T) == \text{SWIZZLE\_SIZE}$  (where:  $T$  array[][NX])
- Allowed values of SWIZZLE\_SIZE: 32, 64, 128

Given  $[y][x]$  index in  $T$  array[][NX]:

1. Compute index of 16-byte chunk within 128-byte segment:

$$i16 := (y * NX + x) * \text{sizeof}(T) / 16$$

$$y16 := i16 / 8$$

$$x16 := i16 \% 8$$

2. Compute swizzled index of 16-byte chunk:

$$x16\_swz := y16 \wedge x16$$

3. Compute swizzled index:

$$x\_swz := x16\_swz * 16 / \text{sizeof}(T) \% NX + x \% (16 / \text{sizeof}(T))$$

Swizzled indices of 16-byte chunks											
32-byte		64-byte				128-byte					
0	1	0	1	2	3	0	1	2	3	4	5
0	1	0	1	2	3	1	0	3	2	5	4
0	1	1	0	3	2	2	3	0	1	6	7
0	1	1	0	3	2	3	2	1	0	7	6
1	0	2	3	0	1	4	5	6	7	0	1
1	0	2	3	0	1	5	4	7	6	1	0
1	0	3	2	1	0	6	7	4	5	2	3
1	0	3	2	1	0	7	6	5	4	3	2



# Asynchronous Data Copies

Segmented sort performance

4194304 sequences of 128 integers (~ 2GB)

Pair-wise sorting network algorithm

H200 experiment

SMEM copies	Time (ms)	BW (GB/s)	% SOL
LDG+STS	1.88	2274	50
LDGSTS	1.66	2600	56
UBLKCP	1.18	3632	79
TMA + swizzle	1.07	4066	87

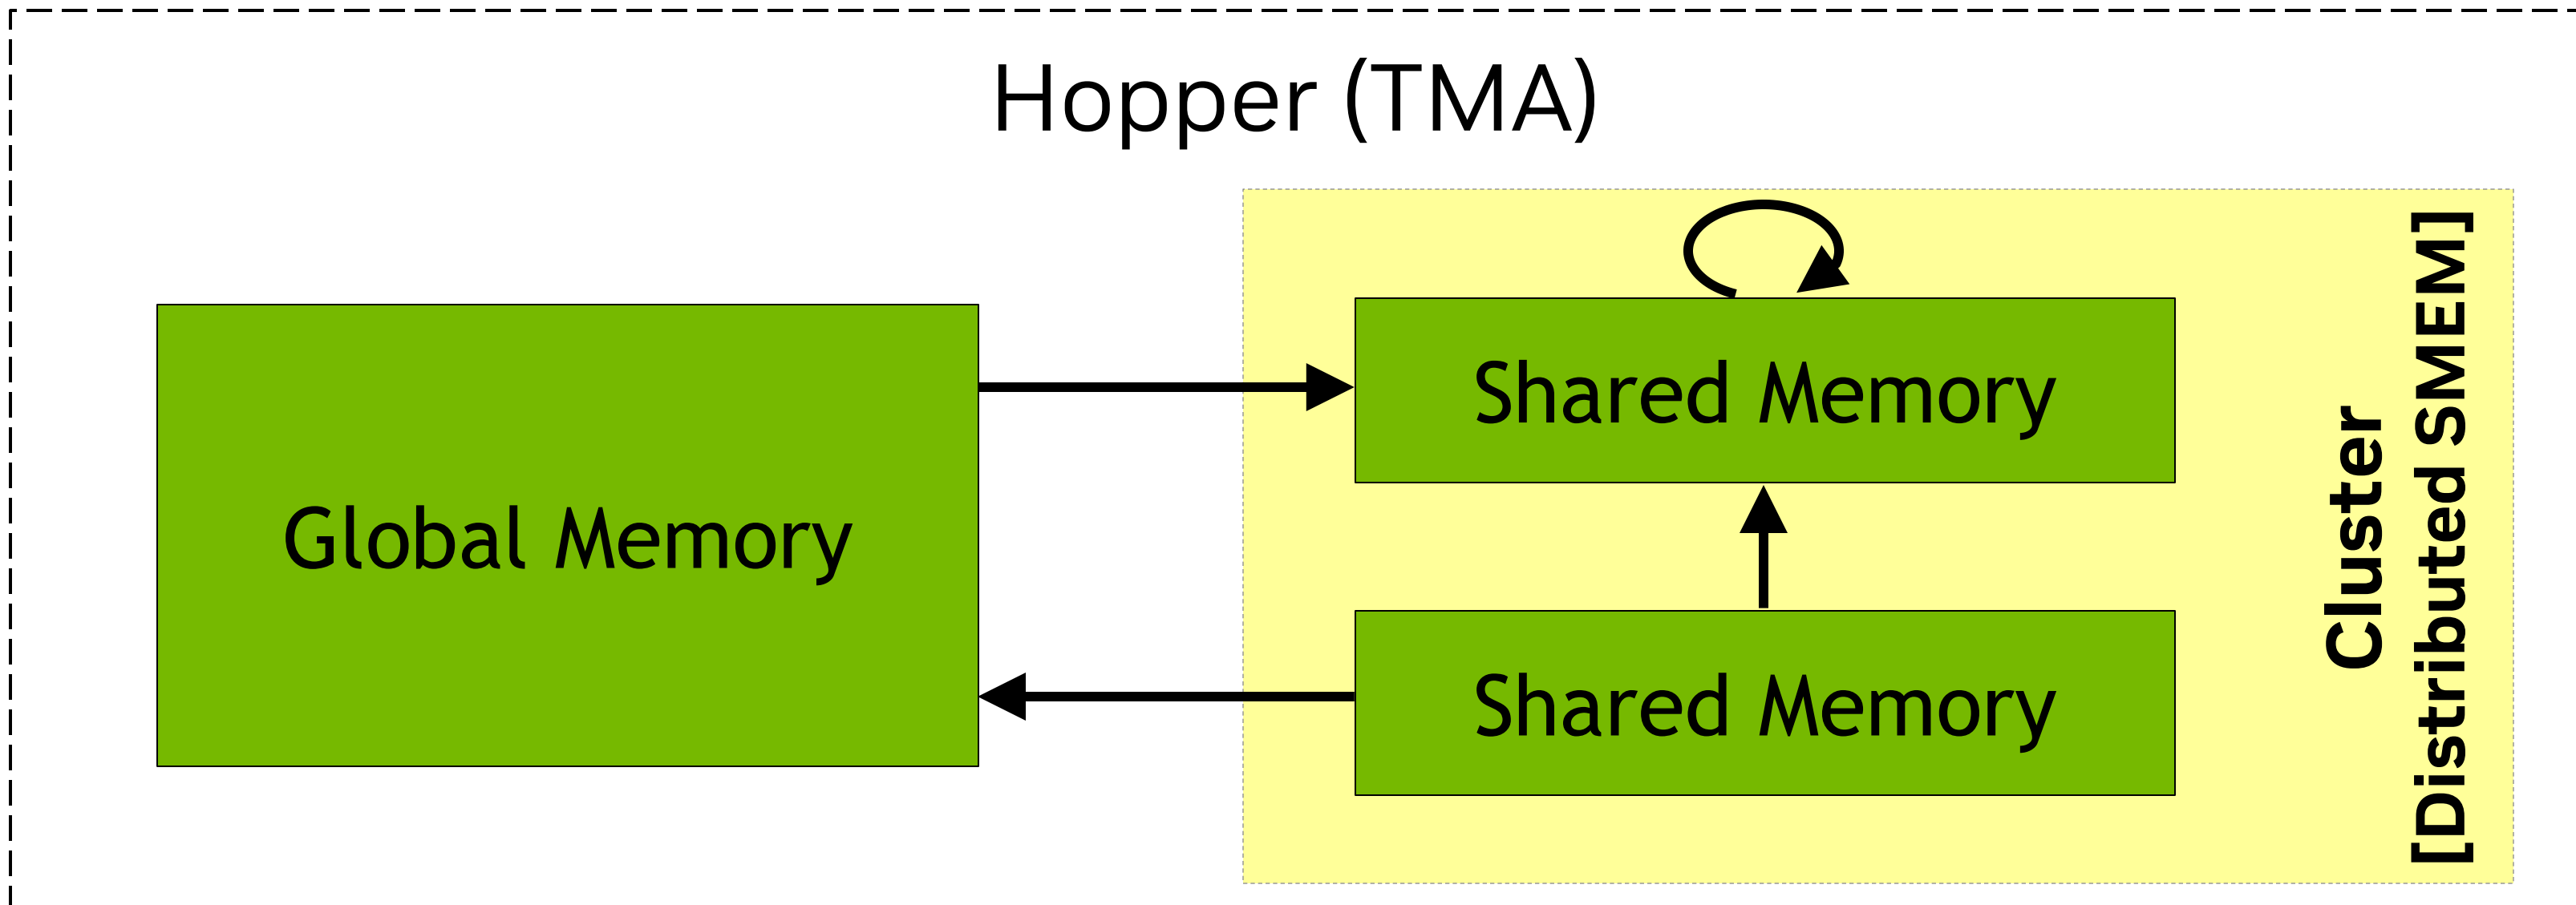
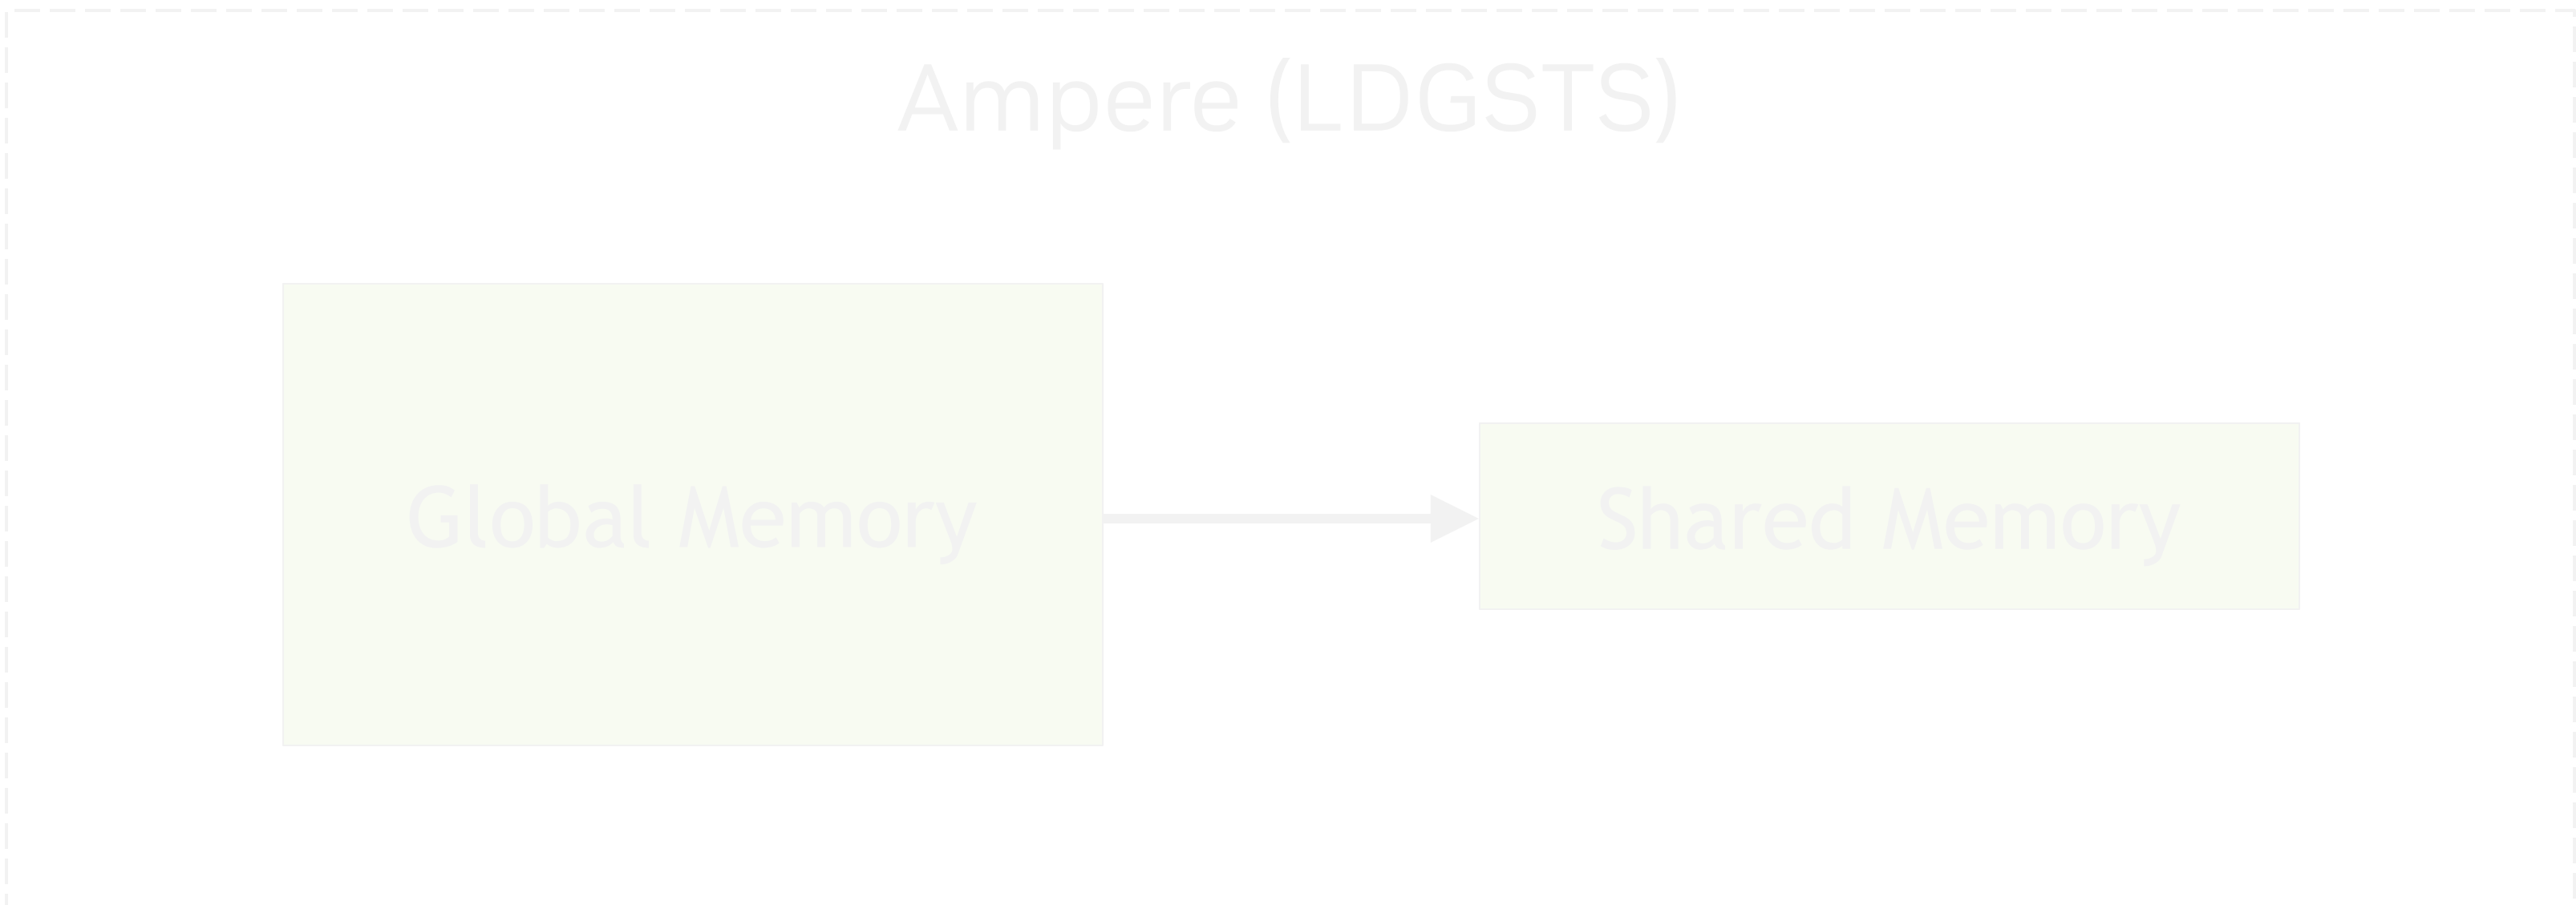
[Test provided by Allard Hendriksen]

# Asynchronous Data Copies

Asynchronous copies – UBLKCP

HW	Direction	Operation	Synchronization	C++	SASS
Hopper	GMEM → SMEM	Asynchronous version of smem[sind] := gmem[gind]	Thread-local ----- SMEM barrier	__pipeline_memcpy_async	LDGSTS
	REG → DSMEM	Asynchronous version of dsmem[sind] := var	DSMEM barrier	cuda::ptx::st_async	STAS
	GMEM → (D)SMEM	1D – 5D block copy Uniform	(D)SMEM barrier	cuda::device::experimental::	UTMALDG
	SMEM → GMEM		Thread-local	cp_async_bulk_tensor_Nd_global_to_shared cp_async_bulk_tensor_Nd_shared_to_global	UTMASTG
	GMEM → SMEM	1D block copy Uniform	SMEM barrier	cuda::device::experimental:: cp_async_bulk_global_to_shared cp_async_bulk_shared_to_global	UBLKCP
	SMEM → GMEM		Thread-local		
	SMEM → (D)SMEM		(D)SMEM barrier	cuda::ptx::cp_async_bulk*	

\* Available in latest github CCCL



# Asynchronous Data Copies

TMA – properties

## UBLKCP ~ async version of memcpy:

- Fired by a single thread
- Uses src/dst pointers and copy size  
(no Tensor descriptors => no bound-box, stride, fill, swizzle, etc. functionality)
- Alignment requirements:
  - 16B alignment of GMEM/SMEM
  - Copy size multiple of 16B
- Same synchronization (barrier or thread-local) as for [TMA nD](#):
  - `mbarrier_arrive_expect_tx(sem_release, scope_cta, space_shared, &bar, SIZE);`
  - `cp_async_bulk_commit_group();`  
`cp_async_bulk_wait_group_read<N>();`





# **Compressible Memory**

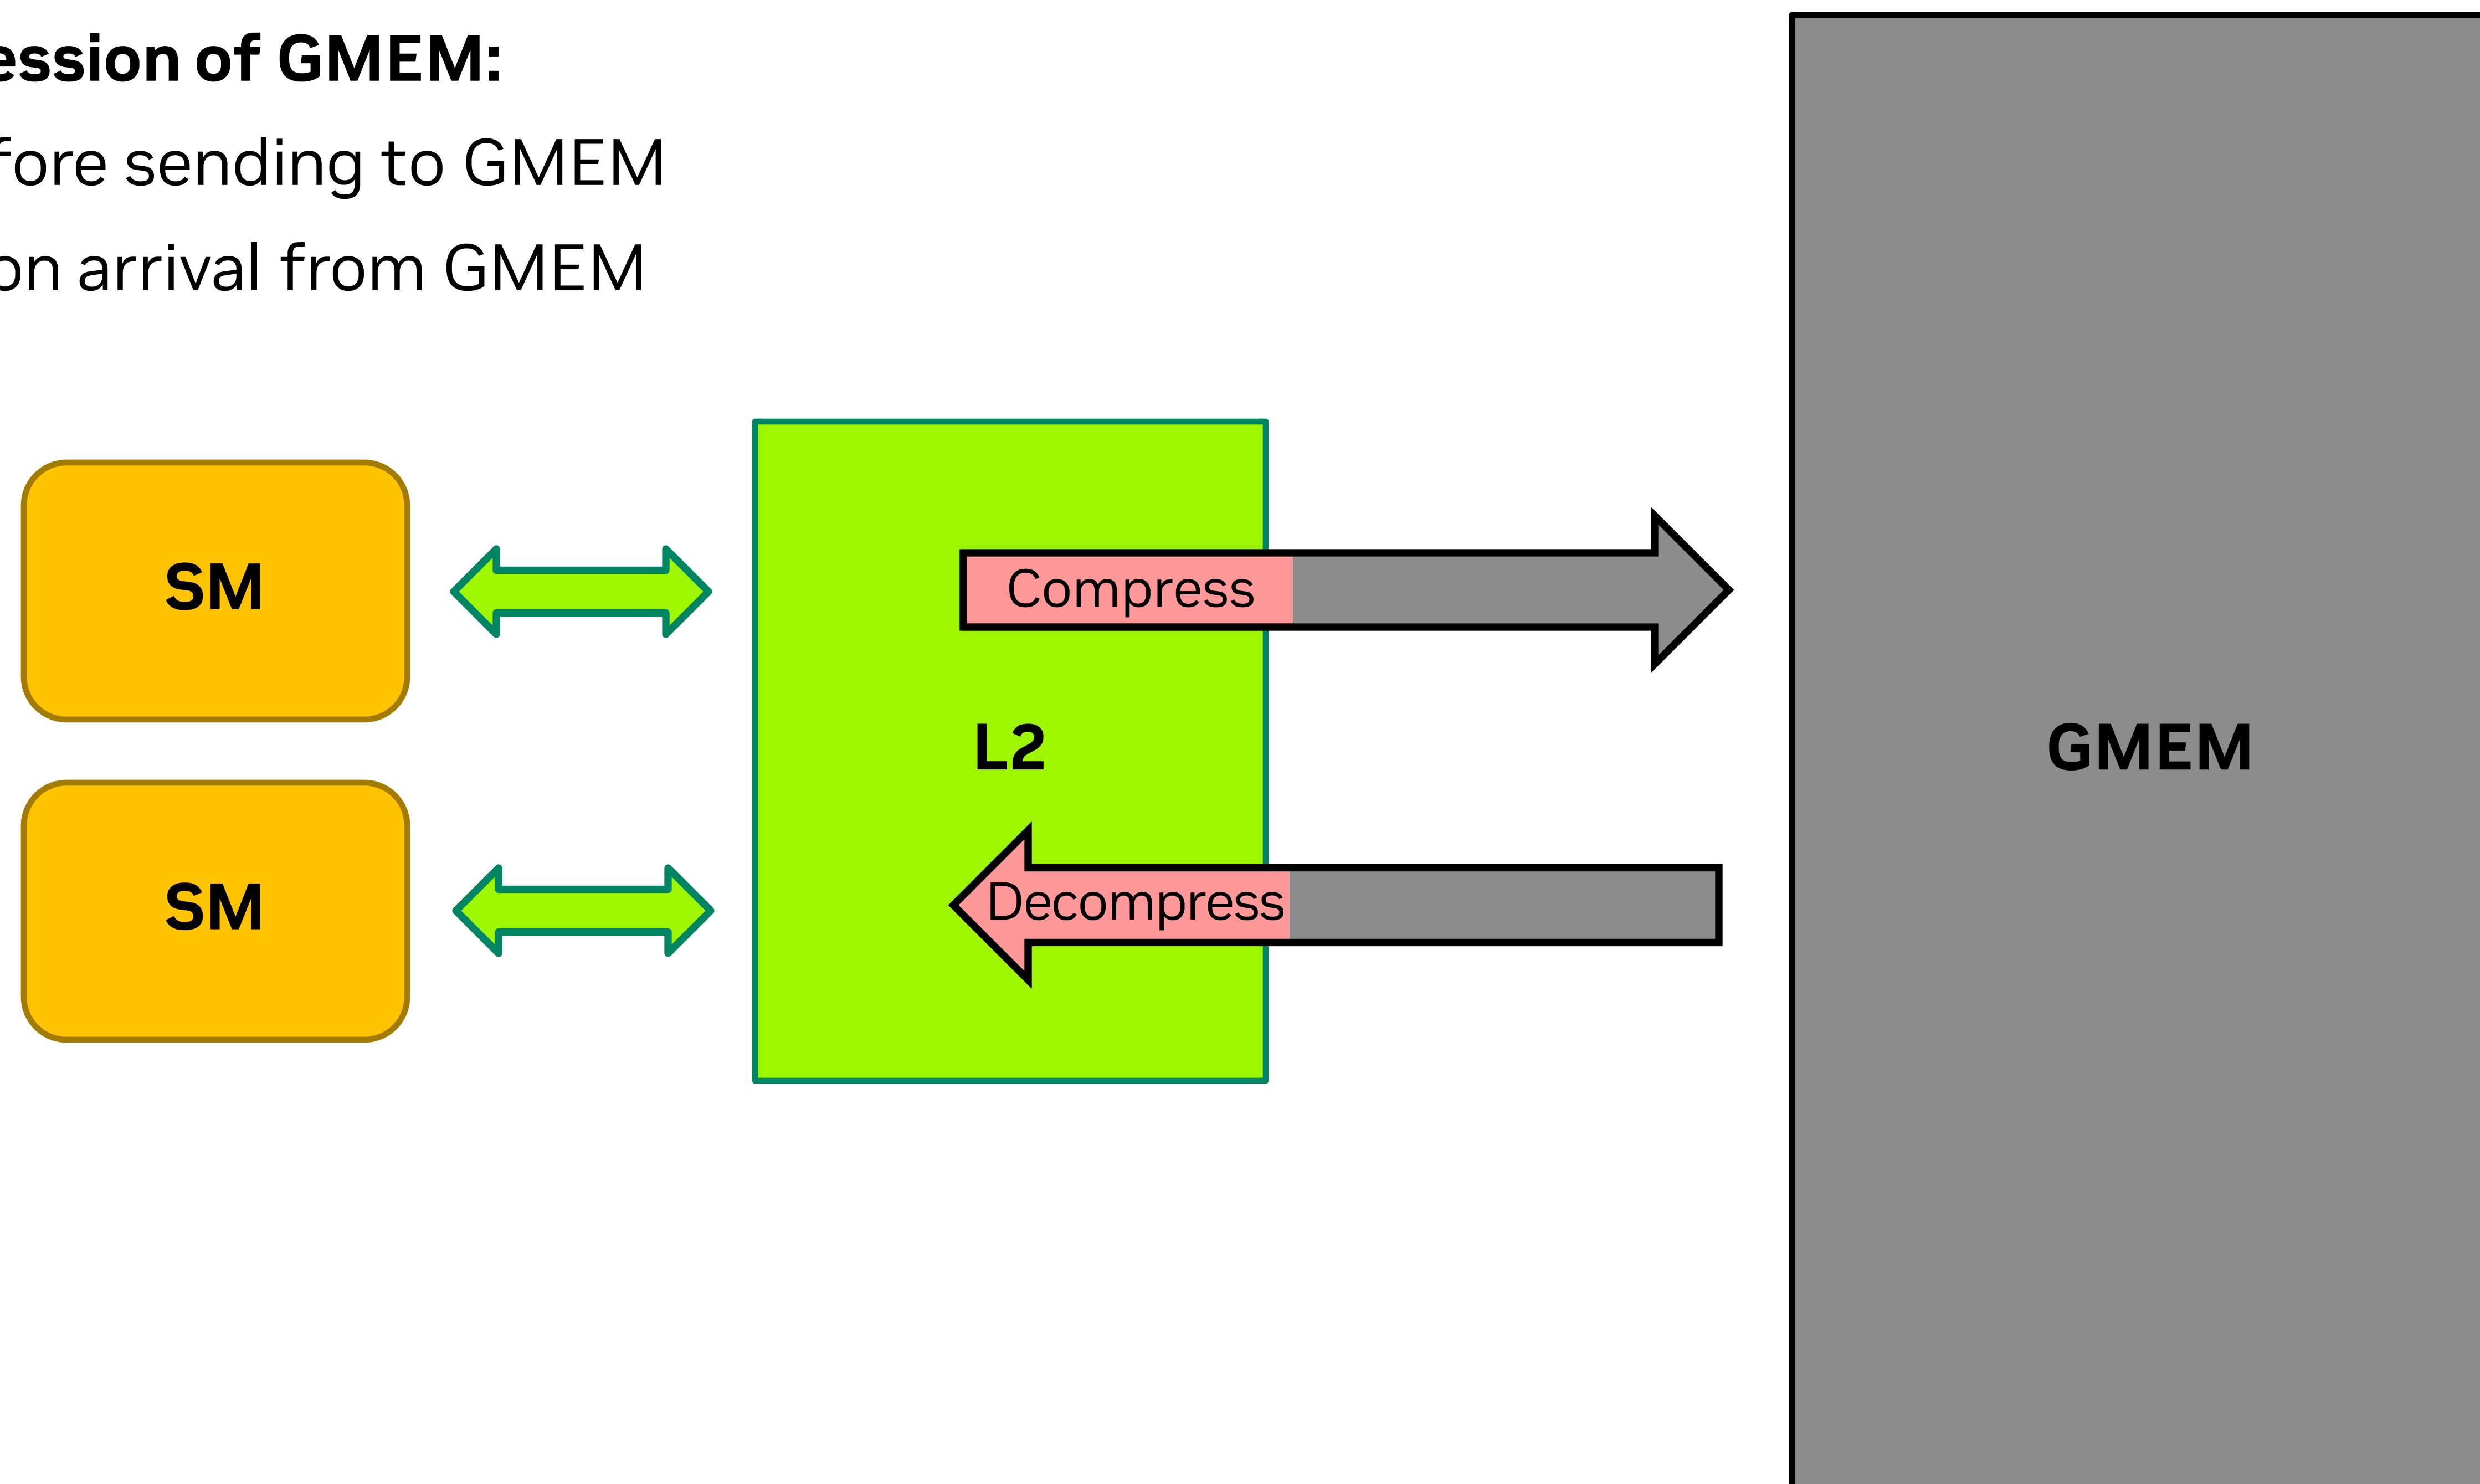


# Compressible Memory

## Introduction

### Hardware compression of GMEM:

- Compresses before sending to GMEM
- Decompresses on arrival from GMEM



# Compressible Memory

## Quick facts

### Compressible memory:

- **Fully transparent** (only change your GMEM allocation calls)
  - “Compresses bandwidth” (GMEM  $\leftrightarrow$  L2)
  - **Doesn't compress storage** (GMEM or L2)
- 128B (cache line) granularity
- **Lossless** (2x, 4x, or 1x if cannot compress)
- Compression is automatic (no user control), not all data will compress
- **Available** on Ada and Hopper
- **May decrease performance** (compression cache misses => large penalties)

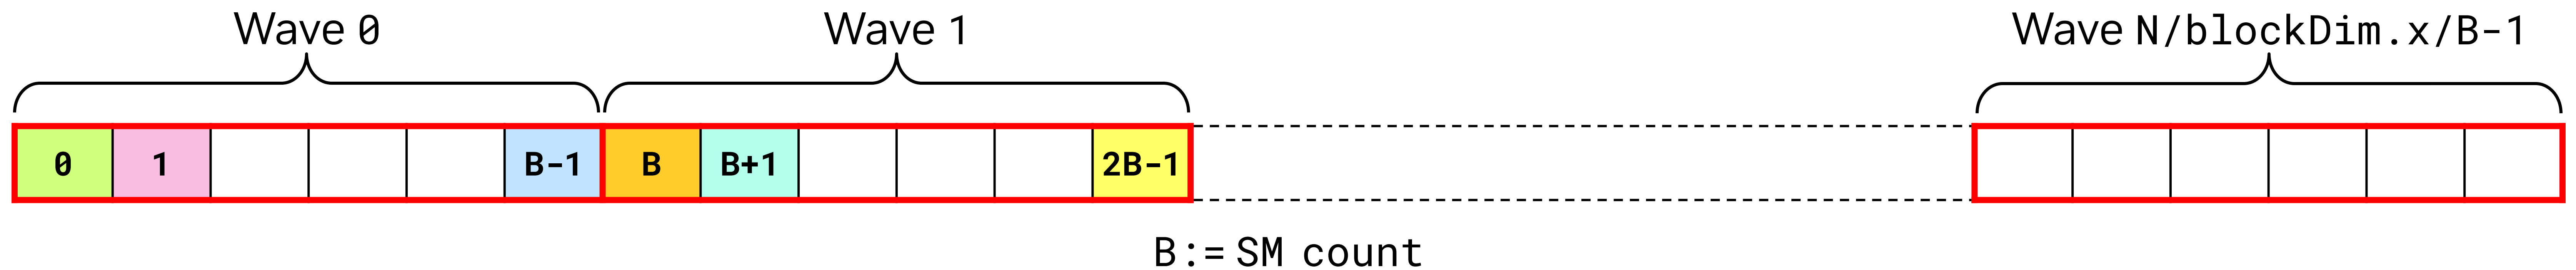


# Compressible Memory

Saxpy example – kernels

“No-loop” kernel:

```
__global__ void saxpy_no_loop(float4 * __restrict__ z,
                             float4 const* __restrict__ x,
                             float4 const* __restrict__ y,
                             float a, int n)
{
    int const i = blockIdx.x * blockDim.x + threadIdx.x;
    if (i < n)
        z[i] = a * x[i] + y[i];
}
```



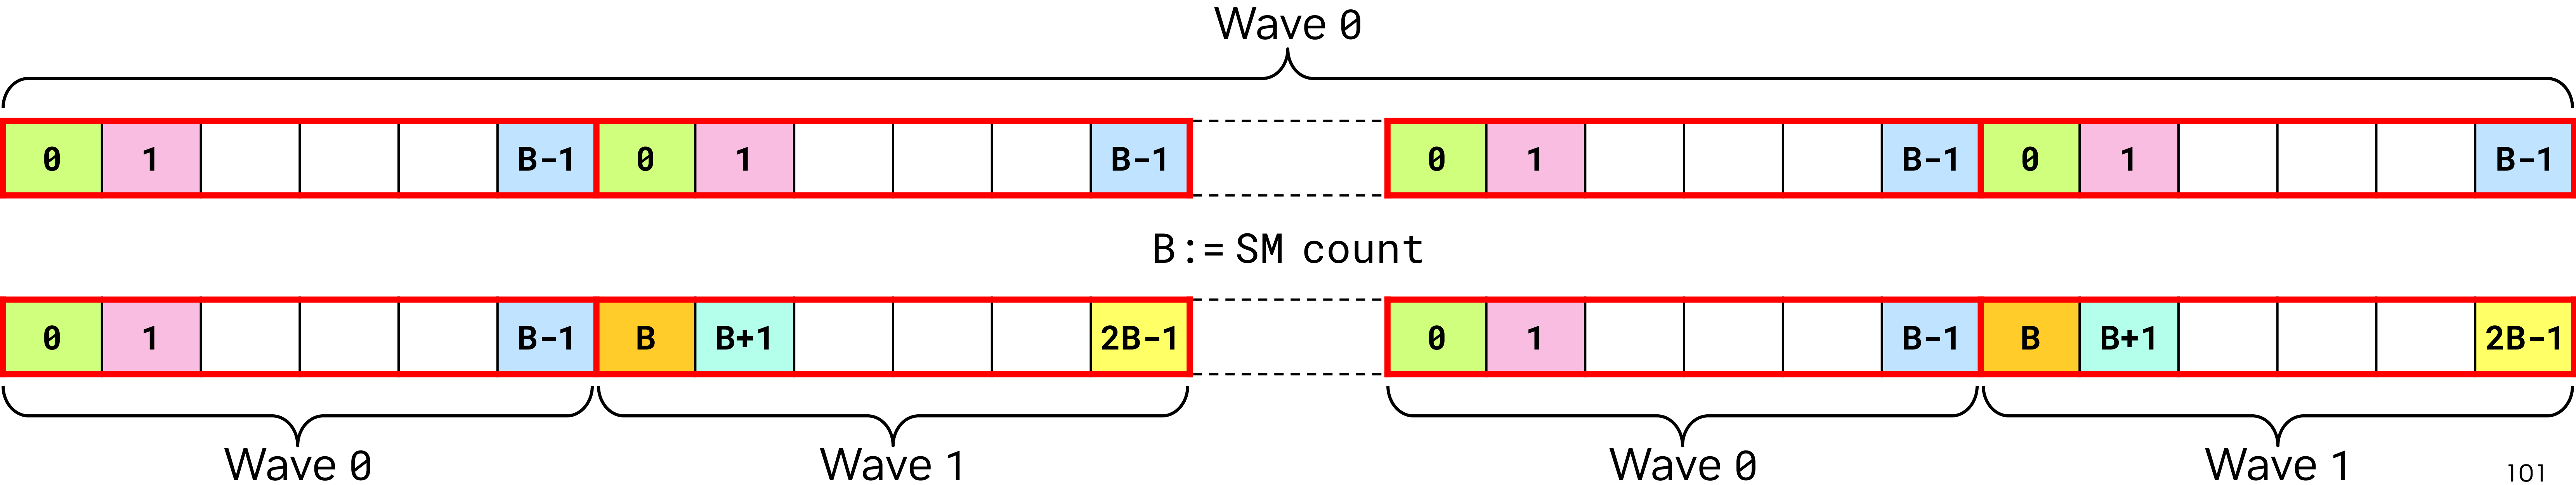
# Compressible Memory

Saxpy example – kernels

“Grid-stride loop” kernel:

```
__global__ void saxpy_grid (float4      * __restrict__ z,  
                           float4 const* __restrict__ x,  
                           float4 const* __restrict__ y,  
                           float a, int n)  
{  
    int const i0 = blockIdx.x * blockDim.x + threadIdx.x;  
    int const step = gridDim.x * blockDim.x;  
    #pragma unroll(4)  
    for (int i = i0; i < n; i += step)  
        z[i] = a * x[i] + y[i];  
}
```

Access locality worsens  
with more waves



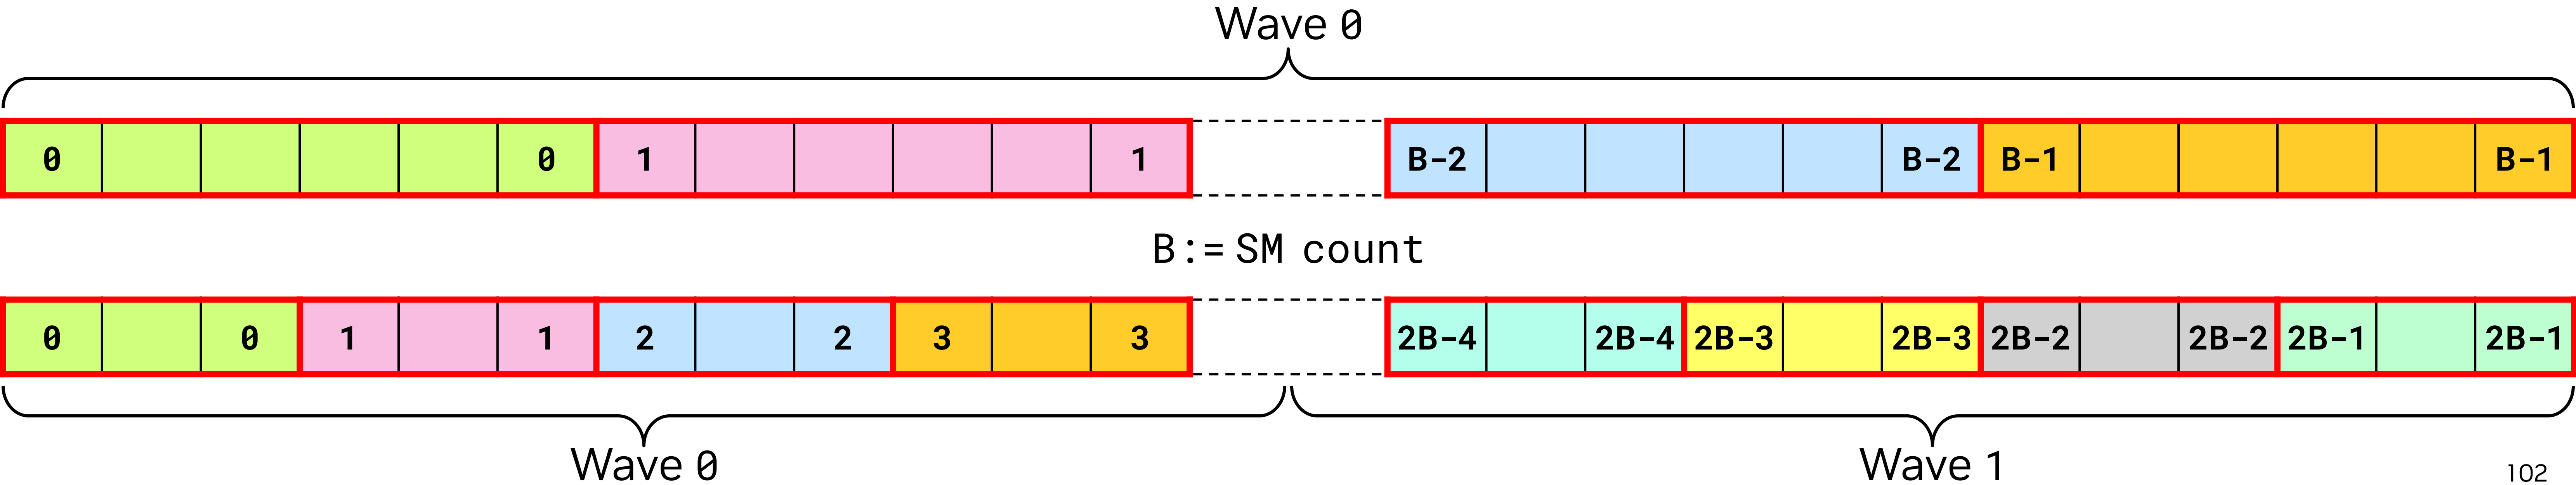
# Compressible Memory

Saxpy example – kernels

“Block-stride loop” kernel:

```
__global__ void saxpy_block (float4      * __restrict__ z,
                             float4 const* __restrict__ x,
                             float4 const* __restrict__ y,
                             float a, int n)
{
    int const i0 = n / gridDim.x * blockIdx.x + threadIdx.x;
    int const i1 = n / gridDim.x * (blockIdx.x + 1) + threadIdx.x;
    #pragma unroll(4)
    for (int i = i0; i < i1; i += blockDim.x)
        z[i] = a * x[i] + y[i];
}
```

Access locality improves  
with more waves





# Compressible Memory

Saxpy example – performance

Performance for  $x[i] = y[i] = \{i, i, i, i\}$  input data

H100 (1.3GB / array)

Kernel	Uncompressed		Compressed	
	Time (ms)	B-width (TB/s)	Time (ms)	B-width (TB/s)
Grid-loop SMx1	1.10	3.53	0.86	4.54
Grid-loop SMx2	1.11	3.50	1.17	3.31
Grid-loop SMx20	1.09	3.59	1.72	2.26
Block-loop SMx1	1.11	3.52	3.00	1.30
Block-loop SMx2	1.10	3.53	2.72	1.43
Block-loop SMx20	1.08	3.59	0.77	5.08
No-loop	1.05	3.71	0.67	5.77

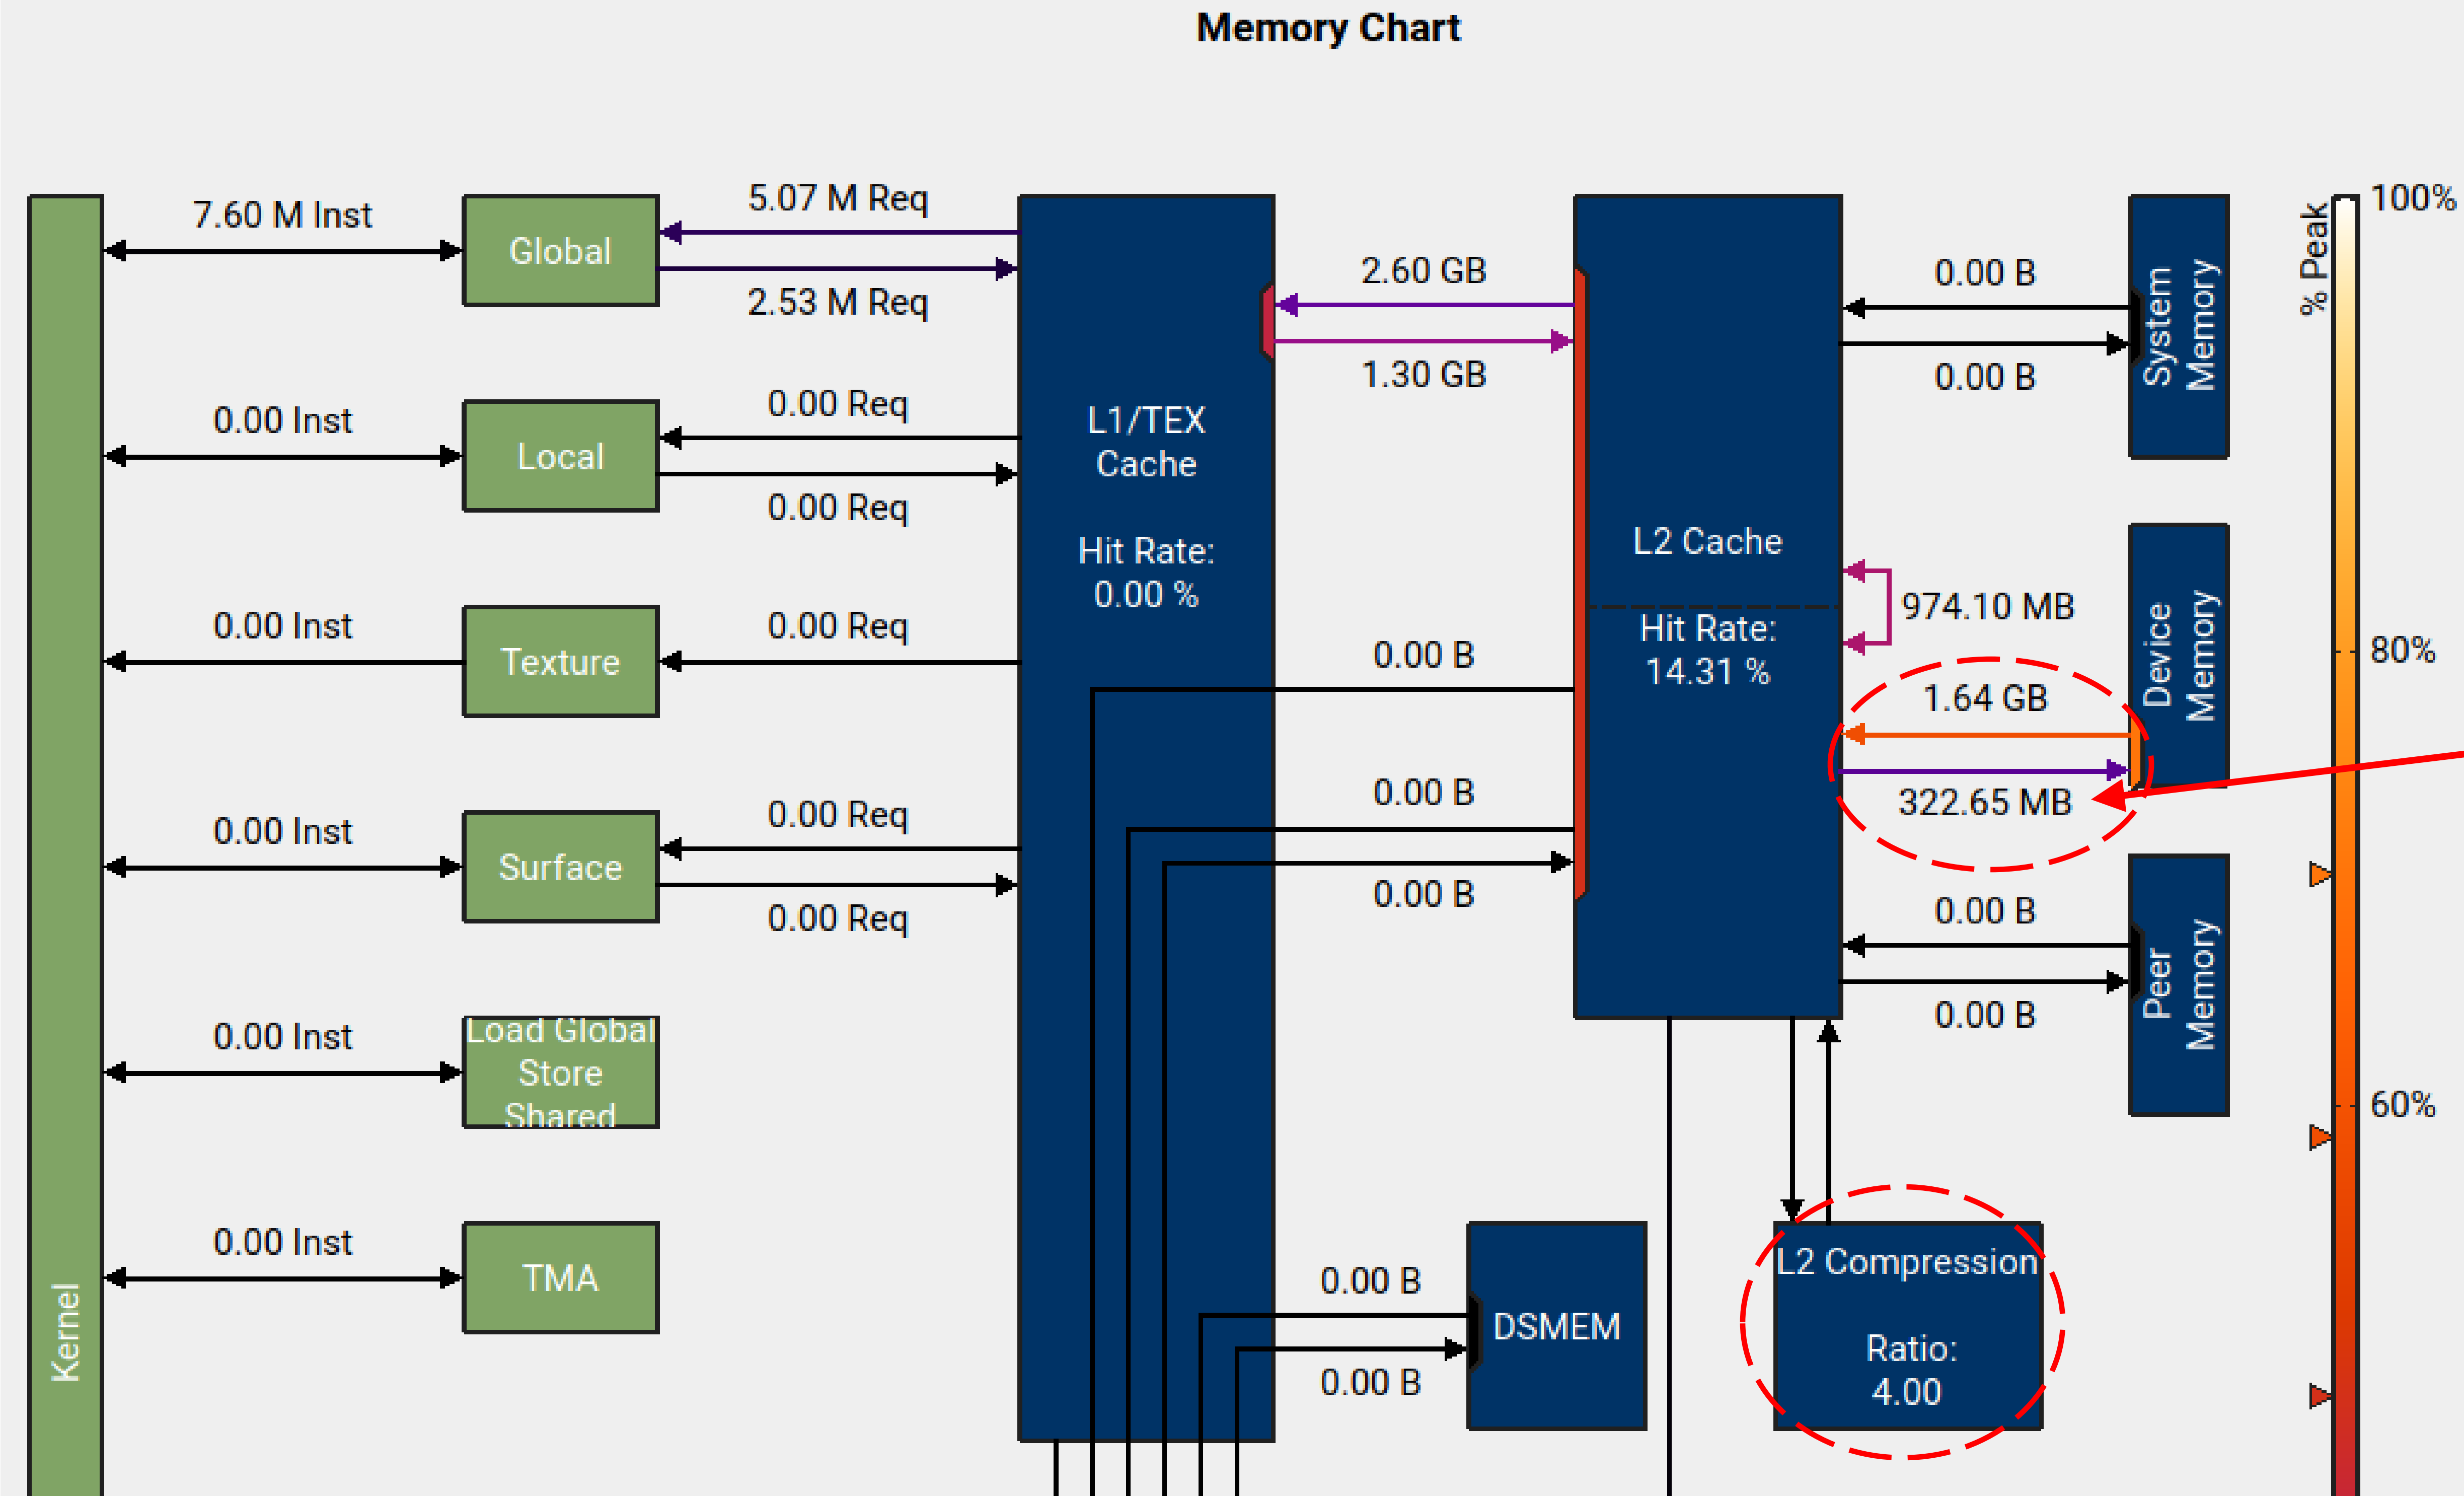
H100 (13GB / array)

Kernel	Uncompressed		Compressed	
	Time (ms)	B-width (TB/s)	Time (ms)	B-width (TB/s)
Grid-loop SMx1	11.0	3.53	26.6	1.46
Grid-loop SMx2	11.1	3.50	39.9	0.98
Grid-loop SMx20	10.9	3.57	44.7	0.87
Block-loop SMx1	10.9	3.57	44.2	0.88
Block-loop SMx2	10.9	3.57	46.5	0.84
Block-loop SMx20	10.9	3.58	30.4	1.28
No-loop	10.4	3.73	6.73	5.78

Locality is crucial!

# Compressible Memory

NCU



Array: 1.3GB / 4 => 322.65 MB

# Compressible Memory


Seismic example

## RTM (3D finite-difference wave propagation):


```
Dim3 block = {BX, BY, 1};  
Dim3 grid = {nx / BX, ny / BY, 1};  
kernel<<<block, grid>>>(...);
```

Kernel z-loop (z is the slowest axis):

```
int zbeg = 0;  
int zend = nz;  
for (int iz = zbeg; iz < zend; ++iz)  
{  
    ...  
}
```



Z-loop





# Compressible Memory



Seismic example – improving locality

## RTM (3D finite-difference wave propagation):

```
int BZ = 64;  
Dim3 block = {BX, BY, 1};  
Dim3 grid = {nx / BX, ny / BY, nz / BZ};  
kernel<<<block, grid>>>(...);
```


Kernel z-loop (z is the slowest axis):

```
int zbeg = blockIdx.z * BZ;  
int zend = min(nz, zbeg + BZ);  
for (int iz = zbeg; iz < zend; ++iz)  
{  
    ...  
}
```




# Compressible Memory

Seismic example – improving locality



Locality  $\approx$  (z-distance between wave slices) \* z-stride



Locality  $\approx$  (z-distance between wave slices) \* z-stride

# Compressible Memory

Seismic example – L40s

**Problem:**

- 8<sup>th</sup> order TTI RTM, single pass kernel
- Domain: 800×800×800 (9 volumes @ 2.2 GB / volume)
- Thread block size: 32×16

**L40s MCells/sec**

	No z-blocking	128 z-blocking	64 z-blocking	48 z-blocking	32 z-blocking
No compression	13888	14145	14574	14610	14661
Compression	1594	2081	4220	5002	51497



# Compressible Memory

Seismic example – L40s

**Problem:**


- 8<sup>th</sup> order TTI RTM, single pass kernel
- Domain: 800×800×800 (9 volumes @ 2.2 GB / volume)
- Thread block size: 32×16

L40s MCells/sec

	No z-blocking	128 z-blocking	64 z-blocking	48 z-blocking	32 z-blocking
No compression	13888	14145	14574	14610	14661
Compression	1594	2081	4220	5002	51497
					43004
					31640
					23422
					21440

– independent of # time-steps  
100 time-steps  
1000 time-steps  
2000 time-steps  
5000 time-steps  
10000 time-steps

Starts with zero wavefields, which become less compressible as they propagate:



# Compressible Memory

Allocating compressible memory

## Via Virtual Memory Management API (CUDA driver):

```
CUmemAllocationProp prop = {};  
prop.type = CU_MEM_ALLOCATION_TYPE_PINNED;  
prop.location.type = CU_MEM_LOCATION_TYPE_DEVICE;  
prop.location.id = device;  
prop.allocFlags.compressionType = CU_MEM_ALLOCATION_COMP_GENERIC;
```

Round-up allocation size  
to granularity

```
size_t granularity = 0;  
cuMemGetAllocationGranularity(&granularity, &prop,  
                             CU_MEM_ALLOC_GRANULARITY_MINIMUM);  
size = ((size + granularity - 1) / granularity) * granularity;
```

Reserve a virtual address range

Allocate physical memory

Map allocated memory to VA range

```
CUdeviceptr ptr;  
CUmemGenericAllocationHandle handle;  
cuMemAddressReserve(&ptr, size, 0, 0, 0);  
cuMemCreate(&handle, size, &prop, 0);  
cuMemMap(ptr, size, 0, handle, 0);
```

Make accessible

```
CUmemAccessDesc desc = {};  
desc.location.type = CU_MEM_LOCATION_TYPE_DEVICE;  
desc.location.id = device;  
desc.flags = CU_MEM_ACCESS_FLAGS_PROT_READWRITE;  
cuMemSetAccess(ptr, size, &desc, 1);
```

```
cuMemRelease(handle);
```





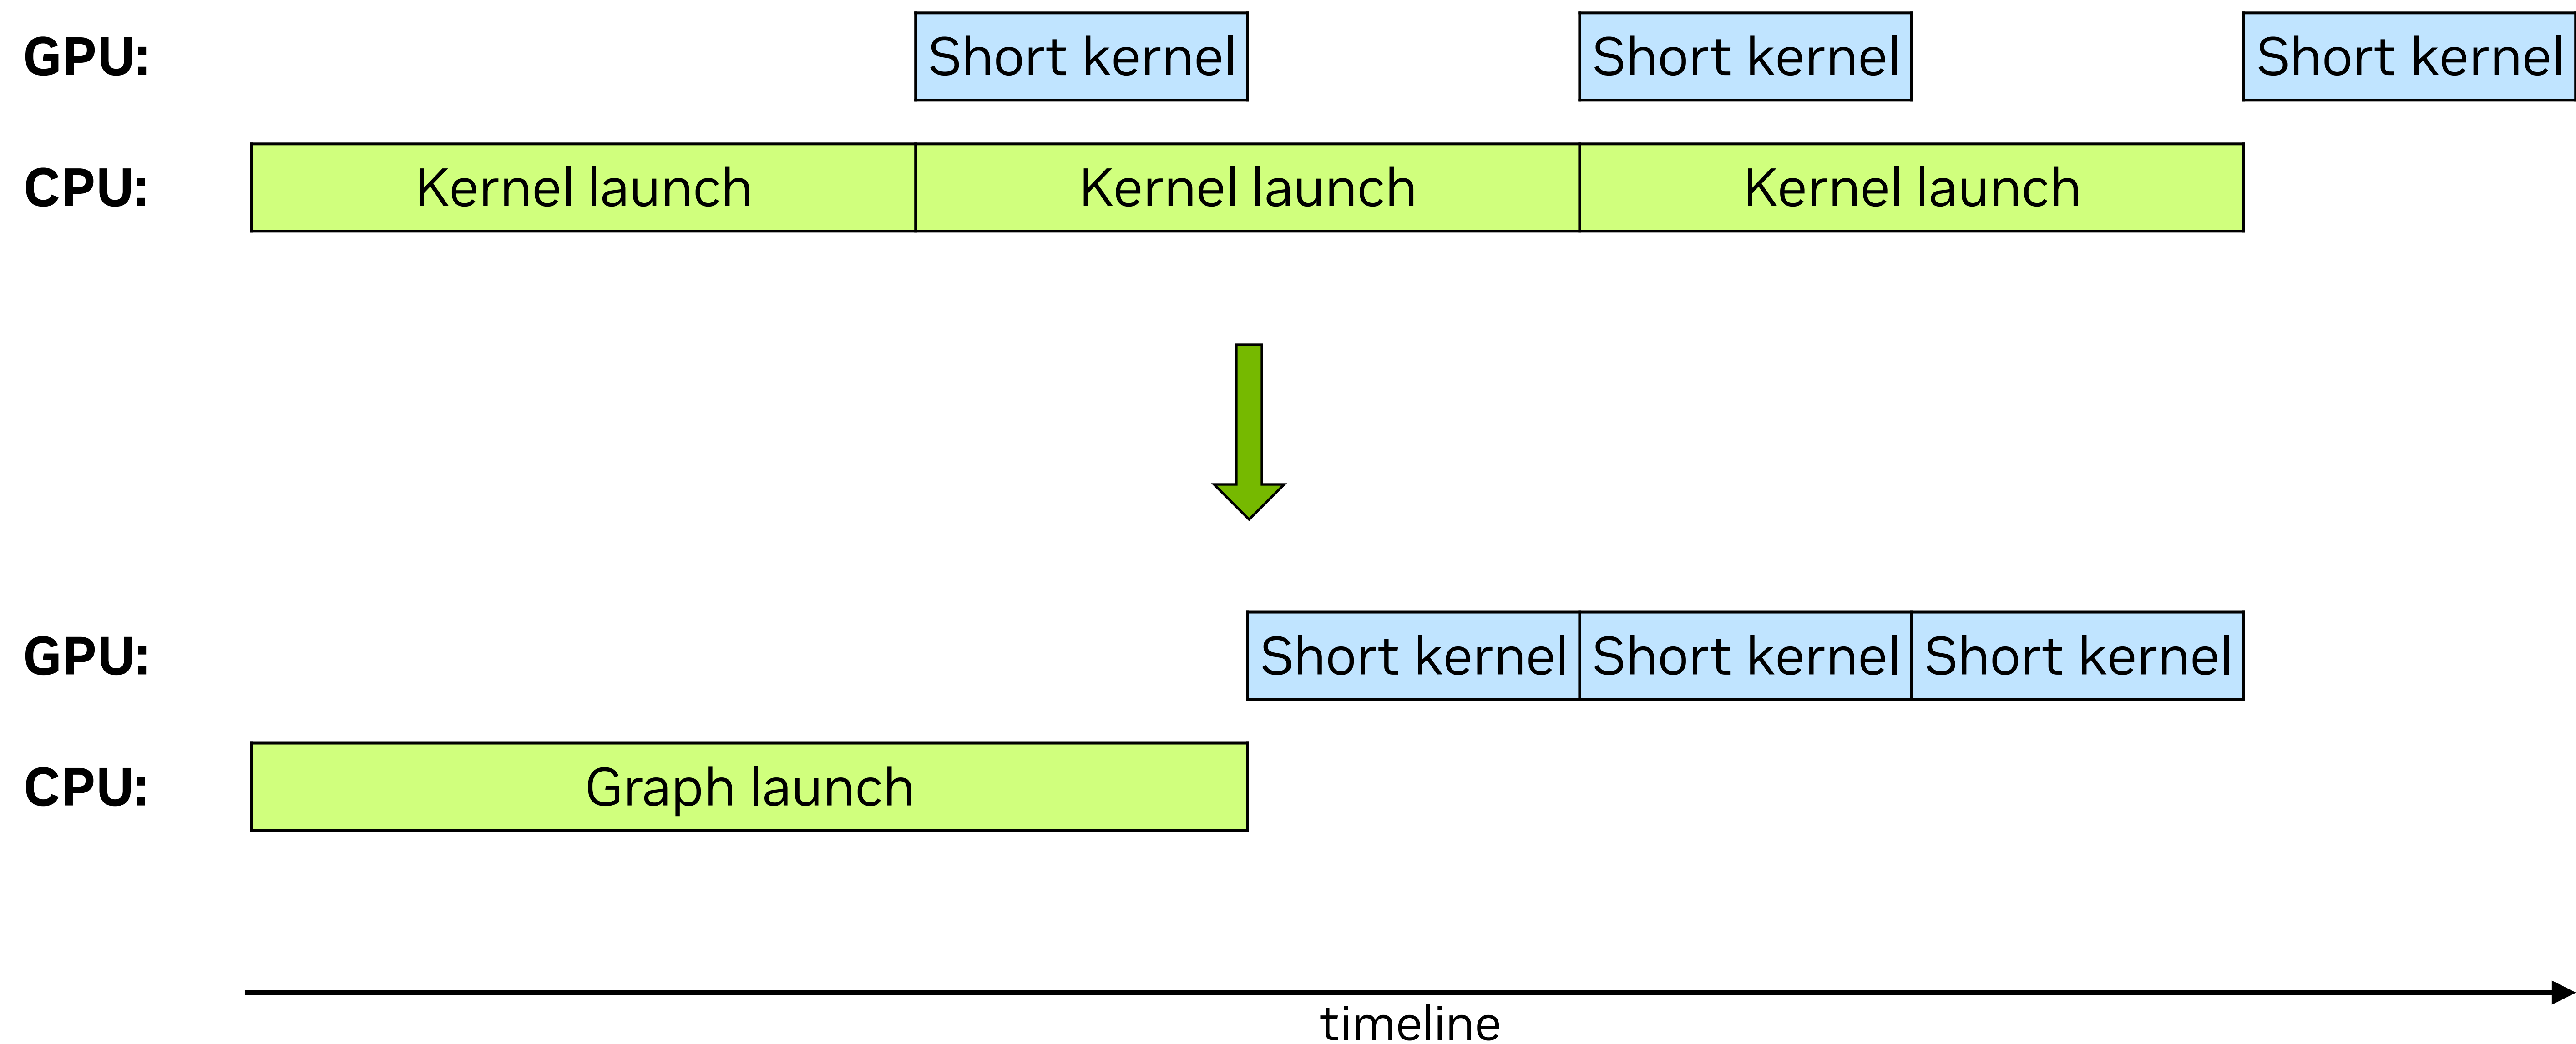
# CUDA Graphs



# CUDA Graphs

Reducing launch overhead


CUDA graphs reduce kernel launch overheads:



# CUDA Graphs

Simple example – capturing

```
kernel_A<<<..., stream_0>>>();  
cudaEventRecord(event_0, stream_0);  
kernel_B<<<..., stream_0>>>();  
cudaStreamWaitEvent(stream_1, event_0);  
kernel_C<<<..., stream_1>>>();  
cudaEventRecord(event_1, stream_1);  
cudaStreamWaitEvent(stream_0, event_1);  
kernel_D<<<..., stream_0>>>();
```





# CUDA Graphs

## Simple example – capturing

```
cudaStreamBeginCapture(stream_0, cudaStreamCaptureModeGlobal);  
kernel_A<<<..., stream_0>>>();  
cudaEventRecord(event_0, stream_0);  
kernel_B<<<..., stream_0>>>();  
cudaStreamWaitEvent(stream_1, event_0);  
kernel_C<<<..., stream_1>>>();  
cudaEventRecord(event_1, stream_1);  
cudaStreamWaitEvent(stream_0, event_1);  
kernel_D<<<..., stream_0>>>();  
cudaGraph_t graph;  
cudaStreamEndCapture(stream_0, &graph);
```

} “dry-run” during capture





# CUDA Graphs

## Simple example – steps

### 1. Capture graph (~ *write code*):

```
cudaStreamBeginCapture(stream_0, cudaStreamCaptureModeGlobal);  
<CAPTURED ACTIONS>  
cudaGraph_t graph;  
cudaStreamEndCapture(stream_0, &graph);
```

### 2. Create executable graph (~ *compile code*):

```
cudaGraphExec_t graph_exec;  
cudaGraphInstantiate(&graph_exec, graph);
```

### 3. Launch executable graph (~ *execute compiled code*):

```
cudaGraphLaunch(graph_exec, stream);
```

Graph encapsulates:

- Stream dependencies and concurrency
  - **Branches execute concurrently even though `cudaGraphLaunch` into a single stream**
  - Stream priorities can be inherited by nodes (requires `cudaGraphInstantiateFlagUseNodePriority` flag)
  - Streams used while capturing can be destroyed after capturing
- Kernel parameters
  - how do you change them?

# CUDA Graphs

## Simple example – graph modifications

Common time-stepping pattern:

```
for (int it = 0; it < nit; ++it) {  
    kernel<<<...>>>(ptr_prev, ptr_next);  
    // more kernels, etc.  
    std::swap(ptr_prev, ptr_next);  
}
```

Several approaches:

- I. Capture even- and odd-iteration graphs and `cudaGraphLaunch(graph_exec[it & 1], stream);`  
(see examples in [https://github.com/NVIDIA/multi-gpu-programming-models/tree/master/nccl\\_graphs](https://github.com/NVIDIA/multi-gpu-programming-models/tree/master/nccl_graphs))
- II. Always capture and `cudaGraphExecUpdate`
- III. Explicit graph construction API:  
`cudaStreamGetCaptureInfo`  
`cudaGraphAddKernelNode` – returns a node handle used by `cudaGraphExecKernelNodeSetParams` during iterations  
`cudaStreamUpdateCaptureDependencies`

**Topology changes typically require graph re-instantiation.**

Minor topology changes may not, e.g., turn node on/off (`cudaGraphNodeSetEnabled`)

# CUDA Graphs

## Other capture restrictions

Restrictions (not comprehensive list):

- Default legacy stream cannot be captured: `libraryCall(cudaStreamDefault) → libraryCall(stream)`
- Synchronous calls cannot be captured: `cudaMemcpy → cudaMemcpyAsync`
- Cannot synchronize: `cudaStreamSynchronize, cudaDeviceSynchronize`
- Host logic: `host_logic → cudaLaunchHostFunc(stream, host_logic)`

**Can capture what can be put in a stream.**

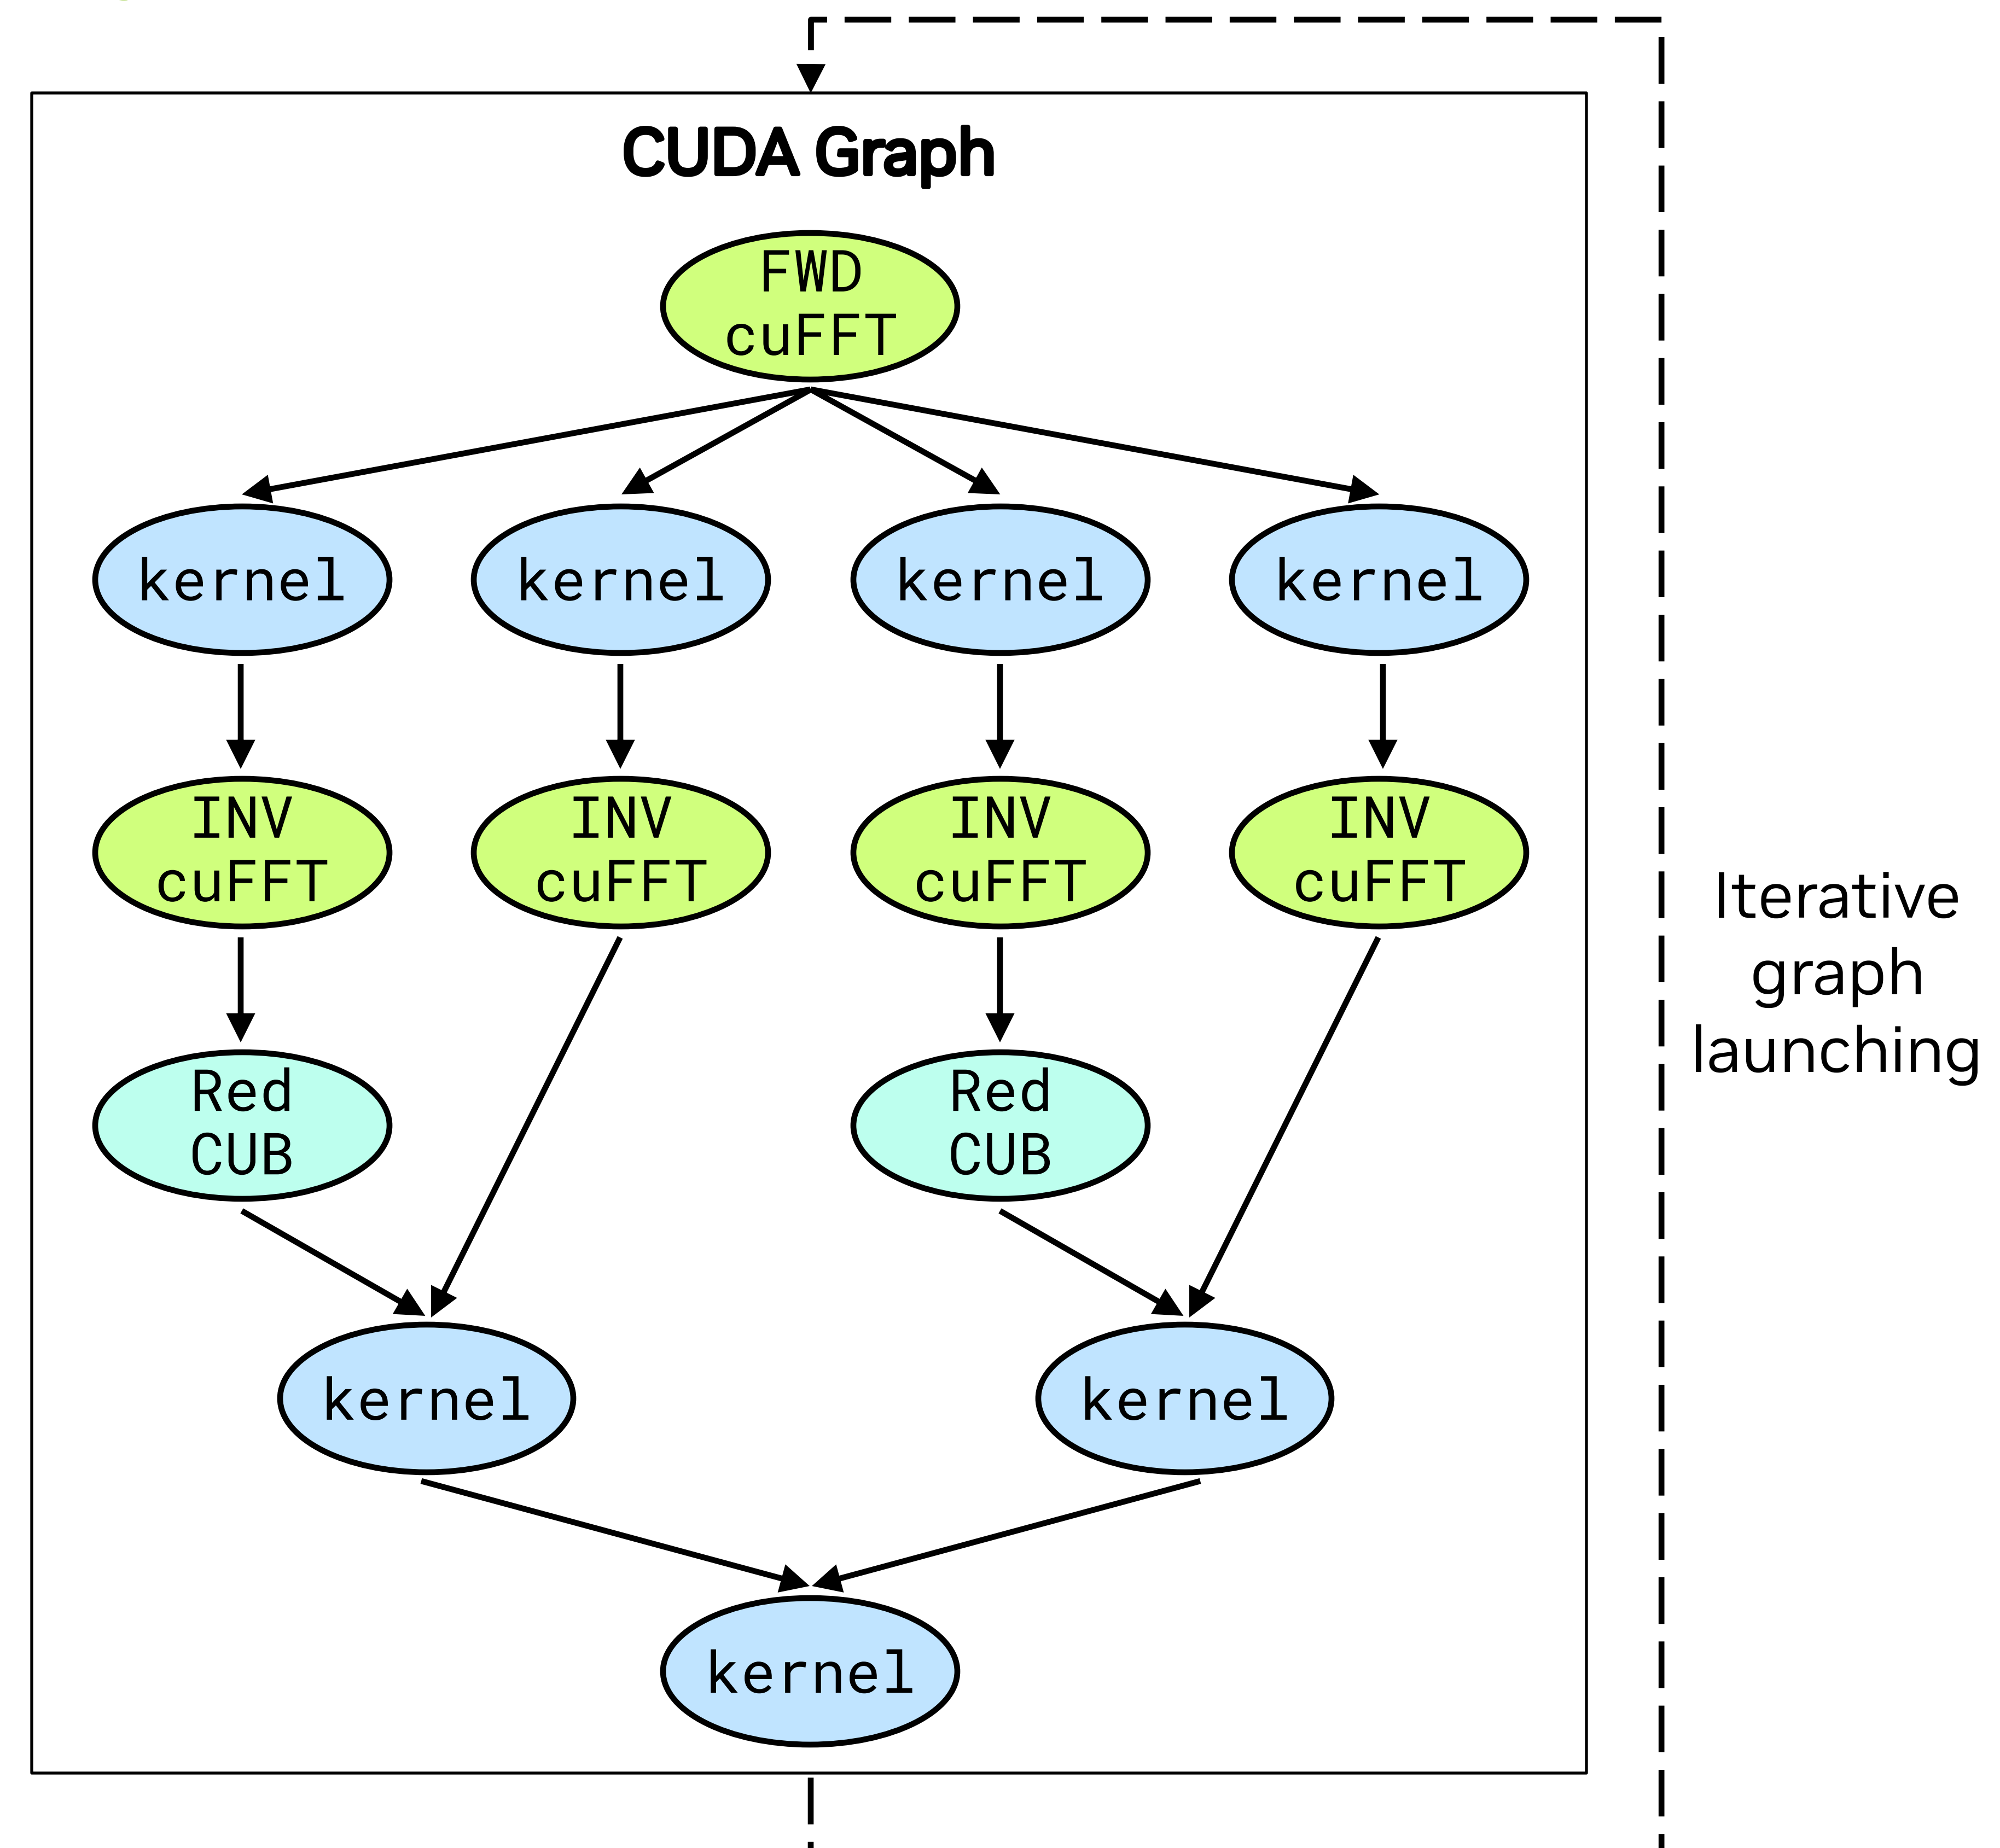


# CUDA Graphs

Example – algo

Workflow based on seismic trace processing:

1. Forward cuFFT
2. Frequency continuations into multiple traces with different frequency bands
3. Inverse cuFFTs of multiple traces (cannot batch due to variable sizes)
4. Reductions (CUB) and combining even/odd
5. Final combining



# CUDA Graphs

## Example – code

```
cudaStreamBeginCapture(f_stream, cudaStreamCaptureModeGlobal);
cufftExecR2C(f_plan, f_data, of_data); // f_plan is associated with f_stream
cudaEventRecord(f_event, f_stream);
for (int i = 0; i < m; ++i) {
    cudaStreamWaitEvent(b_stream[i], f_event);
    kernel_0<<<1, BX, 0, b_stream[i]>>>(b_data[i], of_data, of_n, b_n[i]);
    cufftExecC2R(b_plan[i], b_data[i], ob_data[i]); // b_plan[i] is associated with b_stream[i]
}
for (int i = 0; i < m; i += 2) {
    cub::DeviceReduce::Sum(cub_storage[i].first, cub_storage[i].second, ob_data[i], ocub + i, ob_n[i], b_stream[i]);
    cudaEventRecord(b_event[i], b_stream[i]);
    cudaStreamWaitEvent(b_stream[i + 1], b_event[i]);
    kernel_1<<<1, BX, 0, b_stream[i + 1]>>>(ob_data[i + 1], ocub + i, ob_n[i + 1]);
    cudaEventRecord(b_event[i + 1], b_stream[i + 1]);
    cudaStreamWaitEvent(f_stream, b_event[i + 1]);
}
kernel_2<<<m, BX, 0, f_stream>>>(f_data, ob_data_gpu, ob_n_gpu, f_n);
cudaStreamEndCapture(f_stream, &graph);
```

**No extra effort  
to capture library calls !!  
(into sub-graphs)**

# CUDA Graphs

## Example – performance

Performance:


- 1024 samples input trace.
- 100 output traces 128-8192 sample range.

Graphs provide **5X speedup** (H200 test).

### Graphs: packed and compact



No-graphs: almost useless to employ multiple streams because launch overheads “serialize” kernels execution

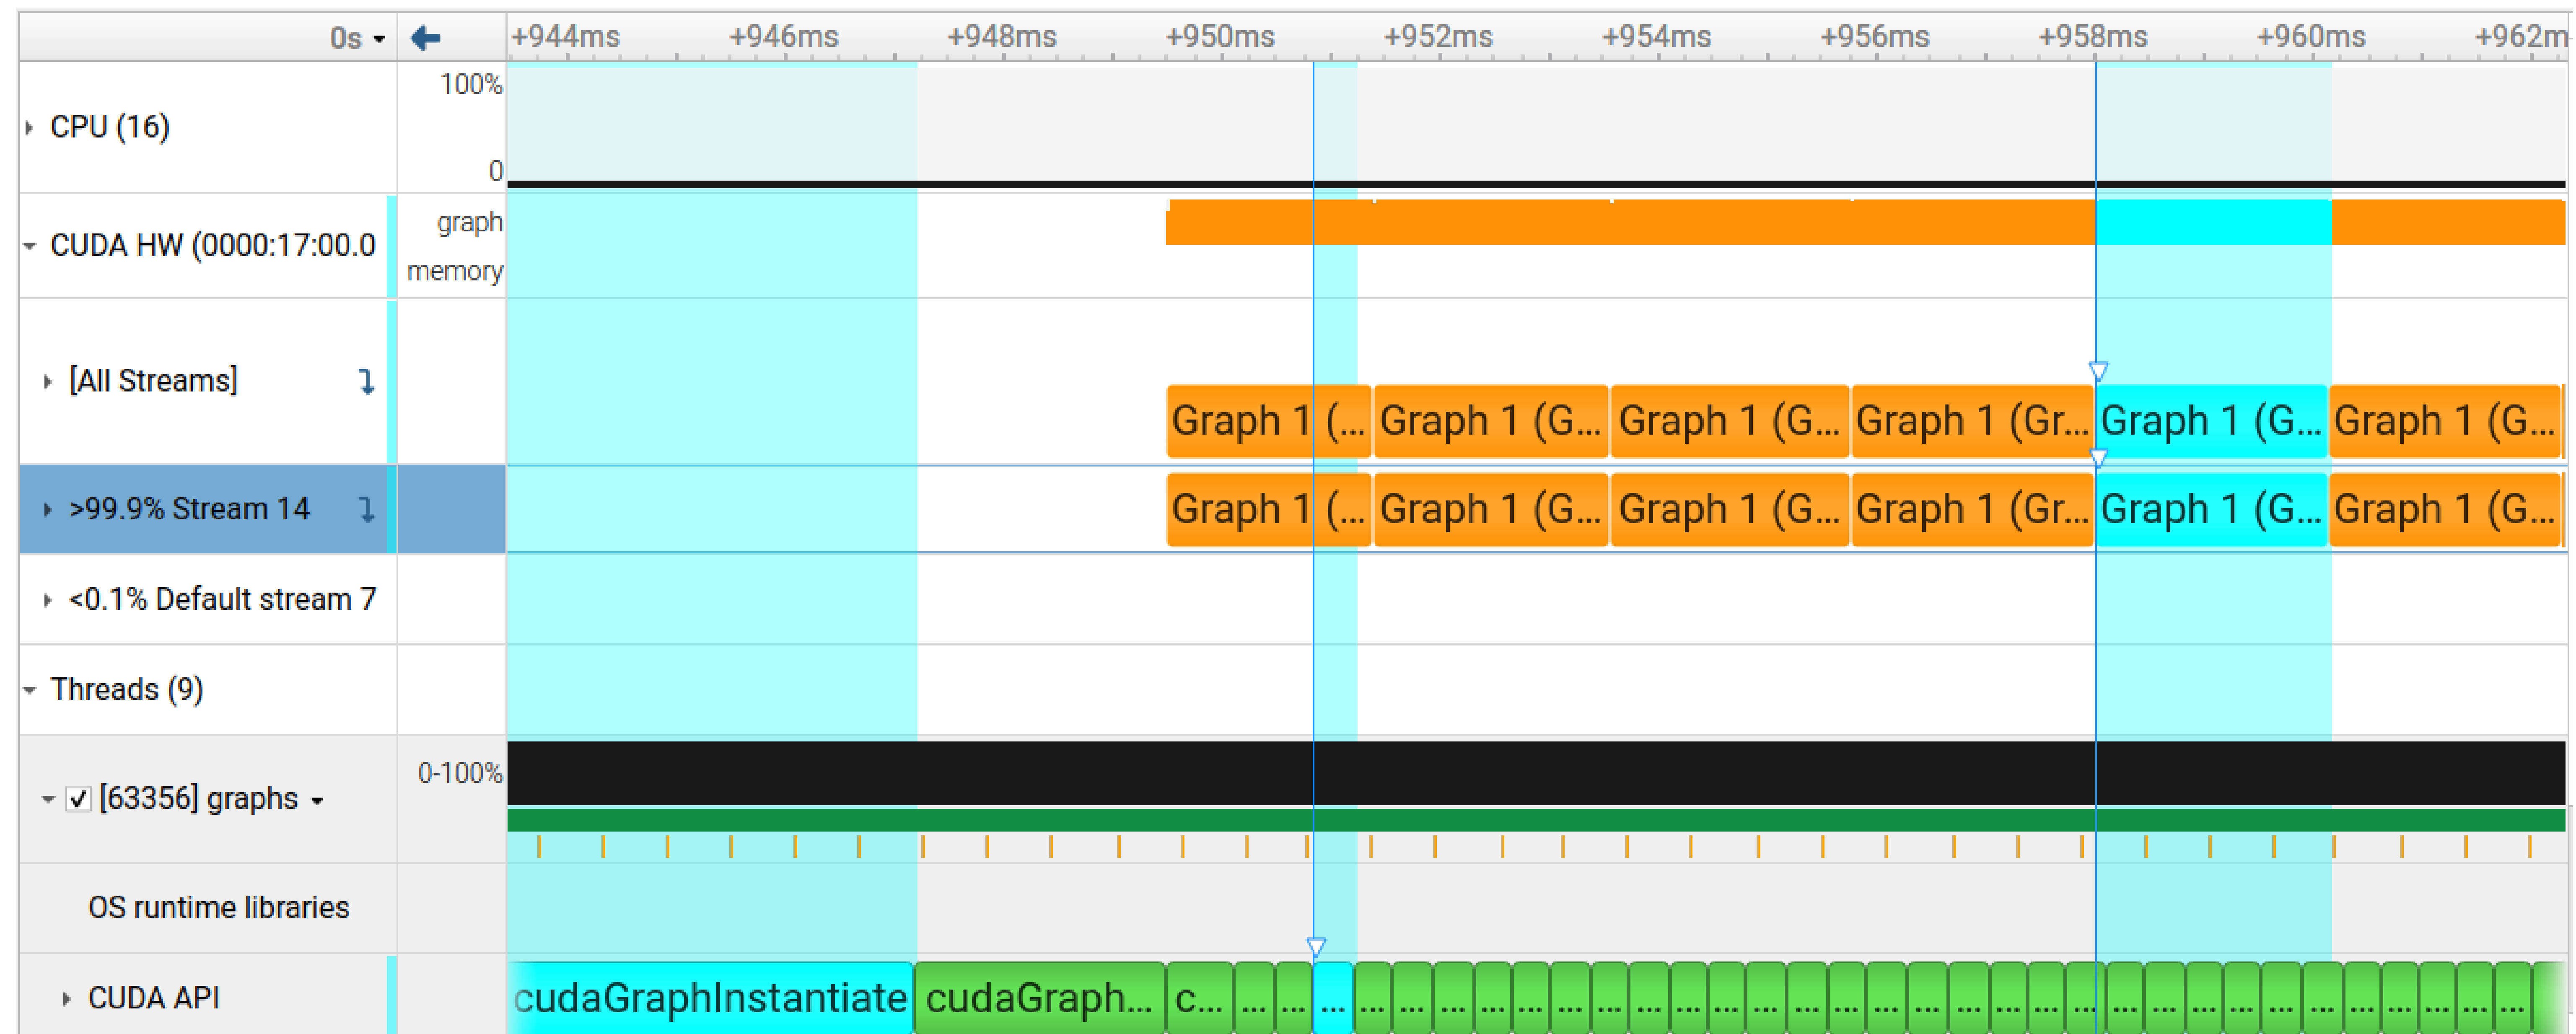




# CUDA Graphs

Nsight systems

- To trace nodes inside graphs, use `--cuda-graph-trace node` nsys option  
Note: has higher profiling overhead than default.
- Default is `--cuda-graph-trace graph` option:



# CUDA Graphs

## Further reading

CUDA graphs API is rich, we just looked at tip of the iceberg.

See also at this GTC:

- [“Multi GPU Programming Models for HPC and AI”](#) by Jiri Kraus [S61339]  
and accompanying examples: [https://github.com/NVIDIA/multi-gpu-programming-models/tree/master/nccl\\_graphs](https://github.com/NVIDIA/multi-gpu-programming-models/tree/master/nccl_graphs)
- [“Accelerating Drug Discovery: Optimizing Dynamic GPU Workflows with CUDA Graphs, Mapped Memory, C++ Coroutines, and More”](#) by Jiqun Tu & Ellery Russell [S61156]

Related resources:

- [GTS’23 \[S51211\]: “CUDA Graphs 101”](#)
- [GTC’21 \[S32082\]: “Effortless CUDA Graphs”](#)
- [Technical Blog: “Constructing CUDA Graphs with Dynamic Parameters”](#)
- [Technical Blog: “Getting Started with CUDA Graphs”](#)
- [CUDA Samples:](#)
  - [simpleCudaGraphs](#)
  - [jacobiCudaGraphs](#)
  - [graphMemoryNodes](#)
  - [graphMemoryFootprint](#)





# Summary



# Summary

## And take-aways

- **Memory model:**
  - What every CUDA programmer should know about memory...
- **Asynchronous memory copies & barriers:**
  - Increase bytes in flight + less pressure (reg, pred, MIO, etc.)
  - Require some programming effort
- **Compressible memory:**
  - Improves effective GMEM bandwidth (if data is compressible)
  - Can be a very low-hanging fruit: just replace cudaMalloc-s
  - Danger: sensitive to access patterns (**can decrease performance**); may require kernel changes to be effective
- **CUDA graphs:**
  - Way to go for short-timed kernels
  - Can be a low-hanging fruit: stream capturing is just a few lines





# **Core Performance Optimization Techniques**



# Core Performance Optimization Techniques at GTC'24

## List of presentations

- Introduction to CUDA Programming and Performance Optimization [S62191]
- **Advanced Performance Optimization in CUDA [S62192]**
- Performance Optimization for Grace CPU Superchip [S62275]
- Grace Hopper Superchip Architecture and Performance Optimizations for Deep Learning Applications [S61159]
- Multi GPU Programming Models for HPC and AI [S61339]
- More Data, Faster: GPU Memory Management Best Practices in Python and C++ [S62550]
- Harnessing Grace Hopper's Capabilities to Accelerate Vector Database Search [S62339]
- From Scratch to Extreme: Boosting Service Throughput by Dozens of Times with Step-by-Step Optimization [S62410]



