Vision AI Demystified

Ayesha Asif, Developer Marketing | GTC March 2024

Agenda

- The History of Computer Vision
- From CNNs to Vision Transformers
- Generative AI: GANs and Diffusion Models
- 3D Reconstruction with NeRFs and Gaussian Splatting
- Emerging Computer Vision Frontiers

A Brief History of Computer Vision

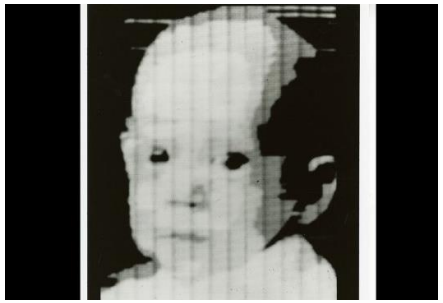

1957

The first digitally scanned photos

Russell Kirsch

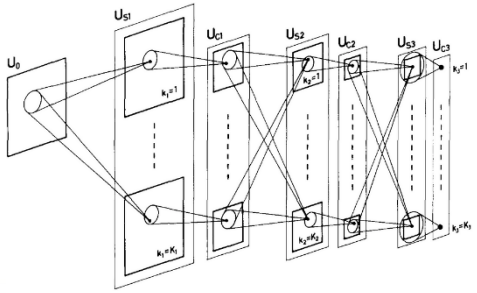

1970

Neocognitron

Kunihiko Fukushima

2001

The Viola-Jones algorithm

Paul Viola and Michael Jones

2012-2024

Deep Learning

Vision Transformers

Diffusion Models

Generative AI

3D Reconstruction

1963

Machine Perception of Three-Dimensional Solids

Lawrence Roberts

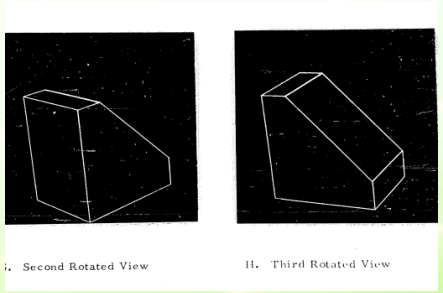

1989

LeNet-5

Yann LeCun

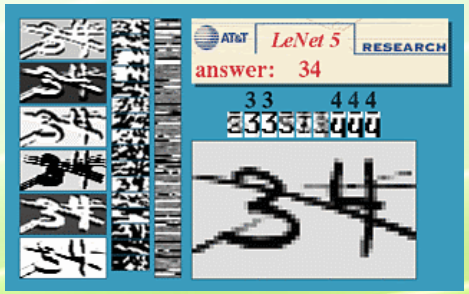

2009

ImageNet Dataset

Fei-Fei Li

Convolutional Neural Networks

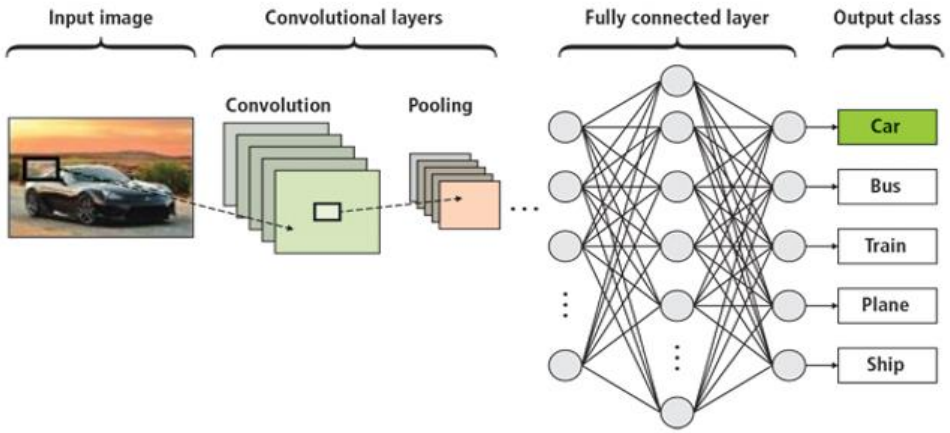

CNN Architecture

Hierarchical feature extraction and classification

Convolutional Neural Networks and Transformers

What are Vision Transformers (ViTs)?

CNN Architecture

Hierarchical feature extraction and classification

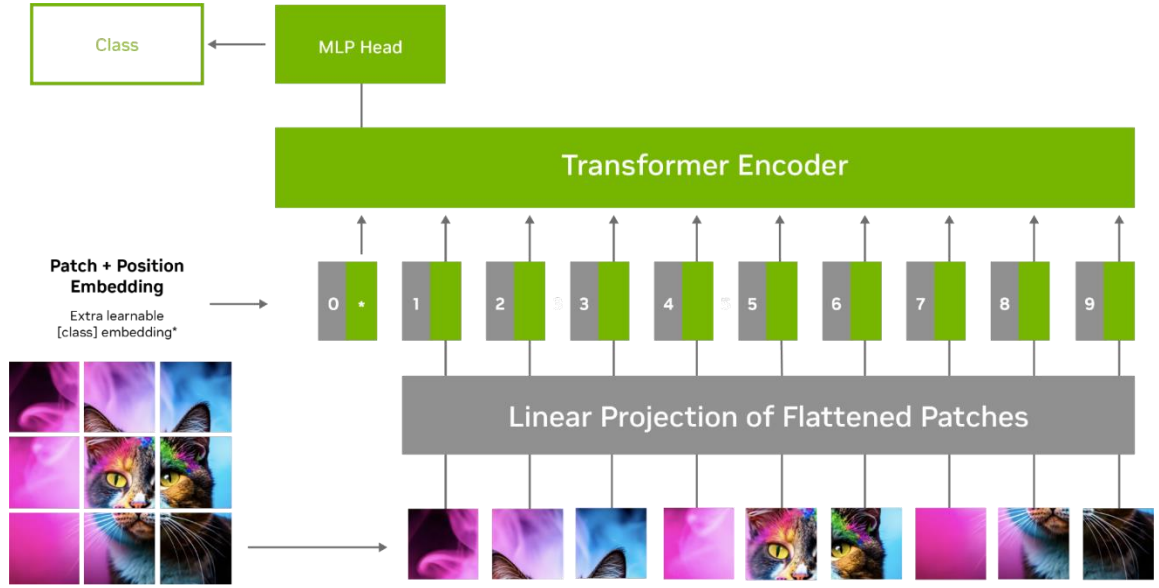

Vision Transformer Architecture

Attention-based patch embedding and classification

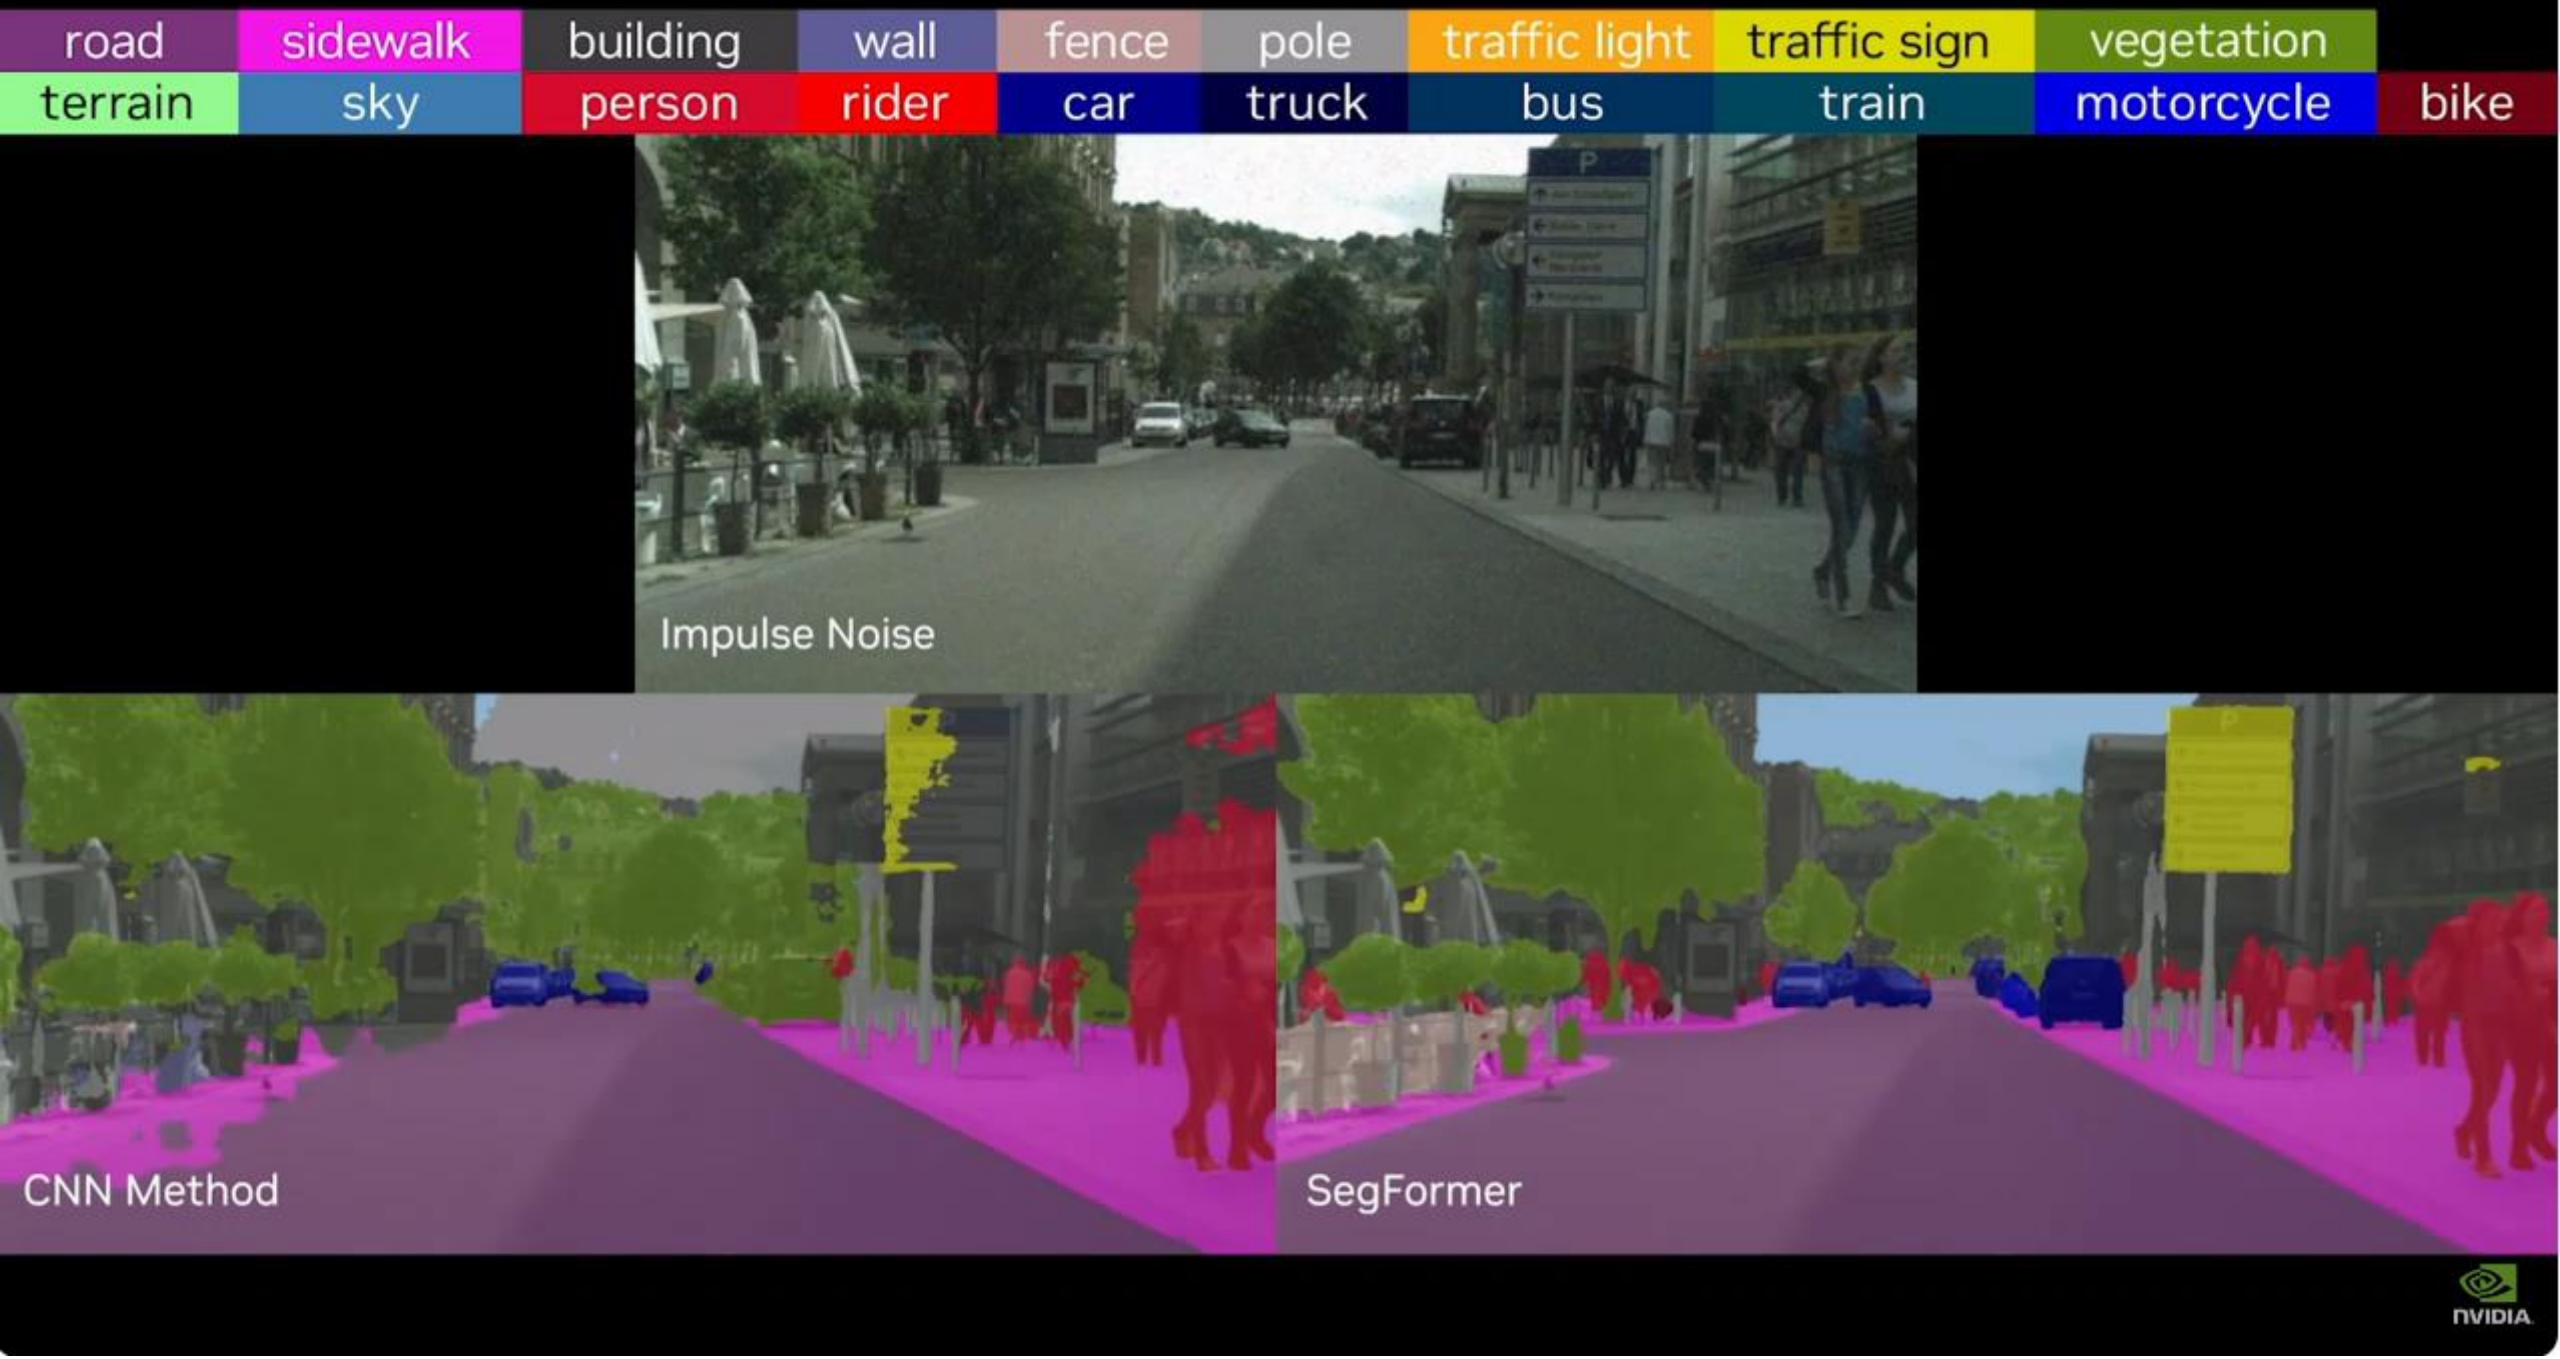
road
sidewalk
building
wall
fence
pole
traffic light
traffic sign
vegetation
terrain
sky
person
rider
car
truck
bus
train
motorcycle
bike
Impulse Noise
CNN Method
SegFormer

Comparison of CNN and ViT Methods

Vision Transformers

self-**DI**stillation with **NO** labels (DINO)

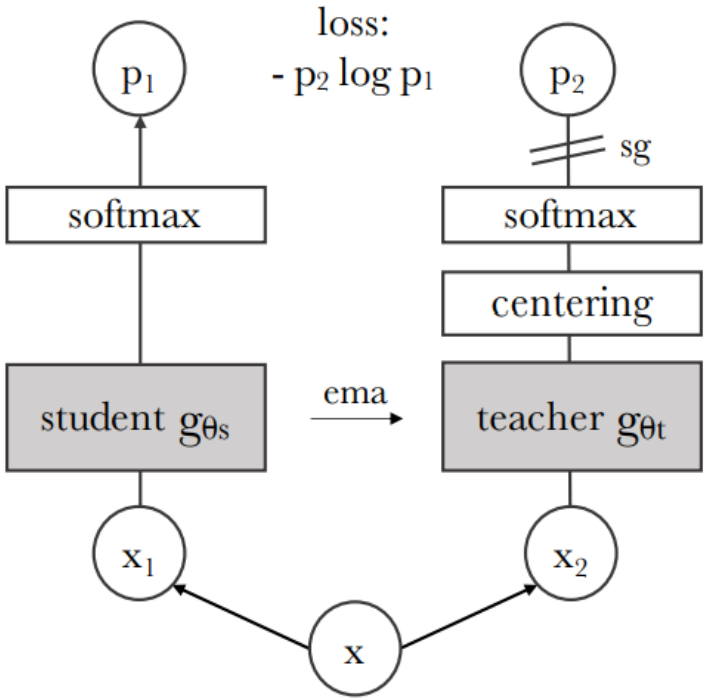

DINO

- Self-supervised learning method
- Same architecture for teacher and student networks
- Student network – matches the output distribution of a teacher network
- Teacher network – learns from the student parameters through exponential moving average (EMA)

Vision Transformers

Segment Anything Model (SAM)

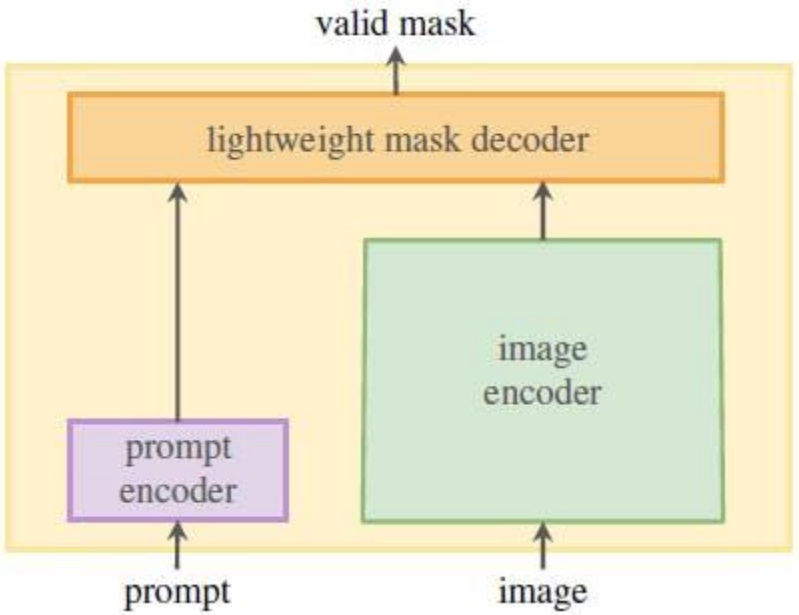

SAM

- SAM was trained on an unprecedented dataset of 11 million images and 1.1 billion segmentation masks
- Strong zero-shot performance
- Promptable segmentation

Vision Transformers

Contrastive Language-Image Pretraining (CLIP)

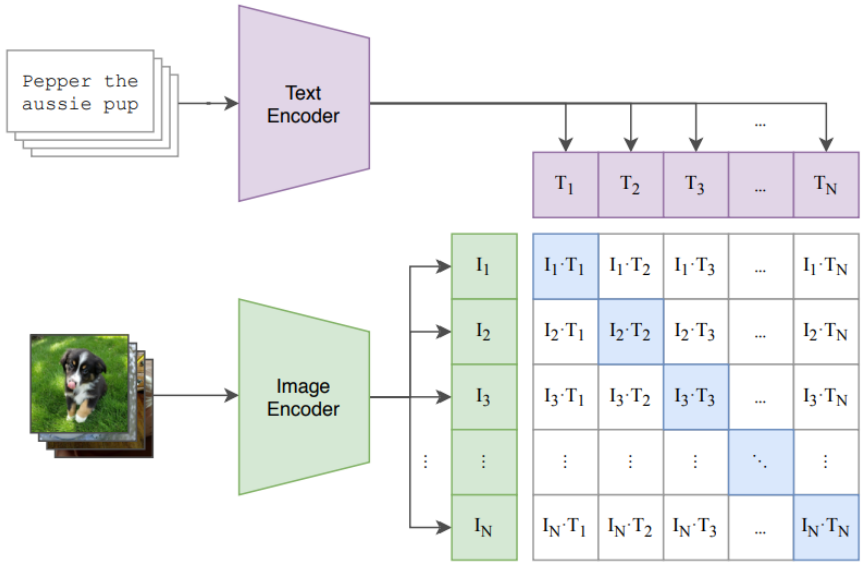

CLIP

- Trained on 400 million (image, text) pairs
- Jointly trains an image and text encoder to predict the correct pairings of a (image, text) examples
- Projected in a shared embedding space to understand the text/image relationships
- Applications – image search, image captioning

Getting Started With Vision Transformers

DINO Notebook

#Installing the TAO launcher

```
!pip3 install nvidia-pyindex
!pip3 install nvidia-tao
```

#Pull pre-trained model from NGC

```
!ngc registry model download-version
nvidia/tao/pretrained_dino_nvimagenet:fan_small_hybri
d_nvimagenet --dest $LOCAL_PROJECT_DIR/dino/
```

#Run Inference of a pre-trained model

```
tao model dino inference -e /path/to/spec.yaml -r
/path/to/results/
inference.checkpoint=/path/to/model.pth
```

Finetuning on Custom Datasets:

- Tao model dino train
- Tao model dino evaluate

Generative Models

Generative Models

Generative Adversarial Networks (GANs)

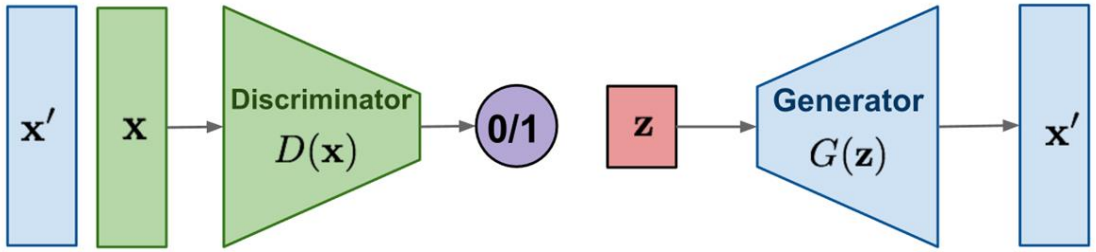

GAN Architecture

Adversarial training

- The generator creates synthetic data, and the discriminator tries to distinguish the generated data from real data
- Training instability and mode collapse issues

Generative Models

From GANs to Diffusion Models

GAN Architecture

Adversarial training

- The generator creates synthetic data, and the discriminator tries to distinguish the generated data from real data
- Training instability and mode collapse issues

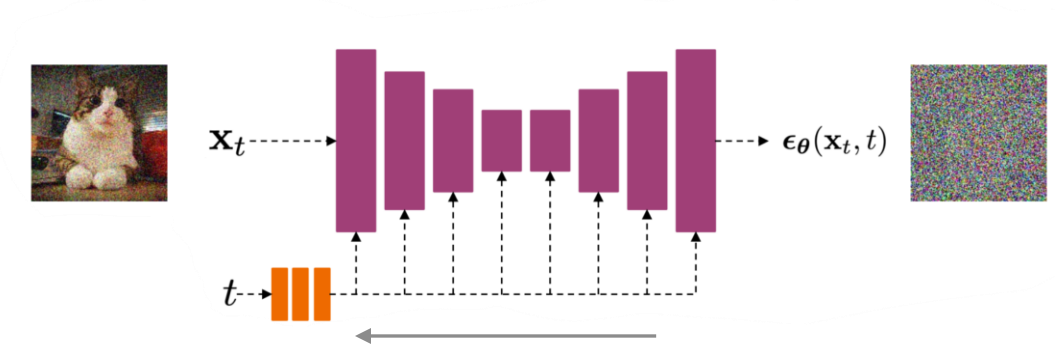

Diffusion Model Architecture

Gradual addition and removal of noise

- Learn to reverse a diffusion process by iteratively removing noise from data
- Can generate high-quality samples across various domains

Generative Models

NVIDIA Edify

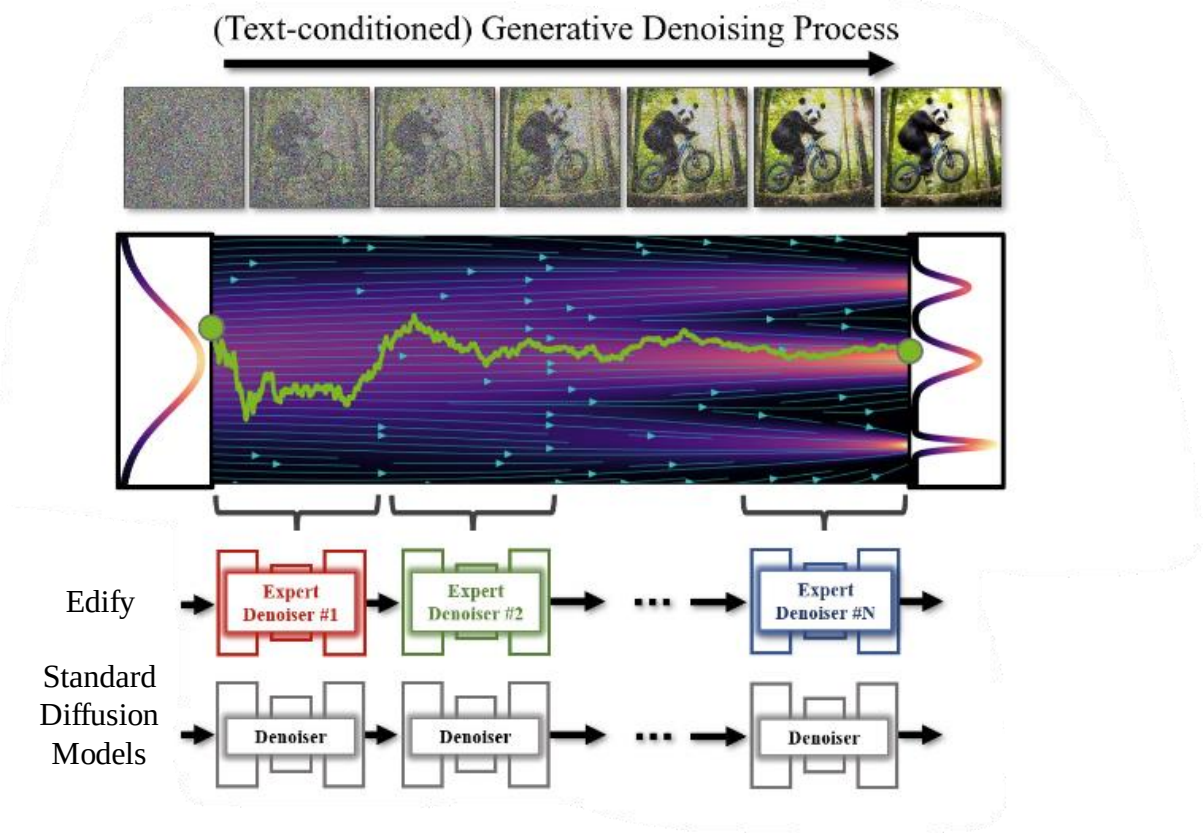

Wildlife photography of wolf, beautiful eyes, golden hour

- Instead of using a single denoiser, Edify trains an ensemble of expert denoising networks
- Edify also combines text embeddings from both T5 and CLIP models instead of using just one text encoder

Generative AI by gettyimages

Built using NVIDIA Picasso

- State-of-the-art NVIDIA Edify model architecture for 4K generative photography
- Commercially safe—trained exclusively on Getty Images licensed data; uncapped indemnification
- Access via Web or iStock.com
- Advanced visual editing APIs for adding subjects, expanding images or replacing specific element

nature photography of a rock arch, a mountain lake with a forest and fog in the background, overcast, majestic




Getting Started With Diffusion Models

A photo of a Shiba Inu dog with a backpack riding a bike

```
import requests

# Call model from NVIDIA NGC
invoke_url = "https://api.nvcf.nvidia.com/v2/nvcf/pexec/functions/89848fb8-549f-41bb-88cb-95d6597044a4"
fetch_url_format = "https://api.nvcf.nvidia.com/v2/nvcf/pexec/status/"

# Provide API key here
headers = {
"Authorization": "Bearer $API_KEY_REQUIRED_IF_EXECUTING_OUTSIDE_NGC",
"Accept": "application/json",
}

# Provide prompts and parameters
payload = {
"prompt": "A photo of a Shiba Inu dog with a backpack riding a bike",
"negative_prompt": "beach",
"inference_steps": 25
}

# re-use connections
session = requests.Session()

response = session.post(invoke_url, headers=headers, json=payload)

while response.status_code == 202:
        response = session.get(fetch_url_format + response.headers.get("NVCF-REQID"), headers=headers)

response.raise_for_status()
response_body = response.json()
print(response_body)
```

catalog.ngc.nvidia.com/orgs/nvidia/teams/ai-foundation/models/sdxl

Getting Started With Diffusion Models

A photo of a Shiba Inu dog with a backpack riding a bike

```
import requests

# Call model from NVIDIA NGC
invoke_url = "https://api.nvcf.nvidia.com/v2/nvcf/pexec/functions/89848fb8-549f-41bb-88cb-95d6597044a4"
fetch_url_format = "https://api.nvcf.nvidia.com/v2/nvcf/pexec/status/"

# Provide API key here
headers = {
"Authorization": "Bearer $API_KEY_REQUIRED_IF_EXECUTING_OUTSIDE_NGC",
"Accept": "application/json",
}

# Provide prompts and parameters
payload = {
"prompt": "A photo of a Shiba Inu dog with a backpack riding a bike",
"negative_prompt": "beach",
"inference_steps": 25
}

# re-use connections
session = requests.Session()

response = session.post(invoke_url, headers=headers, json=payload)

while response.status_code == 202:
        response = session.get(fetch_url_format + response.headers.get("NVCF-REQID"), headers=headers)

response.raise_for_status()
response_body = response.json()
print(response_body)
```

catalog.ngc.nvidia.com/orgs/nvidia/teams/ai-foundation/models/sdxl

Getting Started With Diffusion Models

A photo of a Shiba Inu dog with a backpack riding a bike

```
import requests

# Call model from NVIDIA NGC
invoke_url = "https://api.nvcf.nvidia.com/v2/nvcf/pexec/functions/89848fb8-549f-41bb-88cb-95d6597044a4"
fetch_url_format = "https://api.nvcf.nvidia.com/v2/nvcf/pexec/status/"

# Provide API key here
headers = {
"Authorization": "Bearer $API_KEY_REQUIRED_IF_EXECUTING_OUTSIDE_NGC",
"Accept": "application/json",
}

# Provide prompts and parameters
payload = {
"prompt": "A photo of a Shiba Inu dog with a backpack riding a bike",
"negative_prompt": "beach",
"inference_steps": 25
}

# re-use connections
session = requests.Session()

response = session.post(invoke_url, headers=headers, json=payload)

while response.status_code == 202:
        response = session.get(fetch_url_format + response.headers.get("NVCF-REQID"), headers=headers)

response.raise_for_status()
response_body = response.json()
print(response_body)
```

catalog.ngc.nvidia.com/orgs/nvidia/teams/ai-foundation/models/sdxl

Getting Started With Diffusion Models

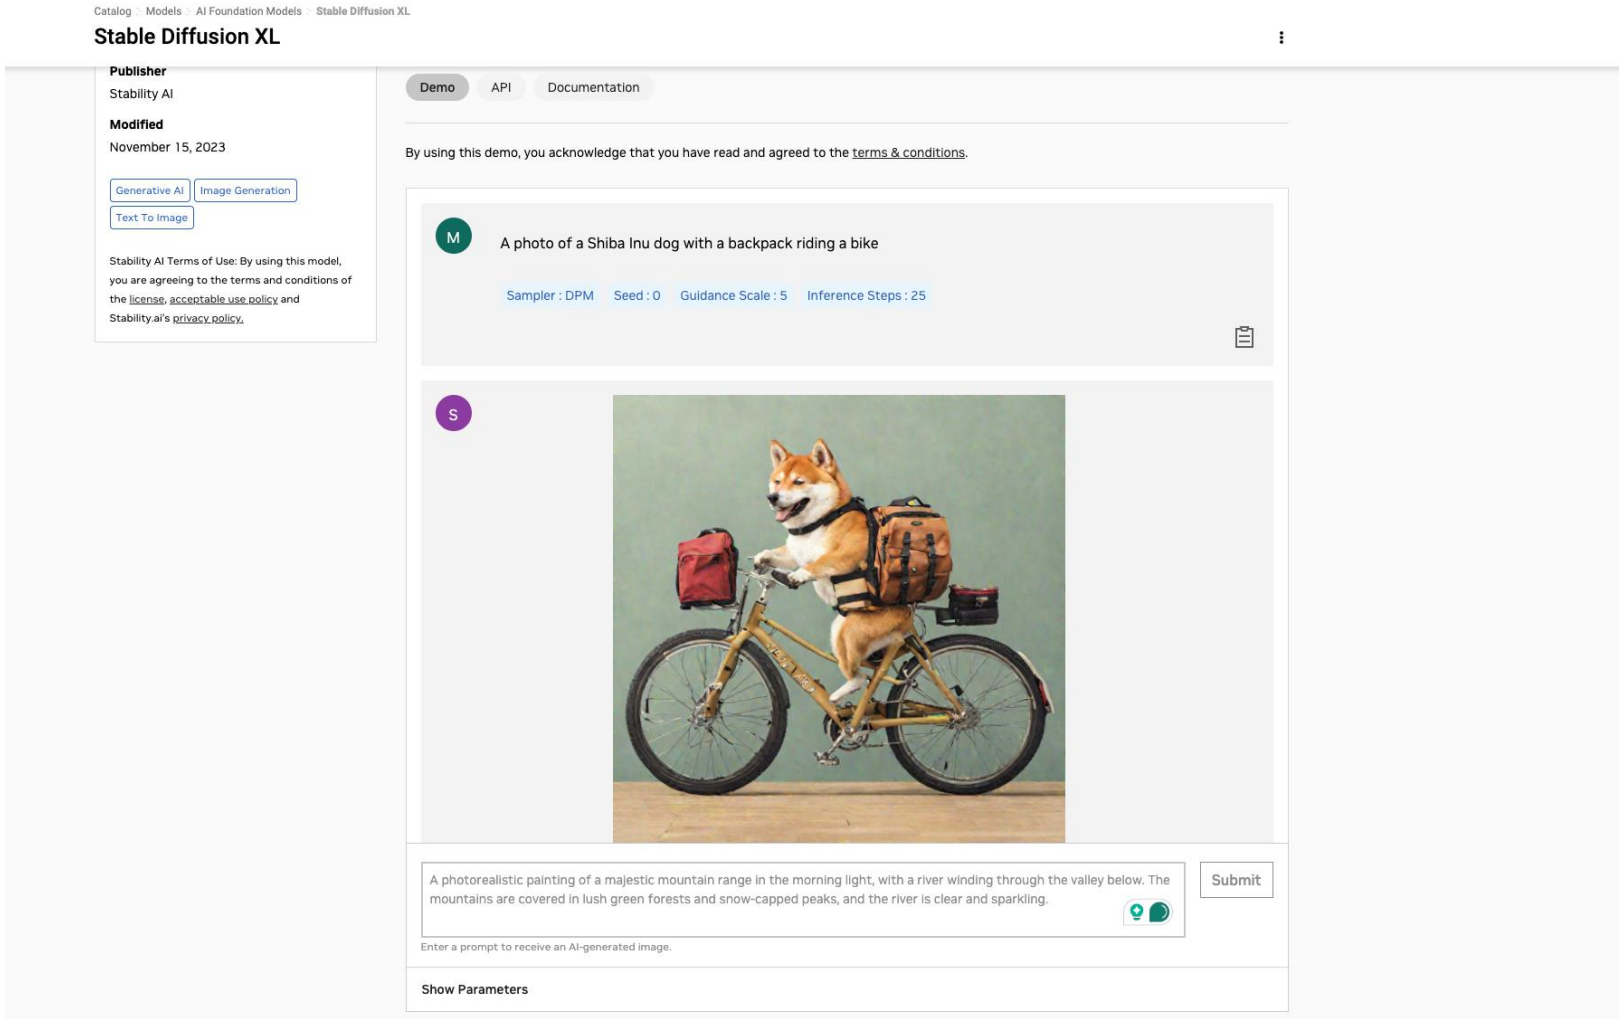

catalog.ngc.nvidia.com/orgs/nvidia/teams/ai-foundation/models/sdxl

3D Reconstruction

3D Reconstruction

Structure from Motion (SfM)

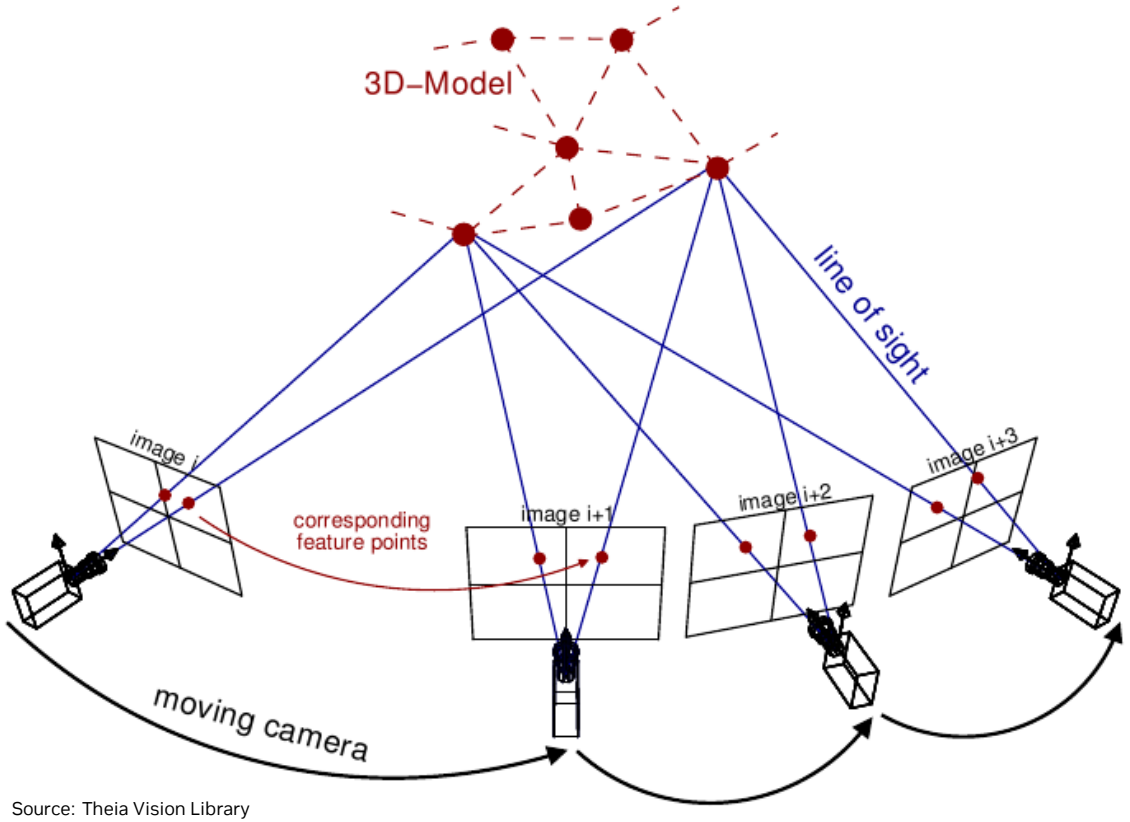

Source: Theia Vision Library

- An imaging technique for estimating 3D structures from 2D image sequences
- SfM tracks features (e.g. corner points, edges) across multiple images to find correspondences.
- This method struggles to capture fine details and complex geometry and struggles with certain materials

3D Reconstruction

Neural Radiance Fields (NeRFs)

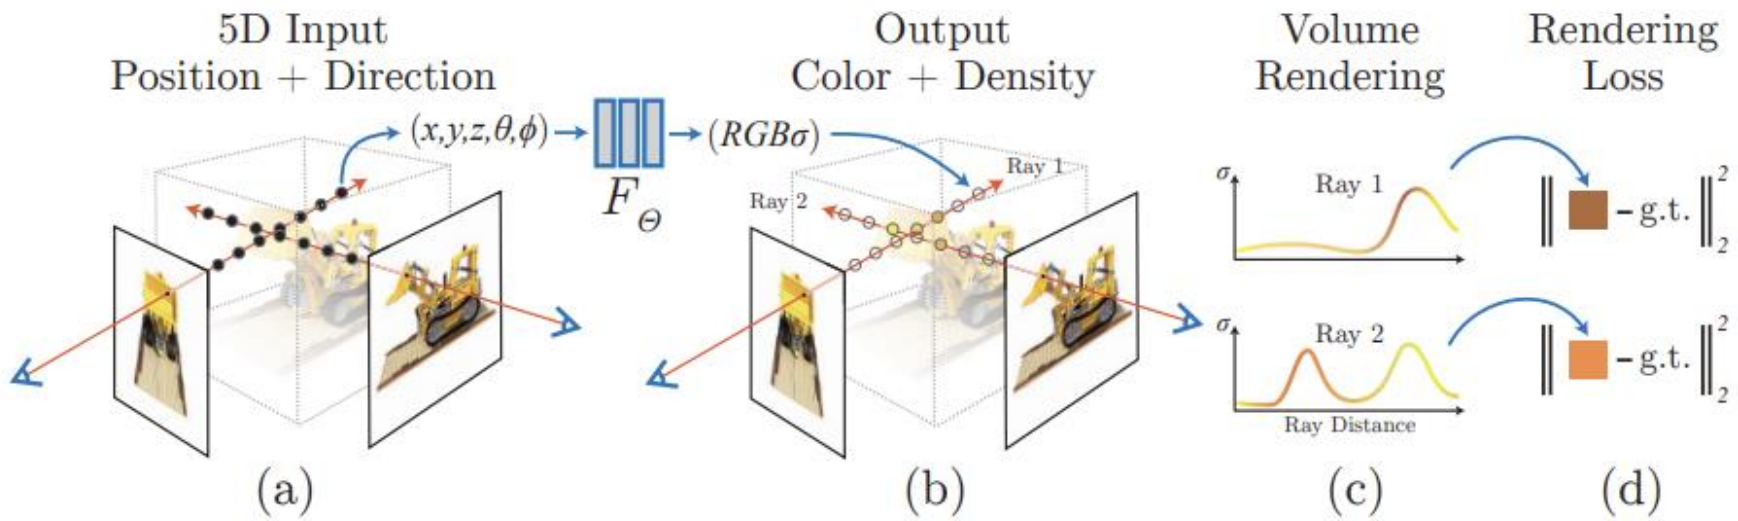

NeRF architecture

- NeRFs can generate novel views of complex 3D scenes, based on a partial set of 2D images
- Several decoupled modules - Sampling, encoding, neural network, and rendering
- Diverse applications spanning virtual/augmented reality, medical imaging, robotics, and autonomous navigation

3D Reconstruction

Neuralangelo

3D Online Shopping Experiences Powered by NeRFs

RECON Labs Inc.

3D Reconstruction

3D Gaussian Splatting

3D Gaussian Splatting for Real-Time Radiance Field Rendering

3D Reconstruction

3D Gaussian Splatting

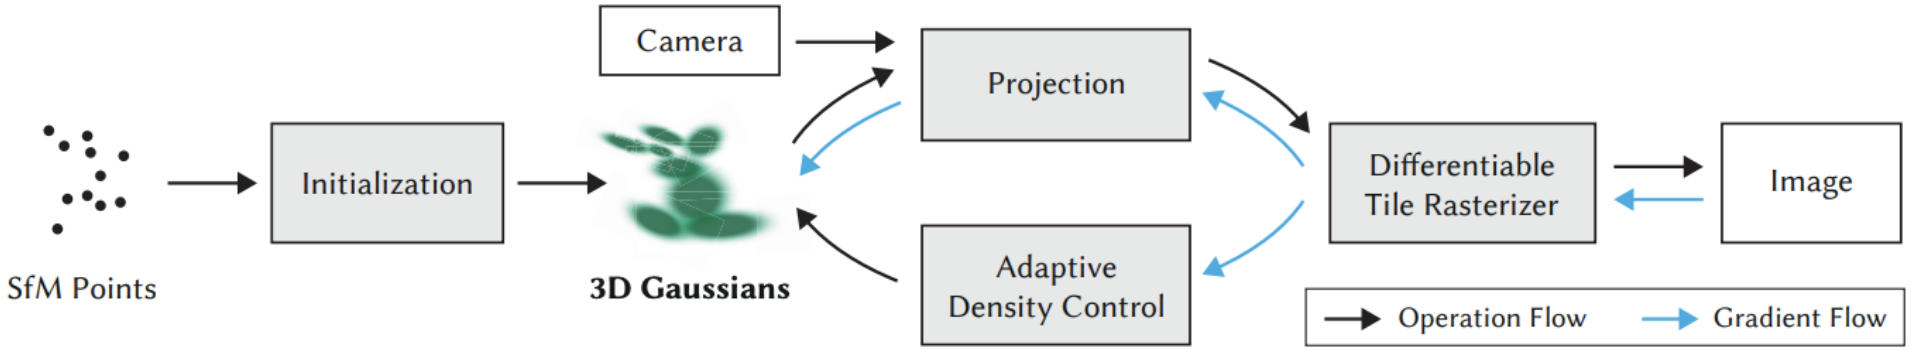

Point Cloud Generation

- 3D sparse point are generated using 3D points

Gaussian Placement

- Each point transformed into a 3D Gaussian function defined by position, spread, color and opacity

Training & Optimization

- Optimizes parameter of each gaussian points (position, spread, color, opacity)

Rendering

- Optimized points are projected on the screen based on camera's viewpoint

3D Gaussian Splatting for Real-Time Radiance Field Rendering

3D Reconstruction

3D Gaussian Splatting

Instant-NGP

Gaussian Splatter

[3D Gaussian Splatting for Real-Time Radiance Field Rendering]

Getting Started With Instant-NGP & 3D Gaussian Splatting

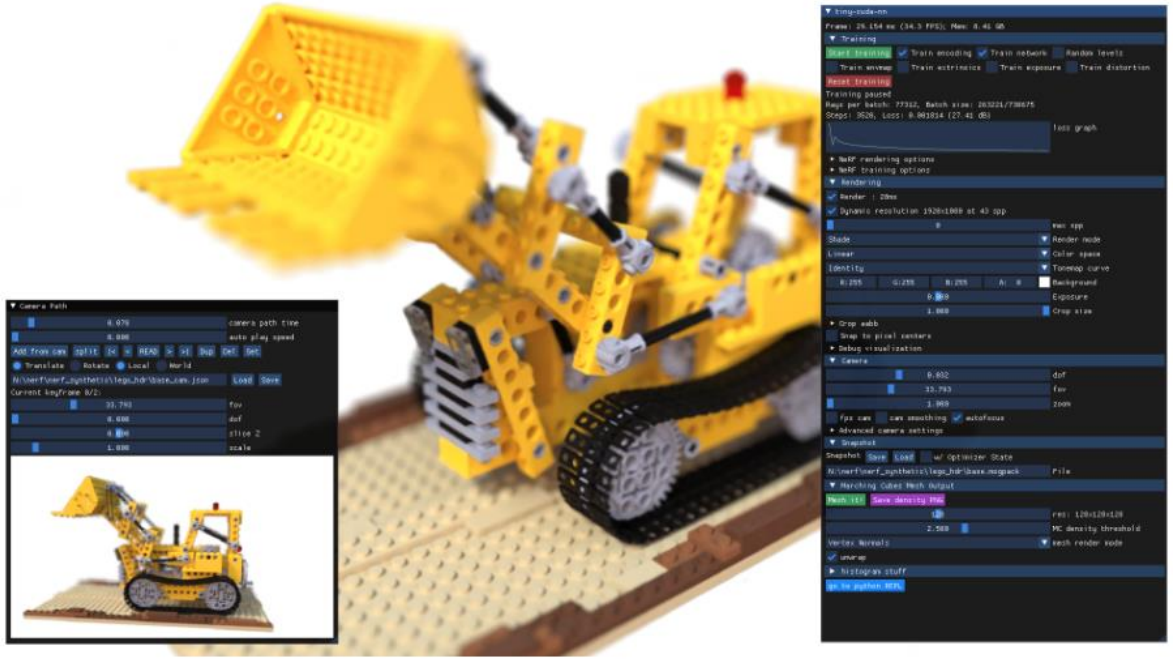

Video Tutorial by Eric Hanes, NVIDIA

Step by Step Instructions in GitHub

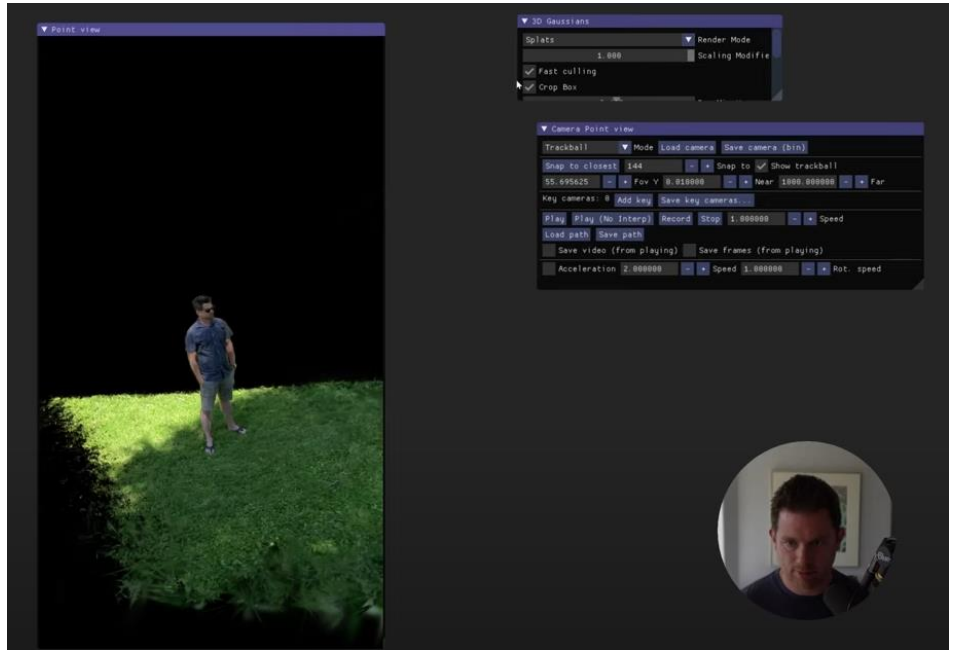

Video Tutorial by The NeRF Guru

Step by Step Instructions in GitHub

Emerging Computer Vision Frontiers

Emerging Computer Vision Frontiers

Text-2-Image with Character Consistency

Emerging Computer Vision Frontiers

Text-to-4D (3D in motion)

https://research.nvidia.com/labs/toronto-ai/AlignYourGaussians/samples.html

Other Relevant GTC Sessions

- [S62724] Revolutionizing Vision AI: From 2D to 3D Worlds
- [S62818] The Visionaries: A Cross-Industry Exploration of Computer Vision
- [CWE63456] Build Accelerated Computer Vision Microservices With NVIDIA Libraries and SDKs
- [S62624] The Vision-AI Revolution powered by DeepStream